



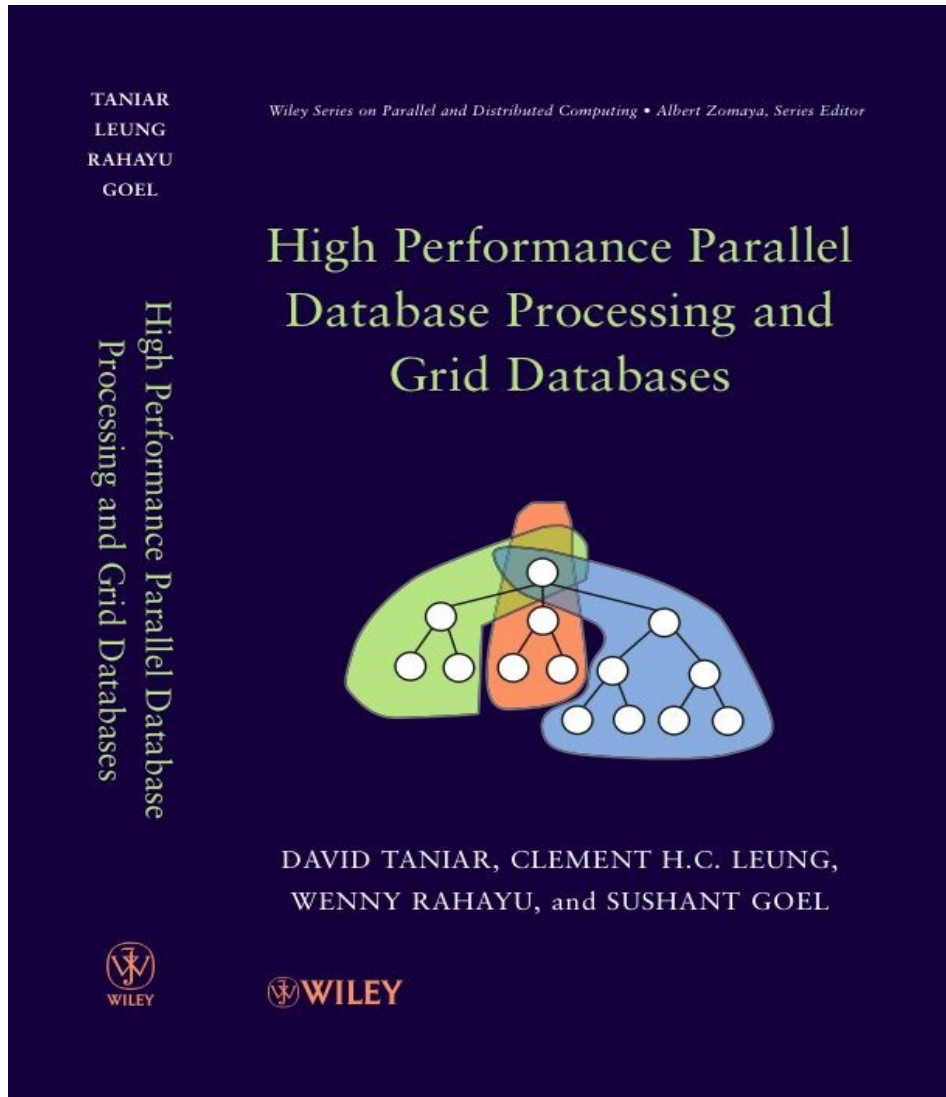
MONASH University

Information Technology

FIT5202 (Volume I - Introduction)

Week 2a – Introduction to Parallel Databases

algorithm distributed systems **database**
systems **computation** knowledge ma
design e-business **model** data mining **int**
distributed systems **database** software
computation knowledge management **an**



Chapter 1 Introduction

- 1.1 A Brief Overview - Parallel Databases and Grid Databases
- 1.2 Parallel Query Processing: Motivations
- 1.3 Parallel Query Processing: Objectives
- 1.4 Forms of Parallelism
- 1.5 Parallel Database Architectures
- 1.6 Grid Database Architecture
- 1.7 Structure of this Book
- 1.8 Summary
- 1.9 Bibliographical Notes
- 1.10 Exercises

Revision

- **Exercise 1 (FLUX Quiz)**

- Using the freeway analogy, number of cars that can pass through the freeway (M1: Monash Freeway) during the morning peak hour from 7 to 9am is called:
 - A. Throughput
 - B. Response Time
 - C. None of the above
 - D. A and B

Revision

- **Exercise 2 (FLUX Quiz)**

- Using the freeway analogy, the duration I take to drive my car to go to work on a freeway (say M1 Monash Freeway) from the Burke Road entrance to the Blackburn Road exit is called:
 - A. Throughput
 - B. Response Time
 - C. None of the above
 - D. A and B

1.3. Objectives (cont'd)

- **Parallel Obstacles**
 - Start-up and Consolidation costs,
 - Interference and Communication, and
 - Skew

1.3. Objectives (cont'd)

- **Start-up and Consolidation**

- Start up: initiation of multiple processes
- Consolidation: the cost for collecting results obtained from each processor by a host processor




Figure 1.3 Serial part vs. parallel part

1.3. Objectives (cont'd)

• Interference and Communication

- Interference: competing to access shared resources
- Communication: one process communicating with other processes, and often one has to wait for others to be ready for communication (i.e. waiting time).




Figure 1.4 Waiting period

• Exercise 3 (Flux Quiz)

- There is a job that will take 1 hour to complete, if this is done by 1 processor.
- The serial part of this job is 10%
- There are 4 processors to use in this job, but each processor will have an overhead of 20% due to waiting time, communication time, etc.
- What type of **speed up** do we get?

1.3. Objectives (cont'd)

- **Skew**

- Unevenness of workload
- Load balancing is one of the critical factors to achieve linear speed up

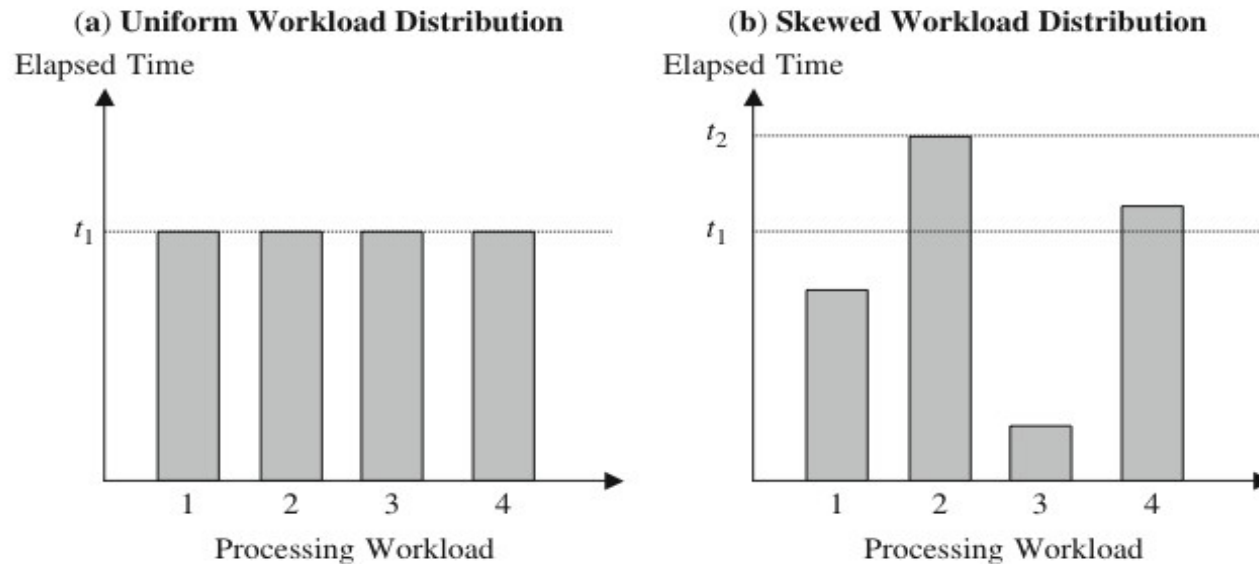


Figure 1.5 Balanced workload vs. unbalanced workload (skewed)

• **Exercise 4 (FLUX Quiz)**

- Pick a number (between 1 and 4)
- A. 1
- B. 2
- C. 3
- D. 4

• **Exercise 5 (FLUX Quiz)**

- Pick a number again (between 1 and 4)
- A. 1
- B. 2
- C. 3
- D. 4

• **Exercise 6 (FLUX Quiz)**

- Pick a number again (between 1 and 4)
- A. 1
- B. 2
- C. 3
- D. 4

1.3. Objectives (cont'd)

• Skew

- *Zipf* distribution model to model skew. Measured in terms of different sizes of fragments allocated to the processors

$$|R_i| = \frac{|R|}{i^\theta \times \sum_{j=1}^N \frac{1}{j^\theta}} \quad \text{where } 0 \leq \theta \leq 1 \quad (2.1)$$

- The symbol θ denotes the degree of skewness, where $\theta = 0$ indicates no skew, and **$\theta = 1$ indicates highly skewed**
- $|R|$ is number of records in the table, **$|R_i|$ is number of records in processor i** , and N is number of processor (j is a loop counter, starting from 1 to N)
- Example: $|R|=100,000$ records, $N=8$ processors

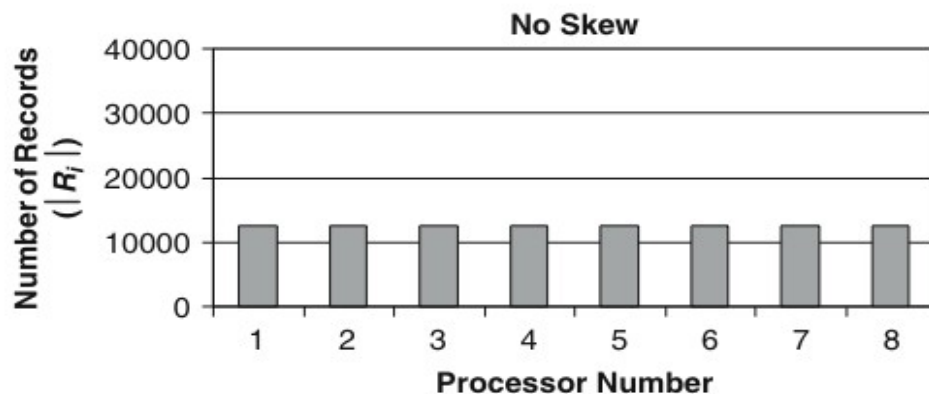


Figure 2.1 Uniform distribution (no skew)




Figure 2.2 Highly skewed distribution

1.3. Objectives (cont'd)

- No skew vs. highly skewed




Figure 2.3 Comparison between highly skewed, less skewed, and no-skew distributions

1.3. Objectives (cont'd)

- No skew vs. highly skewed

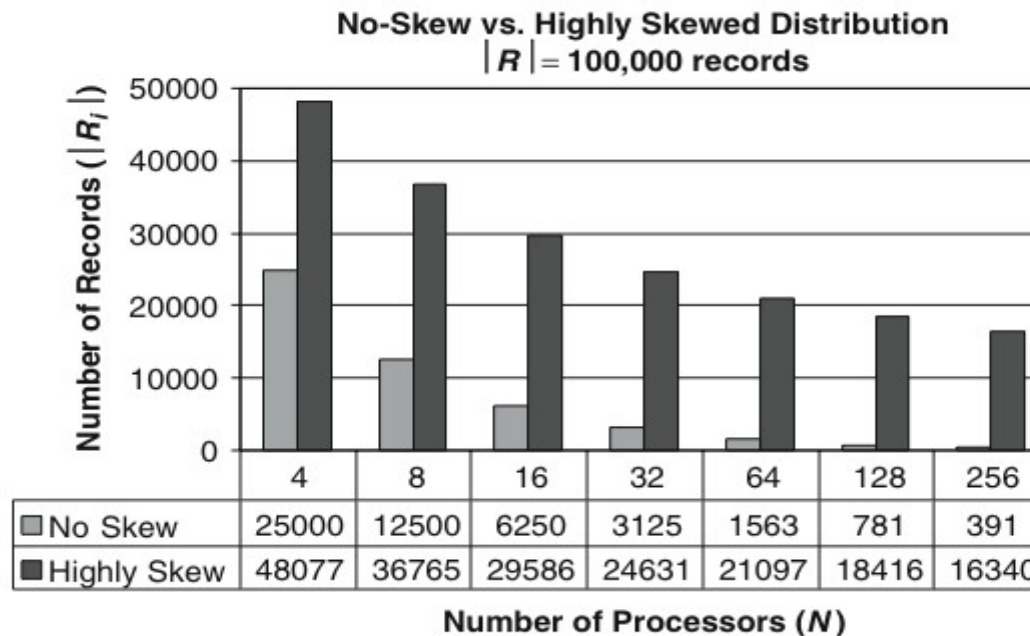


Figure 2.4 Comparison between the heaviest loaded processors using no-skew and highly skewed distributions

1.3. Objectives (cont'd)

- No skew vs. highly skewed

Table 2.2 Divisors (with vs. without skew)

<i>N</i>	4	8	16	32	64	128	256
Divisor without skew	4	8	16	32	64	128	256
Divisor with skew	2.08	2.72	3.38	4.06	4.74	5.43	6.12

• **Exercise 7 (FLUX Quiz)**

- There 100,000 records in the table to be distributed to 32 processors. Assuming that the skewness degree is high ($\theta = 1$), what is the estimated number of records in the heaviest processor?
- A. 48,000 records
- B. 29,000 records
- C. 24,000 records
- D. It is not possible to predict

1.4. Forms of Parallelism

- Forms of parallelism for database processing:
 - Interquery parallelism
 - Intraquery parallelism
 - Interoperation parallelism
 - Intraoperation parallelism
 - Mixed parallelism

1.4. Forms of Parallelism (cont'd)

• Interquery Parallelism

- “Parallelism among queries”
- Different queries or transactions are executed in parallel with one another
- Main aim: scaling up transaction processing systems




Figure 1.6 Interquery parallelism

1.4. Forms of Parallelism (cont'd)

▪ Intraquery Parallelism

- “Parallelism within a query”
- Execution of a single query in parallel on multiple processors and disks
- Main aim: speeding up long-running queries

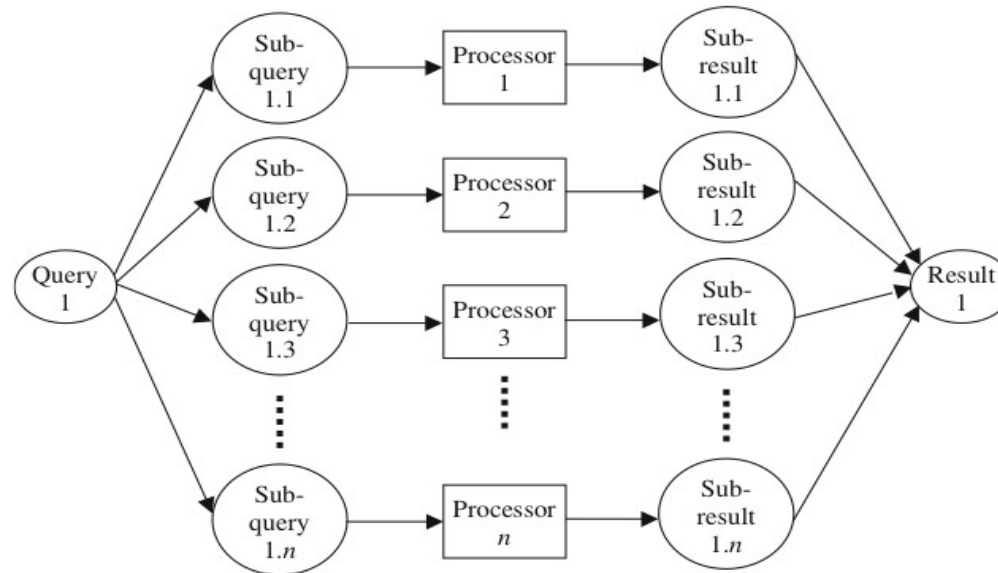


Figure 1.7 Intraquery parallelism

1.4. Forms of Parallelism (cont'd)

- Execution of a single query can be parallelized in two ways:
 - **Intraoperation parallelism**: Speeding up the processing of a query by parallelizing the execution of each individual operation (e.g. parallel sort, parallel search, etc)
 - **Interoperation parallelism**: Speeding up the processing of a query by executing in parallel different operations in a query expression (e.g. simultaneous sorting or searching)

1.4. Forms of Parallelism (cont'd)

▪ Intraoperation Parallelism

- “Partitioned parallelism”
- Parallelism due to the data being partitioned
- Since the number of records in a table can be large, the degree of parallelism is potentially enormous

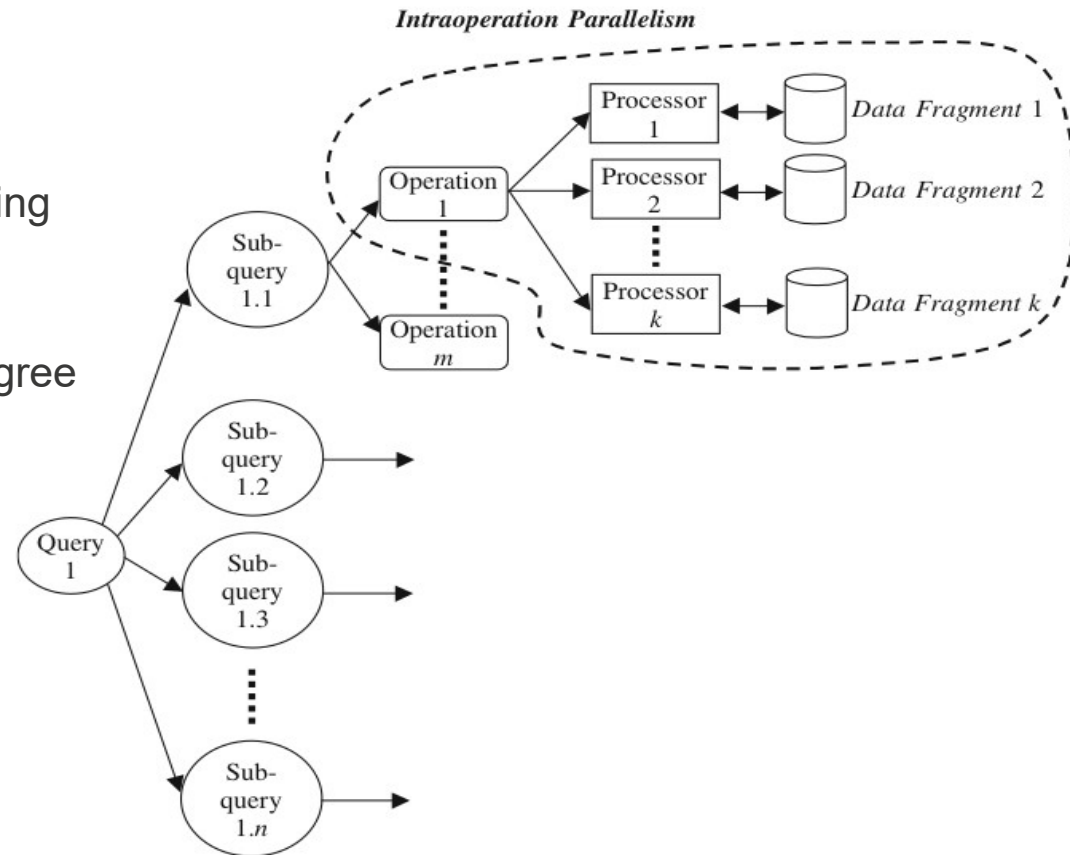


Figure 1.8 Intraoperation parallelism

1.4. Forms of Parallelism (cont'd)

- **Interoperation parallelism:** Parallelism created by concurrently executing different operations within the same query or transaction
 - Pipeline parallelism
 - Independent parallelism

1.4. Forms of Parallelism (cont'd)

• Pipeline Parallelism

- Output record of one operation *A* are consumed by a second operation *B*, even before the first operation has produced the entire set of records in its output
- Multiple operations form some sort of assembly line to manufacture the query results
- Useful with a small number of processors, but does not scale up well

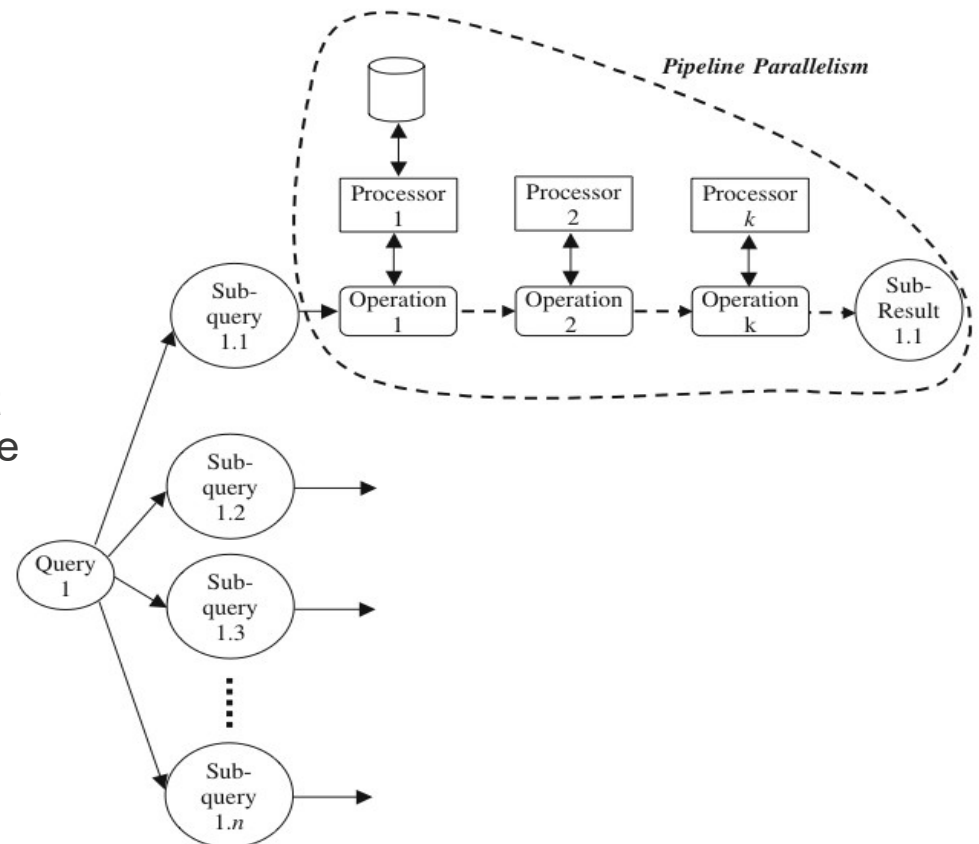


Figure 1.9 Pipeline parallelism

1.4. Forms of Parallelism (cont'd)

- **Independent Parallelism**

- Operations in a query that do not depend on one another are executed in parallel
- Does not provide a high degree of parallelism

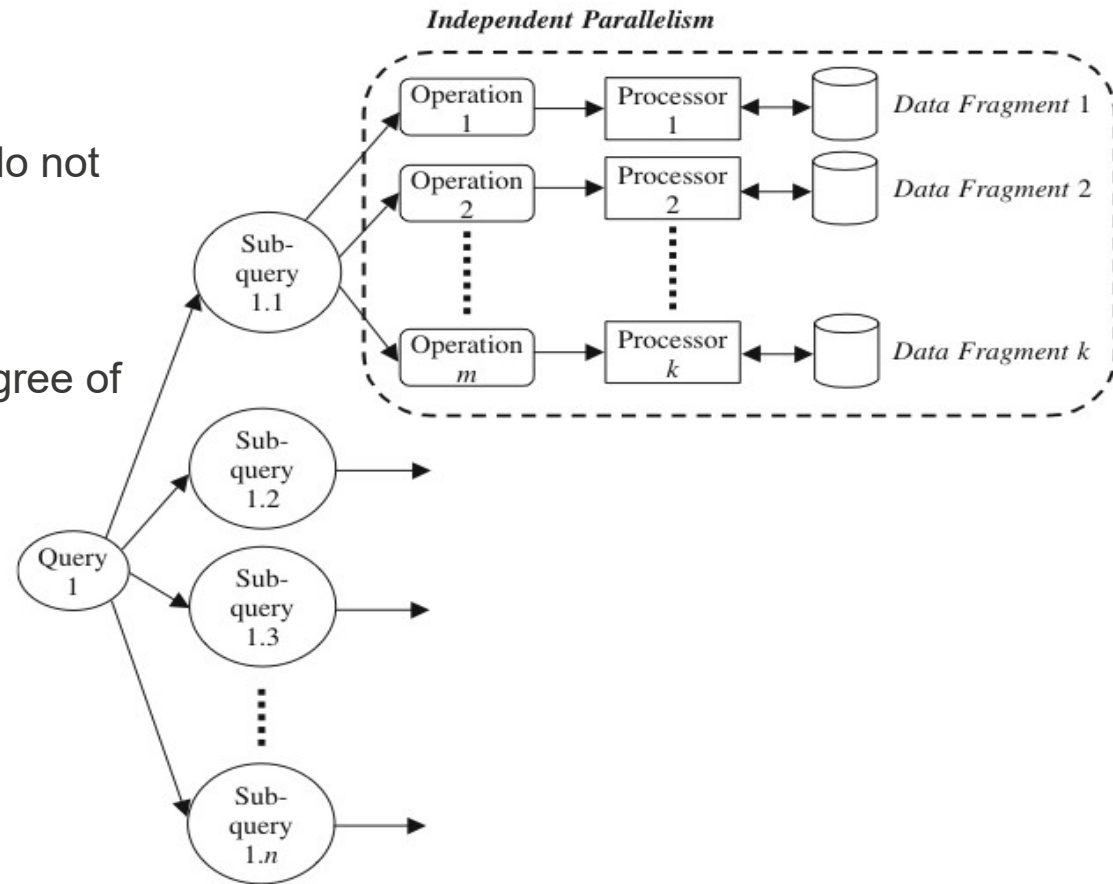


Figure 1.10 Independent parallelism

1.4. Forms of Parallelism (cont'd)

- **Mixed Parallelism**

- In practice, a mixture of all available parallelism forms is used.

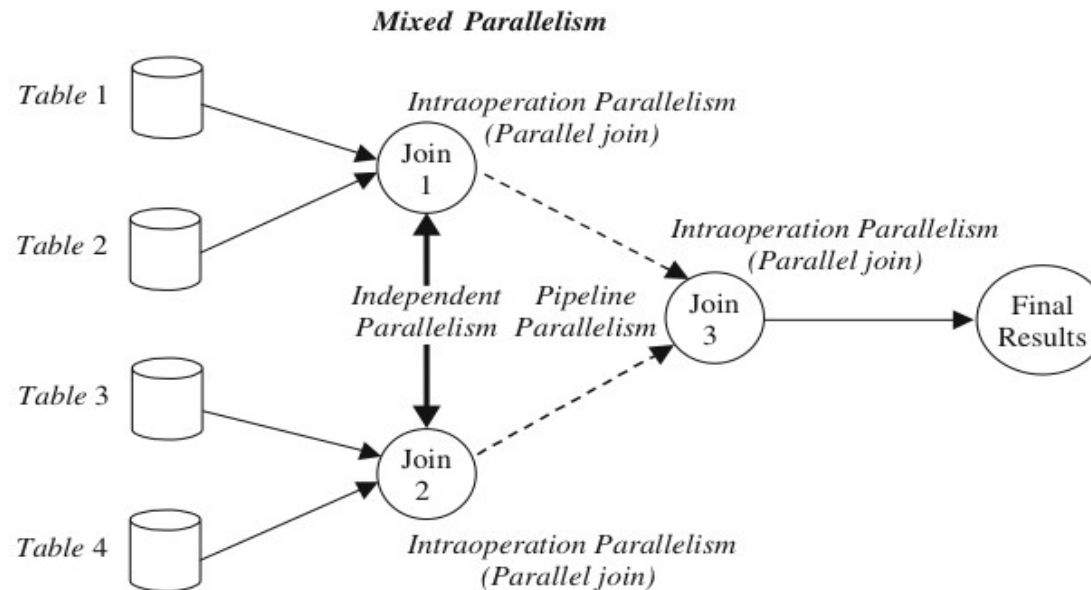


Figure 1.11 Mixed parallelism

1.5. Parallel Database Architectures

- Parallel computers are no longer a monopoly of supercomputers
- Parallel computers are available in many forms:
 - Shared-memory architecture
 - Shared-disk architecture
 - Shared-nothing architecture
 - Shared-something architecture

1.5. Parallel Database Architectures (cont'd)

▪ **Shared-Memory** and **Shared-Disk Architectures**

- Shared-Memory: all processors share a common main memory and secondary memory
- Load balancing is relatively easy to achieve, but suffer from memory and bus contention
- Shared-Disk: all processors, each of which has its own local main memory, share the disks




Figure 1.12 An SMP architecture

1.5. Parallel Database Architectures (cont'd)

- **Shared-Nothing Architecture**

- Each processor has its own local main memory and disks
- Load balancing becomes difficult




Figure 1.13 A shared-nothing architecture

1.5. Parallel Database Architectures (cont'd)

▪ Shared-Something Architecture

- A mixture of shared-memory and shared-nothing architectures
- Each node is a shared-memory architecture connected to an interconnection network aka shared-nothing architecture

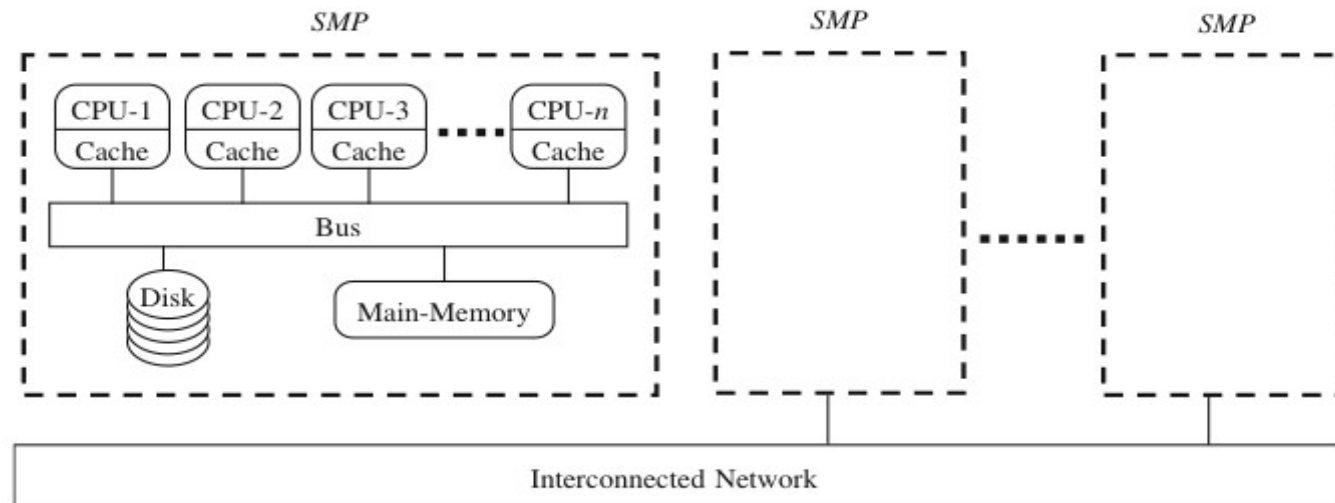


Figure 1.14 Cluster of SMP architectures

1.5. Parallel Database Architectures (cont'd)

• Interconnection Networks

- Bus, Mesh, Hypercube




Figure 1.15 Bus interconnection network





Figure 1.16 Mesh interconnection network

2-dimensional



3-dimensional





Figure 1.17 Hypercube interconnection network

1.7. Exercises (from the textbook)

- **Q1.3:** Highlight the differences between **speed up** and **scale up**.
- **Q1.7:** Skewed workload distribution is generally undesirable. Under what conditions that parallelism (i.e. the workload is divided among all processors) is not desirable?

1.8. Summary

- **Why**, **What**, and **How** of parallel query processing:
 - Why is parallelism necessary in database processing?
 - What can be achieved by parallelism in database processing?
 - How parallelism performed in database processing?
 - What facilities of parallel computing can be used?