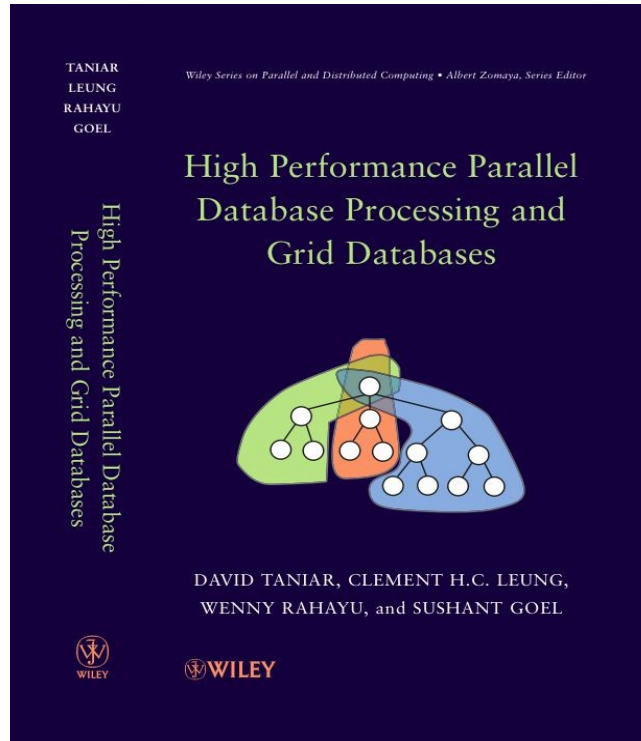


# Machine Learning: Clustering

Prajwol Sangat

Updated by Chee-Ming Ting (15 April 2022)





# Chapter 17

## Parallel Clustering and Classification

- 17.1 Clustering and Classification
- 17.2 Parallel Clustering
- 17.3 Parallel Classification
- 17.4 Summary
- 17.5 Bibliographical Notes
- 17.6 Exercises

# Machine Learning Fundamentals - Revision

- Supervised learning vs. unsupervised learning
- **Supervised learning**: discover patterns in the data that relate to data attributes with a target (class) attribute.
  - These patterns are then utilized to predict the values of the target attribute in future data instances.
- **Unsupervised learning**: The data have no target attribute.
  - Exploring the data to find some intrinsic structures in them.

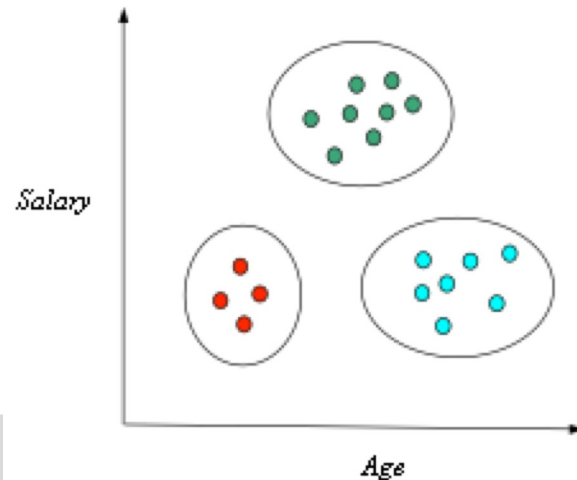
# Clustering: an illustration

- Finds groups (or clusters) of data
- A cluster comprises a number of “similar” objects
- A member is closer to another member within the same group than to a member of a different group (**data points are similar within a cluster, less similar between clusters**)
- Groups have no category or label
- Unsupervised learning

Definition:

**Membership** of a data point

- Indicates which cluster a point belong to



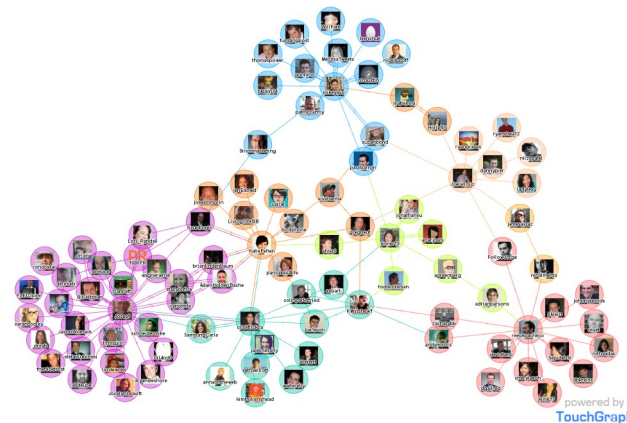
sub	salary	age
1		
2		
3		
4		
⋮		
50		

# What is clustering for?

- Let's see some real-life examples
- **Example 1:** Cluster students based on their examination marks, gender, heights, nationality, etc.
- **Example 2:** In marketing, segment customers according to their similarities
  - To do targeted marketing.

# What is clustering for?

- Clustering is one of the most utilized machine learning techniques.
  - Used in almost every field, e.g., medicine, psychology, botany, sociology, biology, archeology, marketing, insurance, libraries, etc.
  - Most popular applications of clustering are:
    - recommendation engines,
    - market segmentation,
    - social network analysis,
    - image segmentation,
    - anomaly detection



# Some Applications in Digital Health

## Partitioning of Heart sound signals

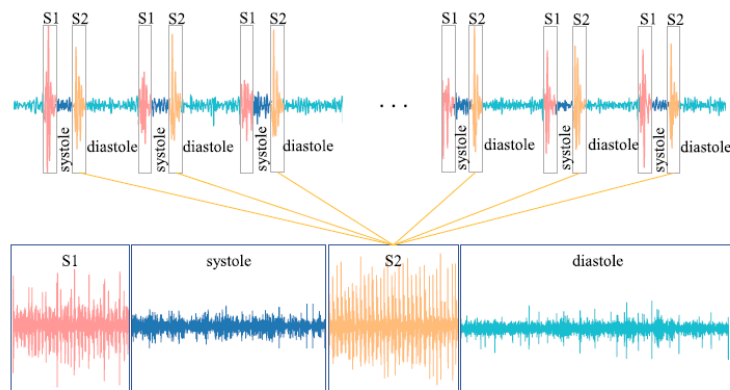
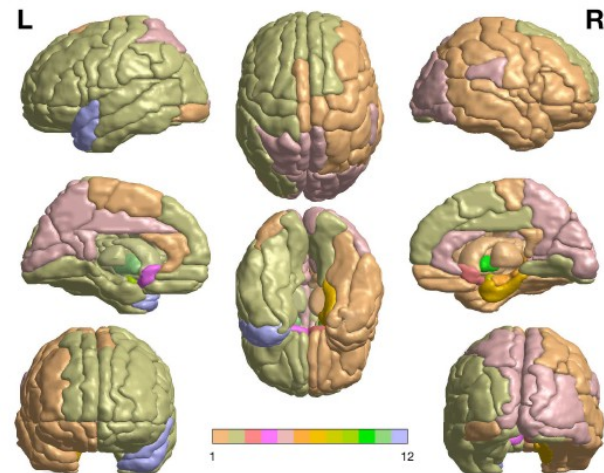


Fig. 2. Dynamic clustering of heart sound into four fundamental components.

Noman, Fuad, Sh-Hussain Salleh, Chee-Ming Ting. "A markov-switching model approach to heart sound segmentation and classification." *IEEE Journal of Biomedical and Health Informatics* 24, no. 3 (2019).

## Partitioning of brain regions into clusters (communities)



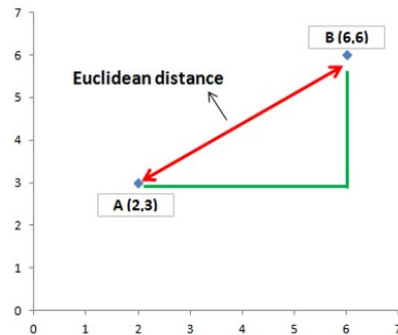
Ting, Chee-Ming, et al. "Detecting Dynamic Community Structure in Functional Brain Networks Across Individuals: A Multilayer Approach." *IEEE Trans Medical Imaging* (2020).

# What is clustering for?

## Similarities Measures

- Key factor in clustering is the similarity measure
- Measure the degree of similarity between two objects
- Distance measure: the shorter the distance, the more similar are the two objects (zero distance means identical objects)
- Euclidean Dis

$$\text{dist}(x_i, x_j) = \sqrt{\sum_{k=1}^h (x_{ik} - x_{jk})^2}$$



$$\text{Euclidean distance } (a, b) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2}$$

$$\begin{aligned} d(x_1, x_2) &= \sqrt{(x_{11} - x_{12})^2 + (x_{21} - x_{22})^2} \\ &= \sqrt{(2 - 6)^2 + (3 - 6)^2} \end{aligned}$$

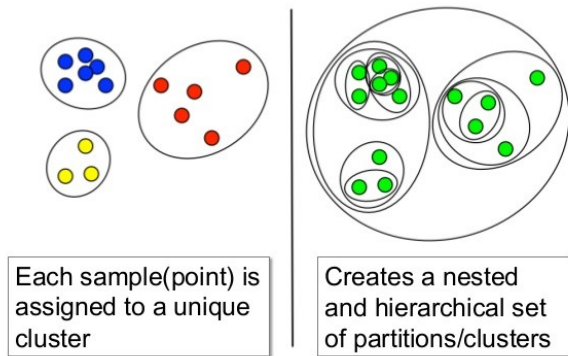
$h$  = number of features (or attributes)



# Clustering Techniques

- **Goal of clustering:**
  - maximize intra-cluster similarity & minimize inter-cluster similarity
- **Hierarchical** clustering (nested clustering)
  - Seeks to build a **hierarchy of clusters** (clusters within clusters)
  - Strategies:
    - *Agglomerative*: Bottom up approach
    - *Divisive*: Top down approach.
- **Partitional** clustering (non-overlapping clustering)
  - Partitions the data objects based on a clustering criterion.
  - Places the data objects into clusters to maximise intra-cluster similarity.
  - So that data in a cluster are more similar to each other than to data in different clusters.

## Partitional vs Hierarchical



# K-Means clustering (Partitional clustering)

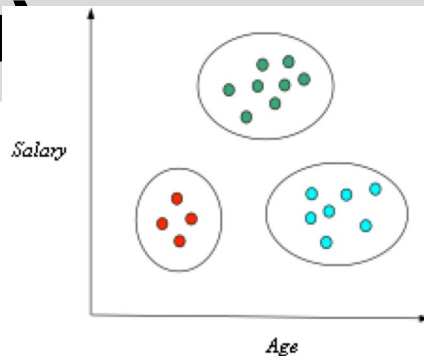
- K-means is a **partitional clustering** algorithm
- Let a set of data points (or instances)  $D$  be

$$\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\},$$

where  $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{ir})$  is a **vector** in a real-valued space  $X \subseteq R^r$ , and  $r$  is the number of attributes (dimensions) in the data.

- The  $k$ -means algorithm partitions the given data into  $k$  clusters.
  - Each cluster has a cluster **center**, called **centroid**.
  - Centroid can be mean/average of member data points in each cluster

$k$  is specified by the user



sub	salary	age
x1		
x2		
x3		
x4		
⋮		
x50		

# K-Means clustering

## ▪ Algorithm k-Means:

- (**Initialization**) Specifies  $k$  number of clusters, and guesses the  $k$  seed cluster centroid
- (**Assignment Step**) Assign each data point to the cluster with the closest centroid
  - Current clusters may receive or loose their members
- (**Update Step**) Each cluster must re-calculate the mean (centroid) based on the newly assigned members
- The process is repeated until the clusters are stable (no change of members)

---

**Algorithm:** k-means

---

Input:

$D = \{x_1, x_2, \dots, x_n\}$  //Data objects

$k$  //Number of desired clusters

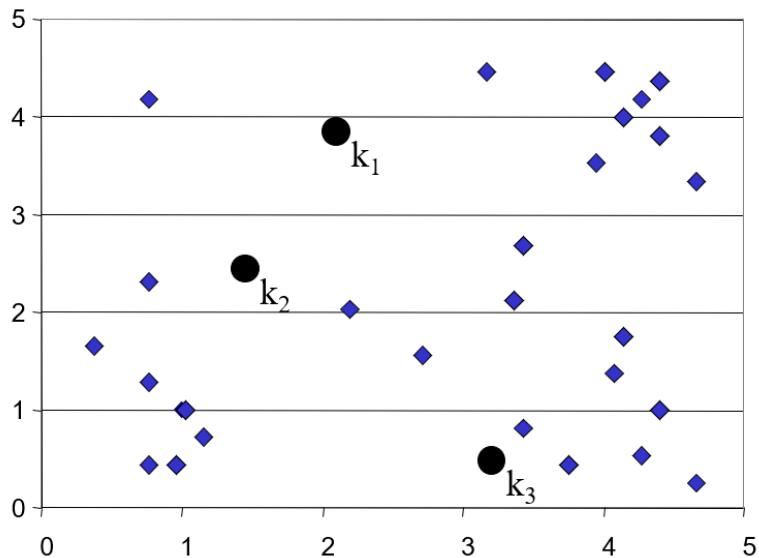
Output:

$K$  //Set of clusters

1. Assign initial values for means  $m_1, m_2, \dots, m_k$
2. Repeat
3. Assign each data object  $x_i$  to the cluster which has the closest mean
4. Calculate new mean for each cluster
5. Until convergence criteria is met

# K-Means Clustering: Step 1

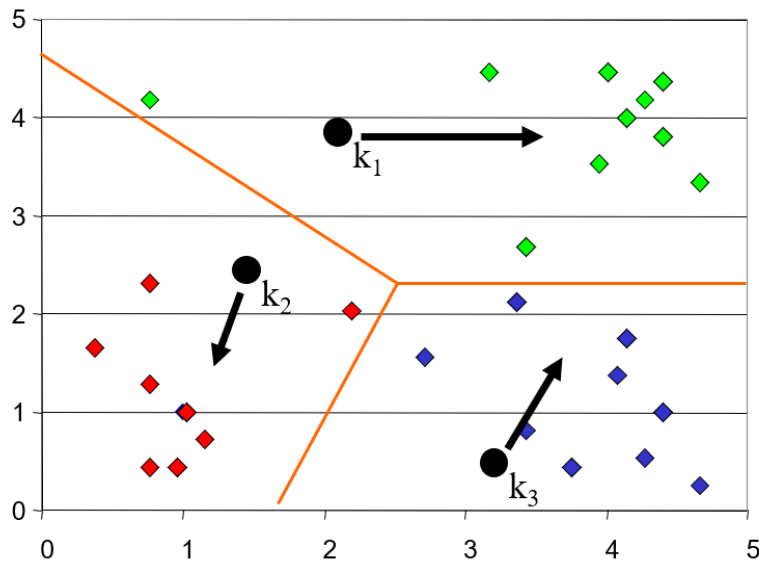
- Algorithm: k-means, Distance Metric: Euclidean Distance



- (**Initialization**)  
Specifies  $k$  number of clusters, and guesses the  $k$  seed cluster centroid

# K-Means Clustering: Step 2

- Algorithm: k-means, Distance Metric: Euclidean Distance

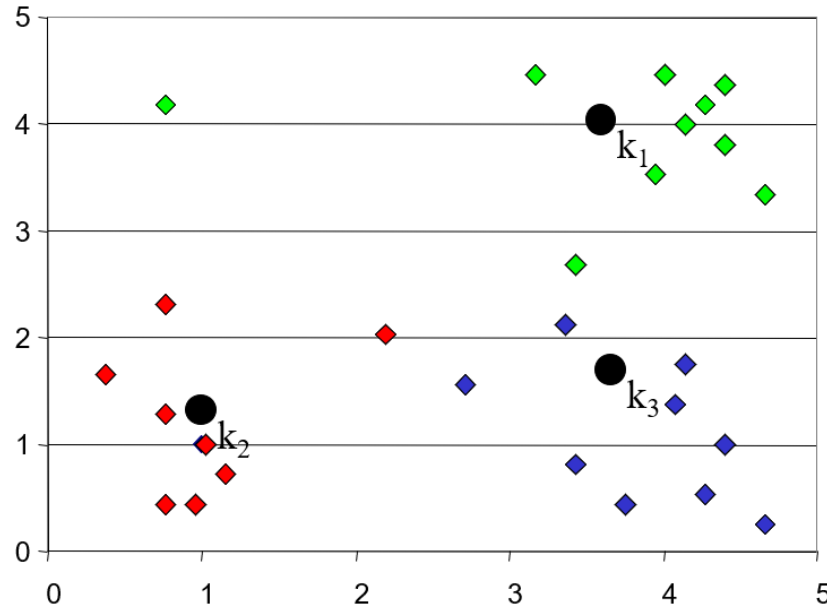


## (Assignment Step)

Iteratively looks at each data point and assigns it to the closest centroid

# K-Means Clustering: Step 3

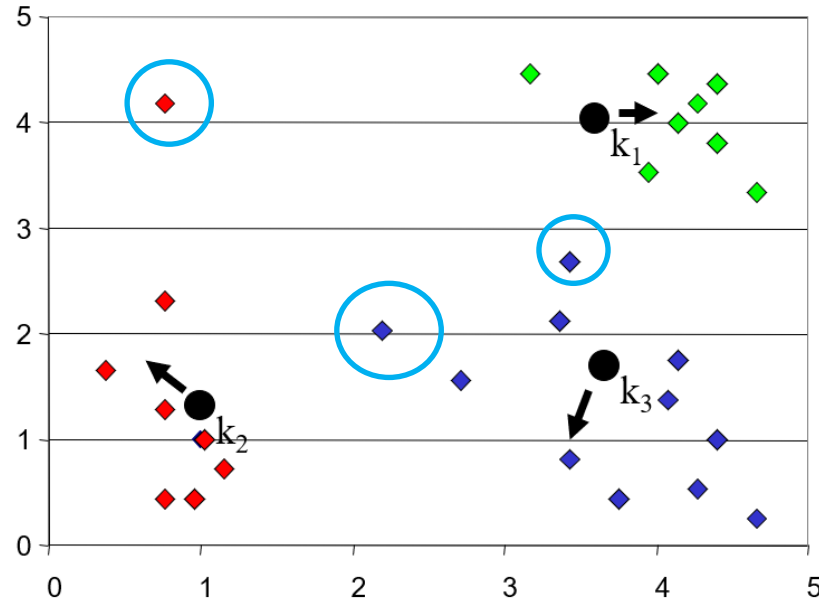
- Algorithm: k-means, Distance Metric: Euclidean Distance



**(Update Step)** Re-calculate the mean (centroid) for each cluster based on the membership of the cluster

# K-Means Clustering: Step 4

- Algorithm: k-means, Distance Metric: Euclidean Distance

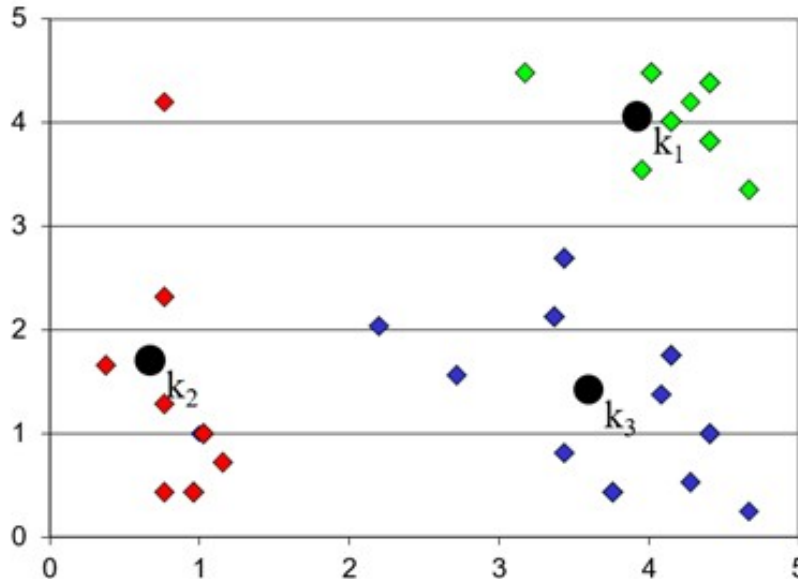


Iteratively looks at each data point and assigns it to the closest centroid,

Current clusters may receive or lose their members

# K-Means Clustering: Step 5

- Algorithm: k-means, Distance Metric: Euclidean Distance



Re-calculate the mean (centroid) for each cluster based on the membership of the cluster



# k-Means: Step-By-Step Example

- Data  $D = \{5, 19, 25, 21, 4, 1, 17, 23, 8, 7, 6, 10, 2, 20, 14, 11, 27, 9, 3, 16\}$
- Number of clusters:  $k = 3$
- Initial centroids:  $m_1=6$ ,  $m_2=7$ , and  $m_3=8$

## First Iteration

### ■ (Assignment Step) Clusters:

- $C_1 = \{1, 2, 3, 4, 5, 6\}$
- $C_2 = \{7\}$
- $C_3 = \{8, 9, 10, 11, 14, 16, 17, 19, 20, 21, 23, 25, 27\}$

### ■ First Iteration: Calculating euclidean distance, determining the cluster membership and calculating new centroid.

D	5	19	25	21	4	1	17	23	8	7	6	10	2	20	14	11	27	9	3	16
$d(m_1, D_i)$	1	13	19	15	2	5	11	17	2	1	0	4	4	14	8	5	21	3	3	10
$d(m_2, D_i)$	2	12	18	14	3	6	10	16	1	0	1	3	5	13	7	4	20	2	4	9
$d(m_3, D_i)$	3	11	17	13	4	7	9	15	0	1	2	2	6	12	6	3	19	1	5	8

# k-Means: Step-By-Step Example

- Clusters:
  - $C_1 = \{1, 2, 3, 4, 5, 6\}$
  - $C_2 = \{7\}$
  - $C_3 = \{8, 9, 10, 11, 14, 16, 17, 19, 20, 21, 23, 25, 27\}$
- New centroids:  $m_1 = 3.5$ ,  $m_2 = 7$ , and  $m_3 = 16.9$

## Second Iteration

- Clusters:
  - $C_1 = \{1, 2, 3, 4, 5\}$
  - $C_2 = \{6, 7, 8, 9, 10, 11\}$
  - $C_3 = \{14, 16, 17, 19, 20, 21, 23, 25, 27\}$
- Re-calculated centroids:  $m_1 = 3$ ,  $m_2 = 8.5$ , and  $m_3 = 20.2$

Second Iteration: Calculating euclidean distance, determining the cluster membership and calculating new centroid.

D	5	19	25	21	4	1	17	23	8	7	6	10	2	20	14	11	27	9	3	16
d(m1, Di)	1.5	15.5	21.5	17.5	0.5	2.5	13.5	19.5	4.5	3.5	2.5	6.5	1.5	16.5	10.5	7.5	23.5	5.5	0.5	12.5
d(m2, Di)	2	12	18	14	3	6	10	16	1	0	1	3	5	13	7	4	20	2	4	9
d(m3, Di)	11.9	2.1	8.1	4.1	12.9	15.9	0.1	6.1	8.9	9.9	10.9	6.9	14.9	3.1	2.9	5.9	10.1	7.9	13.9	0.9



# k-Means: Step-By-Step Example

- Clusters:
  - $C_1 = \{1, 2, 3, 4, 5\}$
  - $C_2 = \{6, 7, 8, 9, 10, 11\}$
  - $C_3 = \{14, 16, 17, 19, 20, 21, 23, 25, 27\}$
- New centroids:  $m_1=3$ ,  $m_2=8.5$ , and  $m_3=20.2$
- **Third Iteration**
  - Clusters:
    - $C_1 = \{1, 2, 3, 4, 5\}$
    - $C_2 = \{6, 7, 8, 9, 10, 11, 14\}$
    - $C_3 = \{16, 17, 19, 20, 21, 23, 25, 27\}$

- Re-calculated centroids:  $m_1=3$ ,  $m_2=9.29$ , and  $m_3=21$

**Third Iteration: Calculating euclidean distance, determining the cluster membership and calculating new centroid.**

D	5	19	25	21	4	1	17	23	8	7	6	10	2	20	14	11	27	9	3	16
d(m1, Di)	2	16	22	18	1	2	14	20	5	4	3	7	1	17	11	8	24	6	0	13
d(m2, Di)	3.5	10.5	16.5	12.5	4.5	7.5	8.5	14.5	0.5	1.5	2.5	1.5	6.5	11.5	5.5	2.5	18.5	0.5	5.5	7.5
d(m3, Di)	15.2	1.2	4.8	0.8	16.2	19.2	3.2	2.8	12.2	13.2	14.2	10.2	18.2	0.2	6.2	9.2	6.8	11.2	17.2	4.2

# k-Means: Step-By-Step Example

- Clusters:
  - $C_1 = \{1, 2, 3, 4, 5\}$
  - $C_2 = \{6, 7, 8, 9, 10, 11, 14\}$
  - $C_3 = \{16, 17, 19, 20, 21, 23, 25, 27\}$
- New centroids:  $m_1=3$ ,  $m_2=9.29$ , and  $m_3=21$
- **Fourth Iteration**
  - Clusters:
    - $C_1 = \{1, 2, 3, 4, 5, 6\}$
    - $C_2 = \{7, 8, 9, 10, 11, 14\}$
    - $C_3 = \{16, 17, 19, 20, 21, 23, 25, 27\}$
  - Re-calculated centroids:  $m_1=3.5$ ,  $m_2=9.83$ , and  $m_3=21$

**Fourth Iteration: Calculating euclidean distance, determining the cluster membership and calculating new centroid.**

D	5	19	25	21	4	1	17	23	8	7	6	10	2	20	14	11	27	9	3	16
d(m1, Di)	2	16	22	18	1	2	14	20	5	4	3	7	1	17	11	8	24	6	0	13
d(m2, Di)	4.3	9.7	15.7	11.7	5.3	8.3	7.7	13.7	1.3	2.3	3.3	0.7	7.3	10.7	4.7	1.7	17.7	0.3	6.3	6.7
d(m3, Di)	16.0	2.0	4.0	0.0	17.0	20.0	4.0	2.0	13.0	14.0	15.0	11.0	19.0	1.0	7.0	10.0	6.0	12.0	18.0	5.0

# k-Means: Step-By-Step Example

- Clusters:
  - $C_1 = \{1, 2, 3, 4, 5, 6\}$
  - $C_2 = \{7, 8, 9, 10, 11, 14\}$
  - $C_3 = \{16, 17, 19, 20, 21, 23, 25, 27\}$
- New centroids:  $m_1 = 3.5$ ,  $m_2 = 9.83$ , and  $m_3 = 21$
- **Fifth Iteration**

- No data movement from clusters (Process Terminated)

$m_1$	$m_2$	$m_3$	$C_1$	$C_2$	$C_3$
6	7	8	1, 2, 3, 4, 5, 6	7	8, 9, 10, 11, 14, 16, 17, 19, 20, 23, 25, 27
3.5	7	16.9	1, 2, 3, 4, 5	6, 7, 8, 9, 10, 11	14, 16, 17, 19, 20, 21, 23, 25, 27
3	8.5	20.2	1, 2, 3, 4, 5	6, 7, 8, 9, 10, 11, 14	16, 17, 19, 20, 21, 23, 25, 27
3	9.29	21	1, 2, 3, 4, 5, 6	7, 8, 9, 10, 11, 14	16, 17, 19, 20, 21, 23, 25, 27
3.5	9.83	21	1, 2, 3, 4, 5, 6	7, 8, 9, 10, 11, 14	16, 17, 19, 20, 21, 23, 25, 27

# Evaluating K-Means Clusters

- One common measure is **sum of squared error (SSE)**

$$SSE = \sum_{i=1}^K \sum_{x \in C_i} dist^2(m_i, x)$$

- $x$  is a data point in cluster  $C_i$
- $m_i$  is the centroid of cluster  $C_i$

Example: How to calculate SSE?

SSE measures how close the assigned data points to centroids

- small value  $\rightarrow$  maximized intra-cluster similarity

Fifth Iteration: Calculating euclidean distance, determining the cluster membership and calculating new centroid.

D	5	19	25	21	4	1	17	23	8	7	6	10	2	20	14	11	27	9	3	16
d(m1, Di)	1.5	15.5	21.5	17.5	0.5	2.5	13.5	19.5	4.5	3.5	2.5	6.5	1.5	16.5	10.5	7.5	23.5	5.5	0.5	12.5
d(m2, Di)	4.8	9.2	15.2	11.2	5.8	8.8	7.2	13.2	1.8	2.8	3.8	0.2	7.8	10.2	4.2	1.2	17.2	0.8	6.8	6.2
d(m3, Di)	16.0	2.0	4.0	0.0	17.0	20.0	4.0	2.0	13.0	14.0	15.0	11.0	19.0	1.0	7.0	10.0	6.0	12.0	18.0	5.0

Clusters:

$C_1 = \{1, 2, 3, 4, 5, 6\}$

$C_2 = \{7, 8, 9, 10, 11, 14\}$

$C_3 = \{16, 17, 19, 20, 21, 23, 25, 27\}$

centroids:  $m_1=3.5$ ,  $m_2=9.83$ , and  $m_3=21$

$$\sum_{x \in C_1} d^2(x, m_1)$$



$$\sum_{x \in C_2} d^2(x, m_2)$$



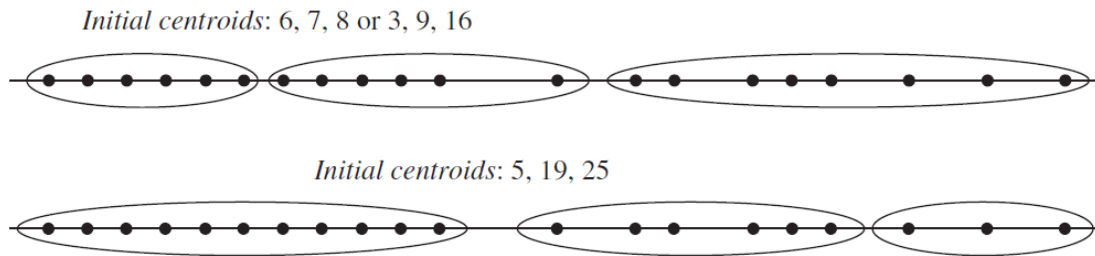
$$\sum_{x \in C_3} d^2(x, m_3)$$



MONASH  
University

# K-Means Clustering

- The number of clusters  $k$  is predefined. The algorithm does not discover the ideal number of clusters. During the process, the number of clusters remains fixed – it does not shrink nor expand.
- The final composition of clusters is very sensitive to the choice of initial centroid values. Different initialisations may result in (




**Figure 17.4** Different clustering results for different initial centroids

# K-Means Clustering: Pros and Cons

## Pros

- Simple and fast for low dimensional data (time complexity of K Means is linear i.e.  $O(n)$ )
- Scales to large data sets
- Easily adapts to new data points

## Cons

-  It will not identify outliers
- Restricted to data which has the notion of a centre (centroid)



# K-means clustering

## ■ Exercise 1

- Data  $D = \{8, 11, 12, 14, 16, 17, 24, 28\}$
- Number of clusters:  $k = 3$
- Initial centroids:  $m_1=11$ ,  $m_2=12$ , and  $m_3=28$
- Use the *k*-means *serial* algorithm to cluster the data in three clusters

# Finding Optimal number of the clusters

- As  $k$  increases, clusters become smaller.
- The neighbouring clusters become less distinct from one another.

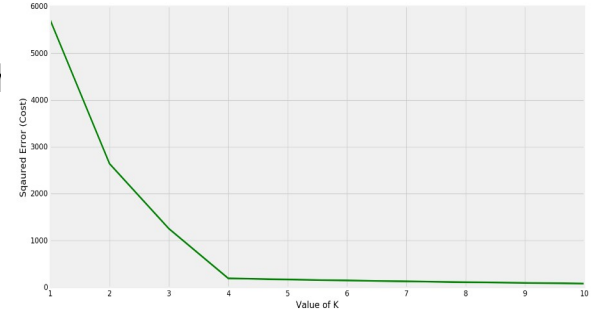
## ■ How to choose an optimal $k$ ?

### - Elbow Method

- Plot sum of squared errors as a function of  $k$  (a scree plot)
- Select the value of  $k$  at the “elbow” ie the point after which the SSE start decreasing in a linear fashion.

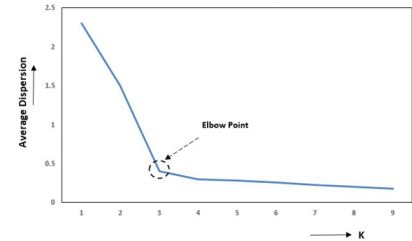
### - Silhouette analysis

- Measure of how close each point in one cluster is compared to points in the neighbouring clusters and provides a way to assess number of clusters.
- If most points have a high silhouette value, then the clustering configuration is appropriate.
- If many points have a low or negative value, then the clustering configuration may have too many or too few clusters.



optimal value for  $k$   
= 4

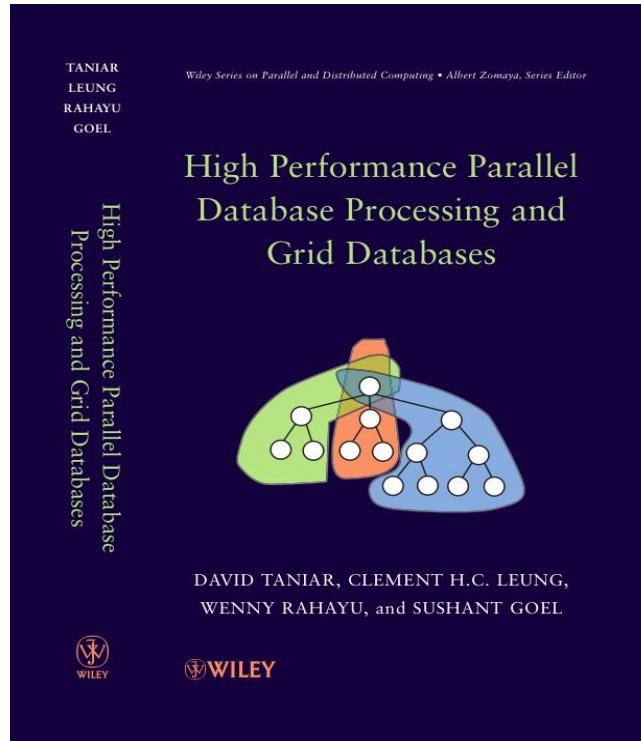
Elbow Method for selection of optimal “K” clusters



```
For n_clusters = 2 The average silhouette_score is : 0.7049787496083262
For n_clusters = 3 The average silhouette_score is : 0.5882004012129721
For n_clusters = 4 The average silhouette_score is : 0.6505186632729437
For n_clusters = 5 The average silhouette_score is : 0.56376469026194
For n_clusters = 6 The average silhouette_score is : 0.4504666294372765
```

# DEMO





# Chapter 17

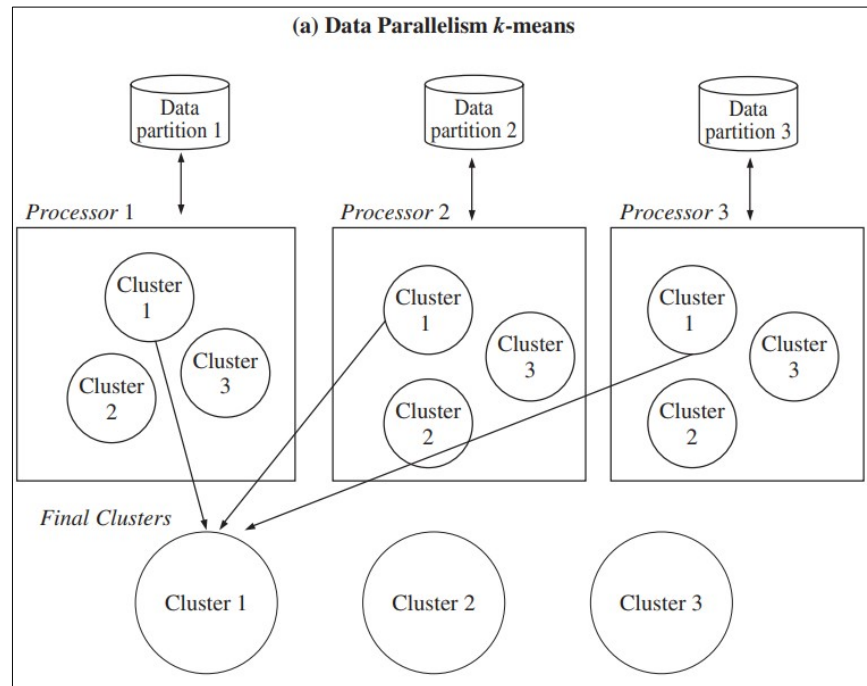
## Parallel Clustering and Classification

- 17.1 Clustering and Classification
- 17.2 Parallel Clustering
- 17.3 Parallel Classification
- 17.4 Summary
- 17.5 Bibliographical Notes
- 17.6 Exercises

# Parallel K-means clustering

## ■ **Data parallelism** of k-means

- ❑ Create parallelism **from the beginning** because of partitioning of the dataset.
- ❑ Data is partitioned into multiple partition
- ❑ Each processor will work independently to create three clusters
- ❑ The final clusters from each processor are respectively united

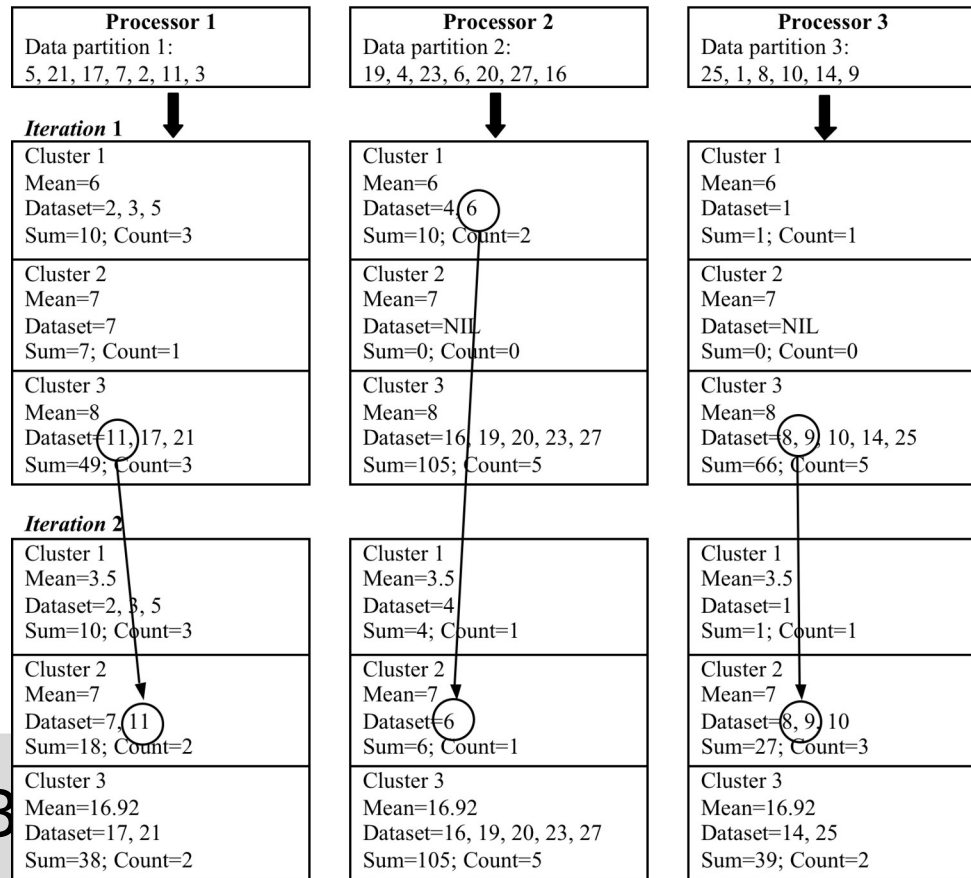


# Parallel K-means

## ■ Data parallelism

- ❑ Example: Data partitioning using round-robin
- ❑ Initial centroids: 6, 7, 8
- ❑ Each processor will run k\_Means locally
- ❑ At the end of each iteration, info about sum & count of data points in each local cluster is shared to calculate new centroid/mean
- ❑ Data does not move among processors (it stays where it was allocated initially)
- ❑ Data move across clusters within same processor

Initial dataset: 5, 19, 25, 21, 4, 1, 17, 23, 8, 7, 6, 10, 2, 20, 14, 11, 27, 9, 3, 16



# Parallel K-means

## ■ Data parallelism

### k-means

Processor 1: Cluster 1 = 2, 3, 5

Cluster 2 = 7, 11

Cluster 3 = 17, 21

Processor 2: Cluster 1 = 4, 6

Cluster 2 = NIL

Cluster 3 = 16, 19, 20, 23, 27

Processor 3: Cluster 1 = 1

Cluster 2 = 8, 9, 10, 14

Cluster 3 = 25

**Cluster 1 = 1, 2, 3, 4, 5, 6**

**Cluster 2 = 7, 8, 9, 10, 11, 14**

**Cluster 3 = 16, 17, 19, 20, 21, 23, 25, 27**

**Initial dataset:** 5, 19, 25, 21, 4, 1, 17, 23, 8, 7, 6, 10, 2, 20, 14, 11, 27, 9, 3, 16

**Processor 1**  
Data partition 1:  
5, 21, 17, 7, 2, 11, 3

**Iteration 1**

Cluster 1  
Mean=6  
Dataset=2, 3, 5  
Sum=10; Count=3

Cluster 2  
Mean=7  
Dataset=7  
Sum=7; Count=1

Cluster 3  
Mean=8  
Dataset=11, 17, 21  
Sum=49; Count=3

**Iteration 2**

Cluster 1  
Mean=3.5  
Dataset=2, 3, 5  
Sum=10; Count=3

Cluster 2  
Mean=7  
Dataset=7, 11  
Sum=18; Count=2

Cluster 3  
Mean=16.92  
Dataset=17, 21  
Sum=38; Count=2

**Processor 2**  
Data partition 2:  
19, 4, 23, 6, 20, 27, 16

Cluster 1  
Mean=6  
Dataset=4, 6  
Sum=10; Count=2

Cluster 2  
Mean=7  
Dataset=NIL  
Sum=0; Count=0

Cluster 3  
Mean=8  
Dataset=16, 19, 20, 23, 27  
Sum=105; Count=5

Cluster 1  
Mean=3.5  
Dataset=4  
Sum=4; Count=1

Cluster 2  
Mean=7  
Dataset=6  
Sum=6; Count=1

Cluster 3  
Mean=16.92  
Dataset=16, 19, 20, 23, 27  
Sum=105; Count=5

**Processor 3**  
Data partition 3:  
25, 1, 8, 10, 14, 9

Cluster 1  
Mean=6  
Dataset=1  
Sum=1; Count=1

Cluster 2  
Mean=7  
Dataset=NIL  
Sum=0; Count=0

Cluster 3  
Mean=8  
Dataset=8, 9, 10, 14, 25  
Sum=66; Count=5

Cluster 1  
Mean=3.5  
Dataset=1  
Sum=1; Count=1

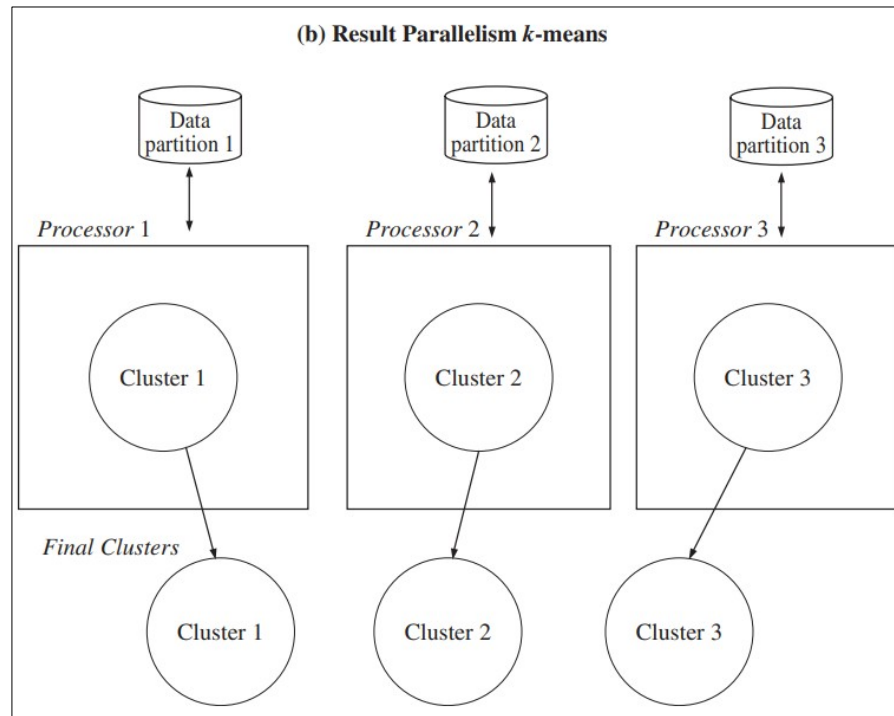
Cluster 2  
Mean=7  
Dataset=8, 9, 10  
Sum=27; Count=3

Cluster 3  
Mean=16.92  
Dataset=14, 25  
Sum=39; Count=2

# Parallel K-means clustering

## ■ **Result Parallelism** of k-means

- ❑ Focuses on clusters partitioning
- ❑ Each processor will work on a particular target cluster
- ❑ For example, from the very beginning, processor 1 will produce only one cluster assigned to it, that is cluster 1.
- ❑ During the iteration, the memberships of cluster can change. -> data movement across processors



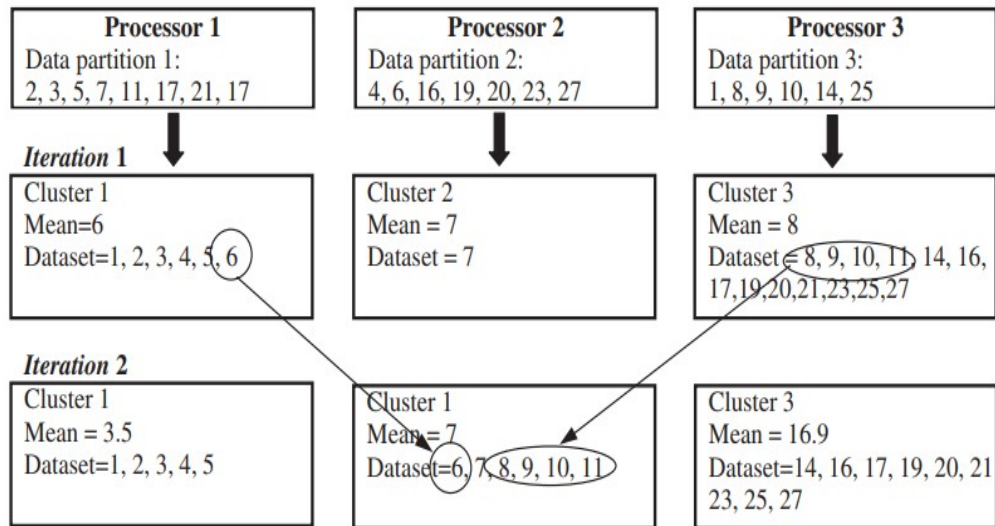


# Parallel K-means

## ■ *Result parallelism* k-means

- ❑ **Example:** Data partitioning using round-robin
- ❑ Each processor is allocated only one cluster.
- ❑ Three initial means are distributed among the three processors,
- ❑ **Data points may move from one processor to another** at each iteration to join a cluster in a different processor
- ❑ Since a cluster is processed by one processor, calculating the mean is straightforward because all the data points within a cluster are located at the same processor

Initial dataset: 5, 19, 25, 21, 4, 1, 17, 23, 8, 7, 6, 10, 2, 20, 14, 11, 27, 9, 3, 16



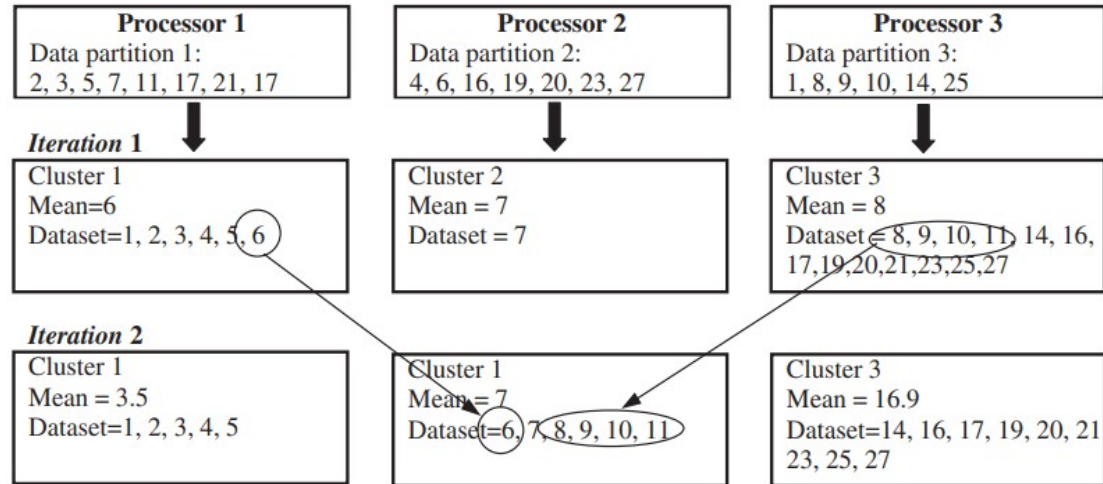
# Parallel K-means

## ■ *Result parallelism* k-means

At the end, the final cluster result is basically the union of all local clusters from each processor.

**Processor 1 cluster 1 = 1, 2, 3, 4, 5, 6**  
**Processor 2 cluster 2 = 7, 8, 9, 10, 11, 14**  
**Processor 3 cluster 3 = 16, 17, 19, 20, 21, 23, 25, 27**

Initial dataset: 5, 19, 25, 21, 4, 1, 17, 23, 8, 7, 6, 10, 2, 20, 14, 11, 27, 9, 3, 16



# Data parallelism vs Result parallelism

## Data parallelism

- Parallelism is created due to the fragmentation of initial input data
- Each processor focuses on its partition of the dataset
- Final results are formed by combining all local results produced by individual processors.

## Result parallelism

- Focuses on the fragmentation of the results, not necessarily the input data.
- Each processor focuses on its target result partition.

# What have we learnt today?

- Partitional (k-means) to attain meaningful groups of data
- Algorithmic examples for clustering of data