

Week 2 Tutorial

- **Week 1 Review**
- **Accumulators**
- **SparkSession vs SparkContext**
- **Data Partitioning**
- **RDD vs DataFrame**
- **Searching in RDDs and DataFrames**
- **Spark SQL**



- **VM Setup and Jupyter Notebooks**
- **RDDs**
 - How to create RDDs?
- **Transformation**
 - Map
 - FlatMap
- **Action**
 - Take
 - Collect (take vs collect)
 - Reduce
 - Count
- **Spark UI (port 4040)**




Fig : Src : [<https://spark.apache.org/docs/2.3.2/running-on-mesos.html>]

Word Count Example Review

step 1: Read the text file twitter.txt

```
rdd = sc.textFile("twitter.txt")
```

step 2: Use a transformation to break the lines to individual words

```
words = rdd.flatMap(lambda line: line.split(" "))
```

step 3: Use a transformation to convert word to a key/value pair of (word, 1)

```
wordCounts = words.map(lambda word: (word, 1))
```

step 4: Use a transformation to reduce the value based on the word

```
finalrdd = wordCounts.reduceByKey(lambda a,b:a +b)
```

step 5: Collect and display the results of the count

```
finalrdd.collect()
```

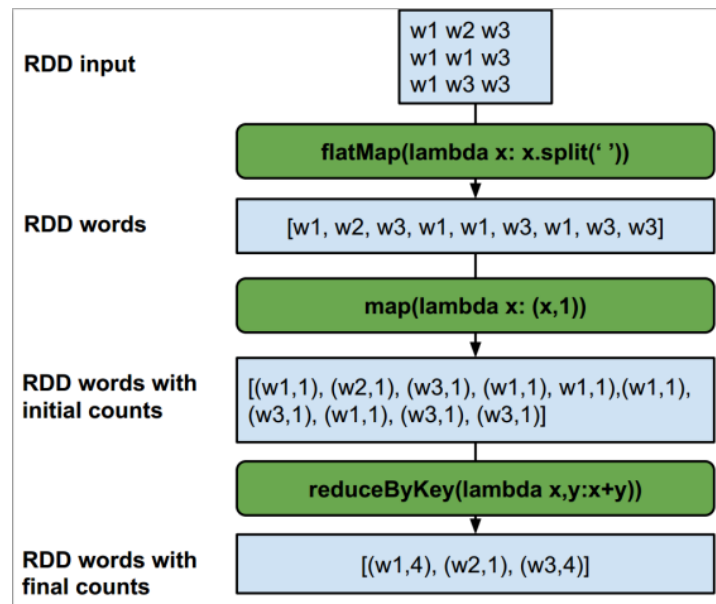


Fig : [\[Source\]](#)

Accumulators

- Accumulators provides a simple syntax for aggregating values from worker nodes back to the driver program.
- They are only “added” through an associative and commutative operation and can therefore be efficiently supported in parallel.
- They can be used to implement counters (as in MapReduce) or sums.

Broadcast Variables

- Broadcast variables allow the program to efficiently send a large, read-only value to all the worker nodes for use in one or more Spark operations.
- Spark automatically sends all variables referenced in your closures to the worker nodes.

Accumulators


```
twitter_rdd = sc.textFile('twitter.txt', 3)
blank_lines = 0 # global variable

def extract_blank_lines(line):
    if line == "":
        blank_lines += 1
    return line.split(" ")

word_rdds = twitter_rdd.flatMap(extract_blank_lines)
word_rdds.collect()

print("Blank lines: %d" %blank_lines)
```

Fails, as blank_lines is not accessible in the executors



Accumulator

```
twitter_rdd = sc.textFile('twitter.txt', 3)
blank_lines = sc.accumulator(0) # Create Accumulator[int] initialized to 0

def extract_blank_lines(line):
    global blank_lines # make the global variable accessible
    lll = {'a':1}
    if line == "":
        print(type(line))
        blank_lines += 1
    return line.split(" ")

word_rdds = twitter_rdd.flatMap(extract_blank_lines)
word_rdds.collect()

print("Blank lines: %d" %blank_lines.value)
```

blank_lines = 0

Spark Driver

Executor

extract_blank_lines()

Executor

extract_blank_lines()

Introducing SparkSession

SparkContext vs SparkSession

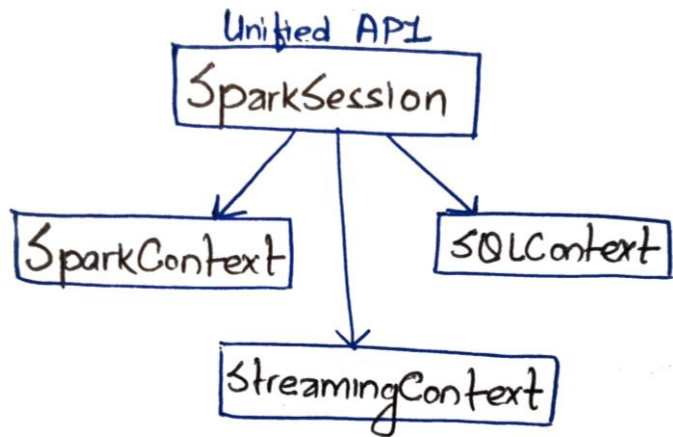
- Unified entry point of Spark application from Spark 2.0

```
# Import SparkConf class into program
from pyspark import SparkConf

# Local[*]: run Spark in local mode with as many working processors as
# If we want Spark to run locally with 'k' worker threads, we can specify
master = "local[*]"
# The `appName` field is a name to be shown on the Spark cluster UI page
app_name = "Parallel Search"
# Setup configuration parameters for Spark
spark_conf = SparkConf().setMaster(master).setAppName(app_name)
```

```
# Import SparkSession
from pyspark.sql import SparkSession # Spark SQL

# Method 1: Using SparkSession
spark = SparkSession.builder.config(conf=spark_conf).getOrCreate()
sc = spark.sparkContext
sc.setLogLevel('ERROR')
```



Data Partitioning Strategies :

1. **Round-robin partitioning** : distribute evenly among processors
2. **Range data partitioning** : partition based on given range
3. **Hash data partitioning** : partition based on a particular attribute using a hash function

DEMO : Partitioning in a RDD!

Data Partitioning in Spark

RDDs are partitioned by Default!

Then why?

Reduce the overhead of the shuffle

Cost Reduction – better utilization of cluster

Avoiding Data Skew??

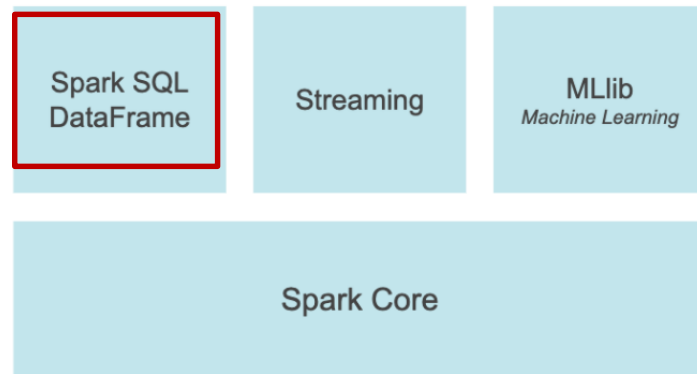


Parallel Search in RDD

- Searching in RDDs using Multiple Conditions
- Finding max/min values of an attribute in RDDs

Spark SQL

- ❑ A Spark module for **structured data processing**.
- ❑ It provides a programming abstraction called **DataFrame** and act as distributed SQL query engine.
- ❑ Unlike the basic Spark RDD API, Spark SQL provide Spark with **more information about the structure of both the data and the computation being performed (used for optimization)**.



Two uses of Spark SQL:

- ❑ **DataFrames APIs:**
 - Dataframe is distributed collection of data organized into **named columns**
- ❑ **To execute SQL queries**

| <i>Id</i> <i>(Int)</i> | <i>First</i> <i>(String)</i> | <i>Last</i> <i>(String)</i> | <i>Url</i> <i>(String)</i> | <i>Published</i> <i>(Date)</i> | <i>Http</i> <i>(Int)</i> |
|---|---|--|---|---|---|
| 1 | Jules | Damji | https://tinyurl.1 | 1/4/2016 | 4535 |
| 2 | Brooke | Wenig | https://tinyurl.2 | 5/5/2018 | 8908 |
| 3 | Denny | Lee | https://tinyurl.3 | 6/7/2019 | 7658 |
| 4 | Tathagata | Das | https://tinyurl.4 | 5/12/2018 | 10566 |

Spark Core

- underlying execution engine that all other functionalities build on it
- Working with an RDD

Partitioning with DataFrames

Round-robin partitioning :

```
df_round = df.repartition(5)
```

repartition()

repartitionByRange()

Range data partitioning :

```
df_range = df.repartitionByRange(5, "balance")
```

Hash data partitioning :

```
column_hash = "education"  
df_hash = df.repartition(column_hash)
```

Searching in Dataframe

- **Filter()** - Perform search
- **Where()** - Alias to Filter()
- **Select()** - Select particular columns to show
- **Show()** - Select rows to show

<https://spark.apache.org/docs/latest/api/python/reference/pyspark.sql/api/pyspark.sql.DataFrame.html#>

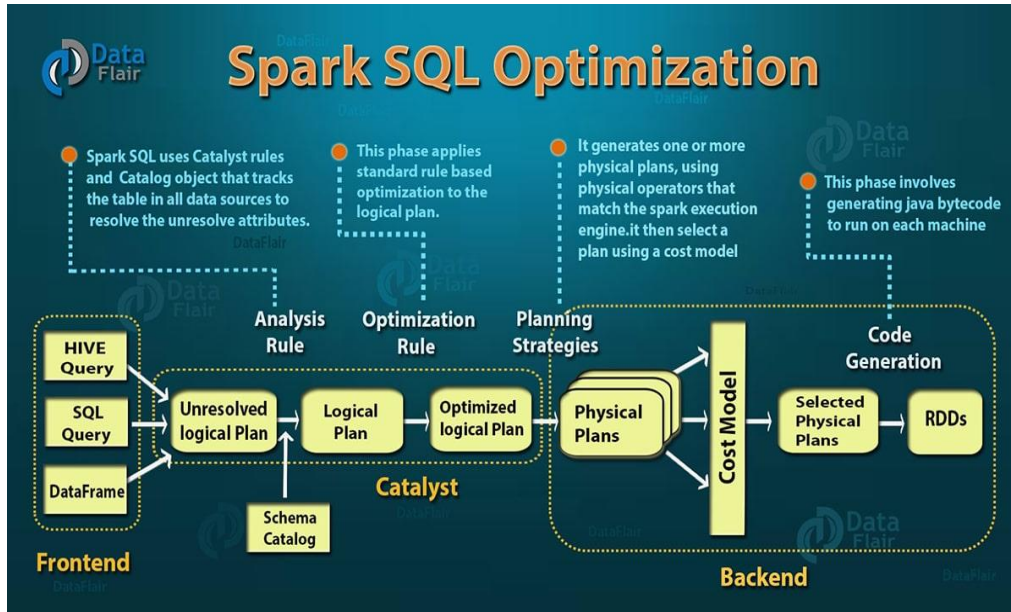
- To execute SQL queries.
- For further reading [link](#)
- Temporary views in Spark SQL

```
df = spark.read.csv("bank.csv",header=True)
# Register the DataFrame as a SQL temporary view
df.createOrReplaceTempView("bank")
```

```
sqlDF = spark.sql("SELECT * FROM bank")
sqlDF.show()
```

```
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
|age|      job| marital|education|default|balance|housing|loan|contact|day|month|duration|campaign|pdays|previous|poutcome|de
posit|
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
| 59|    admin.| married|secondary|    no|   2343|   yes|  no|unknown| 5|  may|   1042|     1|   -1|     0| unknown|
yes|
| 56|    admin.| married|secondary|    no|     45|   no|  no|unknown| 5|  may|   1467|     1|   -1|     0| unknown|
yes|
| 41| technician| married|secondary|    no|   1270|   yes|  no|unknown| 5|  may|   1389|     1|   -1|     0| unknown|
|
```

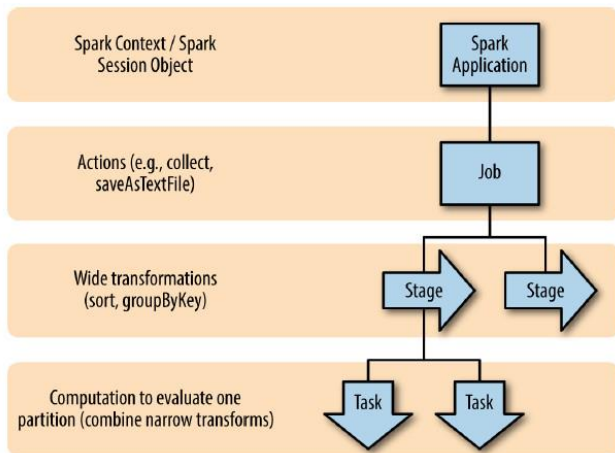
Logical plan & Physical plan



- ❑ **Execution plan** - set of operations executed to **translate a query language statement** (SQL, Spark SQL, Dataframe operations etc.) to a set of **optimized logical and physical operations**.
- ❑ **Logical plan** - Abstract of all data **transformation steps** (specified by the user) that need to be **executed**. No details about the Driver(Master Node) or Executor (Worker Node)
- ❑ **Physical plan** - Contains more **specific description of how execution should happen** (e.g., specific RDD to create, specific choice of algorithms for join or agg., how data partitioned/shuffled)
- ❑ Once the finest Physical Plan is selected, executable code (**DAG of RDDs**) for the query is created

Directed Acyclic Graph (DAG)

DAG is a graph denoting a sequence of operations (transformations & actions) that are performed on the targeted RDDs



Details for Job 8

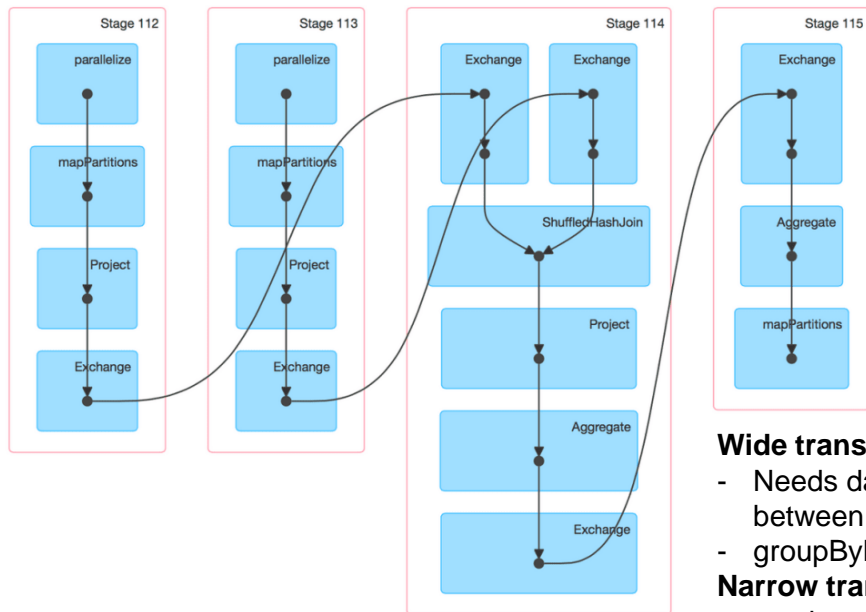
Status: SUCCEEDED

Completed Stages: 4

► Event Timeline

▼ DAG Visualization

The graph is split into stages of tasks for execution in worker nodes



Wide transformation

- Needs data shuffling between nodes
- `groupByKey()`, `join()`...

Narrow transformation

- each partition of output RDD depends only on a **single partition** of input RDD.
- `map()`, `filter()`...

<https://databricks.com/blog/2015/06/22/understanding-your-spark-application-through-visualization.html>

<https://www.linkedin.com/pulse/demystifying-spark-jobs-stages-data-shuffling-shahzad-aslam/>

Thank You!

See you next week.