


# Week 7

## FIT5202 Big Data Processing

K-Means Clustering

Model Selection

Updated by CM Ting – 11 April 2025



# Week 7 Agenda

- Part - A

- Week 6 Review
- K-means Clustering
  - Silhouette Score
- Tutorial Instructions
  - Use case : Identify if 3 hackers were involved

- Part - B

- Model Selection
  - Hyperparameter Tuning
  - Cross Validation
    - K-fold Cross Validation
  - TrainValidationSplit
- Model Persistence
  - Saving and Loading a Model

# Model Selection (a.k.a Hyperparameter Tuning)

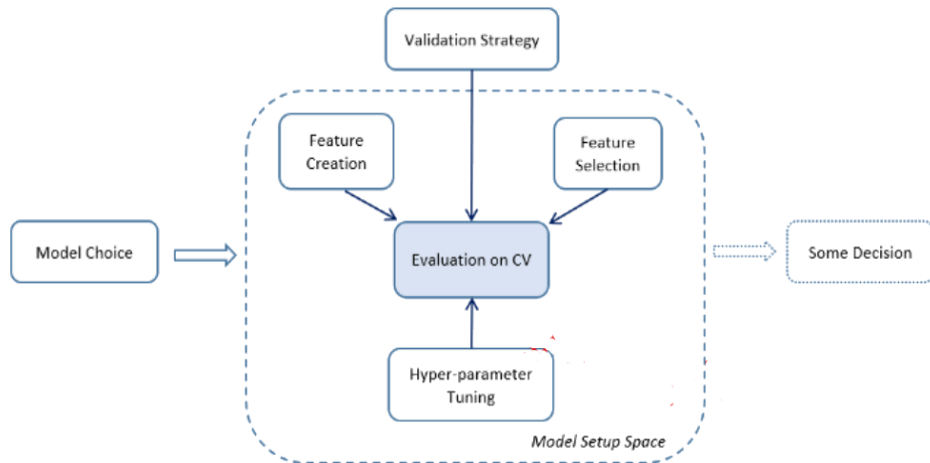
All models are wrong; some are useful (George E.P. Box)

- **HyperParameter Tuning**
- Finding the best model or parameters (e.g. **maxDepth of DT, number of clusters in k-means clustering**)
- Tuning can be done for individual Estimators or the entire Pipeline



Model selection for MLlib has the following tools:

1. CrossValidator
2. TrainValidationSplit




<https://spark.apache.org/docs/latest/ml-tuning.html>

<https://medium.com/pythoneers/hyperparameter-tuning-in-data-science-ad1a03d830b9>

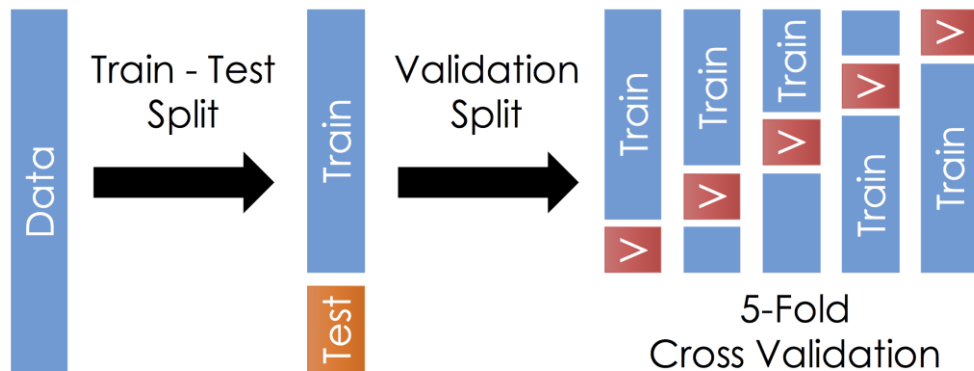
# Hyperparameter Tuning

- Hyper-parameters **are not model parameters** : they cannot be trained from the data
- Hyperparameter tuning : choosing a set of optimal hyperparameters for a learning algorithm
- `model.extractParamMap()` to get the list of hyperparameters for the model



# Cross Validation (K-Fold)

- Splitting dataset into a set of folds, which are used as separate training and test datasets.



## Model Selection:

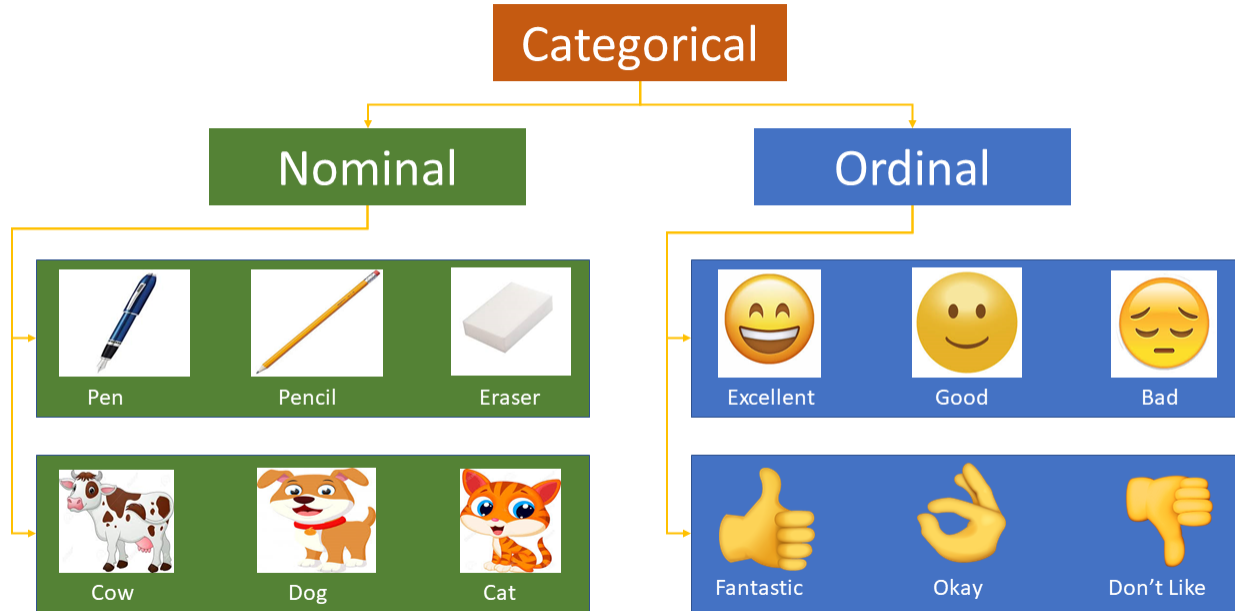
- ❑ Evaluate performance over a range of model hyper-parameters on validation set,
- ❑ Choose the model which give highest performance



Why not just tune hyperparamters on the test set?

# Categorical features

Categorical variables represent types of data which may be divided into groups.



No ordering

The variables have natural,  
ordered categories

# DT Hyperparameter: maxBins

## Continuous features

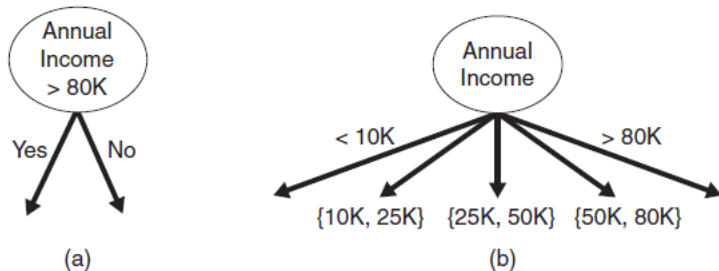


Figure 4.11. Test condition for continuous attributes.

- ❑ The test condition can be expressed as a comparison test ( $A < v$ ) and ( $A > v$ ) with binary outcome, or a **range of outcomes**  $v_i < A < v_{i+1}$  for  $i=1, \dots, k$
- ❑ For binary tree, algorithm will consider all split position  $v$  (splitting point / threshold)

## Example


Consider variable  $X$  with instances [1,3,4,6,2,5,18,10,-3,-5]

We can sort data, and **cluster data into bins to choose splitting point** (e.g., -1, 2.5, 4.5, and 8)

[-5, -3, 1, 2, 3, 4, 5, 6, 10, 18]

[-5, -3], [1, 2], [3, 4], [5, 6], [10, 18]

Maximum number of bins can be specified using **maxBins**.



If **maxBins** is large, more splitting points to consider in building the tree.

# Cross Validation (Decision Tree)

## Grid-based Hyperparameter tuning

```
from pyspark.ml.tuning import ParamGridBuilder, CrossValidator, CrossValidatorModel
from pyspark.ml.evaluation import BinaryClassificationEvaluator
# Create ParamGrid for Cross Validation
dtparamGrid = (ParamGridBuilder()
               .addGrid(dt.maxDepth, [2, 5, 10, 20, 30])
               .addGrid(dt.maxBins, [10, 20, 40, 80, 100])
               .build())
```

```
dtevaluator = BinaryClassificationEvaluator(rawPredictionCol="rawPrediction")
```

```
dtcv = CrossValidator(estimator = pipeline,
                      estimatorParamMaps = dtparamGrid,
                      evaluator = dtevaluator,
                      numFolds = 3)
```

```
dtcvModel = dtcv.fit(train)
```

```
bestModel = dtcvModel.bestModel
```

```
print('Best Param (regParam): ', bestModel.stages[-1]._java_obj.paramMap())
```

```
Best Param for DT: {
  DecisionTreeClassifier_ba35db4d44b0-featuresCol: features,
  DecisionTreeClassifier_ba35db4d44b0-labelCol: label,
  DecisionTreeClassifier_ba35db4d44b0-maxBins: 20,
  DecisionTreeClassifier_ba35db4d44b0-maxDepth: 20
}
```

maxBins

maxDepth

	2	5	10	20	30
10					
20					
40					
80					
100					

- Evaluate performance for each pair of hyperparameters on validation set
- Choose the best set of hyperparameters



# TrainValidationSplit

- Creates a single dataset pair
- Only evaluates each combination of parameter once as opposed to k-times in case of CrossValidator
- Less expensive but not reliable if the training dataset is not large enough


```
: from pyspark.ml.tuning import TrainValidationSplit
  dtparamGrid = (ParamGridBuilder()
    .addGrid(dt.maxDepth, [2, 5, 10, 20, 30])
    .addGrid(dt.maxBins, [10, 20, 40, 80, 100])
    .build())

  dttv = TrainValidationSplit(estimator = pipeline,
    estimatorParamMaps = dtparamGrid,
    evaluator = dtevaluator,
    trainRatio = 0.8)

  model = dttv.fit(train)

: bestModel_tv = model.bestModel

: print(bestModel_tv.stages[-1]._java_obj.paramMap())
```



# K-Means Clustering

Finds groups (or clusters) of data

A cluster comprises a number of “similar” objects

A member is closer to another member within the same group than to a member of a different group


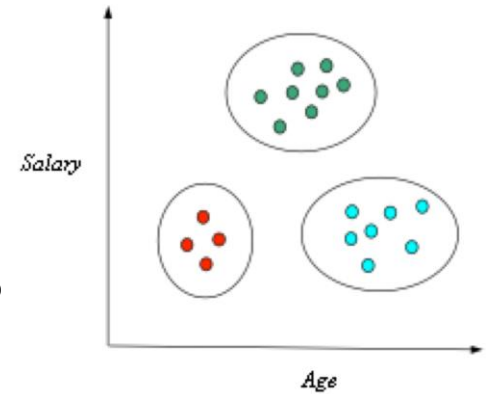
Groups have no category or label

Unsupervised learning

[Animation Demo](#) , [DEMO 2](#)

**Silhouette Score [-1 1]** : calculates the goodness of a clustering technique

- **1** - Clusters are well apart from each other and clearly distinguishes
- **0** - Clusters are not clearly distinguished, the distance between the clusters is not significant (overlapping cluster)
- **-1** – Clusters assigned wrongly

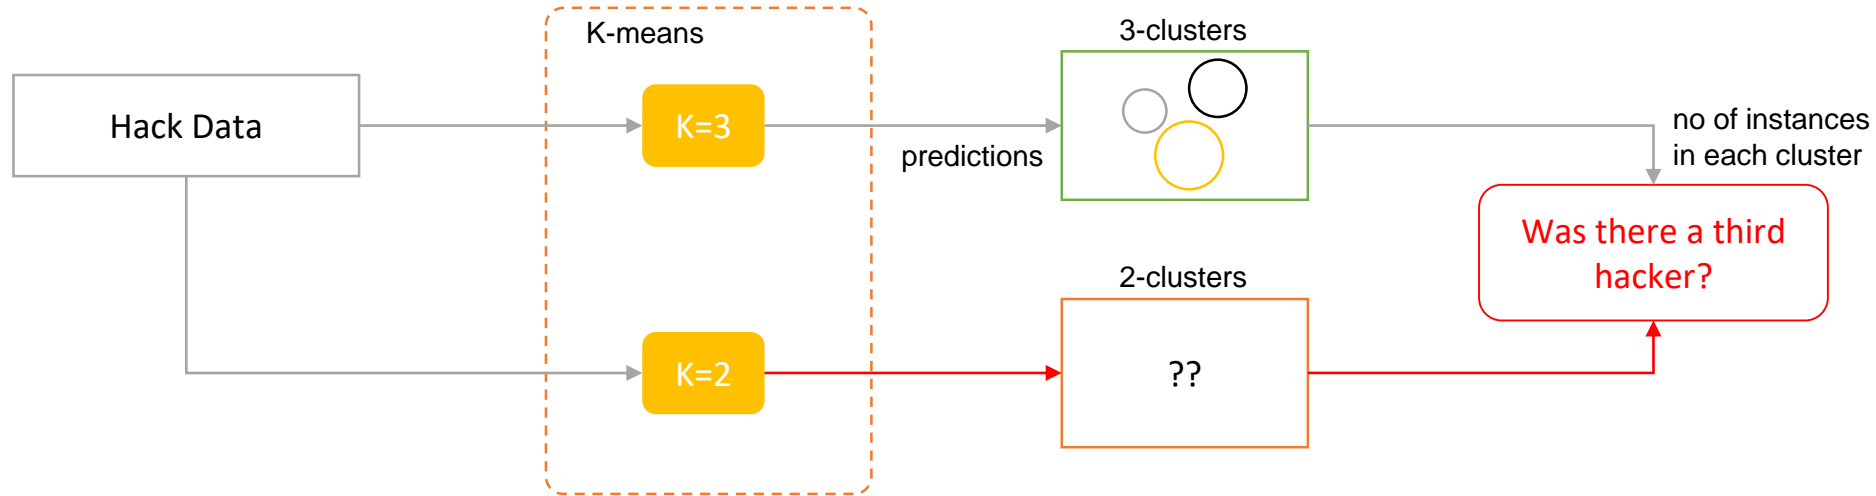


# Use case : Was there a third hacker?

**Assumption:**

Each cluster should have the same number of records

**Assumption :** Hackers trade off attacks equally



Feature transformation:

- (1) Vector Assembler
- (2) StandardScaler (normalizing the features to have mean 0 and variance 1)

**Thank You!**

See you next week.