

COPYRIGHT WARNING: Copyright in these original lectures is owned by Monash University. You may transcribe, take notes, download or stream lectures for the purpose of your research and study only. If used for any other purpose, (excluding exceptions in the Copyright Act 1969 (Cth)) the University may take legal action for infringement of copyright.

Do not share, redistribute, or upload the lecture to a third party without a written permission!

FIT3181/5215 Deep Learning

Week 04: Back-Propagation and Optimization for Deep Learning

Lecturers: Dr Trung Le

A/Prof Zongyuan Ge

Dr Arghya Pal

Email: trunglm@monash.edu



Outline

- Revision of calculus.
- Gradient descent and stochastic gradient descent.
- Backpropagation in feed-forward neural networks.
- Optimizers for deep learning.
- **Further reading recommendations**
 - Deep Learning – Chapter 8
 - Dive into Deep Learning – Chapter 11 (https://d2l.ai/chapter_optimization/index.html)
 - Ruder's blog: <https://ruder.io/optimizing-gradient-descent/index.html>

Our Story

```
class OurFFN(torch.nn.Module):
    def __init__(self, n_features, n_classes):
        super(OurFFN, self).__init__()
        self.n_features = n_features
        self.n_classes = n_classes
        self.fc1 = nn.Linear(n_features, 10)
        self.fc2 = nn.Linear(10, 20)
        self.fc3 = nn.Linear(20, 15)
        self.fc4 = nn.Linear(15, n_classes)

    def forward(self, x): #x is the mini-batch
        x = self.fc1(x)
        x = F.relu(x)
        x = self.fc2(x)
        x = F.relu(x)
        x = self.fc3(x)
        x = F.relu(x)
        x = self.fc4(x)
        return x
```

Declare the model

```
# Loss and optimizer
learning_rate = 0.005
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(ffn_model.parameters(), lr=learning_rate)

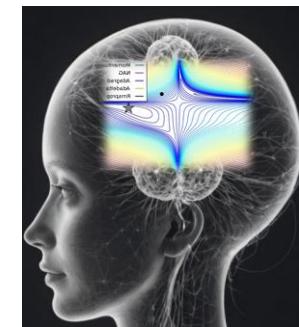
for epoch in range(num_epochs):
    for i, (X, y) in enumerate(train_loader):
        X, y = X.to(device), y.to(device)
        # Forward pass
        outputs = ffn_model(X.type(torch.float32))
        loss = loss_fn(outputs, y.type(torch.long))

        # Backward and optimize
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

Train the model

Declare loss and optimizer

Optimizers



`optimizer.step()`

SINCE 1986

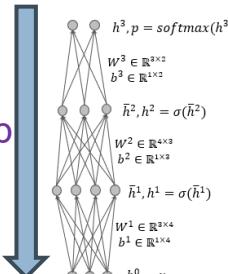


Back-propagation



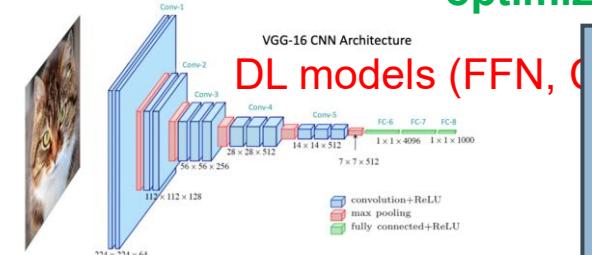
`loss.backward()`

Backprop



Compute gradients

$$\frac{\partial l}{\partial W}, \frac{\partial l}{\partial b}$$



Update

$$W = W - \eta \frac{\partial l}{\partial W}$$

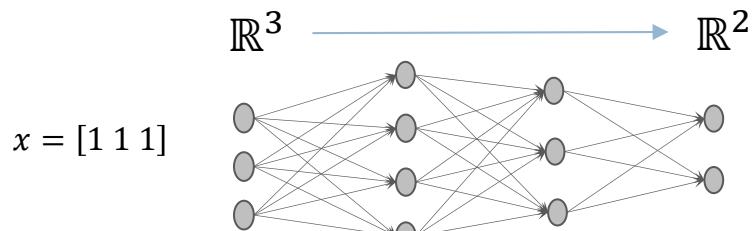
$$b = b - \eta \frac{\partial l}{\partial b}$$

Small Detour to Calculus

A small detour to calculus

- Calculus = **mathematics of change** (very important for deep learning)
- Properties of derivative:
 - $f'(x) = \nabla f(x) = \lim_{h \rightarrow 0} \frac{f(x+h)-f(x)}{h}$
 - $(uv)' = u'v + uv'$
 - $\left(\frac{u}{v}\right)' = \frac{u'v - uv'}{v^2}$
 - $(e^u)' = u'e^u$
 - $(\log u)' = \frac{u'}{u}$
- Multi-variate function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ with $y = f(x) = f(x_1, \dots, x_n)$.
 - Gradient/derivative: $\frac{\partial f}{\partial x}(a) = \nabla_x f(a) = [\nabla_{x_1} f(a), \nabla_{x_2} f(a), \dots, \nabla_{x_n} f(a)]$.
- Chain rule ∞ :
 - $\frac{\partial u}{\partial x} = \frac{\partial u}{\partial v} \times \frac{\partial v}{\partial x}$

Derivative for multi-variate functions



□ Given a function $f: \mathbb{R}^m \rightarrow \mathbb{R}^n$

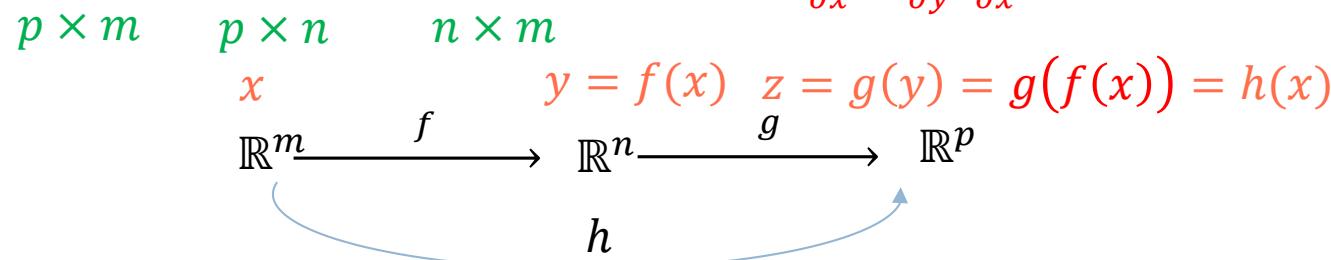
- $f(x) = (f_1(x), \dots, f_n(x))$ where $f_1, \dots, f_n: \mathbb{R}^m \rightarrow \mathbb{R}$ and $x = (x_1, \dots, x_m)$. Let denote $y = f(x)$.
- The **derivative** of f at the point $a \in \mathbb{R}^m$, denoted by $\nabla f(a)$ (**function related** notion) or $\frac{\partial y}{\partial x}(a)$ (**variable related** notion) is a matrix n by m (i.e., the Jacobian matrix).

$$\frac{\partial y}{\partial x}(a) = \nabla f(a) = \begin{bmatrix} \frac{\partial f_1}{\partial x_1}(a) & \cdots & \cdots & \frac{\partial f_1}{\partial x_j}(a) & \cdots & \frac{\partial f_1}{\partial x_m}(a) \\ \cdots & \cdots & \cdots & \cdots & \cdots & \cdots \\ \cdots & \cdots & \cdots & \cdots & \cdots & \cdots \\ \frac{\partial f_i}{\partial x_1}(a) & \cdots & \cdots & \frac{\partial f_i}{\partial x_j}(a) & \cdots & \frac{\partial f_i}{\partial x_m}(a) \\ \cdots & \cdots & \cdots & \cdots & \cdots & \cdots \\ \frac{\partial f_n}{\partial x_1}(a) & \cdots & \cdots & \frac{\partial f_n}{\partial x_j}(a) & \cdots & \frac{\partial f_n}{\partial x_m}(a) \end{bmatrix}_{n \times m}$$

Jacobian matrix

□ Chain rule ∞

- Given a function $f: \mathbb{R}^m \rightarrow \mathbb{R}^n$, $g: \mathbb{R}^n \rightarrow \mathbb{R}^p$, denote $h = g \circ f: \mathbb{R}^m \rightarrow \mathbb{R}^p$, meaning that $h(x) = g(f(x))$. We further define $y = f(x)$ and $z = g(y) = g(f(x)) = h(x)$.
- For $x \in \mathbb{R}^m$, $\nabla h(x) = \nabla g(f(x)) \times \nabla f(x)$ or equivalently $\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \cdot \frac{\partial y}{\partial x}$.



Example

□ $y = f(x) = f(x_1, x_2, x_3) = (x_1^2 + x_2^2, x_2^2 + x_3^2 x_2)$

- $f: \mathbb{R}^3 \rightarrow \mathbb{R}^2$
- $f_1(x) = f_1(x_1, x_2, x_3) = x_1^2 + x_2^2$
- $f_2(x) = f_2(x_1, x_2, x_3) = x_2^2 + x_3^2 x_2$
- $\frac{\partial y}{\partial x} = \nabla f \in \mathbb{R}^{2 \times 3}$

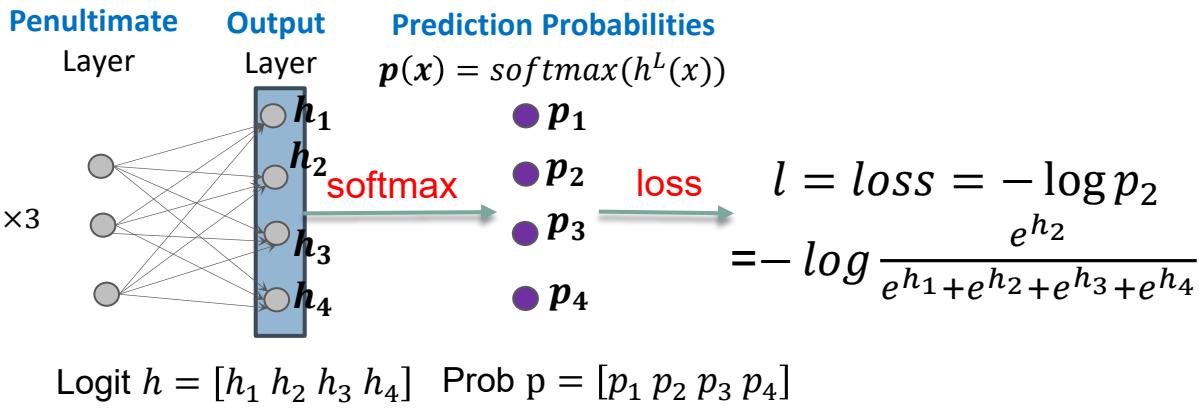
$$\frac{\partial y}{\partial x} = \nabla_x f = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \frac{\partial f_1}{\partial x_3} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \frac{\partial f_2}{\partial x_3} \end{bmatrix} = \begin{bmatrix} 2x_1 & 2x_2 & 0 \\ 0 & 2x_2 + x_3^2 & 2x_2 x_3 \end{bmatrix}$$

Example

Output layer

$$x = [x_1 \ x_2 \ x_3] \in \mathbb{R}^{1 \times 3}$$

$$y = 2$$



$$p_1 = \frac{e^{h_1}}{e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4}}$$

$$p_2 = \frac{e^{h_2}}{e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4}}$$

$$p_3 = \frac{e^{h_3}}{e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4}}$$

$$p_4 = \frac{e^{h_4}}{e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4}}$$

Compute $\frac{\partial l}{\partial h}$?

- $l = -\log \frac{e^{h_2}}{e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4}} = \log \underbrace{(e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4})}_u - h_2$
- $\frac{\partial l}{\partial h_1} = \frac{\nabla_{h_1} u}{u} = \frac{e^{h_1}}{e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4}} = p_1$
- $\frac{\partial l}{\partial h_2} = \frac{\nabla_{h_2} u}{u} - 1 = \frac{e^{h_2}}{e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4}} - 1 = p_2 - 1$
- $\frac{\partial l}{\partial h_3} = \frac{\nabla_{h_3} u}{u} = \frac{e^{h_3}}{e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4}} = p_3$
- $\frac{\partial l}{\partial h_4} = \frac{\nabla_{h_4} u}{u} = \frac{e^{h_4}}{e^{h_1} + e^{h_2} + e^{h_3} + e^{h_4}} = p_4$
- $\frac{\partial l}{\partial h} = [p_1, p_2 - 1, p_3, p_4] = [p_1, p_2, p_3, p_4] - [0, 1, 0, 0] = p - \mathbf{1}_2 = p - \mathbf{1}_y$

Example

Intermediate layer

□ $\bar{h} = xW + b$ and $h = \sigma(\bar{h})$

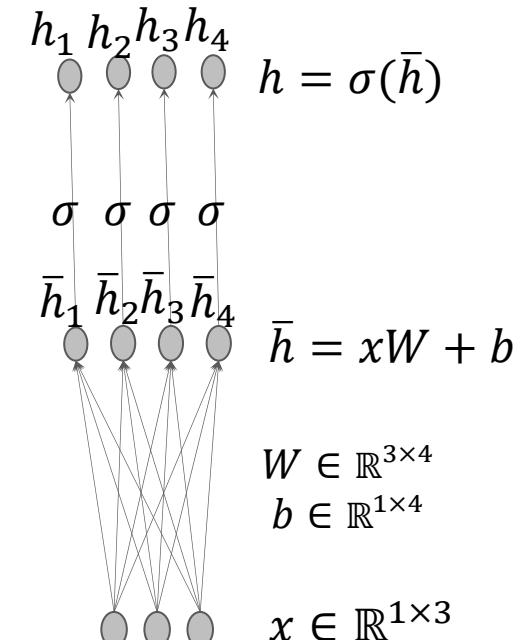
- $h = \sigma(xW + b)$

- σ is the activation function

$$\square \frac{\partial h}{\partial x} = \frac{\partial h}{\partial \bar{h}} \times \frac{\partial \bar{h}}{\partial x} = \text{diag}(\sigma'(\bar{h})) W^T \in \mathbb{R}^{4 \times 3}$$

$$\square \frac{\partial h}{\partial \bar{h}} = \begin{bmatrix} \frac{\partial h_1}{\partial \bar{h}_1} & \frac{\partial h_1}{\partial \bar{h}_2} & \frac{\partial h_1}{\partial \bar{h}_3} & \frac{\partial h_1}{\partial \bar{h}_4} \\ \frac{\partial h_2}{\partial \bar{h}_1} & \frac{\partial h_2}{\partial \bar{h}_2} & \frac{\partial h_2}{\partial \bar{h}_3} & \frac{\partial h_2}{\partial \bar{h}_4} \\ \frac{\partial h_3}{\partial \bar{h}_1} & \frac{\partial h_3}{\partial \bar{h}_2} & \frac{\partial h_3}{\partial \bar{h}_3} & \frac{\partial h_3}{\partial \bar{h}_4} \\ \frac{\partial h_4}{\partial \bar{h}_1} & \frac{\partial h_4}{\partial \bar{h}_2} & \frac{\partial h_4}{\partial \bar{h}_3} & \frac{\partial h_4}{\partial \bar{h}_4} \end{bmatrix} = \begin{bmatrix} \sigma'(\bar{h}_1) & 0 & 0 & 0 \\ 0 & \sigma'(\bar{h}_2) & 0 & 0 \\ 0 & 0 & \sigma'(\bar{h}_3) & 0 \\ 0 & 0 & 0 & \sigma'(\bar{h}_4) \end{bmatrix} = \text{diag}(\sigma'(\bar{h}))$$

$$\square \frac{\partial \bar{h}}{\partial x} = W^T$$



How to code with PyTorch

- $\bar{h} = xW + b$ and $h = \text{sigmoid}(\bar{h})$
 - $h = \text{sigmoid}(xW + b)$
 - $\sigma = \text{sigmoid}$ is the activation function
- $\frac{\partial h}{\partial x} = \frac{\partial h}{\partial \bar{h}} \times \frac{\partial \bar{h}}{\partial x} = \text{diag}(\sigma'(\bar{h}))W^T = \text{diag}(\sigma(\bar{h}) \otimes [1 - \sigma(\bar{h})])W^T = \text{diag}(h \otimes (1 - h))W^T$
 - \otimes is element-wise product

```
x = torch.tensor([1,-1,1], dtype=torch.float32)
print(x)
W = torch.rand(3,4)
b = torch.rand(1,4)
print(W)
print(b)

tensor([ 1., -1.,  1.])
tensor([[0.7266,  0.7925,  0.7952,  0.2159],
        [0.3108,  0.1950,  0.6448,  0.6367],
        [0.0521,  0.2200,  0.5038,  0.9020]])
tensor([[0.3059,  0.2612,  0.9326,  0.5597]])
```

Declare W, x, b

```
hbar = x@W+b
print(hbar)

tensor([[0.7737, 1.0787, 1.5868, 1.0409]])
```

```
h = torch.nn.Sigmoid()(hbar)
print(h)

tensor([[0.6843, 0.7463, 0.8302, 0.7390]])
```

Forward propagation

```
v = h*(1-h)
v = v.squeeze()
print(v)
```

```
tensor([0.2160, 0.1894, 0.1410, 0.1929])
A = torch.diag(v)
print(A)
```

```
tensor([[0.2160, 0.0000, 0.0000, 0.0000],
        [0.0000, 0.1894, 0.0000, 0.0000],
        [0.0000, 0.0000, 0.1410, 0.0000],
        [0.0000, 0.0000, 0.0000, 0.1929]])
```

```
derivative = A@W.T
print(derivative)
```

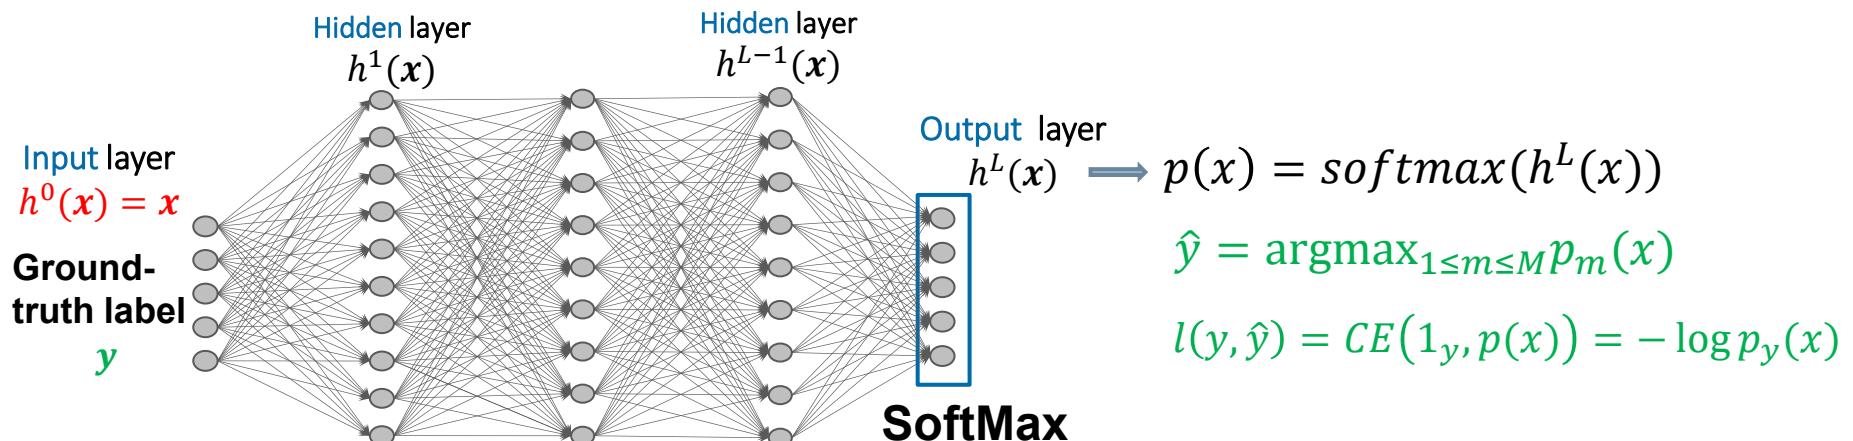
```
tensor([[0.1570, 0.0671, 0.0113],
        [0.1501, 0.0369, 0.0417],
        [0.1121, 0.0909, 0.0710],
        [0.0416, 0.1228, 0.1740]])
```

Backward propagation



Loss Landscape of DL models

Recall optimization problem in deep learning



Training set

$$D = \{(x_1, y_1), \dots, (x_N, y_N)\}$$

Loss function

$$L(D; \theta) := \frac{1}{N} \sum_{i=1}^N CE(1_{y_i}, p(x_i)) = -\frac{1}{N} \sum_{i=1}^N \log p_{y_i}(x_i)$$

□ How to **solve the optimization problem** efficiently ($\theta := \{(W^l, b^l)\}_{l=1}^L$)?

- $\min_{\theta} L(D; \theta) := -\frac{1}{N} \sum_{i=1}^N \log p_{y_i}(x_i) = -\frac{1}{N} \sum_{i=1}^N \log \frac{\exp\{h_{y_i}^L(x_i)\}}{\sum_{m=1}^M \exp\{h_m^L(x_i)\}}$
- Generalize: $\min_{\theta} J(\theta) := \frac{1}{N} \sum_{i=1}^N l(f(x_i; \theta), y_i)$

Optimization problem in ML and DL

- Most of optimization problems (OP) in machine learning (deep learning) has the following form:

$$\min_{\theta} J(\theta) = \underbrace{\Omega(\theta)}_{\text{Regularization term}} + \frac{1}{N} \sum_{i=1}^N l(y_i, f(x_i; \theta))$$

- **Occam's Razor principle:** prefer simplest model that can well predict data.

Regularization term

- $\Omega(\theta) = \lambda \sum_k \sum_{i,j} (W_{i,j}^k)^2 = \lambda \sum_k \|W^k\|_F^2$
- Encourage simple models
- Avoid overfitting

Empirical loss

- Work well on training set

How to efficiently solve this optimization problem?

N is the **training size** and might be very big (e.g., $N \approx 10^6$)

First-order iterative methods (gradient descent, steepest descent)

Use the **gradient** (first derivative) $g = \nabla_{\theta} J(\theta)$ to update parameters

Second-order iterative methods (Newton and quasi Newton methods)

Use the **Hessian** matrix (second derivative) $H = \nabla_{\theta}^2 J(\theta)$ to update parameters

Gradient and Hessian matrix

- Given an **objective function** $J(\theta)$ with $\theta = [\theta_1, \theta_2, \dots, \theta_P]$
 - For DL models
 - θ includes **weight matrices**, **filters**, and **biases** which are trainable model parameters.
 - P is the number of **trainable parameters** (P could be 20×10^6).
 - $J(\theta)$ is the loss function over a training set.
- Gradient $g = \nabla J(\theta)$ is the **first order derivative** and defined as

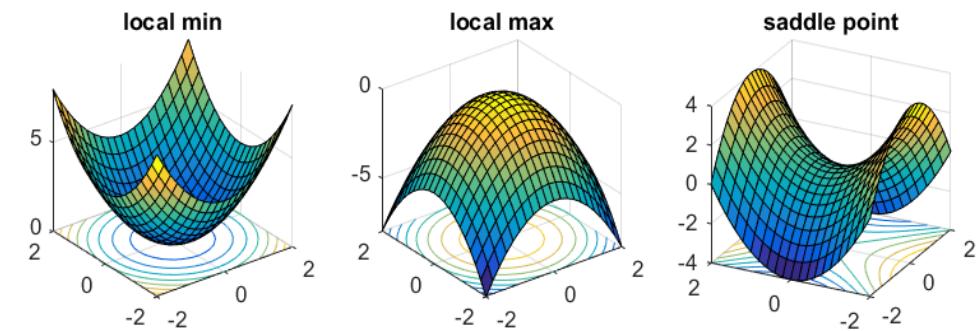
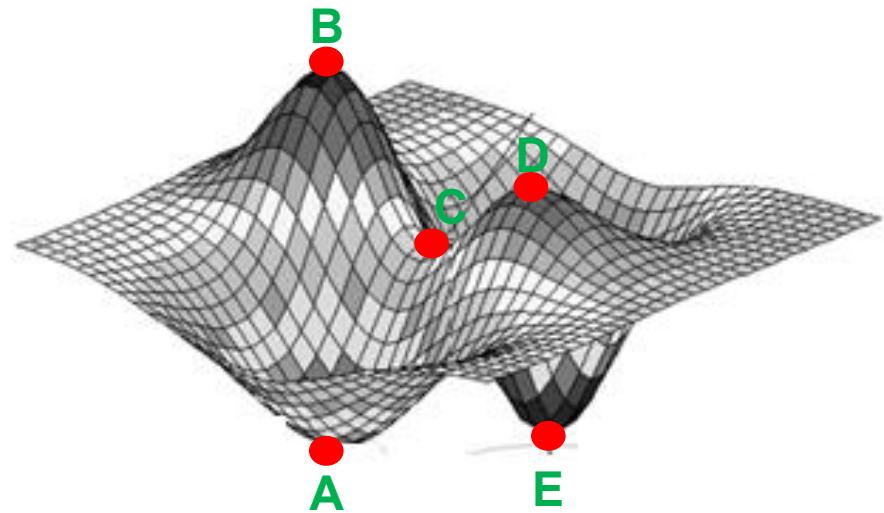
$$\circ \quad \nabla J(\theta) = g = \begin{bmatrix} \frac{\partial J}{\partial \theta_1}(\theta) \\ \cdots \\ \frac{\partial J}{\partial \theta_P}(\theta) \end{bmatrix}$$

- Hessian matrix $H(\theta)$ is the **second order derivative** $\nabla^2 J(\theta)$ and defined as

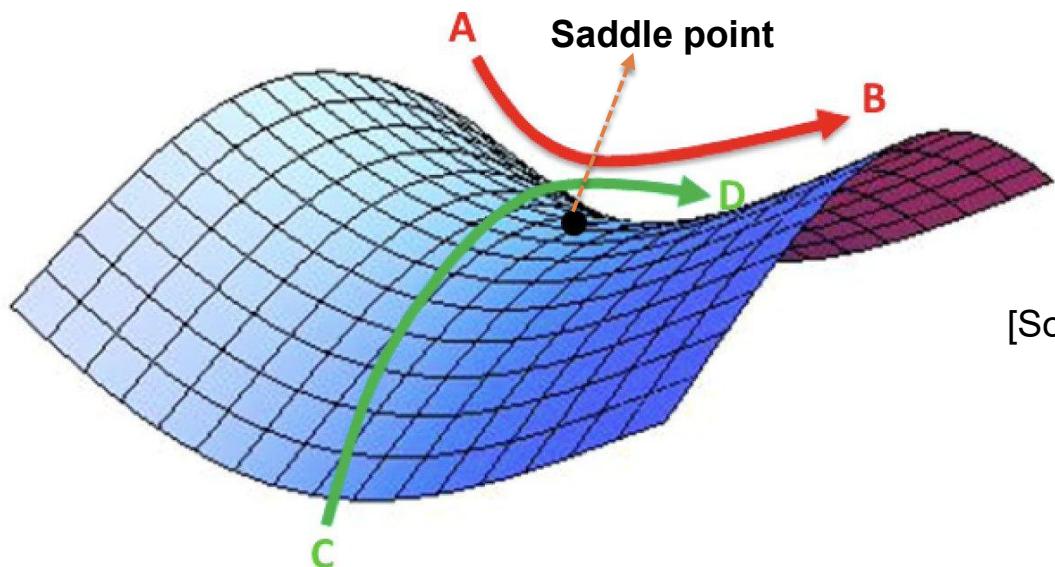
$$\circ \quad \nabla^2 J(\theta) = H(\theta) = \begin{bmatrix} \frac{\partial^2 J}{\partial \theta_1 \partial \theta_1}(\theta) & \cdots & \cdots & \frac{\partial^2 J}{\partial \theta_1 \partial \theta_j}(\theta) & \cdots & \frac{\partial^2 J}{\partial \theta_1 \partial \theta_P}(\theta) \\ \cdots & \cdots & \cdots & \cdots & \cdots & \cdots \\ \cdots & \cdots & \cdots & \cdots & \cdots & \cdots \\ \frac{\partial^2 J}{\partial \theta_i \partial \theta_1}(\theta) & \cdots & \cdots & \frac{\partial^2 J}{\partial \theta_i \partial \theta_j}(\theta) & \cdots & \frac{\partial^2 J}{\partial \theta_i \partial \theta_P}(\theta) \\ \cdots & \cdots & \cdots & \cdots & \cdots & \cdots \\ \frac{\partial^2 J}{\partial \theta_P \partial \theta_1}(\theta) & \cdots & \cdots & \frac{\partial^2 J}{\partial \theta_P \partial \theta_j}(\theta) & \cdots & \frac{\partial^2 J}{\partial \theta_P \partial \theta_P}(\theta) \end{bmatrix}$$

Local minima-maxima and saddle point

- Given an objective function $J(\theta)$ with $\theta = [\theta_1, \theta_2, \dots, \theta_P]$
 - θ is said to be a **critical point** if $\nabla J(\theta) = \mathbf{0}$ (vector $\mathbf{0}$)
- Let us denote the **set of eigenvalues** of Hessian matrix $\nabla^2 J(\theta) = H(\theta)$ by
 - $\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_P$
- Local minima
 - $\nabla J(\theta) = \mathbf{0}$ and $\nabla^2 J(\theta) = H(\theta) > 0$ (positive semi-definite matrix)
 - $0 \leq \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_P$
- Local maxima
 - $\nabla J(\theta) = \mathbf{0}$ and $\nabla^2 J(\theta) = H(\theta) < 0$ (negative semi-definite matrix)
 - $\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_P \leq 0$
- Saddle point
 - $\nabla J(\theta) = \mathbf{0}$ and $\nabla^2 J(\theta) = H(\theta) \prec 0$ (indefinite matrix)
 - $\lambda_1 \leq \lambda_2 \leq \dots < 0 < \dots \leq \lambda_P$



More on saddle point



[Source: Internet]

$$f(\theta) = f(\theta_1, \theta_2) = \theta_1^2 - \theta_2^2$$

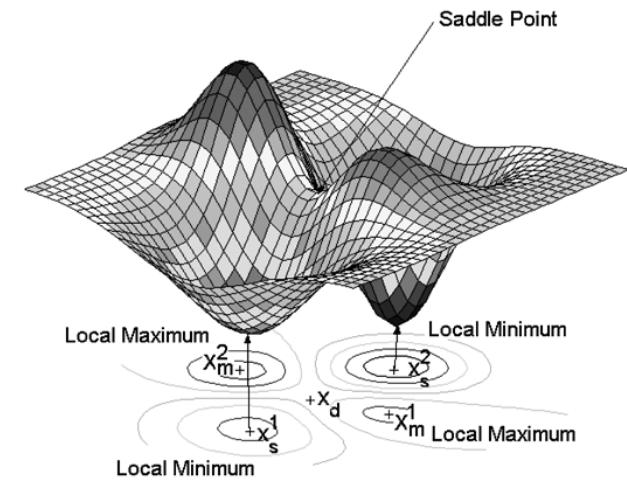
Gradient $\mathbf{g} = \begin{bmatrix} \frac{\partial f}{\partial \theta_1} \\ \frac{\partial f}{\partial \theta_2} \end{bmatrix} = \begin{bmatrix} 2\theta_1 \\ -2\theta_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix} \rightarrow$ a **critical point** $\begin{bmatrix} 0 \\ 0 \end{bmatrix}$.

Hessian matrix is $\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial \theta_1^2} & \frac{\partial^2 f}{\partial \theta_1 \partial \theta_2} \\ \frac{\partial^2 f}{\partial \theta_2 \partial \theta_1} & \frac{\partial^2 f}{\partial \theta_2^2} \end{bmatrix} = \begin{bmatrix} 2 & 0 \\ 0 & -2 \end{bmatrix}$

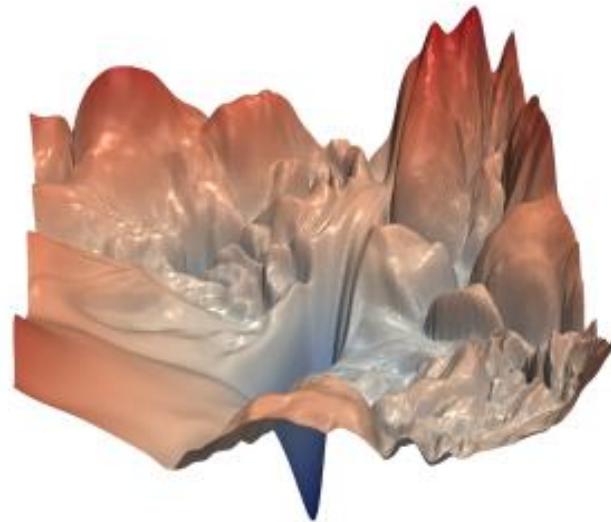
Two eigenvalues $\lambda_1 = -2 < 0 < 2 = \lambda_2 \rightarrow \begin{bmatrix} 0 \\ 0 \end{bmatrix}$ is a **saddle point**.

Numbers of local minima vs saddle points

- We assume to pick randomly a training set
 - The Hessian matrix $H(\theta)$ is a random matrix with **random eigenvalues** $\lambda_1, \lambda_2, \dots, \lambda_P$
 - We assume that $\mathbb{P}(\lambda_1 \geq 0) = \mathbb{P}(\lambda_2 \geq 0) = \dots = \mathbb{P}(\lambda_P \geq 0) = 0.5$
- Therefore, we have
 - $\mathbb{P}(\text{minima}) = \mathbb{P}(\lambda_1 \geq 0)\mathbb{P}(\lambda_2 \geq 0) \dots \mathbb{P}(\lambda_P \geq 0) = 0.5^P$
 - $\mathbb{P}(\text{maxima}) = \mathbb{P}(\lambda_1 \leq 0)\mathbb{P}(\lambda_2 \leq 0) \dots \mathbb{P}(\lambda_P \leq 0) = 0.5^P$
 - $\mathbb{P}(\text{saddle point}) = 1 - \mathbb{P}(\text{minima}) - \mathbb{P}(\text{maxima}) = 1 - 0.5^{P-1}$
- The ratio of #local minima/maxima against #saddle points
 - **#local-minima:#local-maxima:#saddle-point=1: 1: $(2^P - 2)$**
 - Number of saddle points is even **exponentially much more than that of local minima/maxima**



The loss surface of DL optimization problem



Loss surface of a ResNet without skip connection [Hao Li et al., NeurIPS 2017]

- The optimization problem in deep learning:

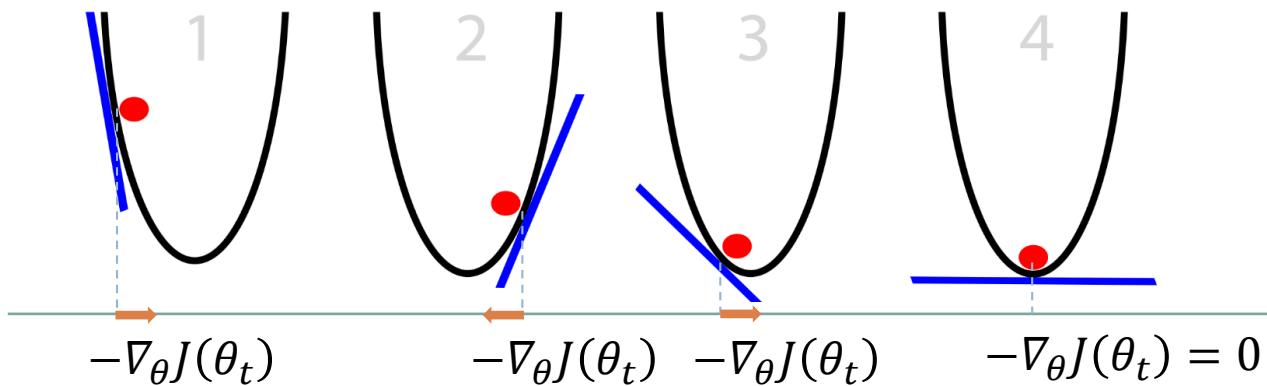
- $\min_{\theta} J(\theta) := L(D; \theta) := \frac{1}{N} \sum_{i=1}^N l(f(x_i; \theta), y_i) = -\frac{1}{N} \sum_{n=1}^N \log \frac{\exp\{h_{y_i}^L(x_i)\}}{\sum_{m=1}^M \exp\{h_m^L(x_i)\}}$

- A very **complex** and **complicated** objective function

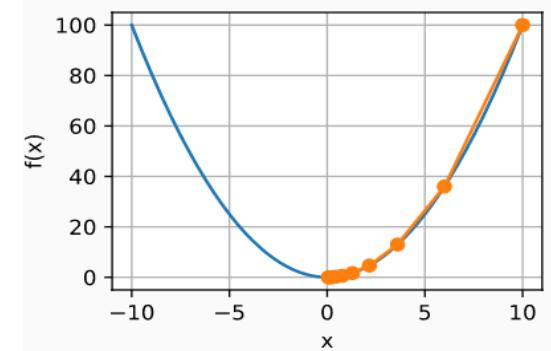
- Highly **non-linear** and **non-convex** function
 - The **loss surface** is very **complex**
 - Many local minima points, but the number of saddle points is even **exponentially** much more

Gradient descent and stochastic gradient descent

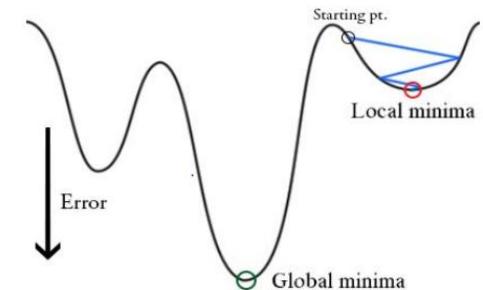
Gradient descend



- We need to solve
 - $\min_{\theta} J(\theta)$
- Follow to the opposite side of the current gradient
 - $\theta_{t+1} = \theta_t - \eta \nabla_{\theta}J(\theta_t)$ where $\eta > 0$ is the learning rate.
- Guarantee to converge to a **global minima** if $J(\cdot)$ is **convex**.
- Get stuck in a **local minima** or **saddle points** if $J(\cdot)$ is non-convex.

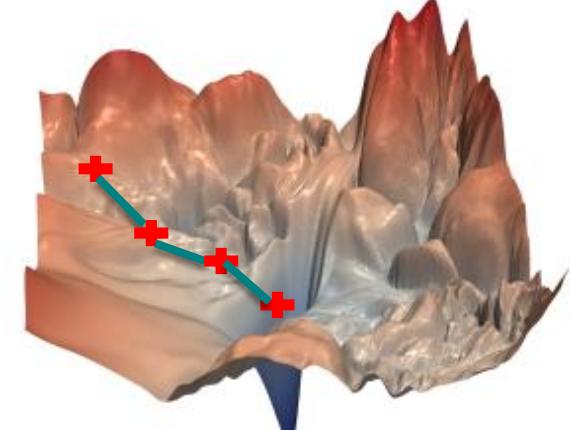


Convex case



(Source: www.cs.ubc.ca)

Non-convex case

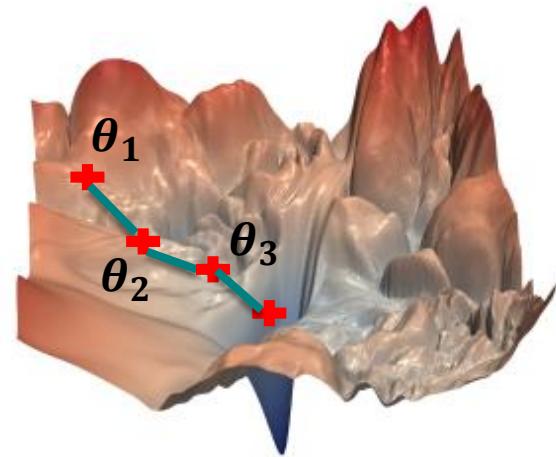


DL case: easy to get stuck in saddle points

Gradient descend

Algorithm

- **Input:** objective function $J(\theta)$
 - **Output:** optimal solution θ^*
1. Initialize parameters θ_0 randomly $\sim N(0, \sigma^2)$.
 2. for $t=1$ to T
 3. Compute gradients $\nabla_{\theta} J(\theta_t) = \frac{\partial J}{\partial \theta}(\theta_t)$
 4. Update $\theta_{t+1} = \theta_t - \eta_t \nabla_{\theta} J(\theta_t)$
 5. Return $\theta^* = \theta_{T+1}$



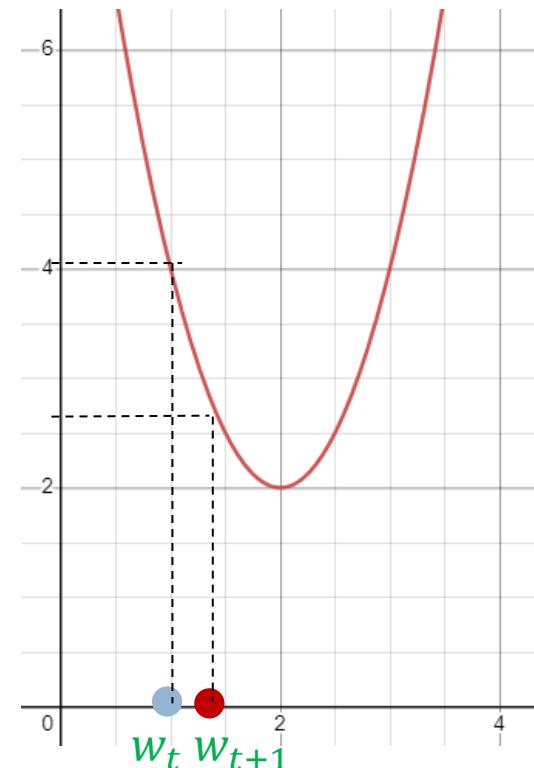
Example of Gradient Descend

- Consider the optimization problem:

$$\min_w [f(w) = (w - 1)^2 + (w - 3)^2]$$

- Currently, we are at $w_t = 1$ with $f(w_t) = f(1) = 0^2 + 2^2 = 4$. What is w_{t+1} if using learning rate $\eta = 0.1$?

- $f'(w) = 2(w - 1) + 2(w - 3) = 4w - 8$
- $f'(w_t) = f'(1) = -4$
- $w_{t+1} = w_t - \eta f'(w_t) = 1 - 0.1(-4) = 1.4$
- $f(w_{t+1}) = 2.72 < f(w_t) = 4$



Gradient descent for deep learning

- For training deep nets, we need to solve

- $\min_{\theta} L(D; \theta) := \frac{1}{N} \sum_{i=1}^N l(x_i, y_i; \theta) = \frac{1}{N} \sum_{i=1}^N l(y_i, f(x_i; \theta))$

where $l(x_i, y_i; \theta) = -\log p(y = y_i | x_i) = -\log \frac{\exp\{h_{y_i}^L(x_i)\}}{\sum_{m=1}^M \exp\{h_m^L(x_i)\}}$ is the loss incurred by (x_i, y_i) .

- Gradient descent update

- $\theta_{t+1} = \theta_t - \eta \nabla_{\theta} L(D; \theta_t) = \theta_t - \frac{\eta}{N} \sum_{i=1}^N \nabla_{\theta} l(x_i, y_i; \theta_t)$ where $\eta > 0$ is a learning rate.
- To compute the gradient $\nabla_{\theta} L(D; \theta_t) = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} l(x_i, y_i; \theta_t)$, we need to go through all data points in $D \rightarrow$ the computational cost is $O(N)$.

- This is very **computationally expensive** for big datasets ($N \approx 10^6$).
- How to **estimate** the gradient $\nabla_{\theta} L(D; \theta_t)$ more efficiently?

Stochastic gradient descent

- The optimization problem in deep learning has the form
 - $\min_{\theta} L(D; \theta) := \frac{1}{N} \sum_{i=1}^N l(x_i, y_i; \theta) = \frac{1}{N} \sum_{i=1}^N l(y_i, f(x_i; \theta))$
- Evaluation of the full gradient is expensive. We want to just estimate this gradient
 - Sample a mini-batch $i_1, i_2, \dots, i_b \sim Uni(\{1, 2, \dots, N\})$ where b is the mini-batch (batch) size.
 - The batch size is usually 32, 64, 128, 256, and so on.
 - Construct $\tilde{L}(\theta) := \frac{1}{b} \sum_{k=1}^b l(x_{i_k}, y_{i_k}; \theta)$ as the average loss of those in the current batch.
 - $E_{i_1, \dots, i_b} [\nabla_{\theta} \tilde{L}(\theta_t)] = \nabla_{\theta} L(D; \theta_t)$
 - $\nabla_{\theta} \tilde{L}(\theta_t) = \frac{1}{b} \sum_{k=1}^b \nabla_{\theta} l(x_{i_k}, y_{i_k}; \theta_t)$ is unbiased estimation of $\nabla_{\theta} L(D; \theta_t)$
 - $O(b)$ compares to $O(N)$.
- The update rule of SGD
 - $\theta_{t+1} = \theta_t - \eta_t \nabla_{\theta} \tilde{L}(\theta_t)$ with learning rate $\eta_t \propto O(\frac{1}{t})$
 - We use $\nabla_{\theta} \tilde{L}(\theta_t)$ as an unbiased estimate of the full gradient $\nabla_{\theta} L(D; \theta)$
 - How to compute $\nabla_{\theta} \tilde{L}(\theta_t)$ efficiently for deep networks?

Example of Stochastic Gradient Descent

- Given the function $f(w) = \frac{1}{1000} \sum_{i=1}^{1000} (w - i)^2$, we need to solve

$$\min_w f(w)$$

using SGD with the learning rate $\eta = 0.1$.

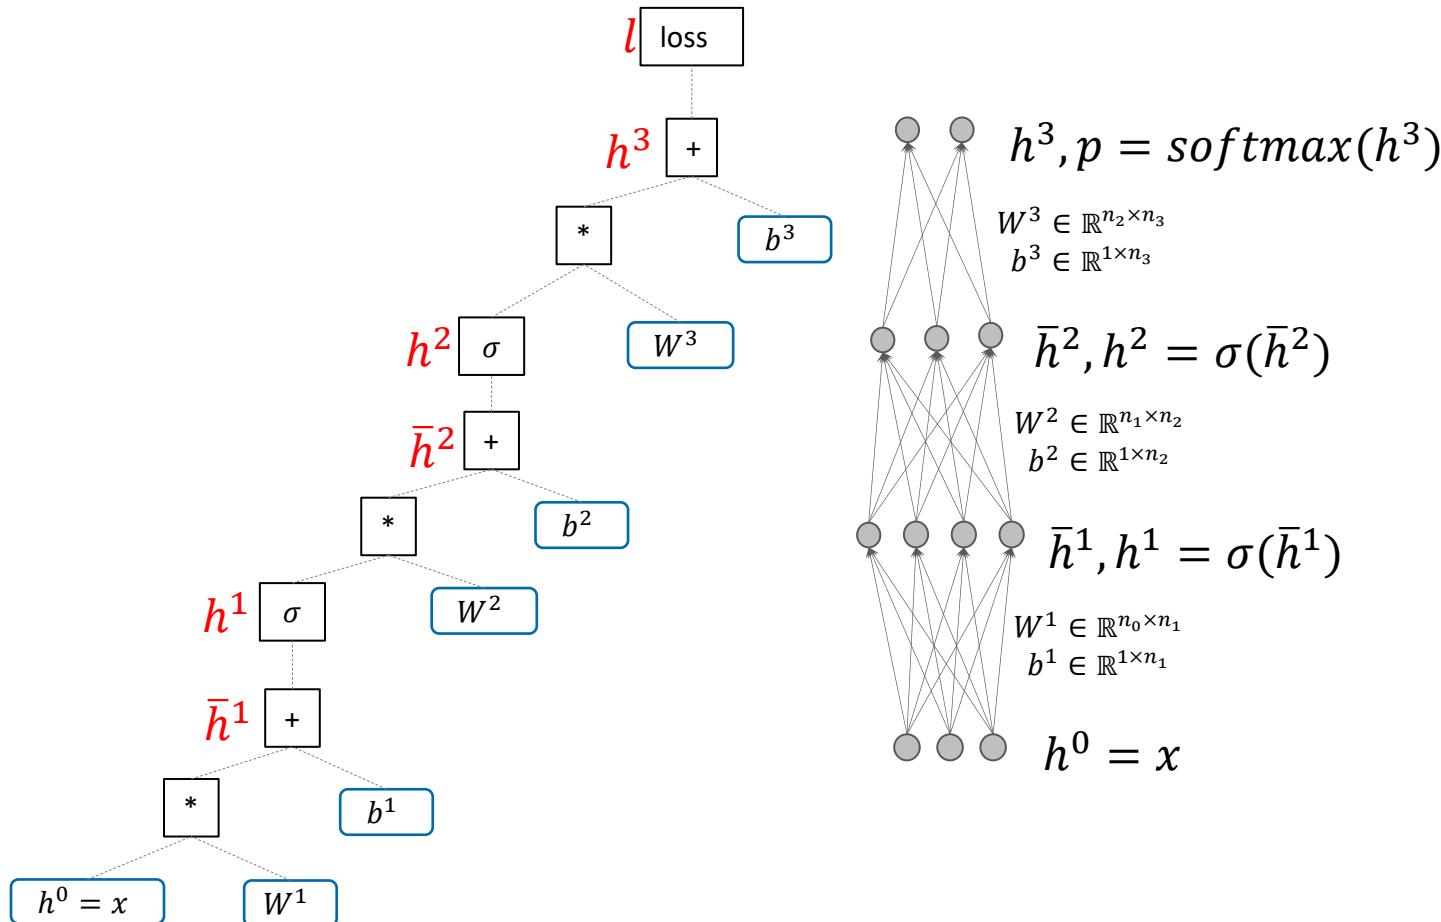
- Assume we sample a batch $i_1 = 1, i_2 = 2, i_3 = 3, i_4 = 4$ of indices. At the iteration t , $w_t = 10$. What is the value of w_{t+1} at the next iteration?

- $\tilde{f}(w) = \frac{1}{4} [(w - 1)^2 + (w - 2)^2 + (w - 3)^2 + (w - 4)^2]$
- $\tilde{f}'(w) = 2w - 5$
- $\tilde{f}'(w_t) = \tilde{f}'(10) = 2 \times 10 - 5 = 15$
- $w_{t+1} = w_t - \eta \tilde{f}'(w_t) = 10 - 0.1 \times 15 = 8.5$.



Back propagation in feed-forward neural networks

Back propagation



- Given a data point and label pair (x, y)
 - $$l(x, y; \theta) = -\log \frac{\exp\{h_y^3(x)\}}{\sum_{m=1}^M \exp\{h_m^3(x)\}}$$
- What are the derivatives?
 - $\nabla_{W^k} l(x, y; \theta)$ and $\nabla_{b^k} l(x, y; \theta)$ for $k = 1, 2, 3$?
 - Using **back propagation** to compute these derivatives conveniently.
- Update the model using SGD with the derivatives.

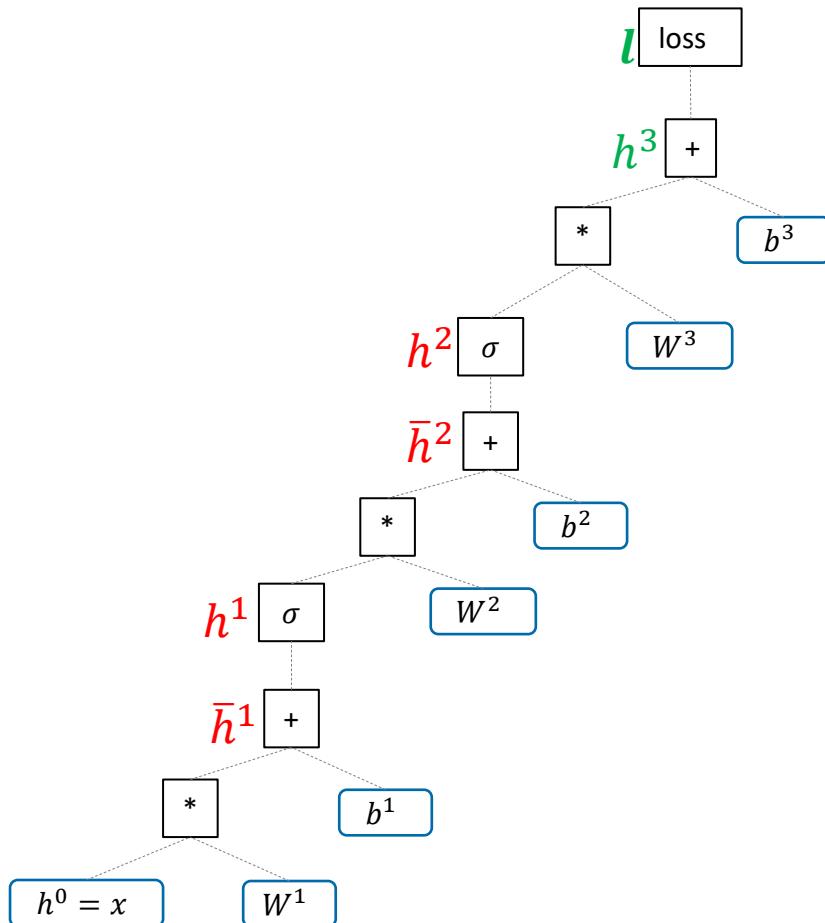
Back propagation

From loss to h^3



$$f(x, y, z) = \log(\exp(x) + \exp(y) + \exp(z))$$

$$\nabla f = [f'_x, f'_y, f'_z] = \text{softmax}([x, y, z])$$



- $$l(x, y; \theta) = -\log \frac{\exp\{h_y^3(x)\}}{\sum_{m=1}^M \exp\{h_m^3(x)\}} =$$

$$-\log\{h_y^3(x)\} + \log[\sum_{m=1}^M \exp\{h_m^3(x)\}] =$$

$$-\sum_{m=1}^M \mathbf{1}_{m=y} h_m^3 + \log[\sum_{m=1}^M \exp\{h_m^3\}]$$

where $\mathbf{1}_{m=y} = \mathbf{1}$ if $m = y$ and $\mathbf{0}$ otherwise.

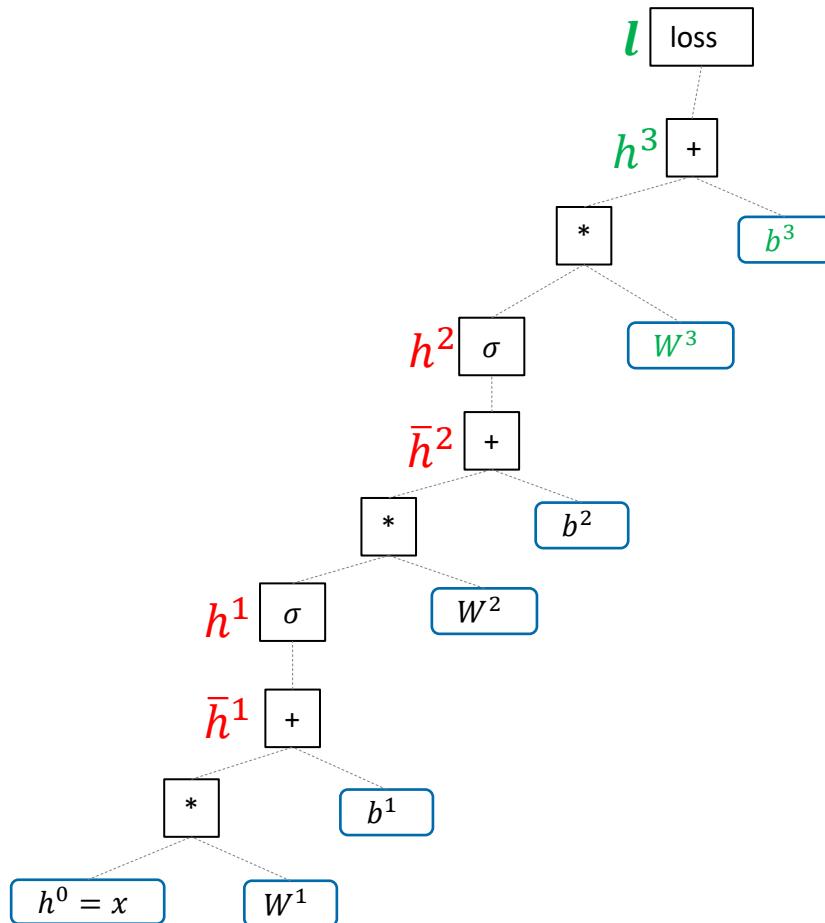
- $$\frac{\partial l}{\partial h_m^3} = -\mathbf{1}_{m=y} + \frac{\exp\{h_m^3\}}{\sum_{k=1}^M \exp\{h_k^3\}}, m = 1, \dots, M$$
- $$g^3 = \frac{\partial l}{\partial h^3} = -\mathbf{1}_y + \text{softmax}(h^3) = p - \mathbf{1}_y$$

where $\mathbf{1}_y$ is the corresponding one-hot vector.

□ g^3 has a shape $[1 \times n_3]$.

Back propagation

From loss to W^3, b^3



□ $h^3 = h^2 W^3 + b^3$

□ $\frac{\partial l}{\partial W^3} = \frac{\partial l}{\partial h^3} \cdot \frac{\partial h^3}{\partial W^3} = (\mathbf{h}^2)^T \mathbf{g}^3$

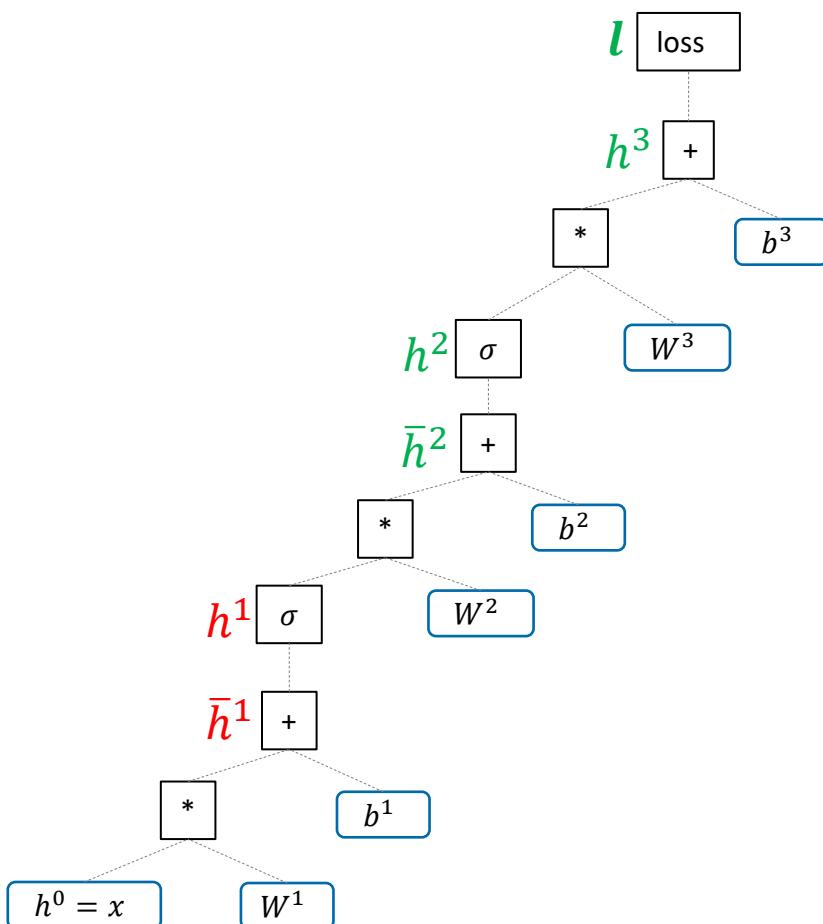
- $[n_2 \times 1] \times [1 \times n_3] \rightarrow [n_2 \times n_3]$

□ $\frac{\partial l}{\partial b^3} = \frac{\partial l}{\partial h^3} \cdot \frac{\partial h^3}{\partial b^3} = \mathbf{g}^3$

- $[1 \times n_3]$

Back propagation

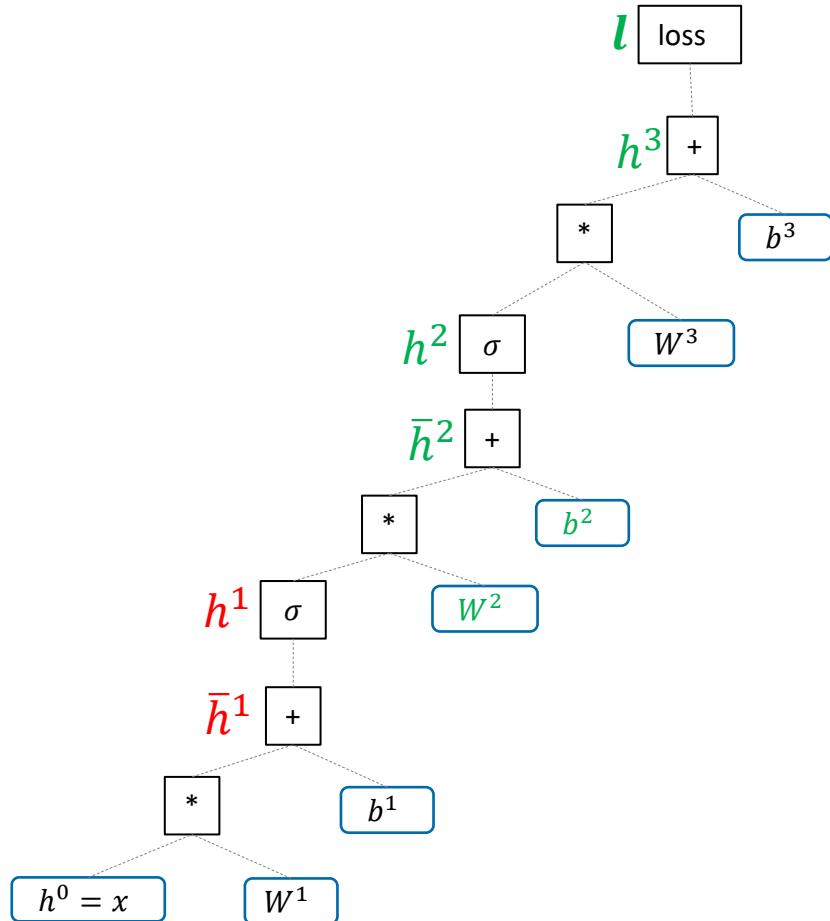
From loss to h^2 and \bar{h}^2



- $h^3 = h^2 W^3 + b^3$
- $g^2 = \frac{\partial l}{\partial h^2} = \frac{\partial l}{\partial h^3} \cdot \frac{\partial h^3}{\partial h^2} = g^3 (W^3)^T$
 - $[1 \times n_3] \times [n_3 \times n_2] \rightarrow [1 \times n_2]$
- $h^2 = \sigma(\bar{h}^2)$ (element-wise activation)
- $\frac{\partial h^2}{\partial \bar{h}^2} = \text{diag}(\sigma'(\bar{h}^2))$
 - $\sigma'(\bar{h}^2)$ is element-wise derivative and $\text{diag}(u)$ is the diagonal matrix corresponding to the vector u (the diagnose is u and others are zeros).
- $\bar{g}^2 = \frac{\partial l}{\partial \bar{h}^2} = \frac{\partial l}{\partial h^2} \cdot \frac{\partial h^2}{\partial \bar{h}^2} = g^2 \text{diag}(\sigma'(\bar{h}^2))$
 - $[1 \times n_2] \times [n_2 \times n_2] \rightarrow [1 \times n_2]$

Back propagation

From loss to W^2 and b^2



□ $\bar{h}^2 = h^1 W^2 + b^2$

□ $\frac{\partial l}{\partial W^2} = \frac{\partial l}{\partial \bar{h}^2} \cdot \frac{\partial \bar{h}^2}{\partial W^2} = (\mathbf{h}^1)^T \bar{g}^2$

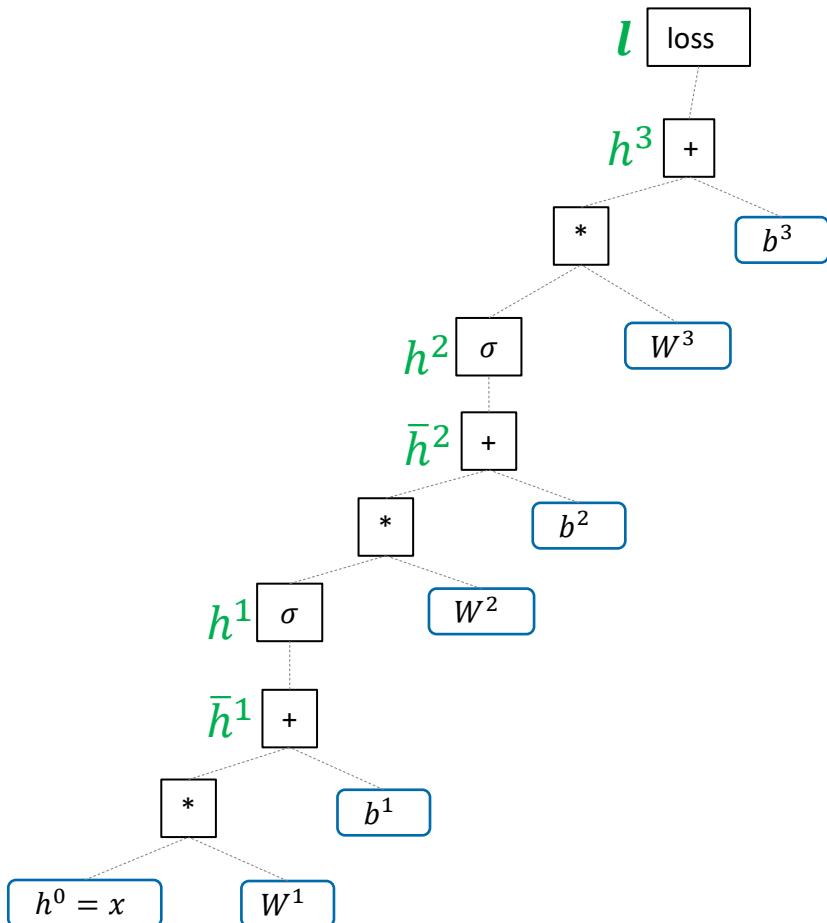
- $[n_1 \times 1] \times [1 \times n_2] \rightarrow [n_1 \times n_2]$

□ $\frac{\partial l}{\partial b^2} = \frac{\partial l}{\partial \bar{h}^2} \cdot \frac{\partial \bar{h}^2}{\partial b^2} = \bar{g}^2$

- $[1 \times n_2]$

Back propagation

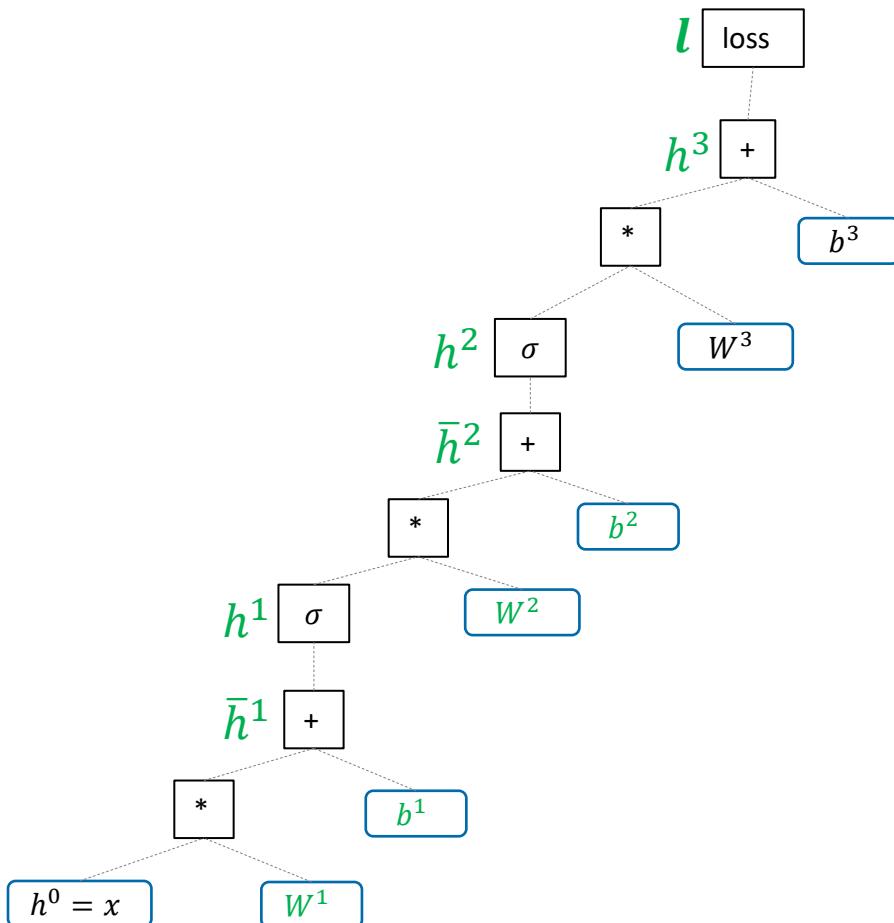
From loss to h^1 and \bar{h}^1



- $\bar{h}^2 = h^1 W^2 + b^2$
- $\mathbf{g}^1 = \frac{\partial l}{\partial h^1} = \frac{\partial l}{\partial \bar{h}^2} \cdot \frac{\partial \bar{h}^2}{\partial h^1} = \bar{g}^2 (W^2)^T$
 - $[1 \times n_2] \times [n_2 \times n_1] \rightarrow [1 \times n_1]$
- $\mathbf{h}^1 = \sigma(\bar{h}^1)$ (element-wise activation)
- $\frac{\partial h^1}{\partial \bar{h}^1} = \text{diag}(\sigma'(\bar{h}^1))$
 - $\sigma'(\bar{h}^1)$ is element-wise derivative and $\text{diag}(u)$ is the diagonal matrix corresponding to the vector u (the diagnose is u and others are zeros).
- $\bar{g}^1 = \frac{\partial l}{\partial \bar{h}^1} = \frac{\partial l}{\partial h^1} \cdot \frac{\partial h^1}{\partial \bar{h}^1} = \mathbf{g}^1 \text{diag}(\sigma'(\bar{h}^1))$
 - $[1 \times n_1] \times [n_1 \times n_1] \rightarrow [1 \times n_1]$

Back propagation

From loss to W^1 and b^1



- $\bar{h}^1 = h^0 W^1 + b^1 \quad (h^0 = x)$

- $\frac{\partial l}{\partial W^1} = \frac{\partial l}{\partial \bar{h}^1} \cdot \frac{\partial \bar{h}^1}{\partial W^1} = (\mathbf{h}^0)^T \bar{g}^1$

- $[n_0 \times 1] \times [1 \times n_1] \rightarrow [d = n_0 \times n_1]$

- $\frac{\partial l}{\partial b^1} = \frac{\partial l}{\partial \bar{h}^1} \cdot \frac{\partial \bar{h}^1}{\partial b^1} = \bar{g}^1$

- $[1 \times n_1]$



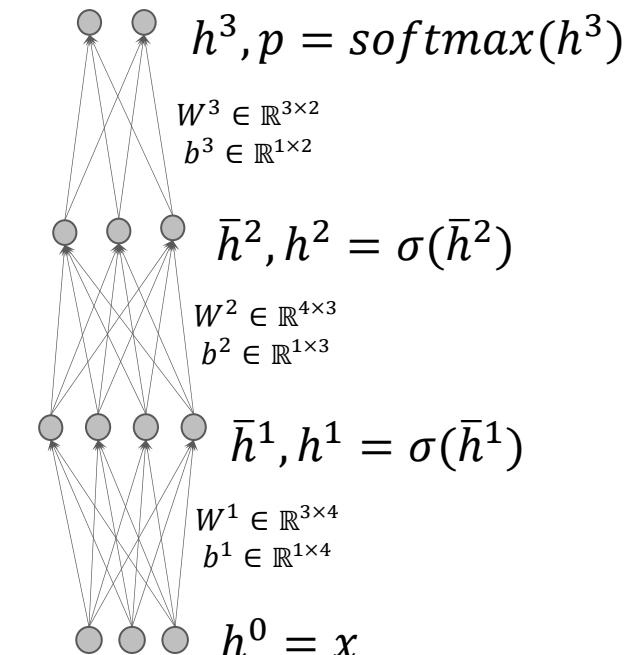
Exercise: How to compute $\frac{\partial l}{\partial x}$?

SGD for deep learning

```

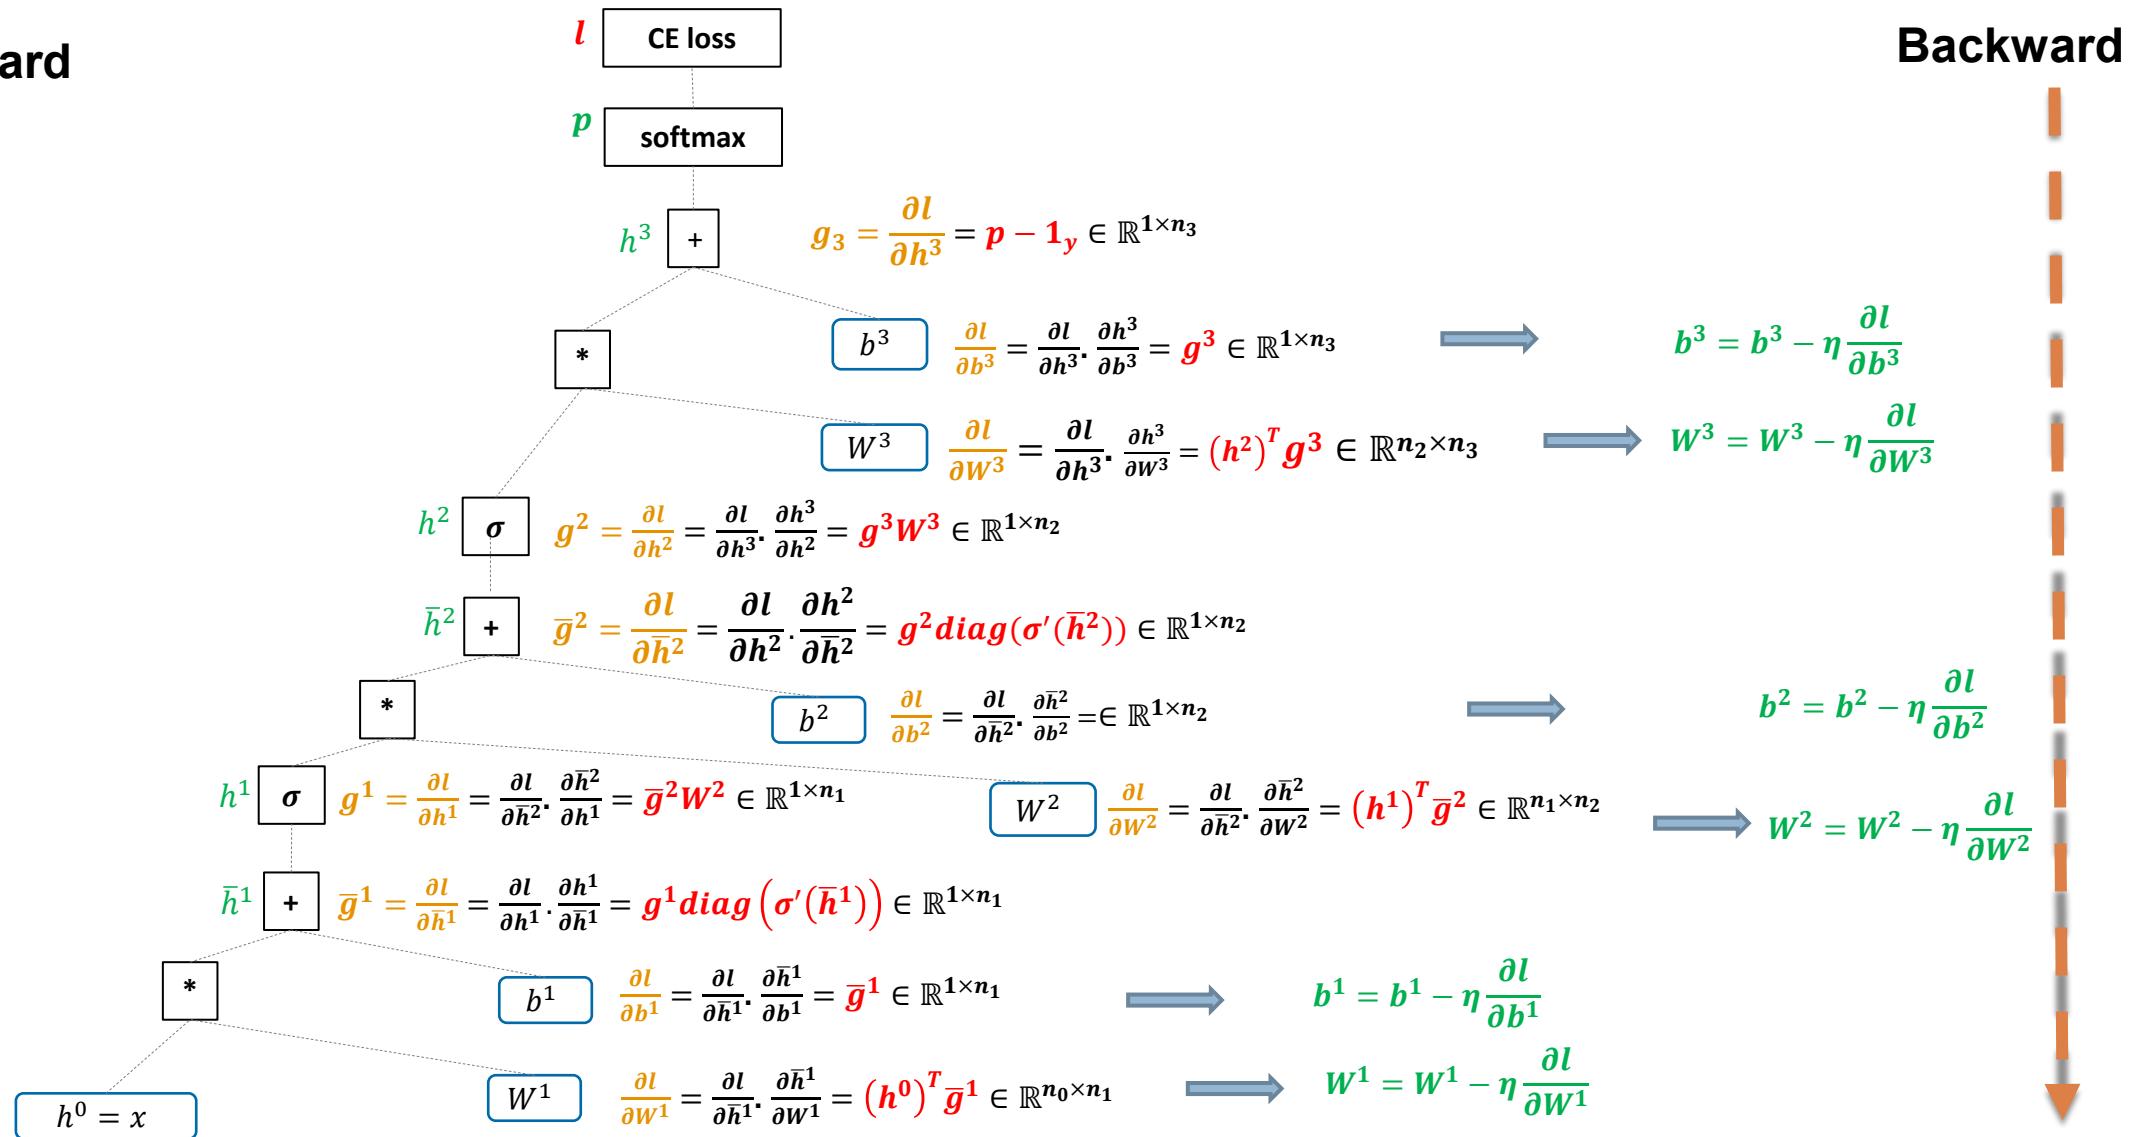
b = 32           //batch size
iter_per_epoch = N/b //epoch means one round going
                     through all data points
n_epoch = 50      //number of epochs
for epoch=1 to n_epoch do
  for i=1 to iter_per_epoch do
    Sample a minibatch  $B = \{(x_{ij}, y_{ij})\}_{j=1}^b$  from the training set
    Do forward propagation for B
    Do back propagation to compute  $\left(\frac{\partial l}{\partial W^k}, \frac{\partial l}{\partial b^k}\right)_{k=1}^L$ 
  for k=1 to L do
     $W_k = W_k - \eta \frac{\partial l}{\partial W^k}$ 
     $b_k = b_k - \eta \frac{\partial l}{\partial b^k}$ 

```



Forward – Backward Propagations

Forward



Mini-batch feed-forward

- **Input**
 - Batch $X: [b, n_0 = d = 5]$ (b is the batch size)
- **Hidden layer 1**
 - $h^1 = \sigma(XW^1 + b^1)$
 - Tensor $[b, n_1 = 7]$
- **Hidden layer 2**
 - $h^2 = \sigma(h^1W^2 + b^2)$
 - Tensor $[b, n_2 = 5]$
- **Output layer**
 - $h^3 = h^2W^3 + b^3$
 - $P = \text{softmax}(h^3, \text{dim} = 1)$
 - Tensor $P: [b, n_L = M = 4]$
- **The loss of the batch**
 - $\frac{1}{b} \sum_{i=1}^b CE(1_{y_i}, p_i) = -\frac{1}{b} \sum_{i=1}^b \log p_{y_i}$
 - Update weight matrices and biases to **minimize** the batch loss.

Backward propagation

One-hot labels

$$y_1 = 3, M = 4$$

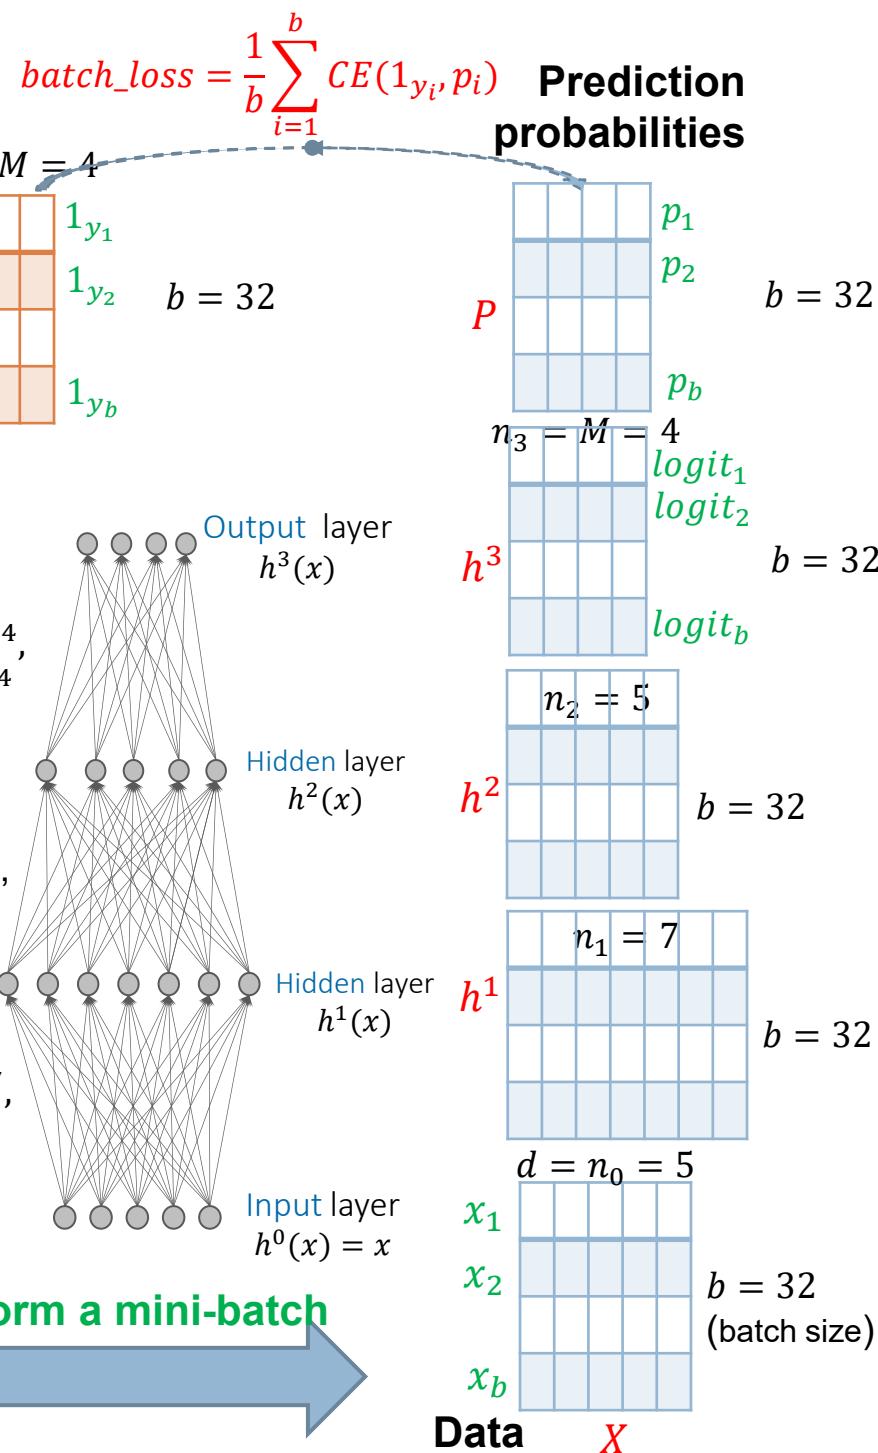
$$1_{y_1} = [0, 0, 1, 0]$$

Forward propagation

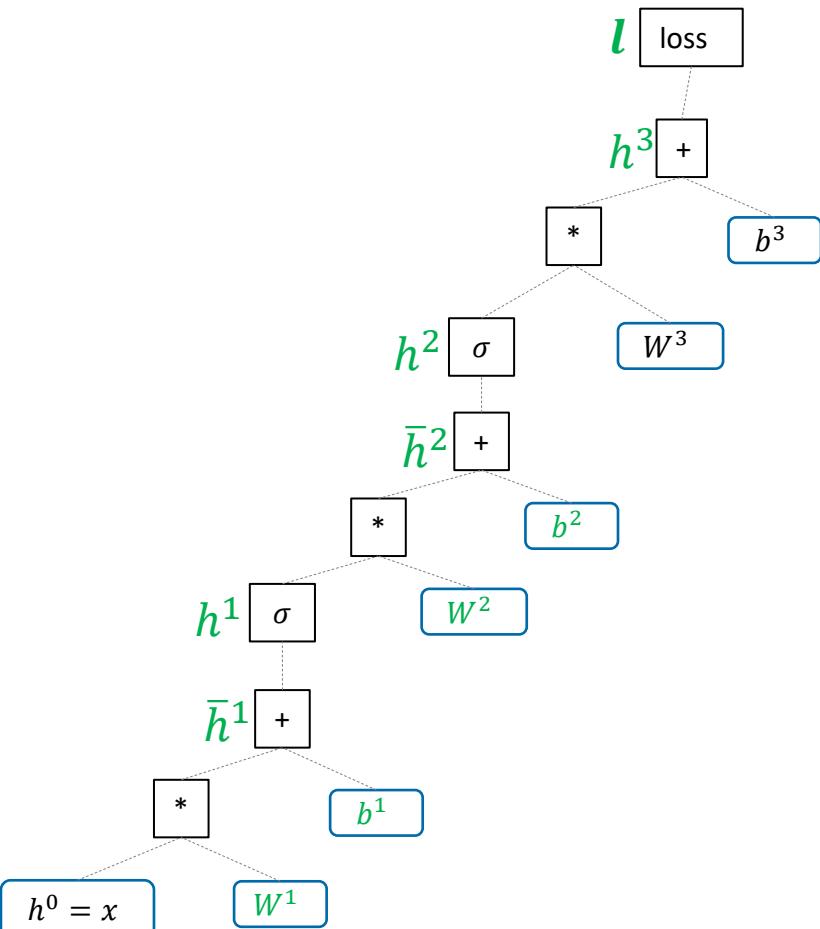
Training set

Mini-batches

$$(x_1, y_1), \dots, (x_b, y_b), \\ (x_{b+1}, y_{b+1}), \dots, (x_{2b}, y_{2b}) \\ \dots \\ (x_*, y_*), \dots, (x_N, y_N)$$



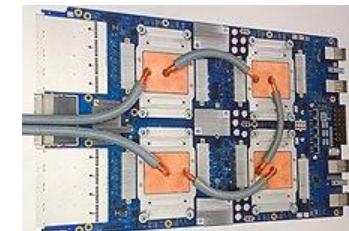
Why does deep learning need GPU and TPU?



- Let consider
$$\frac{\partial l}{\partial W^1} = \frac{\partial l}{\partial h^3} \cdot \frac{\partial h^3}{\partial h^2} \cdot \frac{\partial h^2}{\partial \bar{h}^2} \cdot \frac{\partial \bar{h}^2}{\partial h^1} \cdot \frac{\partial h^1}{\partial \bar{h}^1} \cdot \frac{\partial \bar{h}^1}{\partial W^1}$$
$$= (h^0)^T (p - 1_y) (W^3)^T \text{diag}(\sigma'(\bar{h}^2)) (W^2)^T \text{diag}(\sigma'(\bar{h}^1))$$
- For a deep net, this **back propagation** requires many **matrix multiplications**
 - We need specific hardware that can parallel and significantly speed up matrix multiplication operation
 - **GPU** (Graphic Processing Unit) and **TPU** (Tensor Processing Unit)



GPU (Source: HelloTech)

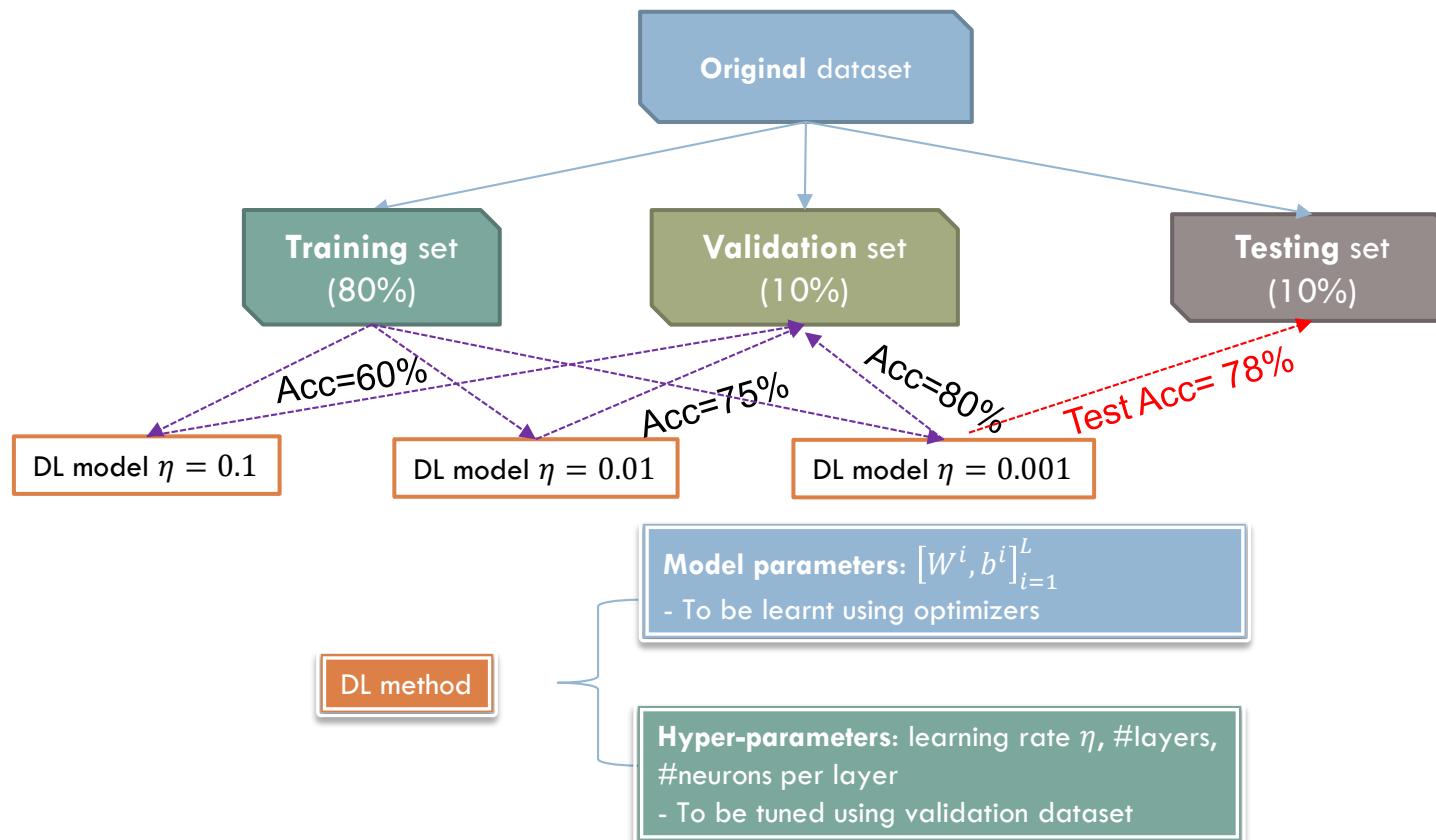


TPU (Source: Wikipedia)

Deep learning pipeline

Tuning hyper-parameters

- We want to train our DL model on a **training set** such that the **trained model** can predict well **unseen data** in a **separate testing set**.





Optimizers for deep learning

Challenges of optimization for Deep Learning

- The optimization problem in deep learning:

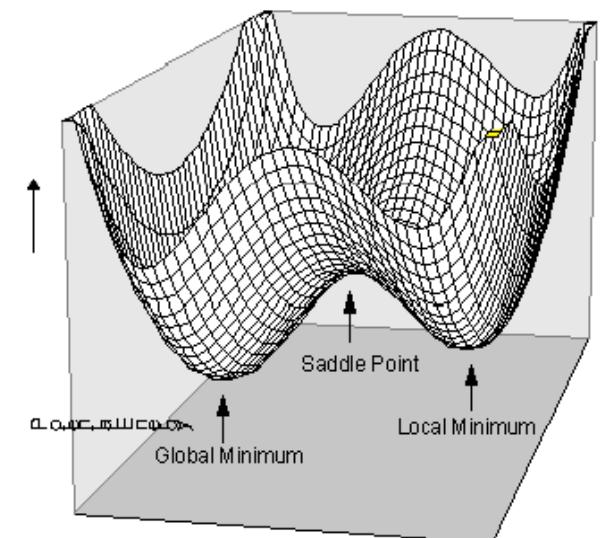
- $\circ \min_{\theta} J(\theta) := L(D; \theta) := \frac{1}{N} \sum_{i=1}^N l(x_i, y_i; \theta) = -\frac{1}{N} \sum_{i=1}^N \log \frac{\exp\{h_{y_i}^L(x_i)\}}{\sum_{m=1}^M \exp\{h_m^L(x_i)\}}$

- A very **complex** and **complicated** objective function

- \circ Highly **non-linear** and **non-convex** function
- \circ The **loss surface** is very **complex**
- \circ Many local minima points, but the number of saddle points is even **exponentially much more**

- Need **efficient optimizers** to solve

- \circ SGD with momentum, Adagrad, Adadelta, RMSProp, Adam, and Nadam
- \circ They are **built-in optimizers** of PyTorch.



(Source: Jan Jakubik)

SGD and SGD with momentum

SGD

- Input: $\eta > 0$ and *initial model* θ

while stopping criterion not met do

 Sample a mini-batch $\{(x^1, y^1), \dots, (x^b, y^b)\}$

 Compute $\mathbf{g} = \frac{1}{b} \sum_{i=1}^b \nabla_{\theta} l(f(x^i, \theta), y^i)$

 Update $\theta = \theta - \eta \mathbf{g}$

end while

- SGD uses only the **gradient of the mini-batch** to update the model
- It is fast at first several epochs and becomes **much slower** later.



(Source: Sebastian Ruder)

SGD with momentum

- Input: $\eta > 0, \alpha \in [0,1]$ and *initial model* θ

while stopping criterion not met do

 Sample a mini-batch $\{(x^1, y^1), \dots, (x^b, y^b)\}$

 Compute $\mathbf{g} = \frac{1}{b} \sum_{i=1}^b \nabla_{\theta} l(f(x^i, \theta), y^i)$

 Compute $\mathbf{v} = \alpha \mathbf{g} + (1 - \alpha) \mathbf{v}$ //velocity v

 Update $\theta = \theta - \eta \mathbf{v}$

end while

□ SGD with momentum uses a **velocity vector v** which **stores the past gradients** together with the **current gradient** to speed up SGD

- α is a hyper-parameter that indicates how quickly the contributions of previous gradients. In practice, this is usually set to 0.5, 0.9, and 0.99.
- The momentum primarily solves 2 problems: **poor conditioning** of the Hessian matrix and **variance** in the stochastic gradient.



AdaGrad

AdaGrad

- Input: $\eta > 0$, $\epsilon > 0$ (10^{-6}), and *initial model* θ

while stopping criterion not met **do**

 Sample a mini-batch $\{(x^1, y^1), \dots, (x^b, y^b)\}$

 Compute $\mathbf{g} = \frac{1}{b} \sum_{i=1}^b \nabla_{\theta} l(f(x^i, \theta), y^i)$

 Accumulate the square gradient: $\gamma = \gamma + \mathbf{g} \odot \mathbf{g}$

 Update $\theta = \theta - \frac{\eta}{\sqrt{\epsilon + \gamma}} \odot \mathbf{g}$

end while

Note: \odot means element-wise product

\mathbf{g}^1	g_1^1	g_2^1	...	g_P^1
\mathbf{g}^2	g_1^2	g_2^2	...	g_P^2
...
\mathbf{g}^t	g_1^t	g_2^t	...	g_P^t
γ	$\sum_{i=1}^t (g_1^i)^2$	$\sum_{i=1}^t (g_2^i)^2$...	$\sum_{i=1}^t (g_P^i)^2$

- Learning rates are scaled by the square root of the cumulative sum of squared gradients
- Direction with large partial derivatives
 - Thus, rapid decrease in their learning rates
- Direction with small partial derivatives
 - Hence relatively small decrease in their learning rates
- Weakness: always decrease the learning rate!
 - Excellent for convex problem, but not so good for DL (with non-convex problems)

RMSProp

RMSProp

- **Input:** $\eta > 0, \epsilon > 0 (10^{-6}), \beta \in [0,1]$ and *initial model* θ

while stopping criterion not met do

Sample a mini-batch $\{(x^1, y^1), \dots, (x^b, y^b)\}$

Compute $\mathbf{g} = \frac{1}{b} \sum_{i=1}^b \nabla_{\theta} l(f(x^i, \theta), y^i)$

Accumulate the square gradient: $\gamma = \beta\gamma + (1 - \beta)\mathbf{g} \odot \mathbf{g}$

Update $\theta = \theta - \frac{\eta}{\sqrt{\epsilon + \gamma}} \odot \mathbf{g}$

end while

Note: \odot means element-wise product

- A modification of AdaGrad to work better for **non-convex** setting.
- Instead of cumulative sum, use exponential moving/smoothing average.
- RMSProp has been shown to be an effective and practical optimization algorithm for DNN.
 - Currently one of the go-to optimization methods being employed routinely by DL applications.

Adam

- The best variant that essentially combines RMSProp with momentum

- Suggested default values:
 $\eta = 0.001$, $\alpha_1 = 0.9$, $\alpha_2 = 0.999$ and $\epsilon = 10^{-6}$.

Adam

○ **Input:** $\eta > 0, \epsilon > 0 (10^{-6}), \beta_1, \beta_2 \in [0,1]$ and *initial model* θ

$t = 1$

while stopping criterion not met do

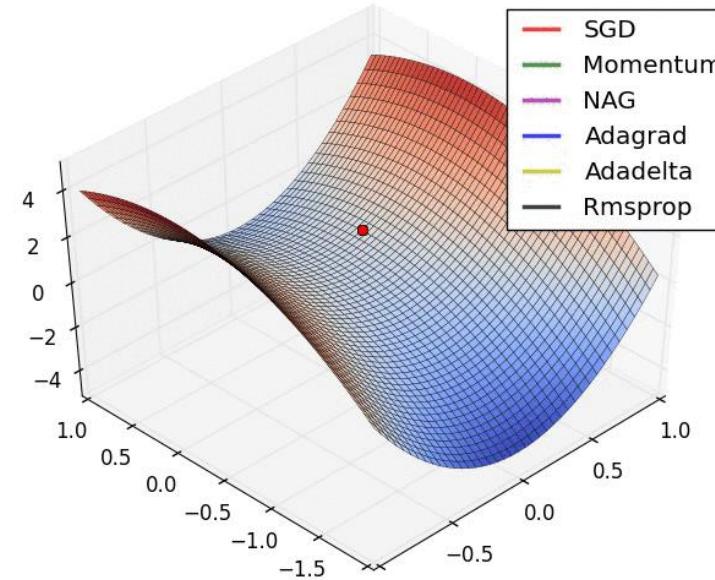
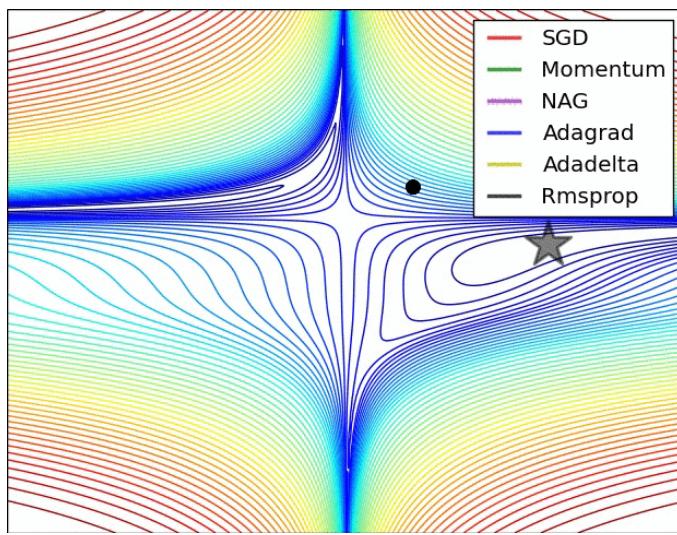
- Sample a mini-batch $\{(x^1, y^1), \dots, (x^b, y^b)\}$
- Compute $\mathbf{g} = \frac{1}{b} \sum_{i=1}^b \nabla_{\theta} l(f(x^i, \theta), y^i)$
- Accumulate the gradient: $\mathbf{s} = \beta_1 \mathbf{s} + (1 - \beta_1) \mathbf{g}$
- Accumulate the square gradient: $\mathbf{\gamma} = \beta_2 \mathbf{\gamma} + (1 - \beta_2) \mathbf{g} \odot \mathbf{g}$
- Correct s : $\hat{s} = \frac{s}{1 - \beta_1^t}$ # t is the current iteration
- Correct γ : $\hat{\gamma} = \frac{\gamma}{1 - \beta_2^t}$
- Update $\theta = \theta - \frac{\eta}{\sqrt{\epsilon + \hat{\gamma}}} \odot \hat{s}$

$t = t + 1$

end while

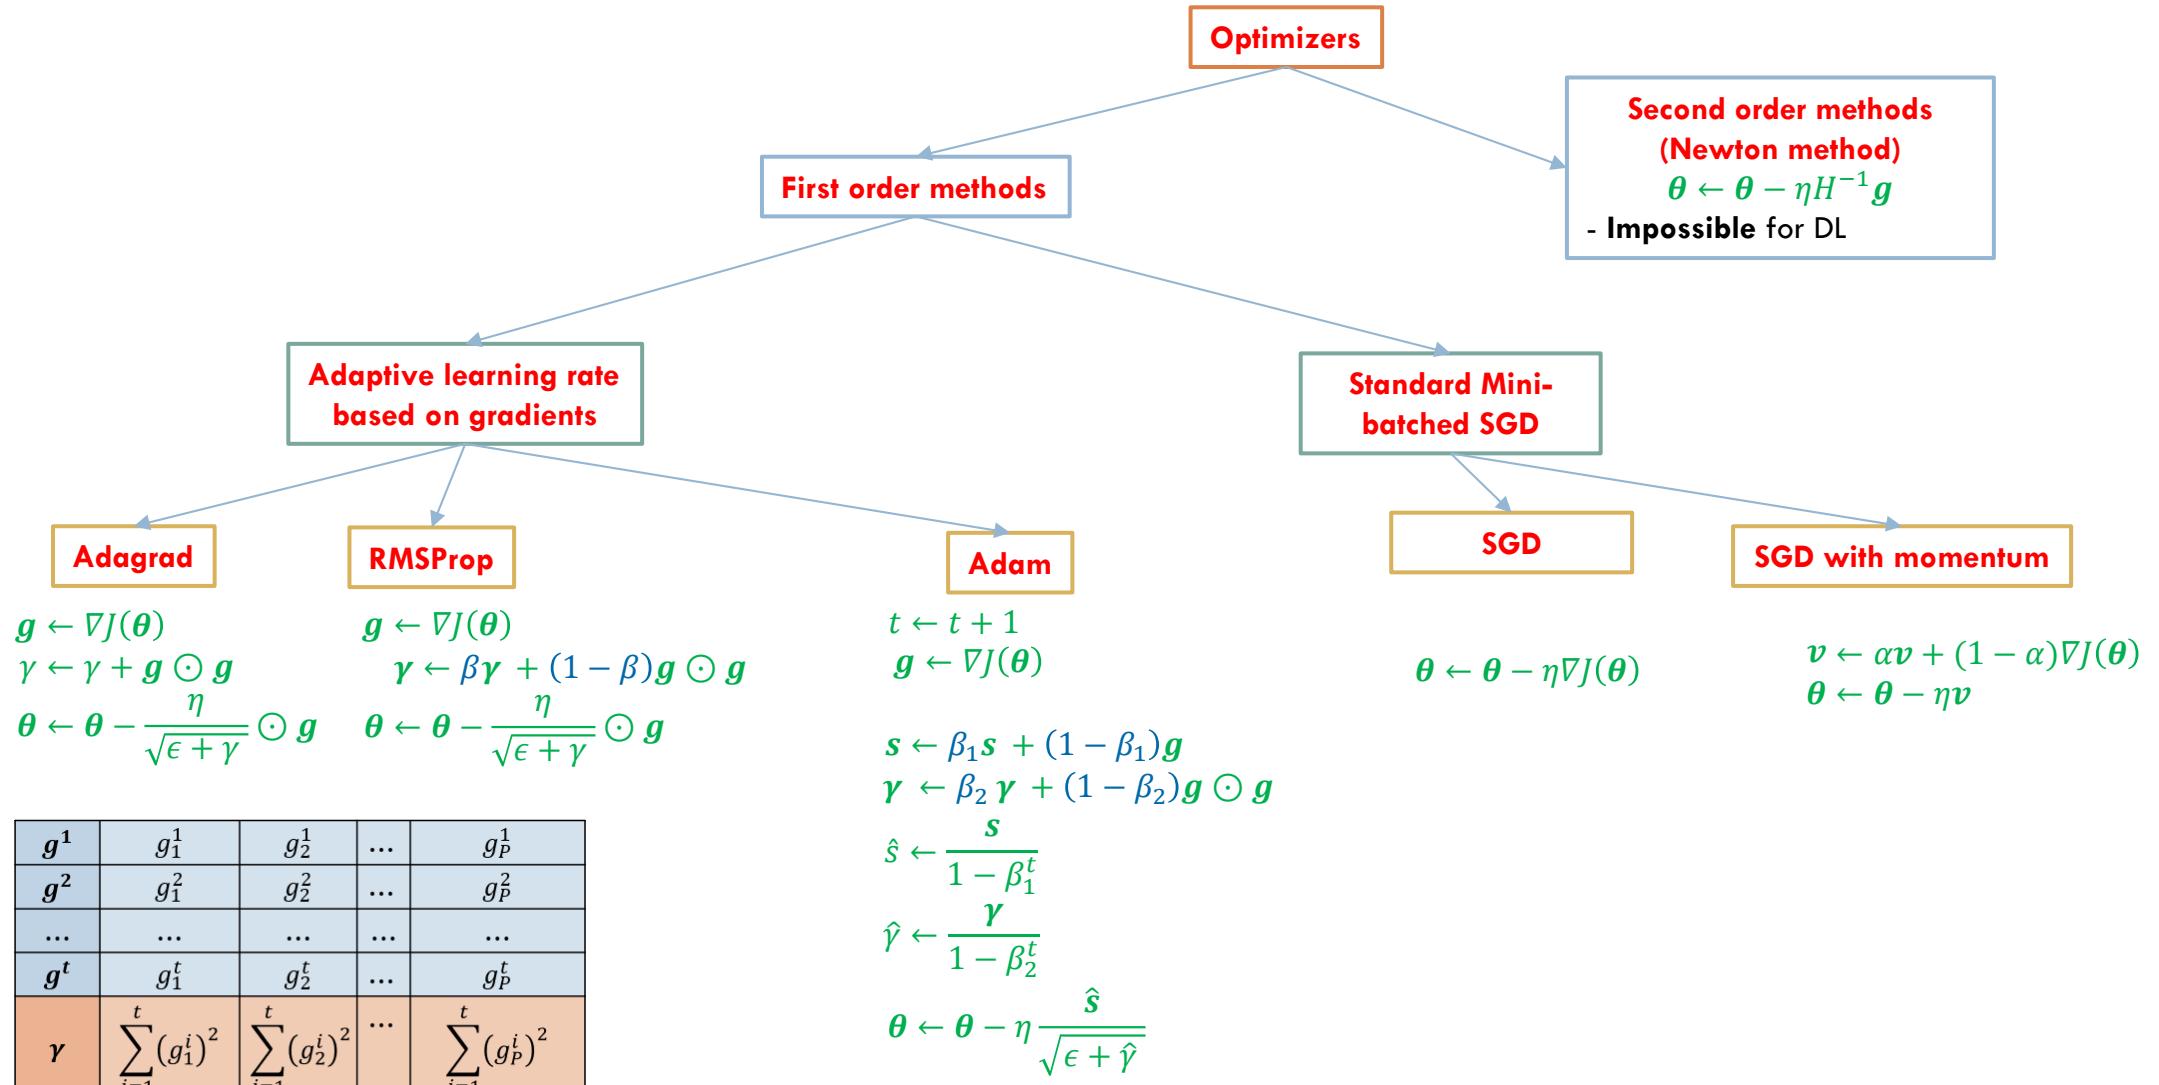
Note: \odot means element-wise product

Visual comparison of all optimizers



[Source: Sebastian Ruder]

Optimizers in deep learning



Summary

- Optimization problem in DL and ML
 - Regularization term + Empirical loss term
- Gradient descent
- Stochastic gradient descent
- Backward propagation
- Other optimizers in DL
 - SGD with momentum, Adagrad, RMSProp, and Adam
- First order methods and second order methods

Thanks for your attention!

