

COPYRIGHT WARNING: Copyright in these original lectures is owned by Monash University. You may transcribe, take notes, download or stream lectures for the purpose of your research and study only. If used for any other purpose, (excluding exceptions in the Copyright Act 1969 (Cth)) the University may take legal action for infringement of copyright.

Do not share, redistribute, or upload the lecture to a third party without a written permission!

FIT3181/5215 Deep Learning

Week 05: Practical skills in deep learning

Lecturer: Trung Le

Email: trunglm@monash.edu


Outline

- Setting of a machine learning problem
 - General loss versus empirical loss
- Gradient vanishing/exploding and network initialization.
- Overfitting and underfitting
- Recipe for overfitting
 - Use regularization term, dropout, batch norm, data augmentation, transfer learning
 - Label smoothing, data mix-up, cut-mix

□ Further reading recommendation

- [Deep Learning, Sections 4.1-4.3, 8.1 -8.5, 11.3, 11.4].
- [Dive into Deep Learning, Chapters 5 and 11].

Machine learning setting



Machine learning setting

Empirical data/label distribution \tilde{p}_{data}

(x, y)	(x_1, y_1)	(x_2, y_2)	...	(x_N, y_N)
p	$1/N$	$1/N$...	$1/N$

Data-label generative process

$$(x', y') \sim p_{data}(x, y) = p_{data}(x)p_{data}(y | x)$$


$y' = 1$ or cat

$$y' \sim p_{data}(y | x')$$

Cat (1)	Dog (2)	Lion (3)	Panda (4)
0.7	0.2	0.05	0.05

$$p_{data}(y | x')$$

Ground-truth labelling mechanism



Data space x
 $x' \sim p_{data}(x)$

Data distribution


Dataset collecting process

$$D = \{(x_1, y_1), \dots, (x_N, y_N)\}$$

$$(x_i, y_i) \sim p_{data}(x, y) \xleftarrow{N \rightarrow \infty} \tilde{p}_{data}(x, y) = \frac{1}{N} \sum_{i=1}^N \delta_{(x_i, y_i)}(x, y)$$

$$y_1 \sim p_{data}(y | x_1)$$

$$y_N \sim p_{data}(y | x_N)$$



Data space x
 $x \sim p_{data}(x)$

Data distribution

Not in assessment

Machine learning setting

Generalization (general) loss

(loss on data/label distribution)

$$\mathcal{L}_{gen}(\theta) = \mathbb{E}_{p_{data}} [l(f(x; \theta), y)]$$

$N \rightarrow \infty$
Law of large numbers


Ideal: $\theta^* = \operatorname{argmin}_{\theta} \mathcal{L}_{gen}(\theta)$

$$(x', y') \sim p_{data}(x, y) = p_{data}(x)p_{data}(y | x)$$

$$y' \sim p_{data}(y | x')$$

$$p_{data}(y | x')$$

Labelling mechanism



Data space x
 $x' \sim p_{data}(x)$

Data distribution

Empirical loss

(loss on a collected training set D)

$$\mathcal{L}_{emp}(\theta) = \mathbb{E}_{\tilde{p}_{data}} [l(f(x; \theta), y)] = \frac{1}{N} \sum_{i=1}^N l(f(x_i; \theta), y_i)$$


Reality: $\theta^* = \operatorname{argmin}_{\theta} \mathcal{L}_{emp}(\theta)$

$$D = \{(x_1, y_1), \dots, (x_N, y_N)\}$$

$$(x_i, y_i) \sim p_{data}(x, y) \xleftarrow{N \rightarrow \infty} \tilde{p}_{data}(x, y) = \frac{1}{N} \sum_{i=1}^N \delta_{(x_i, y_i)}(x, y)$$

$$y_1 \sim p_{data}(y | x_1)$$

$$y_N \sim p_{data}(y | x_N)$$



Data space x
 $x \sim p_{data}(x)$

Data distribution

How Learning Differs from Pure Optimization?

□ Some important notations:

- $p_{data}(x, y)$: existed, but unknown, distribution of data and label
- $\tilde{p}_{data}(x, y) = \frac{1}{N} \sum_{i=1}^N \delta_{(x_i, y_i)}(x, y)$: empirical data/label distribution - this is what we observed from training data $D = \{(x_i, y_i)\}_{i=1}^N$ where $(x_i, y_i) \stackrel{\text{iid}}{\sim} p_{data}(x, y)$.
- Per-sample loss: $l(f(x; \theta), y)$

□ Empirical loss minimisation (pure optimisation):

$$\mathcal{L}_{emp}(\theta) = \tilde{p}_{data} \mathbb{E} [l(f(x; \theta), y)] = \frac{1}{N} \sum_{i=1}^N l(f(x_i; \theta), y_i)$$

Empirical loss (maths)

VS.

Generalisation loss (ML)

□ But what ML wants is true generalisation loss:

$$\mathcal{L}_{gen}(\theta) = p_{data} \mathbb{E} [l(f(x; \theta), y)]$$

How to achieve this when we only have access to empirical data?

Optimization Problem in ML and DL

- Most of **optimization problems** in machine learning (deep learning) has the following form:

$$\min_{\theta} J(\theta) = \Omega(\theta) + \frac{1}{N} \sum_{i=1}^N l(y_i, f(x_i; \theta))$$

Regularization term

- $\Omega(\theta) = \lambda \sum_k \sum_{i,j} (W_{i,j}^k)^2 = \lambda \sum_k \|W^k\|_F^2$
- Encourage **simple models**
- Avoid **overfitting**

Empirical loss

- Work well on training set

Guiding principles:

1. **Occam Razer:** prefer the simplest model that can do well.

How to efficiently solve this optimization problem?
 N is the **training size** and might be very big (e.g., $N \approx 10^6$)

Works well for Deep Learning (non-convex)

First-order iterative methods

gradient descent, steepest descent

Use the **gradient** (first derivative) $g = \nabla_{\theta} J(\theta)$ to update parameters:

$$w = w - \text{learning rate} \times \text{gradient}$$

Works well for convex problems, but not DL


Second-order iterative methods

Newton and quasi Newton methods

Use the Hessian matrix (second derivative) $H = \nabla_{\theta}^2 J(\theta)$ to update parameters

Optimization Problem in ML and DL

- Most of optimization problems in machine learning (deep learning) has the following form:

$$\min_{\theta} J(\theta) = \Omega(\theta) + \frac{1}{N} \sum_{i=1}^N l(y_i, f(x_i; \theta))$$



What **loss** function should one use?

Some typical loss functions used in the classification and regression problems

- 0-1 loss, Hinge loss, Logistic loss (binary classification)
- L1 loss, L2 loss, ϵ –insensitive loss (regression)
- Popular surrogate loss: cross-entropy loss (multi-class classification, deep learning)

How Learning Differs from Pure Optimization?

- Achieving **generalisation capacity** is the holy-grail of machine learning.
 - Empirical risk is prone to **over-fitting**
 - Some times, it is not really feasible if the loss function **does not have useful derivatives** (e.g., 0-1 loss), hence we usually resort to **surrogate loss function**, e.g., cross-entropy, Hinge loss, etc.
- In most cases, DL algorithm **doesn't halt at local minimum**
 - It halts when **certain convergence criterion is met** (e.g., based on **early stopping** when overfitting start to occur, reach **certain budgets, number of epochs**, etc).
 - For training DL models, it **might stop** when the loss function still has **large derivates**, which is **different** from a pure optimisation setting.





Challenges in deep learning optimization: gradient vanishing, gradient exploding


Gradient vanishing

Gradient Vanishing


- Gradients get **smaller and smaller** as the algorithm progresses down to the **lower layers**.
 - SGD update leaves the lower layer connection weights **virtually unchanged**, and training **never converges** to a good solution.

Some activation functions are easy to get saturated


- Sigmoid or tanh



$$\sigma(z) = s(z) = \frac{1}{1 + \exp\{-z\}}$$



$$\begin{aligned}\frac{\partial l}{\partial W^1} &= \frac{\partial l}{\partial h^3} \cdot \frac{\partial h^3}{\partial h^2} \cdot \frac{\partial h^2}{\partial \bar{h}^2} \cdot \frac{\partial \bar{h}^2}{\partial h^1} \cdot \frac{\partial h^1}{\partial \bar{h}^1} \cdot \frac{\partial \bar{h}^1}{\partial W^1} \\ &= \left[(p^T - 1_y) W^3 \text{diag}(\sigma'(\bar{h}^2)) W^2 \text{diag}(\sigma'(\bar{h}^1)) \right]^T (h^0)^T\end{aligned}$$



$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$

Gradient vanishing

● Gradient Vanishing


- Gradients get **smaller and smaller** as the algorithm progresses down to the **lower layers**.
 - SGD update leaves the lower layer connection weights **virtually unchanged**, and training **never converges** to a good solution.

● Some activation functions are easy to **get saturated**

- Sigmoid or tanh

● Recipe

- Activation function plays an important role! Common practice:
 - Avoid sigmoid or saturated activation function
 - ReLU is a common good choice
- Good weight initialization is critical!



$$\begin{aligned}\frac{\partial l}{\partial W^1} &= \frac{\partial l}{\partial h^3} \cdot \frac{\partial h^3}{\partial h^2} \cdot \frac{\partial h^2}{\partial \bar{h}^2} \cdot \frac{\partial \bar{h}^2}{\partial h^1} \cdot \frac{\partial h^1}{\partial \bar{h}^1} \cdot \frac{\partial \bar{h}^1}{\partial W^1} \\ &= \left[(p^T - 1_y) W^3 \text{diag}(\sigma'(\bar{h}^2)) W^2 \text{diag}(\sigma'(\bar{h}^1)) \right]^T (h^0)^T\end{aligned}$$

Gradient exploding


- **Gradient Exploding**

- The gradients can grow **bigger and bigger**, so many layers get **insanely large weight** updates, and the training **diverges**.
- Mostly being encountered in **recurrent neural networks**.
- Often happen for **recursive models** for example **Recurrent Neural Network (RNN)**, **Bidirectional RNN**



Tip:

- Let simplify W as a **scalar** in the real-valued set
 - $W^m \rightarrow 0$ if $|W| < 1$.
 - $W^m \rightarrow \infty$ if $|W| > 1$.



$$\frac{\partial l_T}{\partial h_0} = \frac{\partial l_T}{\partial h_T} \times \frac{\partial h_T}{\partial h_{T-1}} \times \cdots \times \frac{\partial h_1}{\partial h_0}$$

W W

Multiplication of many matrices W

Gradient Clipping

- One way to lessen the **exploding gradients problem** is to simply **clip the gradients** during backpropagation so that they never **exceed** some threshold
 - Either clip by values (direction might change)
 - or clip by norms (keep direction but rescale the magnitude)
- Widely applied to **Recurrent Neural Networks**
- Widely used for NLP tasks, not so much used for CNNs.

```

import torch
import torch.nn as nn
import torch.optim as optim

# Example model
class SimpleModel(nn.Module):
    def __init__(self):
        super(SimpleModel, self).__init__()
        self.fc = nn.Linear(10, 1)

    def forward(self, x):
        return self.fc(x)

# Initialize the model, loss function, and optimizer
model = SimpleModel()
criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)

# Example training loop
for epoch in range(100):
    optimizer.zero_grad()
    # Example input and target tensors
    inputs = torch.randn(32, 10) # Batch of 32, input size 10
    targets = torch.randn(32, 1) # Batch of 32, target size 1
    # Forward pass
    outputs = model(inputs)
    loss = criterion(outputs, targets)
    # Backward pass
    loss.backward()
    # Gradient clipping by norm
    torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=2.0)
    # Update model parameters
    optimizer.step()
    # Print loss
    print(f'Epoch {epoch+1}, Loss: {loss.item()}')

```

Observing Gradient Vanishing and Exploding

Log the histogram of gradients to TensorBoard

```

def on_epoch_end(self, params):
    epoch = params['epoch']

    # Log scalar values
    if self.log Scalars:
        self.writer.add_scalar('Loss/train', params['train_loss'], epoch)
        self.writer.add_scalar('Loss/validation', params['val_loss'], epoch)
        self.writer.add_scalar('Accuracy/train', params['train_accuracy'], epoch)
        self.writer.add_scalar('Accuracy/validation', params['val_accuracy'], epoch)


    # Log a batch of images (example with first batch from trainloader)
    if self.log Images:
        dataiter = iter(self.train_loader)
        images, labels = next(dataiter)
        img_grid = vutils.make_grid(images)
        self.writer.add_image('images_batch', img_grid, epoch)

    # Log histogram of model parameters
    if self.log Histograms:
        #log the model parameters
        for name, param in self.model.named_parameters():
            self.writer.add_histogram(name, param, epoch)

#log the gradients w.r.t. the model parameters
num_selected = 100
subset, _ = torch.utils.data.random_split(val_subset, [num_selected, len(val_subset) - num_selected])
subset_loader = torch.utils.data.DataLoader(subset, batch_size=num_selected) # Create a dataloader for
inputs, labels = next(iter(subset_loader)) # Get a batch of data
inputs, labels = inputs.to(device), labels.to(device) # Move data to device
outputs = self.model(inputs) # Use self.model instead of models
loss = self.criterion(outputs, labels)
loss.backward()
for name, param in self.model.named_parameters():
    self.writer.add_histogram(name + '_grad', param.grad, epoch)

```

Using TensorBoard



Weight initialization


He and Xavier weight initialization

- Initialisation in deep learning training is **crucial**:

- Some optimizers can be **theoretically guaranteed to converge** regardless of initializations
- Deep learning algorithms do not have these luxuries:
 - It is iterative, but optimization deep neural networks is **not yet well understood**
 - Initial point is extremely important:** it can determine if the algorithm converges, or with some **bad initialisation**, it becomes unstable and fails together

- What is a **good weight/filter initialization**?

- Break the ‘symmetry’** of the network: two hidden nodes with the **same input** should have **different weights**.
- The **gradient signal to flow well** in both directions and **don’t want** the signal to die out or to **explode and saturate**.
- Large initial weights** has **better symmetry breaking effect**, help avoiding **losing signals** and **redundant units**, but could **result in exploding values** during back-ward and forward passes, especially in Recurrent Neural Networks.



Initialization is important for training DL models.


He and Xavier weight initialization

● Xavier initialization

- Try to ensure the **variance of the outputs** of each layer equal to the **variance of its inputs**
- Also need the gradients to have **equal variance** before and **after flowing through a layer** in the **reverse direction**
- Good for **sigmoid** and **tanh** functions
- Not good for ReLU

Gaussian version

$$w_{Xa} \sim N\left(0, \sqrt{\frac{2}{n_{in} + n_{out}}}\right)$$



Uniform alternative

$$w_{Xa} \sim \text{Uniform}\left(-\sqrt{\frac{6}{n_{in} + n_{out}}}, \sqrt{\frac{6}{n_{in} + n_{out}}}\right)$$

$$\begin{aligned} n_{in} &= 3 \\ n_{out} &= 2 \end{aligned}$$

Why $\sqrt{\frac{2}{n_{in}+n_{out}}}$?

Understanding the difficulty of training deep feedforward neural networks

Xavier Glorot

DIRO, Université de Montréal, Montréal, Québec, Canada

Yoshua Bengio

Paper link: <https://proceedings.mlr.press/v9/glorot10a/glorot10a.pdf>

He and Xavier weight initialization

□ Xavier initialization

- Ensure the **variance of the outputs** of each layer **equal** to the **variance of its inputs**
- Also need the **gradients** to have **equal variance before and after** flowing through a layer in the reverse direction
- Good for **sigmoid** and **tanh** activation functions
- **Not** good for **ReLU**

□ He initialization

- A **variant of Xavier initialization** where $\alpha = 1$
- Works better for ReLU.

$$w_{Xa} \sim N\left(0, \sqrt{\frac{2}{n_{in} + n_{out}}}\right)$$

$n_{out} = 2$
 $n_{in} = 3$

$$w_{He} \sim N\left(0, \alpha \times \sqrt{\frac{2}{n_{in} + n_{out}}}\right) \quad \alpha = \begin{cases} 1 & \text{if sigmoid} \\ 4 & \text{if tanh} \\ \sqrt{2} & \text{if ReLU} \end{cases}$$



This ICCV paper is the Open Access version, provided by the Computer Vision Foundation.
Except for this watermark, it is identical to the version available on IEEE Xplore.

Delving Deep into Rectifiers:
Surpassing Human-Level Performance on ImageNet Classification

Kaiming He Xiangyu Zhang Shaoqing Ren Jian Sun
Microsoft Research

Paper link: https://www.cv-foundation.org/openaccess/content_iccv_2015/papers/He_Delving_Deep_into_ICCV_2015_paper.pdf




Where do we go from here?

- Challenges in deep learning optimisation and how to address them:
 - Local minima, saddle points and complex loss surfaces
 - Gradient vanishing and exploding
 - What we **don't** cover in this unit (see DL section 8.2):
 - Ill-conditioning problem
 - Long-term dependencies
 - Poor correspondence between local and global structures
 - Theoretical limits of optimisation (but they usually have little use in practice of deep learning)
- Initialization Strategies
- Regularization in deep learning
 - Parameter norm penalty: l1, l2 regularization
 - Early stopping
 - Dropout
 - Batch normalization
- Choice of optimizers:
 - Basic algorithms: SGD, Momentum, Nesterov Momentum
 - Algorithms with adaptive learning rate: AdaGrad, RMSProp, Adam

Overfitting and Regularization in Deep Learning

Deep Learning Pipeline

- We want to train our DL model on a **training set** such that the **trained model** can predict well **unseen data** in a separate testing set.




Model parameters: weight matrices, biases, filters which will be learnt

Hyper-parameters to consider: learning rate, #layers, #neurons which need to be tuned

Deep Learning Pipeline

- We want to train our DL model on a **training set** such that the **trained model** can predict well **unseen data** in a separate testing set.




- A) Extremely high
- B) Too high
- C) Too small
- D) Just right

Hyper-parameters to consider: learning rate, #layers, #neurons

Deep Learning Pipeline

- We want to train our DL model on a **training set** such that the **trained model** can predict well **unseen data** in a separate testing set.




Hyper-parameters to consider: learning rate, #layers, #neurons

Regularization and Overfitting

- Three elements in ML and DL
 - **Data:** training data, testing data, validation data
 - **Model:** a mathematical function $f(x; \theta)$ that maps an input instance x to outcome y
 - **Evaluation:** a performance metric to quantitatively measure how well $f(x; \theta)$ is
- Machine learning as an **optimization process**. Learning from data by **optimizing** its loss
 - Define a measure of loss via a **loss function**: $l(f(x; \theta), y)$
 - Compute its **loss over all training data** $J(\theta) = N^{-1} \sum_{i=1}^N l(f(x_i, \theta), y_i)$
 - Learning = **finding θ^* that minimizes the loss**: $\theta^* = \arg \min_{\theta} J(\theta)$
- What might **go wrong** with this formulation?
 - The choice of **learning function** $f(x; \theta)$ is too hard to learn
 - The choice of **loss function** $l(f(x), y)$ is inadequate
 - The model **does well** on training data, but **perform poorly** on unseen test data: **overfitting** problem!

Overfitting & Underfitting



Underfitting

- The model is **too simple** to characterise a training set
 - Use linear model to learn from non-linear data


Overfitting

- The model performs **very well** on the **training set**, but **cannot generalise** to perform well on a separate **testing set**
- This is the most common problem in DL since deep networks are very powerful!

Overfitting in Deep NNs

Overfitting

Overfit: tendency of the network to “memorize” all training samples, leading to poor generalization



[Source: <https://www.ibm.com/cloud/learn/overfitting>]

- **Overfitting** occurs when your network models the training data *too well* and fails to generalize to your test (validation) data.
 - Performance measured by errors on “unseen” data.
 - **Minimize error alone on training data is not enough**
 - Causes: **too many layers, too many hidden nodes, and overtrained**.

Hyperparameter	Increases capacity when...	Reason	Caveats
Number of hidden units	increased	Increasing the number of hidden units increases the representational capacity of the model.	Increasing the number of hidden units increases both the time and memory cost of essentially every operation on the model.
Learning rate	tuned optimally	An improper learning rate, whether too high or too low, results in a model with low effective capacity due to optimization failure	
Convolution kernel width	increased	Increasing the kernel width increases the number of parameters in the model	A wider kernel results in a narrower output dimension, reducing model capacity unless you use implicit zero padding to reduce this effect. Wider kernels require more memory for parameter storage and increase runtime, but a narrower output reduces memory cost.
Implicit zero padding	increased	Adding implicit zeros before convolution keeps the representation size large	Increased time and memory cost of most operations.
Weight decay coefficient	decreased	Decreasing the weight decay coefficient frees the model parameters to become larger	
Dropout rate	decreased	Dropping units less often gives the units more opportunities to “conspire” with each other to fit the training set	

We can also experiment with model capacity itself in parallel.





Table 11.1: The effect of various hyperparameters on model capacity.

Recipe for Overfitting

- ❑ **Early stopping**
 - ❑ Stopping the training on time before it becomes overtrained and overly complex.
- ❑ **Train with more data**
 - ❑ Expanding the training set to include more data
- ❑ **Data augmentation**
 - ❑ Creating many variations of clean data to make the model to generalize better to unseen examples.
- ❑ **Regularization**
 - ❑ Overfitting occurs when a model is too complex comparing with data.
 - ❑ Using regularization terms (L1, L2), batch norm, dropout, data mix-up, label smoothing, VAT (virtual adversarial training).
- ❑ **Ensemble methods (not cover in this lecture)**
 - ❑ Ensemble learning methods are made up of a set of classifiers and their predictions are aggregated to identify the most popular result.
 - ❑ The most well-known ensemble methods are **bagging** and **boosting**.

Early stopping



- At first, the train and valid losses **gradually drop**, but the model is **not good enough** to characterise data
 - **Underfitting** is happening
- At a certain point, the train loss **still decreases**, while the valid loss **starts increasing**
 - **Overfitting** starts happening and we need to do **early stopping** at this point

Add Regularization

Reduce Overfitting

- Optimization problem

$$\min_{\theta} J(\theta) = \Omega(\theta) + \frac{1}{N} \sum_{i=1}^N l(y_i, f(x_i; \theta))$$

- L2 regularization

$$\Omega(\theta) = \lambda \sum_k \sum_{i,j} (W_{i,j}^k)^2 = \lambda \sum_k \|W^k\|_F^2$$


- $\lambda > 0$ is regularization parameter
- Gradient: $\nabla_{W^k} \Omega(\theta) = 2W^k$
- Apply on weights (W) only, not on biases (b)

- L1 regularization

$$\Omega(\theta) = \lambda \sum_k \sum_{i,j} |W_{i,j}^k|$$


- $\lambda > 0$ is regularization parameter
- Optimization is now much harder – subgradient.
- Apply on weights (W) only, not on biases (b)

Training Vs. Test Set Error




Dropout


Reduce Overfitting



Without dropout



With dropout



(Source: chatbotslife)

- This is a **cheap technique** to reduce model capacity
 - Reduce overfitting
- In each iteration, at each layer, **randomly choose** some neurons and **drop all connections** from these neurons
 - `dropout_rate = 1 - keep_prob`


Dropout

Reduce Overfitting

- Computationally efficient
- Can be considered a **bagged ensemble** of an exponential number (2^N) of neural networks.
- **Training**

$$\begin{aligned} \mathbf{r} &\sim \text{Bernoulli}(\mu) \\ \tilde{\mathbf{h}}^l &= \mathbf{h}^l \odot \mathbf{r} \\ \mathbf{h}^{l+1} &= \sigma(W^{(l)\top} \tilde{\mathbf{h}}^l + \mathbf{b}^l) \end{aligned}$$

$$\begin{matrix} h^{l+1} \\ \uparrow \\ h^l \\ \uparrow \\ h^{l-1} \end{matrix}$$



- **Testing**
 - **No dropout** (`dropout_rate = 0`).


Internal covariate shift problem

Input batch

x_1^1	x_2^1	x_3^1	x_4^1
x_1^2	x_2^2	x_3^2	x_4^2
x_1^3	x_2^3	x_3^3	x_4^3
x_1^4	x_2^4	x_3^4	x_4^4

Representation batch

z_1^1	z_2^1	z_3^1
z_1^2	z_2^2	z_3^2
z_1^3	z_2^3	z_3^3
z_1^4	z_2^4	z_3^4



Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift

Sergey Ioffe
Christian Szegedy

Google, 1600 Amphitheatre Pkwy, Mountain View, CA 94043


SIOFFE@GOOGLE.COM
SZEGEDY@GOOGLE.COM

Paper link: <http://proceedings.mlr.press/v37/ioffe15.pdf>

Internal covariate shift problem

Input batch 1


x_1^1	x_2^1	x_3^1	x_4^1
x_1^2	x_2^2	x_3^2	x_4^2
x_1^3	x_2^3	x_3^3	x_4^3
x_1^4	x_2^4	x_3^4	x_4^4



difference

Input batch 2

x_1^5	x_2^5	x_3^5	x_4^5
x_1^6	x_2^6	x_3^6	x_4^6
x_1^7	x_2^7	x_3^7	x_4^7
x_1^8	x_2^8	x_3^8	x_4^8



difference


Input batch 3

x_1^9	x_2^9	x_3^9	x_4^9
x_1^{10}	x_2^{10}	x_3^{10}	x_4^{10}
x_1^{11}	x_2^{11}	x_3^{11}	x_4^{11}
x_1^{12}	x_2^{12}	x_3^{12}	x_4^{12}



Representation batch 1

z_1^1	z_2^1	z_3^1
z_1^2	z_2^2	z_3^2
z_1^3	z_2^3	z_3^3
z_1^4	z_2^4	z_3^4



Representation batch 2

z_1^5	z_2^5	z_3^5
z_1^6	z_2^6	z_3^6
z_1^7	z_2^7	z_3^7
z_1^8	z_2^8	z_3^8




Representation batch 3

z_1^9	z_2^9	z_3^9
z_1^{10}	z_2^{10}	z_3^{10}
z_1^{11}	z_2^{11}	z_3^{11}
z_1^{12}	z_2^{12}	z_3^{12}



□ **Internal covariate shift problem:**

- Significant difference in statistics of input batches and also representation batches.
- Statistical differences among mini-batches make the classifier harder to learn from data.



Feature extractor

Classifier

Batch Normalization

Input batch 1

x_1^1	x_2^1	x_3^1	x_4^1
x_1^2	x_2^2	x_3^2	x_4^2
x_1^3	x_2^3	x_3^3	x_4^3
x_1^4	x_2^4	x_3^4	x_4^4

Input batch 2

x_1^5	x_2^5	x_3^5	x_4^5
x_1^6	x_2^6	x_3^6	x_4^6
x_1^7	x_2^7	x_3^7	x_4^7
x_1^8	x_2^8	x_3^8	x_4^8

Input batch 3

x_1^9	x_2^9	x_3^9	x_4^9
x_1^{10}	x_2^{10}	x_3^{10}	x_4^{10}
x_1^{11}	x_2^{11}	x_3^{11}	x_4^{11}
x_1^{12}	x_2^{12}	x_3^{12}	x_4^{12}

Representation batch 1

z_1^1	z_2^1	z_3^1
z_1^2	z_2^2	z_3^2
z_1^3	z_2^3	z_3^3
z_1^4	z_2^4	z_3^4

Normalize to $N(0, I)$

\hat{z}_1^1	\hat{z}_2^1	\hat{z}_3^1
\hat{z}_1^2	\hat{z}_2^2	\hat{z}_3^2
\hat{z}_1^3	\hat{z}_2^3	\hat{z}_3^3
\hat{z}_1^4	\hat{z}_2^4	\hat{z}_3^4

Normalized representation 1

difference

difference

difference

Representation batch 2

z_1^5	z_2^5	z_3^5
z_1^6	z_2^6	z_3^6
z_1^7	z_2^7	z_3^7
z_1^8	z_2^8	z_3^8

Normalize to $N(0, I)$

\hat{z}_1^5	\hat{z}_2^5	\hat{z}_3^5
\hat{z}_1^6	\hat{z}_2^6	\hat{z}_3^6
\hat{z}_1^7	\hat{z}_2^7	\hat{z}_3^7
\hat{z}_1^8	\hat{z}_2^8	\hat{z}_3^8

Normalized representation 2


Representation batch 3

z_1^9	z_2^9	z_3^9
z_1^{10}	z_2^{10}	z_3^{10}
z_1^{11}	z_2^{11}	z_3^{11}
z_1^{12}	z_2^{12}	z_3^{12}

Normalize to $N(0, I)$

\hat{z}_1^9	\hat{z}_2^9	\hat{z}_3^9
\hat{z}_1^{10}	\hat{z}_2^{10}	\hat{z}_3^{10}
\hat{z}_1^{11}	\hat{z}_2^{11}	\hat{z}_3^{11}
\hat{z}_1^{12}	\hat{z}_2^{12}	\hat{z}_3^{12}



Normalized representation 3



Feature extractor

Classifier


Batch Normalization



Batch Normalization

1. Cope with internal covariate shift
2. Reduce gradient vanishing/exploding
3. Reduce overfitting
4. Make training more stable
5. Converge faster
 1. Allow us to train with bigger learning rate



- Let $z = W^k h^k + b^k$ be the mini-batch before activation. We compute the normalized \hat{z} as
 - $\hat{z} = \frac{z - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$ where ϵ is a small value such as $1e^{-7}$
 - $\mu_B = \frac{1}{b} \sum_{i=1}^b z_i$ is the empirical mean
 - $\sigma_B^2 = \frac{1}{b} \sum_{i=1}^b (z_i - \mu_B)^2$ is the empirical variance
- We scale the normalized \hat{z}
 - $z_{BN} = \gamma \hat{z} + \beta$ where $\gamma, \beta > 0$ are two learnable parameters (i.e., scale and shift parameters)
- We then apply the activation to obtain the next layer value
 - $h^{k+1} = \sigma(z_{BN})$



Batch Normalization


Testing Phase

- At testing time, let's say we only want to test on a single input x
- Hence, we don't have a set of mini-batch samples to compute mean and standard deviation.
- So, we replace the mini-batch μ_B and σ_B with running averages of $\tilde{\mu}_B$ and $\tilde{\sigma}_B$ computed during the training process.
 - $\tilde{\mu}_B = \theta\tilde{\mu}_B + (1 - \theta)\mu_B$
 - $\tilde{\sigma}_B = \theta\tilde{\sigma}_B + (1 - \theta)\sigma_B$
 - $0 < \theta < 1$ is the momentum decay (i.e., $\theta = 0.9$).




(Source: medium.com)

Data Augmentation



Reality

- You have a **tiny dataset of 10K images**, and you need to train a **good deep net**.



What are the **criteria of a qualified training set?**

Quantity?

- Collect **more and more data**?




Prof Dinh Phung

Create many variants

Quality?

- More **diverse** data?
- More **qualified** data?



Original	Flip	Rotation	Random crop

- Image without any modification

- Flipped with respect to an axis for which the meaning of the image is preserved

- Rotation with a slight angle
- Simulates incorrect horizon calibration

- Random focus on one part of the image
- Several random crops can be done in a row

Color shift	Noise addition	Information loss	Contrast change

- Nuances of RGB is slightly changed
- Captures noise that can occur with light exposure

- Addition of noise
- More tolerance to quality variation of inputs

- Parts of image ignored
- Mimics potential loss of parts of image

- Luminosity changes
- Controls difference in exposition due to time of day

Data Augmentation

- Use **simple transformations** to augment data examples. Models will be **challenged** with **diverse data examples** which might be **encountered** in the testing set
 - Rotation, Width Shifting, Height Shifting, Brightness, Shear Intensity, Zoom, Channel Shift, Horizontal Flip, Vertical Flip
- This will **reduce overfitting**, making this a **regularization technique**. The trick is to generate **realistic training instances**

Original	Flip	Rotation	Random crop
			
<ul style="list-style-type: none"> • Image without any modification 	<ul style="list-style-type: none"> • Flipped with respect to an axis for which the meaning of the image is preserved 	<ul style="list-style-type: none"> • Rotation with a slight angle • Simulates incorrect horizon calibration 	<ul style="list-style-type: none"> • Random focus on one part of the image • Several random crops can be done in a row
Color shift	Noise addition	Information loss	Contrast change
			
<ul style="list-style-type: none"> • Nuances of RGB is slightly changed • Captures noise that can occur with light exposure 	<ul style="list-style-type: none"> • Addition of noise • More tolerance to quality variation of inputs 	<ul style="list-style-type: none"> • Parts of image ignored • Mimics potential loss of parts of image 	<ul style="list-style-type: none"> • Luminosity changes • Controls difference in exposition due to time of day

[Source: Stanford CS 230 Deep Learning]

Data Augmentation in PyTorch

```
test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)), # Normalize the images, each R,G,B value is normalized with mean=0.5 and std=0.5
    transforms.Resize((32,32)), #resizes the image so it can be perfect for our model.
])

train_transform = transforms.Compose([
    transforms.Resize((32,32)), #resizes the image so it can be perfect for our model.
    transforms.RandomHorizontalFlip(), # Flips the image w.r.t horizontal axis
    #transforms.RandomRotation(4), #Rotates the image to a specified angel
    #transforms.RandomAffine(0, shear=10, scale=(0.8,1.2)), #Performs actions like zooms, change shear angles.
    transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2), # Set the color params
    transforms.ToTensor(), # convert the image to tensor so that it can work with torch
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)), # Normalize the images, each R,G,B value is normalized with mean=0.5 and std=0.5
])

full_train_set = torchvision.datasets.CIFAR10("./data", download=True, transform=train_transform) # Apply train_transform to train set
full_valid_set = torchvision.datasets.CIFAR10("./data", download=True, transform=test_transform) # Apply test_transform to generate valid set
full_test_set = torchvision.datasets.CIFAR10("./data", download=True, train=False, transform=test_transform)

n_train, n_valid, n_test = 5000, 5000, 5000

total_num_train = len(full_train_set)
total_num_test = len(full_test_set)
train_valid_idx = torch.randperm(total_num_train)
train_set = torch.utils.data.Subset(full_train_set, train_valid_idx[:n_train])
valid_set = torch.utils.data.Subset(full_valid_set, train_valid_idx[n_train:n_train+n_valid])

test_idx = torch.randperm(total_num_test)
test_set = torch.utils.data.Subset(full_test_set, test_idx[:n_test])

print("Train set\nNumber of samples:{}\nShape of 1 sample:{}\n".format(len(train_set), list(train_set[0][0].shape)))
print("Valid set\nNumber of samples:{}\nShape of 1 sample:{}\n".format(len(valid_set), list(valid_set[0][0].shape)))
print("Test set\nNumber of samples:{}\nShape of 1 sample:{}\n".format(len(test_set), list(test_set[0][0].shape)))


train_loader = torch.utils.data.DataLoader(train_set, batch_size=32, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_set, batch_size=32)
valid_loader = torch.utils.data.DataLoader(valid_set, batch_size=32)

img_train = [train_set[idx][0].numpy().transpose((1,2,0)) for idx in range(32)]
label_train = [train_set[idx][1] for idx in range(32)]
```

Without augmentation



With augmentation






Some additional regularization techniques

Label smoothing

When Does Label Smoothing Help?

Rafael Müller*, Simon Kornblith, Geoffrey Hinton
Google Brain
Toronto
rafaelmuller@google.com

Paper link: <https://papers.nips.cc/paper/2019/file/f1748d6b0fd9d439f71450117eba2725-Paper.pdf>





- Given **data instance** (x, y) with the label $y \in \{1, \dots, M\}$, we compute the **CE loss** between the **prediction probabilities** $p(x)$ and the **smooth label**

- $l(y, \hat{y}) = CE(\tilde{\mathbf{1}}_y, p(x))$ with $\tilde{\mathbf{1}}_y = (1 - \alpha) \times \mathbf{1}_y + \frac{\alpha}{M} \times \mathbf{1}$ where $\mathbf{1}$ is a vector of all 1 and $0 < \alpha < 1$.

Data mix-up

mixup: BEYOND EMPIRICAL RISK MINIMIZATION


 $(x_1, \mathbf{1}_{y_1})$


$\lambda \sim \text{Beta}(\alpha, \alpha)$

Blended image $\tilde{x} = \lambda \times x_1 + (1 - \lambda) \times x_2$

Blended label $\tilde{y} = \lambda \times \mathbf{1}_{y_1} + (1 - \lambda) \times \mathbf{1}_{y_2}$

$\min CE(\tilde{y}, p(\tilde{x}))$

Hongyi Zhang
MIT

Moustapha Cisse, Yann N. Dauphin, David Lopez-Paz*
FAIR

Paper link: <https://openreview.net/pdf?id=r1Ddp1-Rb>



[Source: <https://medium.com/>]

□ for $(x_1, y_1), (x_2, y_2)$ in `zip(batch1, batch 2)`

1. $\lambda \sim \text{Beta}(\alpha, \alpha)$
2. $\tilde{x} = \lambda \times x_1 + (1 - \lambda) \times x_2$
3. $\tilde{y} = \lambda \times \mathbf{1}_{y_1} + (1 - \lambda) \times \mathbf{1}_{y_2}$
4. Update optimizer to minimize $CE(\tilde{y}, p(\tilde{x}))$

Cut-mix

[Source: <https://encord.com/blog/data-augmentation-guide/>]



$$(x_1, \mathbf{1}_{y_1}), x_1 \in \mathbb{R}^{C \times W \times H}$$

$$\lambda \sim \text{Beta}(\alpha, \alpha)$$

$$\tilde{x} = M \odot x_1 + (1 - M) \odot x_2, M \in \{0,1\}^{H \times W}$$

$$\tilde{y} = \lambda \times \mathbf{1}_{y_1} + (1 - \lambda) \times \mathbf{1}_{y_2}$$

$$\min CE(\tilde{y}, p(\tilde{x}))$$

for $(x_1, y_1), (x_2, y_2)$ in zip(batch1, batch 2)

1. $\lambda \sim \text{Beta}(\alpha, \alpha)$

2. Sample a bounding box $B = [r_x, r_y, r_w, r_h]$

1. $r_x \sim \text{Uni}[0, W], r_w \sim W\sqrt{1 - \lambda}$

2. $r_y \sim \text{Uni}[0, H], r_h \sim H\sqrt{1 - \lambda}$

3. $M \in \{0, 1\}^{W \times H}$ by filling 1 within B and 0 otherwise

4. $\tilde{x} = M \odot x_1 + (1 - M) \odot x_2$

5. Update optimizer to minimize $CE(\tilde{y}, p(\tilde{x}))$

CutMix: Regularization Strategy to Train Strong Classifiers with Localizable Features

Sangdoo Yun¹

Dongyoон Han¹
Junsuk Choe^{1,3}

Seong Joon Oh²
Youngjoon Yoo¹

Sanghyuk Chun¹

¹Clova AI Research, NAVER Corp.

²Clova AI Research, LINE Plus Corp.

³Yonsei University

[Source: <https://arxiv.org/pdf/1905.04899>]

M

1	1	1	1	1	1	1
1	0	0	0	1		
1	0	0	0	1		
1	0	0	0	1		
1	0	0	0	1		
1	1	1	1	1	1	1

$$B = [r_x, r_y, r_w, r_h]$$

$$\frac{\text{area}(B)}{\text{area(image)}} = \frac{WH(1 - \lambda)}{WH} = 1 - \lambda$$



$$\odot M =$$



$$\odot (1 - M) =$$





Transfer learning and fine-tuning

Transfer Learning


Motivation

Large-scale dataset (ImageNet)

- over 1.2 million images and 1,000 possible object categories



printer housing animal weight
offspring teacher computer drop headquarters egg white
register gallery court key structure light date spread
king fire place church press market lighter
restaurant counter cup concert market lighter
hotel road Paper side site door pack
sport screen wall means fan hill can camp fish coffee
sky plant house school stock film
bread weapon table top man car fly study bird
cloud cover range leash van suite mirror seat
spring fruit dog shop kit roll bar watch
bed shop train tea overall sleeve goal
kitchen camera box memory sieve cell bar
engine chain boat center step
dinner stone child case student
apple girl flat flag bank home room office
flag bank home room office rule hall club
radio valley cross chair mine castle Support level line street golf
beach library stage video food building
base material player machine security call clock
tool hospital match equipment cell phone mountain telephone
short circuit bridge scale equipment gas pedal microphone recording crowd
football hospital match equipment cell phone mountain telephone



Transfer learning
+ Fine tune



Your small-scale dataset

(Flower-17)

- 17 category dataset with 80 images per class


Not enough data to train
a good model.
HOW?

Powerful pretrained model

- VGG16, VGG19
- Inception V3
- Xception
- ResNet



Low-level filters



Mid-level filters



High-level filters


Fits the target
dataset

Transfer Learning

How to Do That?

- Remove FC layers from the pretrained model


- Replace them with a brand-new FC head.
 - These new FC layers can then be fine-tuned to the specific dataset
 - The old FC layers are no longer used



Transfer Learning

How to Do That?


- **Freeze** all CONV layers in the network
 - Only allow the gradient to backpropagate through the FC layers
 - Doing this allows our network to **warm up**
- Training data is forward propagated through the network
 - However, the backpropagation is stopped after the FC layers
 - Allows these layers to start to learn patterns from the highly discriminative CONV layers



Transfer Learning

How to Do That?

- After the FC layers have had a chance to warm up, we may choose to **unfreeze** all layers in the network
 - Allow each of them to be **fine-tuned**.
 - Continue training the entire network, *but with a very small learning rate*
 - We do not want to **deviate our CONV filters** dramatically. Training is then allowed to continue until sufficient accuracy is obtained.



Model zoo supported by PyTorch

Model	Acc@1	Acc@5			
AlexNet	56.522	79.066	Densenet-169	75.600	92.806
VGG-11	69.020	88.628	Densenet-201	76.896	93.370
VGG-13	69.928	89.246	Densenet-161	77.138	93.560
VGG-16	71.592	90.382	Inception v3	77.294	93.450
VGG-19	72.376	90.876	GoogleNet	69.778	89.530
VGG-11 with batch normalization	70.370	89.810	ShuffleNet V2 x1.0	69.362	88.316
VGG-13 with batch normalization	71.586	90.374	ShuffleNet V2 x0.5	60.552	81.746
VGG-16 with batch normalization	73.360	91.516	MobileNet V2	71.878	90.286
VGG-19 with batch normalization	74.218	91.842	MobileNet V3 Large	74.042	91.340
ResNet-18	69.758	89.078	MobileNet V3 Small	67.668	87.402
ResNet-34	73.314	91.420	ResNeXt-50-32x4d	77.618	93.698
ResNet-50	76.130	92.862	ResNeXt-101-32x8d	79.312	94.526
ResNet-101	77.374	93.546	Wide ResNet-50-2	78.468	94.086
ResNet-152	78.312	94.046	Wide ResNet-101-2	78.848	94.284
SqueezeNet 1.0	58.092	80.420	MNASNet 1.0	73.456	91.510
SqueezeNet 1.1	58.178	80.624	MNASNet 0.5	67.734	87.490
Densenet-121	74.434	91.972			

Transfer learning with PyTorch

```
# Load pretrained VGG19 model
model = models.vgg19(pretrained=True)
model = model.to(device)
summary(model, (3, 224, 224))

/usr/local/lib/python3.10/dist-packages/torch
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/mo
100%[██████████] 548M/548M [00:07<00:00, 75.8]
```

Load pretrained VGG19

```
VGG(
    (features): Sequential(
        (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (3): ReLU(inplace=True)
        (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (6): ReLU(inplace=True)
        (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (8): ReLU(inplace=True)
        (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (11): ReLU(inplace=True)
        (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (13): ReLU(inplace=True)
        (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (15): ReLU(inplace=True)
        (16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (17): ReLU(inplace=True)
        (18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (20): ReLU(inplace=True)
        (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (22): ReLU(inplace=True)
        (23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (24): ReLU(inplace=True)
        (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (26): ReLU(inplace=True)
        (27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (29): ReLU(inplace=True)
        (30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (31): ReLU(inplace=True)
        (32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (33): ReLU(inplace=True)
        (34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (35): ReLU(inplace=True)
        (36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
    (classifier): Sequential(
        (0): Linear(in_features=25088, out_features=4096, bias=True)
        (1): ReLU(inplace=True)
        (2): Dropout(p=0.5, inplace=False)
        (3): Linear(in_features=4096, out_features=4096, bias=True)
        (4): ReLU(inplace=True)
        (5): Dropout(p=0.5, inplace=False)
        (6): Linear(in_features=4096, out_features=1000, bias=True)
    )
)
```

Will be replaced by a fresh new linear layer

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 224, 224]	1,792
ReLU-2	[-1, 64, 224, 224]	0
Conv2d-3	[-1, 64, 224, 224]	36,928
ReLU-4	[-1, 64, 224, 224]	0
MaxPool2d-5	[-1, 64, 112, 112]	0
Conv2d-6	[-1, 128, 112, 112]	73,856
ReLU-7	[-1, 128, 112, 112]	0
Conv2d-8	[-1, 128, 112, 112]	147,584
ReLU-9	[-1, 128, 112, 112]	0
MaxPool2d-10	[-1, 128, 56, 56]	0
Conv2d-11	[-1, 256, 56, 56]	295,168
ReLU-12	[-1, 256, 56, 56]	0
Conv2d-13	[-1, 256, 56, 56]	590,080
ReLU-14	[-1, 256, 56, 56]	0
Conv2d-15	[-1, 256, 56, 56]	590,080
ReLU-16	[-1, 256, 56, 56]	0
Conv2d-17	[-1, 256, 56, 56]	590,080
ReLU-18	[-1, 256, 56, 56]	0
MaxPool2d-19	[-1, 256, 28, 28]	0
Conv2d-20	[-1, 512, 28, 28]	1,180,160
ReLU-21	[-1, 512, 28, 28]	0
Conv2d-22	[-1, 512, 28, 28]	2,359,808
ReLU-23	[-1, 512, 28, 28]	0
Conv2d-24	[-1, 512, 28, 28]	2,359,808
ReLU-25	[-1, 512, 28, 28]	0
Conv2d-26	[-1, 512, 28, 28]	2,359,808
ReLU-27	[-1, 512, 28, 28]	0
MaxPool2d-28	[-1, 512, 14, 14]	0
Conv2d-29	[-1, 512, 14, 14]	2,359,808
ReLU-30	[-1, 512, 14, 14]	0
Conv2d-31	[-1, 512, 14, 14]	2,359,808
ReLU-32	[-1, 512, 14, 14]	0
Conv2d-33	[-1, 512, 14, 14]	2,359,808
ReLU-34	[-1, 512, 14, 14]	0
Conv2d-35	[-1, 512, 14, 14]	2,359,808
ReLU-36	[-1, 512, 14, 14]	0
MaxPool2d-37	[-1, 512, 7, 7]	0
AdaptiveAvgPool2d-38	[-1, 512, 7, 7]	0
Linear-39	[-1, 4096]	102,764,544
ReLU-40	[-1, 4096]	0
Dropout-41	[-1, 4096]	0
Linear-42	[-1, 4096]	16,781,312
ReLU-43	[-1, 4096]	0
Dropout-44	[-1, 4096]	0
Linear-45	[-1, 1000]	4,097,000

Total params: 143,667,240

Trainable params: 143,667,240

Transfer learning with PyTorch

Warm-up the fresh new linear layer

```
# Freeze all layers
for param in model.parameters():
    param.requires_grad = False

# Modify the last fully connected layer for Flower102
num_features = model.classifier[6].in_features
model.classifier[6] = nn.Linear(num_features, 102)
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=1e-3, momentum=0.9)
trainer = BaseTrainer(model, criterion, optimizer, train_loader, val_loader)
trainer.fit(num_epochs=10)

Epoch 1/10
train_loss: 0.9840 - train_accuracy: 0.7027 - val_loss: 0.4316 - top1_acc: 0.8501 - top5_acc: 1.0000
Epoch 2/10
train_loss: 0.4744 - train_accuracy: 0.8287 - val_loss: 0.3686 - top1_acc: 0.8706 - top5_acc: 1.0000
Epoch 3/10
train_loss: 0.4272 - train_accuracy: 0.8505 - val_loss: 0.3627 - top1_acc: 0.8760 - top5_acc: 1.0000
Epoch 4/10
train_loss: 0.3894 - train_accuracy: 0.8535 - val_loss: 0.3507 - top1_acc: 0.8787 - top5_acc: 1.0000
Epoch 5/10
train_loss: 0.3738 - train_accuracy: 0.8631 - val_loss: 0.3283 - top1_acc: 0.8937 - top5_acc: 1.0000
Epoch 6/10
train_loss: 0.3593 - train_accuracy: 0.8719 - val_loss: 0.3252 - top1_acc: 0.8815 - top5_acc: 1.0000
Epoch 7/10
train_loss: 0.3309 - train_accuracy: 0.8886 - val_loss: 0.3134 - top1_acc: 0.8910 - top5_acc: 1.0000
Epoch 8/10
train_loss: 0.3180 - train_accuracy: 0.8903 - val_loss: 0.3070 - top1_acc: 0.8924 - top5_acc: 1.0000
Epoch 9/10
train_loss: 0.3062 - train_accuracy: 0.8917 - val_loss: 0.3121 - top1_acc: 0.8978 - top5_acc: 1.0000
Epoch 10/10
train_loss: 0.3139 - train_accuracy: 0.8849 - val_loss: 0.3082 - top1_acc: 0.8924 - top5_acc: 1.0000
```

Fine-tune the entire model

```
# Unfreeze the last convolutional block
# for param in model.features[-7:].parameters():
#     param.requires_grad = True

# Unfreeze model parameter
for param in model.parameters():
    param.requires_grad = True

criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=1e-4, momentum=0.9)
trainer = BaseTrainer(model, criterion, optimizer, train_loader, val_loader)
trainer.fit(num_epochs=10)

Epoch 1/10
train_loss: 0.2559 - train_accuracy: 0.9114 - val_loss: 0.2631 - top1_acc: 0.9101 - top5_acc: 1.0000
Epoch 2/10
train_loss: 0.1983 - train_accuracy: 0.9305 - val_loss: 0.2489 - top1_acc: 0.9196 - top5_acc: 1.0000
Epoch 3/10
train_loss: 0.1561 - train_accuracy: 0.9472 - val_loss: 0.2435 - top1_acc: 0.9196 - top5_acc: 1.0000
Epoch 4/10
train_loss: 0.1417 - train_accuracy: 0.9513 - val_loss: 0.2349 - top1_acc: 0.9210 - top5_acc: 1.0000
Epoch 5/10
train_loss: 0.1089 - train_accuracy: 0.9632 - val_loss: 0.2379 - top1_acc: 0.9223 - top5_acc: 1.0000
Epoch 6/10
train_loss: 0.0947 - train_accuracy: 0.9690 - val_loss: 0.2370 - top1_acc: 0.9223 - top5_acc: 1.0000
Epoch 7/10
train_loss: 0.0865 - train_accuracy: 0.9690 - val_loss: 0.2356 - top1_acc: 0.9319 - top5_acc: 1.0000
Epoch 8/10
train_loss: 0.0644 - train_accuracy: 0.9816 - val_loss: 0.2403 - top1_acc: 0.9319 - top5_acc: 1.0000
Epoch 9/10
train_loss: 0.0516 - train_accuracy: 0.9816 - val_loss: 0.2477 - top1_acc: 0.9319 - top5_acc: 1.0000
Epoch 10/10
train_loss: 0.0530 - train_accuracy: 0.9837 - val_loss: 0.2477 - top1_acc: 0.9251 - top5_acc: 1.0000
```

Summary

- Setting of a machine learning problem
 - General loss versus empirical loss
- Gradient vanishing/exploding and network initialization.
- Overfitting and underfitting
- Recipe for overfitting
 - Use regularization term
 - Dropout, batch norm
 - Data augmentation
 - Transfer learning
 - Label smoothing, data mix-up

Thanks for your attention!
Question time

