

**COPYRIGHT WARNING:** Copyright in these original lectures is owned by Monash University. You may transcribe, take notes, download or stream lectures for the purpose of your research and study only. If used for any other purpose, (excluding exceptions in the Copyright Act 1969 (Cth)) the University may take legal action for infringement of copyright.

Do not share, redistribute, or upload the lecture to a third party without a written permission!

## **FIT3181/5215 Deep Learning**

Week 09: Deep Learning for Sequential Data (II):  
Seq2Seq and Transformers

**Lecturer: Trung Le**

Email: [trunglm@monash.edu](mailto:trunglm@monash.edu)






# Outline

- Encoder-Decoder models
- Sequence to sequence models
  - Machine Translation and Image Captioning
- Attention mechanism
  - Global attention and Local attention
- Transformer and BERT
  - Self-Attention Mechanism & Multi-Head Self-Attention Mechanism
  - Pre-trained Language Models: BERT

# Encoder-Decoder Models

- Image Captioning
- Machine Translation

# Encoder-Decoder Models: Application Motivations



**Machine translation**

## Many more applications

- ❖ C++ programs to Java programs, texts to images, question answering (questions to answers), and etc.



# Sequence to Sequence Models

# Introduction

- Problem statement of seq2seq
  - Source sequence:  $x = (x_1, x_2, \dots, x_{T_x})$  where  $x_i \in V_x$
  - Target sequence:  $y = (y_1, y_2, \dots, y_{T_y})$  where  $y_i \in V_y$
  - Training set:
    - $\mathcal{D} = \{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(N)}, y^{(N)})\} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$
    - $S_x$ : set of possible source sequences ( $\{x^{(i)}\} \subseteq S_x$ )
    - $S_y$ : set of possible target sequences ( $\{y^{(i)}\} \subseteq S_y$ )
  - Task
    - Learn a function  $f: S_x \rightarrow S_y$
- Applications
  - Machine Translation
  - Image Captioning

# Machine translation

## Sequence to Sequence Learning with Neural Networks


Paper: [paper link](#)

Ilya Sutskever  
Google  
ilyasu@google.com



Oriol Vinyals  
Google  
vinyals@google.com

Quoc V. Le  
Google  
qvl@google.com

- Translate a sentence in a language to another language
  - Input: He loved to eat. (English)
  - Output: Il aimait manger. (French)



### Encoder



# Encoder-decoder model for seq2seq

## Fixed context vector

### Encoder


- Produces the context vector  $\mathbf{c} = \mathbf{h}_{T_x}$  of the input sequence
- Context vector  $\mathbf{c}$  summarizes input sequence  $[\mathbf{x}_1, \dots, \mathbf{x}_{T_x}]$ .

### Decoder


- Decodes the encode  $\mathbf{c}$  to the output sequence

### Special symbols

- $\text{<EOS>}$  signifies the end of a sequence
- $\text{<BOS>}$  signifies the beginning of a sequence



# Training of seq2seq



- We need to maximize the log-likelihood:

$$\max_{\theta} J(\theta) = \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}} \log P(\mathbf{y} | \mathbf{x}, \theta)$$

where  $\theta = [\theta_e, \theta_d]$  and  $\theta_e, \theta_d$  are encoding and decoding parameters respectively.

- Product rule:

$$\begin{aligned} P(\mathbf{y} | \mathbf{x}, \theta) &= P(\mathbf{y}_{1:T_y} | \mathbf{x}_{1:T_x}, \theta) = P(\mathbf{y}_{1:T_y} | \mathbf{c}, \theta) \\ &= P(\mathbf{y}_1 | \mathbf{c}, \theta) P(\mathbf{y}_2 | \mathbf{y}_1, \mathbf{c}, \theta) \dots P(\mathbf{y}_j | \mathbf{y}_{1:j-1}, \mathbf{c}, \theta) \dots P(\mathbf{y}_{T_y} | \mathbf{y}_{1:T_y-1}, \mathbf{c}, \theta) = \prod_{j=1}^{T_y} P(\mathbf{y}_j | \mathbf{y}_{1:j-1}, \mathbf{c}, \theta) \end{aligned}$$

$$\log P(\mathbf{y} | \mathbf{x}, \theta) = \log P(\mathbf{y} | \mathbf{c}, \theta) = \sum_{j=1}^{T_y} \log P(\mathbf{y}_j | \mathbf{y}_{1:j-1}, \mathbf{c}, \theta) = \sum_{j=1}^{T_y} \log P(\mathbf{y}_j | \mathbf{q}_{j-1}, \mathbf{c}, \theta)$$

- We can compute  $P(\mathbf{y}_j | \mathbf{q}_{j-1}, \mathbf{c}) = g(\mathbf{y}_j, \mathbf{q}_{j-1}, \mathbf{c})$  where  $g$  is a nonlinear, potentially multi-layered NN that outputs the probability of  $y_j$ .
- Pay attention on how  $c$  is used in every step during decoding

# Training of seq2seq

- We need to maximize the log-likelihood:

$$\max_{\theta} J(\theta) = \sum_{(x,y) \in \mathcal{D}} \log P(y|x, \theta)$$


where  $\theta = [\theta_e, \theta_d]$  and  $\theta_e, \theta_d$  are encoding and decoding parameters, respectively.

- Product rule:

$$\begin{aligned} \log P(y|x, \theta) &= \log P(y|\textcolor{red}{c}, \theta) = \\ \sum_{j=1}^{T_y} \log P(y_j|y_{1:j-1}, \textcolor{red}{c}, \theta) &= \\ \sum_{j=1}^{T_y} \log P(y_j|\textcolor{green}{q}_{j-1}, \textcolor{red}{c}, \theta) \end{aligned}$$


- $\textcolor{green}{q}_{j-1}$  contains the information of  $\textcolor{red}{c}$ .
- (1):  $P(y_j|\textcolor{green}{q}_{j-1}, \textcolor{red}{c}, \theta) = P(y_j|\textcolor{green}{q}_{j-1}, \theta)$   
 $q_{j-1} = f(q_{j-2}, y_{j-1})$
- (2):  $P(y_j|\textcolor{green}{q}_{j-1}, \textcolor{red}{c}, \theta) = P(y_j|\textcolor{green}{q}_{j-1}, \theta)$   
 $q_{j-1} = f(q_{j-2}, \text{concat}(c, y_{j-1}))$
- $f$  is a memory cell (e.g., LSTM or GRU)

(1) [Sutskever et al. 2014]




**Big drawback:** the context vector  $c$  is fixed across timesteps.

# Inference



We know  $P(y_1 = \circ | c)$  We know  $P(y_2 = \circ | c)$




- Given a **trained model** and **input sequence  $x$** 
  - We need to infer the corresponding **output sequence  $y$**
- Two common strategies
  - Greedy Decoding and Beam Search Decoding

# Inference

## Greedy decoding

### □ Greedy Decoding


- Given  $\mathbf{x}$ , find word  $\mathbf{y}_1$  with highest probability
- Given  $\mathbf{y}_1$  and  $\mathbf{x}$ , find word  $\mathbf{y}_2$  with highest probability
- ...
- Stop when seeing  $\langle EOS \rangle$



# Inference


## Beam search

- Beam Search Decoding with **beam width  $k$** 
  - Given  $x$ , find  $k$  candidates for  $y_1$  with highest probability
  - Given  $x$ , for each candidate  $y_1$ , find  $k$  candidates for word  $y_2$  with highest probability
  - Pick top- $k$  sequences  $y_1y_2$  with highest joint probability
  - For each  $y_1y_2$ , find  $k$  candidates for word  $y_3$  with highest probability
  - Pick top- $k$  sequences  $y_1y_2y_3$  with highest joint probability
  - .....
  - Stop when see  $\langle EOS \rangle$  on each beam
  - Finally, pick 1 sequence with highest probability from top- $k$  sequences



# Inference

## Beam search




**Beam width = 3**

- We always choose three sentences with highest joint probabilities

# Inference

## Beam search



**Joint Probability**

$$\begin{aligned} p_1 p_{11} \dots \\ p_1 p_{12} \dots \\ p_1 p_{13} \dots \\ p_2 p_{11} \dots \\ \dots \end{aligned}$$

**Beam width = 3**

- We **always** choose three sentences with **highest joint probabilities**

# Drawback of fixed context


- Fixed context vector  $c$  is easily overwhelmed by long inputs or long outputs.
- At a specific timestep  $j$ , some words or items in the input sequence might possibly contribute more to the generation of next item or word in the output sequence.
  - I want to see you every day → Je veux te voir chaque jour
  - I want to see **you** every day → Je veux te ? (voir) ....
- How to timely adapt the context vector  $c_j$ ?
  - $c_j = \alpha(h_1, \dots, h_{T_x}, q_{j-1})$
  - Computed using attention mechanism

## NEURAL MACHINE TRANSLATION BY JOINTLY LEARNING TO ALIGN AND TRANSLATE

Dzmitry Bahdanau  
Jacobs University Bremen, Germany

KyungHyun Cho Yoshua Bengio\*  
Université de Montréal

Paper: [paper link](#)






# Attention mechanism

# Attention mechanism

- So far, the input sequence is **summarised** by a **single** context  $c$  !
- Fixed-length context** could be problematic as it is **easily overwhelmed** by long inputs or long outputs
  - The **fixed-length** context  $c$  might **not be powerful enough** to capture long input sequences
- Some **specific items** in an input sequence might be **more relevant** and **contributing in generating** a given item in output sequence.



- Gratefully acknowledge the excellent visualizations used from Jay Alammar blog at:
  - <https://jalammar.github.io/>
  - <https://jalammar.github.io/visualizing-neural-machine-translation-mechanics-of-seq2seq-models-with-attention/>

# Attention mechanism


Published as a conference paper at ICLR 2015

## NEURAL MACHINE TRANSLATION BY JOINTLY LEARNING TO ALIGN AND TRANSLATE

Dzmitry Bahdanau  
Jacobs University Bremen, Germany

KyungHyun Cho Yoshua Bengio\*  
Université de Montréal


Bahdanau, Cho, Bengio, [Neural Machine Translation by Jointly Learning to Align and Translate](#), ICLR 2015



## Effective Approaches to Attention-based Neural Machine Translation

Minh-Thang Luong Hieu Pham Christopher D. Manning  
Computer Science Department, Stanford University, Stanford, CA 94305  
[{lmthang, hyhieu, manning}@stanford.edu](mailto:{lmthang, hyhieu, manning}@stanford.edu)

Luong, Pham, Manning, [Effective Approach Attention-based Neural Machine Translation](#), EMNLP, 2015



- Attention mechanism allows the decoding network to refer to the input.
  - Global attention
    - Use all input hidden states of the encoder when deriving the context  $c_t$ .
  - Local attention
    - Use a selective window of input hidden states of the encoder when deriving the context  $c_t$ .

# Global attention

## □ Main idea

- Consider **all input hidden states** of the encoder when deriving the context:  $c_t = \sum_{s=1}^3 a_t(s) h_s$


## □ Alignment weights:

$$a_t(s) = \text{align}(q_t, h_s) = \frac{\exp(\text{score}(q_t, h_s))}{\sum_{s'} \exp(\text{score}(q_t, h_{s'}))}$$

where  $q_t$  is current target state,  $h_s$  is each source states.

## □ Alignment score function:

$$\text{score}(q_t, h_s) = \begin{cases} q_t^T h_s & \text{dot product} \\ q_t^T W_a h_s & \text{general metric} \\ v_a^T \tanh(W_a[q_t; h_s]) & \text{concat} \end{cases}$$



# Global attention

- Context vector:


$$c_t = \sum_{s=1}^3 a_t(s) h_s$$

- Attentional hidden state:

$$\tilde{q}_t = \tanh(W_c[c_t; q_t])$$

- Predictive distribution:

$$p(y_t | y_{<t}, x) = \text{softmax}(W_s \tilde{q}_t)$$



# Global attention

## Example

- Convert into alignment weights.

$$a_t(s) = \frac{\exp(score(q_t, h_s))}{\sum_{s'} \exp(score(q_t, h_{s'}))}$$

- Build context vector: weighted average


$$c_t = \sum_s a_t(s) h_s$$

- Compute the next hidden state.

$$\tilde{q}_t = \tanh(W_c[c_t; q_t])$$

- Predictive distribution:

$$p(y_t | y_{<t}, x) = \text{softmax}(W_s \tilde{q}_t)$$




# Global attention

## Drawback

- **Drawback:** employ **all items** on the source side for deriving each target item.
  - expensive computation
  - impractical to translate longer sequences





Local attention




# Local attention

- Main idea
  - Selectively focuses on a **small window** of context and is differentiable.
- $c_t$  is then derived as a **weighted average** over the set of source hidden states within the window  $[p_t - D; p_t + D]$ ,  $D$  is empirically selected.



# Local attention

- Main idea
  - Selectively focuses on a **small window** of context and is differentiable
- For current target word, predict aligned position  $p_t$ 
  - $p_t = S \cdot \text{sigmoid}(\nu_p^T \tanh(W_p q_t))$   
where  $W_p$  and  $\nu_p$  are the model parameters,  $S$  is the source sentence length, and hence  $0 \leq p_t \leq S$
- Alignment weights:
  - $a_t(s) = \text{align}(q_t, h_s) \exp\left(-\frac{(s-p_t)^2}{2\sigma^2}\right)$  with  $\sigma = D/2$
  - $\text{align}(q_t, h_s) = \frac{\exp(\text{score}(q_t, h_s))}{\sum_{s'} \exp(\text{score}(q_t, h_{s'}))}$  (**temporary alignment weights**)
- Context is now **computed** as before but **within a windows of size D**, centered at  $p_t$



# Other types of attention mechanism

- Self-attention/Multiple Heads/Transformer Networks
  - "Attention is all you need". Vaswani, Ashish, et al. ." *Advances in neural information processing systems*, 2017.
- Pyramidal encoders
  - "Listen Attend and Spell". William Chan, Navdeep Jaitly, Quoc Le, Oriol Vinyals, ICASSP 2015.
- Hierarchical Attention
  - "Learning efficient algorithms with hierarchical attentive memory". Andrychowicz, Marcin, and Karol Kurach, arXiv:1602.03218, 2016.
- Soft/Hard Attention
  - "Show, attend and tell: Neural image caption generation with visual attention". Xu, Kelvin, et al, ICML 2015.



# Transformer and BERT

Question

What does **GPT** stands for in **ChatGPT**?

# Transformer and BERT

---

## Attention Is All You Need

---

Ashish Vaswani\*  
Google Brain  
avaswani@google.com

Noam Shazeer\*  
Google Brain  
noam@google.com

Niki Parmar\*  
Google Research  
nikip@google.com

Jakob Uszkoreit\*  
Google Research  
usz@google.com

Llion Jones\*  
Google Research  
llion@google.com

Aidan N. Gomez\* †  
University of Toronto  
aidan@cs.toronto.edu

Lukasz Kaiser\*  
Google Brain  
lukasz.kaiser@google.com

Illia Polosukhin\* ‡  
illia.polosukhin@gmail.com

Paper: [paper link](#)

## BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin Ming-Wei Chang Kenton Lee Kristina Toutanova  
Google AI Language  
{jacobdevlin, mingweichang, kentonl, kristout}@google.com


Paper: [paper link](#)

- Transformer proposed **self-attention** mechanism
- BERT was **developed** based on Transformer
  - BERT currently achieves **SOTA performances** in most tasks of NLP.
- Driving force of generative AI revolution


# Transformer Motivations

- Problem using RNNs for Seq2Seq models
  - Slow due to sequential nature
  - Poor long-range dependency modelling
  
- Transformer
  - Enable **parallel computation**, hence exploit the power of GPU. How?
  - Remove the **dependency** between words
  - But **position/location** matters → resolve by **positional encoding**
  - Extremely good at capturing **long-range dependency**

# Transformer – General Architecture



Model architecture






[Source: <https://jalammar.github.io/illustrated-transformer/>]

## Some building blocks

- Positional encoding
- Self-attention and Multi-Head Self-attention Attention
- Masked Multi-Head Attention
- Residual add and layer norm


# Overview

## Architecture of Transformer Encoder




# Transformer – Self Attention

[Source: <https://jalammar.github.io/illustrated-transformer/>]



| Input     | Thinking | Machines |       |
|-----------|----------|----------|-------|
| Embedding | $x_1$    | $x_2$    |       |
| Queries   | $q_1$    | $q_2$    | $W^Q$ |
| Keys      | $k_1$    | $k_2$    | $W^K$ |
| Values    | $v_1$    | $v_2$    | $W^V$ |

Multiplying  $x_1, x_2$  with  $W^Q, W^K, W^V$  to gain **queries, keys, and values.**



- **Self attention** among a sequence of **items/tokens**
  - **Keys** and **queries** for computing self attention weights
  - $W^Q, W^K, W^V$  are **learnable matrices**.

# Transformer

## Self-Attention

- For each token  $x_i$ , calculate its **query**, **key**, and **value**

$$q_i = x_i W^Q, k_i = x_i W^K, v_i = x_i W^V$$

- Matrix/stacked** version ( $X = \begin{bmatrix} x_1 \\ \dots \\ x_L \end{bmatrix}, L = seq\_len$ )

$$Q = XW^Q, K = XW^K, V = XW^V$$


- Calculate the **attention probabilities** between query and key

- Scaled dot-product attention






$$A = softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)$$

- Take a **weighted sum** of values

$$Z = AV$$



# Self-Attention




[batch\_size, seq\_len]

# Transformer


## Multi-Head Self-Attention

- Main idea:
  - Perform self-attention multiple times in parallel and combine results
- Multiple attention heads through multiple Q, K, V matrices
- Each attention head performs attention independently
  - Allow attending to parts of the sequence differently

Scaled Dot-Product Attention




Multi-Head Attention



# Transformer – Self Attention/Multi-Head Attention

- 1) This is our input sentence\*  $X$
- 2) We embed each word\*  $R$
- 3) Split into 8 heads. We multiply  $X$  or  $R$  with weight matrices  $W_0^Q, W_0^K, W_0^V$
- 4) Calculate attention using the resulting  $Q/K/V$  matrices
- 5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^O$  to produce the output of the layer




[Source: <https://jalammar.github.io/illustrated-transformer/>]

# Transformer – Self Attention/Multi-Head Attention

## Matrix calculation for One-Head



$$\begin{aligned}
 X & \times W_Q = Q \\
 X & \times W_K = K \\
 X & \times W_V = V \\
 \text{softmax} \left( \frac{Q \times K^T}{\sqrt{d_k}} \right) & \times V = Z
 \end{aligned}$$

## Matrix calculation for Multi-Head



[Source: <https://jalammar.github.io/illustrated-transformer/>]

# Transformer – Residual Connection/Layer Norm





$$Y_m = \text{mean}(Y, \text{axis} = \text{embed\_size}, \text{keep\_dims} = \text{True})$$

$$Y_\sigma = \text{std}(Y, \text{axis} = \text{embed\_size}, \text{keep\_dims} = \text{True})$$


$$\text{LayerNorm}(Y) = \gamma \frac{Y - Y_m}{Y_\sigma} + \beta$$

- **Layer normalization (LN)** is the same as **batch normalization** except that LN **normalizes across the feature dimension**.
  - Batch normalization is usually empirically less effective than layer normalization in natural language processing tasks, whose inputs are often variable-length sequences.

# Transformer- Pointwise Feed-forward NN





# Transformer – Positional Encoding



- Use to capture **position** of a word in a sentence
  - Transformer has **no recurrence and no convolution**.
  - For the model to make use of the **order of the sequence**, some information about the **relative or absolute position** of the tokens in the sequence **needs to be injected**

$$PE_{(pos,2i)} = \sin(pos/1000^{2i/embed\_size})$$


$$PE_{(pos,2i+1)} = \cos(pos/1000^{2i/embed\_size})$$



All words or word embeddings are inputted simultaneously to the Transformer


# Transformer Decoder

Putting them together




# Transformer Decoder

Putting them together




# Transformer Decoder

## Putting them together




# Masked Self-Attention




# Transformer Decoder

## Putting them together




# Cross-Attention



# Pretrained Language Models

# Bidirectional Encoder Representations from Transformers (BERT)




- Developed based on the architecture of **Transformer**
- **Two steps** in the framework
  - **Pretrain and fine tuning**
- **Pretrain**
  - The model is trained on unlabeled data over different pre-training tasks
  - **Masked word prediction** and **next sentence prediction**
- **Fine tuning**
  - The **model for a downstream task** is first initialized with the pre-trained parameters, and all of the parameters are fine-tuned using labeled data from the downstream tasks.
  - Some downstream tasks, for example, **sequence-level classification**, predicting the **answer text span in the passage**, and **sentence-pair completion**

# BERT - Pretrain


- Task 1: Masked words prediction

- 15% of the words are masked at random
- Not all tokens were masked in the same way. Given a masked word, it happens
  - With an 80% of chance, this word is replaced by the <MASK> token
  - With a 10% of chance, this word is replaced by a random word
  - With a 10% of chance, this word is left intact
- Predict the indices of masked words on top of representations of those words.



- Task 2: Next sentence prediction


- When choosing the sentences A and B for each pretraining example, 50% of the time B is the actual next sentence following A (labeled as **IsNext**), and 50% of the time it is a random sentence from the corpus (labeled as **NotNext**).
- The final hidden vector of the special [CLS] token as  $C \in \mathbb{R}^H$  is used to predict two labels: **IsNext** and **NotNext**.



# BERT – Fine Tuning

Sequence level classification

- Input a **single sentence** to pretrain BERT.
- Observe the representations  $C, T_1, \dots, T_N \in \mathbb{R}^h$ 
  - $C \in \mathbb{R}^h$  aggregates the information of tokens/words  $1, 2, \dots, N$ , hence can be viewed as a **sentence representation**.
- On top of the **sentence representation**  $C$ , we predict sentence label
  - $W \in \mathbb{R}^{M \times h}$  is an **additional weight matrix** where  $M$  is the number of classes.
  - Prediction probabilities  $P = \text{softmax}(WC)$
- Fine tuning  $W$  and **pretrain BERT parameters**.



# BERT – Fine Tuning

## Predicting the answer text span

### Task description

- Input Question:


*Where do water droplets collide with ice crystals to form precipitation?*

- Input Paragraph:

*.... Precipitation forms as smaller droplets coalesce via collision with other raindrops or ice crystals **within a cloud**. ....*

- Output Answer:

*within a cloud*



- ❖ Represent the **input question** and **input paragraph** as a **single packed sequence**
- ❖ The **question** uses the **A embedding** and the **paragraph** uses the **B embedding**
- ❖ **New parameters** to be learned in fine-tuning are **start vector  $S \in \mathbb{R}^h$**  and **end vector  $E \in \mathbb{R}^h$**
- ❖ The probability of **word  $i$**  being **the start** of the answer span
  - ❖ 
$$P_i = \frac{\exp\{S^T T_i\}}{\sum_{j \in \text{input paragraph}} \exp\{S^T T_j\}}$$
- ❖ The probability **word  $i$**  being **the start** and **word  $j$**  being **the end** of the answer span
  - ❖ 
$$P_{i,j} = \frac{\exp\{S^T T_i + E^T T_j\}}{\sum_{(i,j) : j > i} \exp\{S^T T_{i'} + E^T T_{j'}\}}$$
- ❖ Train  $S, E$  and pretrain **BERT parameters** by maximizing log-likelihood.

# Why does pretraining-then-finetuning work?

|   | Model         | Model Size (Parameters)        | Training Data (# Tokens)                  | Notes  |
|---|---------------|--------------------------------|---|--|
|    | GPT-3 (175B)  | 175 billion                    | 300 billion tokens                        | Trained on a diverse dataset including Common Crawl, Wikipedia, and various books      |
|    | GPT-4         | Unknown (estimated 1 trillion) | Trained on hundreds of billions of tokens | Exact size not disclosed, improved multimodal capabilities, and larger context window  |
|    | BERT (Base)   | 110 million                    | 3.3 billion tokens                        | Pretrained on BooksCorpus and English Wikipedia  |
|    | BERT (Large)  | 340 million                    | 3.3 billion tokens                        | Same data as BERT Base, but with more layers and parameters                            |
|    | T5 (Base)     | 220 million                    | 1 trillion tokens                         | Trained on the Colossal Clean Crawled Corpus (C4)                                      |
|   | T5 (Large)    | 770 million                    | 1 trillion tokens                         | Same dataset as T5 Base, scaled up   |
|  | T5 (XXL)      | 11 billion                     | 1 trillion tokens                         | Same dataset as T5 Base, scaled to a massive number of parameters                      |
|  | PaLM          | 540 billion                    | 780 billion tokens                        | Trained on a multilingual dataset, including data from the web, books, Wikipedia       |
|  | LLaMA (7B)    | 7 billion                      | 1.0 trillion tokens                       | Trained on publicly available datasets and academic sources                            |
|  | LLaMA 2 (13B) | 13 billion                     | 2.0 trillion tokens                       | Enhanced version of LLaMA with more training tokens and improved training architecture |
|  | LLaMA 2 (70B) | 70 billion                     | 2.0 trillion tokens                       | Largest version of LLaMA 2 with significant improvements in token efficiency           |
|  | Chinchilla    | 70 billion                     | 1.4 trillion tokens                       | Developed by DeepMind, trained on more tokens with optimal compute scaling laws        |
|  | Gopher        | 280 billion                    | 300 billion tokens                        | Trained on diverse sources such as books, articles, and websites                       |
|  | BLOOM (176B)  | 176 billion                    | 366 billion tokens                        | Trained on multilingual data, including a variety of web sources and datasets          |
|  | Claude 2      | 70 billion                     | Unknown (hundreds of billions estimated)  | Built by Anthropic, focused on safety and interpretability                             |
|  | Mistral (7B)  | 7 billion                      | 2.0 trillion tokens                       | A dense model trained for efficiency and performance in NLP tasks                      |

1. Transformer architecture
2. Large scale training data
3. Large number of parameters

# Further Reading Recommendation

- Seq2Seq
  - <https://jalammar.github.io/>
  - <https://jalammar.github.io/visualizing-neural-machine-translation-mechanics-of-seq2seq-models-with-attention/>
  - <https://trungtran.io/2019/02/27/neural-machine-translation-with-tensorflow-model-creation/>
- Transformer
  - <https://ai.plainenglish.io/transformers-visual-guide-153af370693f>
  - <https://jalammar.github.io/illustrated-transformer/>
- BERT
  - <https://jalammar.github.io/illustrated-bert/>
  - <https://towardsdatascience.com/bert-explained-state-of-the-art-language-model-for-nlp-f8b21a9b6270>
- ViT
  - <https://arxiv.org/pdf/2010.11929.pdf>
- Swin
  - <https://arxiv.org/pdf/2103.14030.pdf>
- Dive into Deep Learning: Chapter 10

# Summary

- Encoder-Decoder models
- Sequence to sequence models
- Attention mechanism
  - Motivation
  - Global attention
  - Local attention
- Transformer and BERT

Thanks for your attention!  
Question time

