







*“... Picture says a
thousand words ...”*

Convolutional Neural Network



Arghya Pal
FIT5215



NOTE



Neocognitron 1982, Fukushima



Simple Cell (S):

Extract features, filter

Complex Cell (C):

Shift tolerance, pooling

References

1. <https://ml4a.github.io/ml4a/convnets/>
- 2.




Image Processing Zone

how to process the pixels.



Computer Vision Zone

Extract meaning from visual data.



Deep Learning Zone

Decision




Neocognitron *how to process the pixels*


Tutorial 3a

Making of Complex Cell

NOTE



A. Shift → Blur



B. Size → Pooling




References

1. <https://ml4a.github.io/ml4a/convnets/>
- 2.



Neocognitron: Extract *meaning* from visual data

猫
बिली
بلی
chat
Katze
বিড়াল






References

Neocognitron: Extract *meaning* from visual data




References

Neocognitron: Extract *meaning* from visual data





References

Neocognitron: Decision



References

From Neocognitron to Convolutional Neural Network






Image Processing Zone

how to process the pixels.



Computer Vision Zone

Extract meaning from visual data.



Deep Learning Zone

Decision





Convolutional Neural Network: *how* to process the pixels

Tutorial 3a

Instead of S and C cells we have Convolutional Operation

NOTE




What to do when it is not perfect fit?


size = 1x14

filter = 1x2

stride = 1



Convolution layer with zero padding



kernel size = (3,3)
 W' (filter or kernel)


1	1	1
1	1	1
1	1	1

$H_i = 7$ $f_w = 3$

feature map			
3	8	19	7
3	16	30	10
6	11	22	5
2	-2	8	3

$H_o = 4$ $W_o = 4$

- The **sliding window** moves from **left to right, top to bottom** with strides.
- We **convolve** the **filter** and the **sliding windows** to work out the neurons on the **feature map**.



Convolution layer with zero padding

x (input tensor (7,7))							strides = (2,2)	
0	0	0	0	0	0	0		
0	1	-2	-1	5	3	2	1	0
0	1	3	-1	4	3	3	1	0
0	1	-2	1	6	3	3	2	0
0	2	-2	2	5	2	1	0	0
0	0	3	-2	5	4	1	2	0
0	1	2	-3	1	1	2	-1	0
0	1	-2	-1	1	2	1	1	0
0	0	0	0	0	0	0	0	0

$W_i = 7$ $p = 1$ zero padding

kernel size = (3,3)
 W (filter or kernel)

1	1	1
1	1	1
1	1	1


$H_i = 7$ $f_h = 3$
 $f_w = 3$

feature map

3	8	19	7
3	15	30	10
6	11	22	5
2	-2	9	3

$H_o = 4$ $W_o = 4$


- The sliding window moves from left to right, top to bottom with strides.
- We convolve the filter and the sliding windows to work out the neurons on the feature map.



Padding


↗ no padding

“Valid”: $n \times n \times f \times f \rightarrow n-f+1 \times n-f+1$



“Same”: Pad so that output size is the same as the input size.

$$\begin{aligned} & n + 2p - f + 1 \times n + 2p - f + 1 \\ & n + 2p - f + 1 = n \Rightarrow p = \frac{f-1}{2} \end{aligned}$$



Padding

```
m = nn.CircularPad2d(2)
input = torch.arange(9,
dtype=torch.float).reshape(1, 1, 3, 3)
input
tensor([[[[0., 1., 2.],
          [3., 4., 5.],
          [6., 7., 8.]]]])
m(input)
tensor([[[[4., 5., 3., 4., 5., 3., 4.],
          [7., 8., 6., 7., 8., 6., 7.],
          [1., 2., 0., 1., 2., 0., 1.],
          [4., 5., 3., 4., 5., 3., 4.],
          [7., 8., 6., 7., 8., 6., 7.],
          [1., 2., 0., 1., 2., 0., 1.],
          [4., 5., 3., 4., 5., 3., 4.]]]])
# using different paddings for different sides
m = nn.CircularPad2d((1, 1, 2, 0))
m(input)
tensor([[[[5., 3., 4., 5., 3.],
          [8., 6., 7., 8., 6.],
          [2., 0., 1., 2., 0.],
          [5., 3., 4., 5., 3.],
          [8., 6., 7., 8., 6.]]]])
```

```
m = nn.ConstantPad2d(2, 3.5)
input = torch.randn(1, 2, 2)
input
tensor([[ 1.6585,  0.4320],
        [-0.8701, -0.4649]]])
m(input)
tensor([[[ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  1.6585,  0.4320,  3.5000,  3.5000],
        [ 3.5000,  3.5000, -0.8701, -0.4649,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000]]])
# using different paddings for different sides
m = nn.ConstantPad2d((3, 0, 2, 1), 3.5)
m(input)
tensor([[[ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  3.5000,  1.6585,  0.4320],
        [ 3.5000,  3.5000,  3.5000, -0.8701, -0.4649],
        [ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000]]])
```

padding_left, padding_right, Padding_top, padding_bottom

Padding

```
m = nn.ReflectionPad2d(2)
input = torch.arange(9,
dtype=torch.float).reshape(1, 1, 3, 3)
input
tensor([[[[0., 1., 2.],
          [3., 4., 5.],
          [6., 7., 8.]]]])]
m(input)
tensor([[[[8., 7., 6., 7., 8., 7., 6.],
          [5., 4., 3., 4., 5., 4., 3.],
          [2., 1., 0., 1., 2., 1., 0.],
          [5., 4., 3., 4., 5., 4., 3.],
          [8., 7., 6., 7., 8., 7., 6.],
          [5., 4., 3., 4., 5., 4., 3.],
          [2., 1., 0., 1., 2., 1., 0.]]]])
# using different paddings for different
sides
m = nn.ReflectionPad2d((1, 1, 2, 0))
m(input)
tensor([[[[7., 6., 7., 8., 7.],
          [4., 3., 4., 5., 4.],
          [1., 0., 1., 2., 1.],
          [4., 3., 4., 5., 4.],
          [7., 6., 7., 8., 7.]]]])
```

```
m = nn.ReplicationPad2d(2)
input = torch.arange(9,
dtype=torch.float).reshape(1, 1, 3, 3)
input
tensor([[[[0., 1., 2.],
          [3., 4., 5.],
          [6., 7., 8.]]]])]
m(input)
tensor([[[[0., 0., 0., 1., 2., 2., 2.],
          [0., 0., 0., 1., 2., 2., 2.],
          [0., 0., 0., 1., 2., 2., 2.],
          [3., 3., 3., 4., 5., 5., 5.],
          [6., 6., 6., 7., 8., 8., 8.],
          [6., 6., 6., 7., 8., 8., 8.],
          [6., 6., 6., 7., 8., 8., 8.]]]])
# using different paddings for different
sides
m = nn.ReplicationPad2d((1, 1, 2, 0))
m(input)
tensor([[[[0., 0., 1., 2., 2.],
          [0., 0., 1., 2., 2.],
          [0., 0., 1., 2., 2.],
          [3., 3., 4., 5., 5.],
          [6., 6., 7., 8., 8.]]]])
```

padding_left, padding_right, Padding_top, padding_bottom

Convolution 2D without zero padding

1	-2	-1	5	3	2	1	1
1	3	-1	4	3	3	1	-1
1	-2	1	6	3	3	2	2
2	-2	2	5	2	1	0	-1
0	3	-2	5	4	1	2	1
1	2	-3	1	1	2	-1	2
1	-2	-1	1	2	1	1	3

$$x = \text{tensor}(7, 8) \quad W_i = 8$$

strides $s = (s_w, s_h) = (2, 2)$



- W_i, H_i : The width and height of the input image
- W_o, H_o : The width and height of the output image (feature map)

$$W_o = \left\lceil \frac{W_i + 2p - f_w}{s_w} \right\rceil + 1 \text{ and } H_o = \left\lceil \frac{H_i + 2p - f_h}{s_h} \right\rceil + 1$$


$$\text{□ Our case: } W_o = \left\lceil \frac{8+0-3}{2} \right\rceil + 1 = 3 \text{ and } H_o = \left\lceil \frac{7+0-3}{2} \right\rceil + 1 = 3$$

Convolution 2D with zero padding



0	0	0	0	0	0	0	0	0
0	1	-2	-1	5	3	2	1	0
0	1	3	-1	4	3	3	1	0
0	1	-2	1	6	3	3	2	0
0	2	-2	2	5	2	1	0	0
0	0	3	-2	5	4	1	2	0
0	1	2	-3	1	1	2	-1	0
0	1	-2	-1	1	2	1	1	0
0	0	0	0	0	0	0	0	0

$x = \text{tensor}(7,7)$
strides $s = (s_w, s_h) = (2,2)$
zero-padding $p = 1$

padding = 1



} padding = 1



- W_i, H_i : The width and height of the input image
- W_o, H_o : The width and height of the output image (feature map)

$$W_o = \left\lfloor \frac{W_i + 2p - f_w}{s_w} \right\rfloor + 1 \text{ and } H_o = \left\lfloor \frac{H_i + 2p - f_h}{s_h} \right\rfloor + 1$$

$$\text{Our case: } W_o = \left\lfloor \frac{7+2 \times 1 - 3}{2} \right\rfloor + 1 = 4 \text{ and } H_o = \left\lfloor \frac{7+2 \times 1 - 3}{2} \right\rfloor + 1 = 4$$





Image Processing Zone

how to process the pixels.




Computer Vision Zone

Extract meaning from visual data.



Deep Learning Zone


Decision



Effects of filters/kernels to images




```
smallBlur = np.ones((7, 7), dtype="float") * (1.0 / (7 * 7))
```




```
largeBlur = np.ones((21, 21), dtype="float") * (1.0 / (21 * 21))
```





```
# construct a sharpening filter  
sharpen = np.array([[0, -1, 0],  
                   [-1, 5, -1],  
                   [0, -1, 0]], dtype="int")
```



```
# construct the Laplacian kernel used to detect edge-like  
# regions of an image  
laplacian = np.array([[0, 1, 0],  
                     [1, -4, 1],  
                     [0, 1, 0]], dtype="int")
```



```
# construct an emboss kernel  
emboss = np.array([[1, 2, 1, 0, 0],  
                  [0, 1, 0, 0, 0],  
                  [-1, -2, -1, 0, 0],  
                  [0, 0, 0, 0, 0],  
                  [0, 0, 0, 0, 0]], dtype="int")
```



Example of pooling with PyTorch

Max pooling

```
output = torch.nn.functional.max_pool2d(input = batch_tensor,
                                         kernel_size=(2,2), stride= (2,2))
output = output.numpy()
print(batch[0].shape)
print(output[0].shape)
plot_image(batch[0].transpose((1,2,0)).astype(np.int32), scale=True) # plot the 1st original image
plot_image(output[0].transpose((1,2,0)).astype(np.int32), scale=True) # plot the output for the 1st image
(3, 280, 320)
(3, 140, 160)
```



Average pooling

```
output = torch.nn.functional.avg_pool2d(input = batch_tensor,
                                         kernel_size=(2,2), stride= (2,2))
output = output.numpy()
print(batch[0].shape)
print(output[0].shape)
plot_image(batch[0].transpose((1,2,0)).astype(np.int32), scale=True) # plot the 1st original image
plot_image(output[0].transpose((1,2,0)).astype(np.int32), scale=True) # plot the output for the 1st image
(3, 280, 320)
(3, 140, 160)
```






Image Processing Zone

how to process the pixels.



Computer Vision Zone


Extract meaning from visual data.




Deep Learning Zone

Decision





CNN in Operation



$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$


```
torch.nn.Conv2d(  
    in_channels,  
    out_channels,  
    kernel_size,  
    stride=1,  
    padding=0,  
    dilation=1,  
    groups=1,  
    bias=True,  
    padding_mode='zeros',  
    device=None,  
    dtype=None  
)
```

$$\begin{bmatrix} 1 & 2 & 0 \\ 0 & 1 & 3 \\ 2 & 1 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 1 & 2 \\ 1 & 0 & 1 \\ 3 & 1 & 2 \end{bmatrix}$$

Tricks


NOTE




Normalization

Standardization

Collapse inputs to be between 0 and 1.



Make mean 0 and variance 1.





References

1. <https://www.youtube.com/watch?v=yXOMHOpb0n8>

Batch Normalization

1. Cope with internal covariate shift
2. Reduce gradient vanishing/exploding
3. Reduce overfitting
4. Make training more stable
5. Converge faster
 1. Allow us to train with bigger learning rate


- Let $z = W^k h^k + b^k$ be the mini-batch before activation. We compute the normalized \hat{z} as
 - $\hat{z} = \frac{z - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$ where ϵ is a small value such as $1e^{-7}$
 - $\mu_B = \frac{1}{b} \sum_{i=1}^b z_i$ is the empirical mean
 - $\sigma_B^2 = \frac{1}{b} \sum_{i=1}^b (z_i - \mu_B)^2$ is the empirical variance
- We scale the normalized \hat{z}
 - $z_{BN} = \gamma \hat{z} + \beta$ where $\gamma, \beta > 0$ are two learnable parameters (i.e., scale and shift parameters)
- We then apply the activation to obtain the next layer value
 - $h^{k+1} = \sigma(z_{BN})$



Dropout



Reduce Overfitting

Without dropout




(Source: Analytics Vidhya)

With dropout




- This is a **cheap technique** to reduce model capacity
 - Reduce overfitting
- In each iteration, at each layer, **randomly choose** some neurons and **drop all connections** from these neurons
 - `dropout_rate = 1 - keep_prob`



Network Architectures




AlexNet



VGG16



CONV block



IDENTITY block

ResNet

MiniVGG for Cifar10

Our Tutorial

Layer Type	Output Size	Filter Size / Stride
INPUT IMAGE	$32 \times 32 \times 3$	
CONV	$32 \times 32 \times 32$	$3 \times 3, K = 32$
ACT	$32 \times 32 \times 32$	
BN	$32 \times 32 \times 32$	
CONV	$32 \times 32 \times 32$	$3 \times 3, K = 32$
ACT	$32 \times 32 \times 32$	
BN	$32 \times 32 \times 32$	
POOL	$16 \times 16 \times 32$	2×2
DROPOUT	$16 \times 16 \times 32$	
CONV	$16 \times 16 \times 64$	$3 \times 3, K = 64$
ACT	$16 \times 16 \times 64$	
BN	$16 \times 16 \times 64$	
CONV	$16 \times 16 \times 64$	$3 \times 3, K = 64$
ACT	$16 \times 16 \times 64$	
BN	$16 \times 16 \times 64$	
POOL	$8 \times 8 \times 64$	2×2
DROPOUT	$8 \times 8 \times 64$	
FC	512	
ACT	512	
BN	512	
DROPOUT	512	
FC	10	
SOFTMAX	10	

```
def create_vgg():
    vgg_model = nn.Sequential(
        #nn.LazyConv2d(32, kernel_size=3, padding=1),
        nn.Conv2d(3, 32, kernel_size=3, padding= 1), #[32,32,32]
        nn.BatchNorm2d(32),
        nn.ReLU(),
        #nn.LazyConv2d(32, kernel_size=3, padding=1),
        nn.Conv2d(32, 32, kernel_size=3, padding=1), #[32,32,32]
        nn.BatchNorm2d(32),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2), #down-sample by two #[32,16,16]
        nn.Dropout(p=0.25),

        #nn.LazyConv2d(64, kernel_size=3, padding=1),
        nn.Conv2d(32, 64 , kernel_size=3, padding=1), #[64,16,16]
        nn.BatchNorm2d(64, momentum=0.1),
        nn.ReLU(),
        #nn.LazyConv2d(64, kernel_size=3, padding=1)
        nn.Conv2d(64, 64, kernel_size=3, padding=1), #[64,16,16]
        nn.BatchNorm2d(64, momentum=0.1),
        nn.ReLU(),
        nn.LazyConv2d(64, kernel_size=3, padding=1),
        nn.BatchNorm2d(64, momentum=0.1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2), #down-sample by two [64,8,8]
        nn.Dropout(p=0.25),

        nn.Flatten(1), #64x8x8
        #nn.Linear(64x8x8, 512)
        nn.LazyLinear(512),
        nn.ReLU(),
        #nn.LazyLinear(10)
        nn.Linear(512, 10),
    )
    return vgg_model
```



60,000 images
10 classes

