

# Deep Learning for Sequential Data (II)

## Seq2Seq and Transformers

### Learning Outcomes

By the end of this week, you'll be able to:

- Understand Seq2Seq and how they might be applied to solve problems such as machine translation, language modeling, and image captioning.
- Understand local and global attention mechanisms in deep models.
- Understand Transformer and BERT, two latest SOTA models for NLP.
- Understand the self-attention mechanism and how it is used in Transformer.

# Use Case

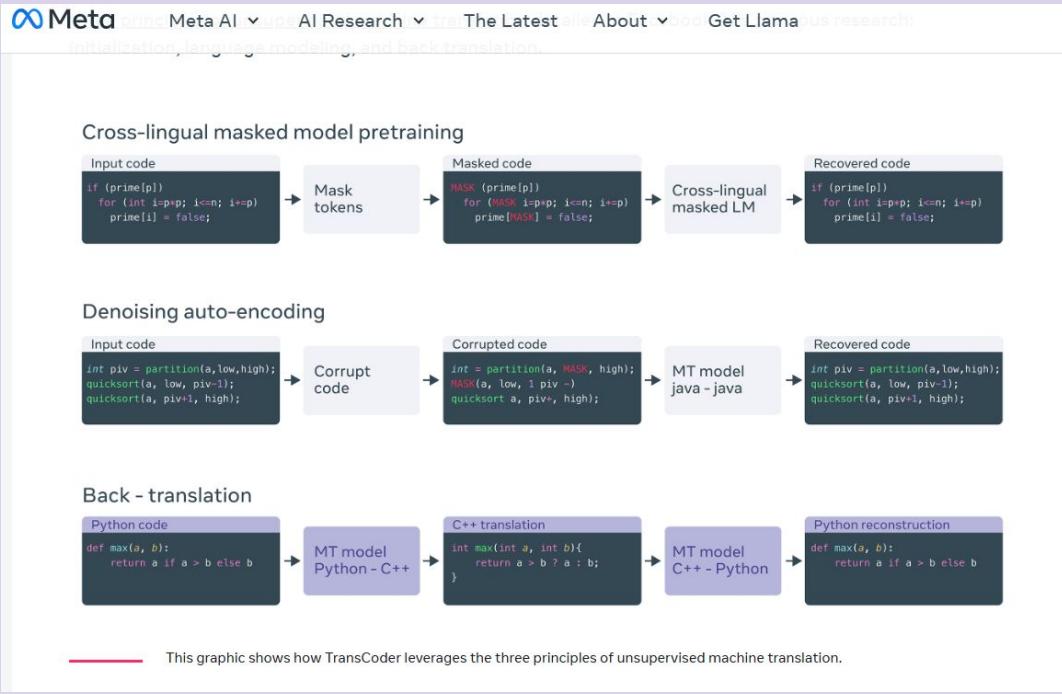


Screenshot of the Viber Translate interface. The top bar shows "ENGLISH" and "RUSSIAN". The main area has a placeholder "Start typing or insert a link". To the right, there is a "Translating photo" section with a "Drag file here or select" button and a "ctrl + v" keyboard shortcut. The bottom navigation bar includes links for "Ask about English", "Explain the Present Perfect", and "When do I use 'the' before nouns?".

Screenshot of the DeepL AI Labs interface. The top bar shows "DeepL AI Labs" and a "Start free trial" button. The main area features tabs for "Translate text" (selected), "Translate files", and "DeepL Write". It includes a "Detect language" dropdown set to "Chinese (simplified)". The right sidebar contains sections for "Editing tools", "Customizations", and "Powered by".

Screenshot of the Google Translate interface. The top bar shows "Text", "Images", "Documents", and "Websites". The main area has a "Detect language" dropdown set to "English" and a "From" dropdown set to "English". The right sidebar lists "Spanish", "French", "Arabic", and "Translation". At the bottom right is a "Send feedback" button.

# Use Case



Python input

```

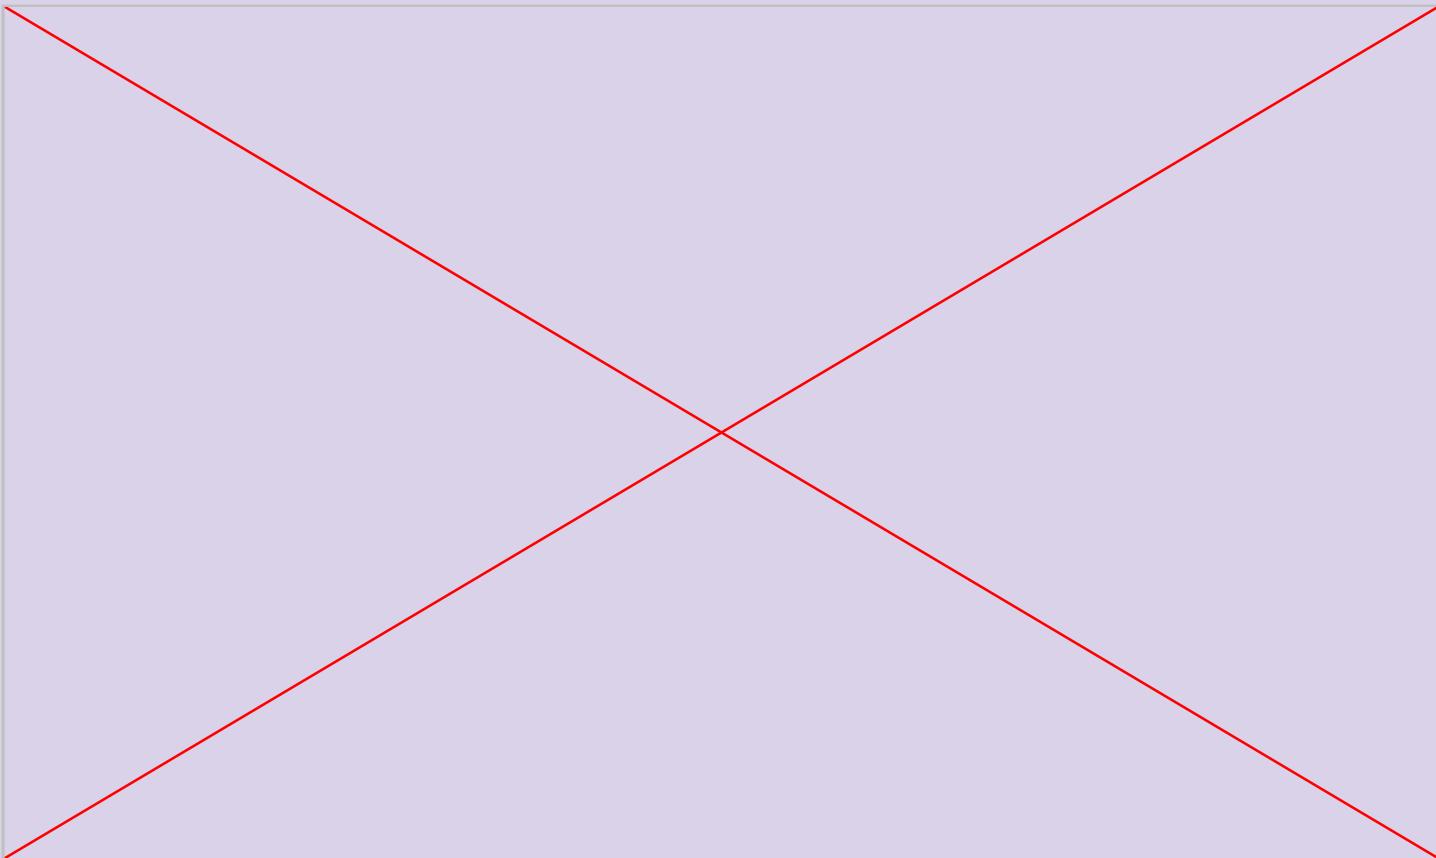
def sumOfksubArray(arr, n, k):
    Sum = 0
    S = deque()
    G = deque()
    for i in range(k):
        while len(S) > 0 and arr[S[-1]] >= arr[i]:
            S.pop()
        while len(G) > 0 and arr[G[-1]] <= arr[i]:
            G.pop()
        G.append(i)
        S.append(i)
    for i in range(k, n):
        Sum += arr[S[0]] + arr[G[0]]
        while len(S) > 0 and S[0] <= i - k:
            S.popleft()
        while len(G) > 0 and G[0] <= i - k:
            G.popleft()
        while len(S) > 0 and arr[S[-1]] >= arr[i]:
            S.pop()
        while len(G) > 0 and arr[G[-1]] <= arr[i]:
            G.pop()
        G.append(i)
        S.append(i)
    Sum += arr[S[0]] + arr[G[0]]
    return Sum
  
```

C++ unsupervised translation

```

int sumOfksubArray(int arr[], int n, int k){
    int Sum = 0;
    deque<int> S;
    deque<int> G;
    for(int i = 0; i < k; i++){
        while((int)S.size() > 0 && arr[S.back()] >= arr[i])
            S.pop_back();
        while((int)G.size() > 0 && arr[G.back()] <= arr[i])
            G.pop_back();
        G.push_back(i);
        S.push_back(i);
    }
    for(int i = k; i < n; i++){
        Sum += arr[S.front()] + arr[G.front()];
        while((int)S.size() > 0 && S.front() <= i - k)
            S.pop_front();
        while((int)G.size() > 0 && G.front() <= i - k)
            G.pop_front();
        while((int)S.size() > 0 && arr[S.back()] >= arr[i])
            S.pop_back();
        while((int)G.size() > 0 && arr[G.back()] <= arr[i])
            G.pop_back();
        G.push_back(i);
        S.push_back(i);
    }
    Sum += arr[S.front()] + arr[G.front()];
    return Sum;
}
  
```

# Sequence and Sequence



# Sequence and Sequence

Source sequence:

Я видел котю на мате <eos>  
"I" "saw" "cat" "on" "mat"

Target sequence:

I saw a **cat** on a mat <eos>

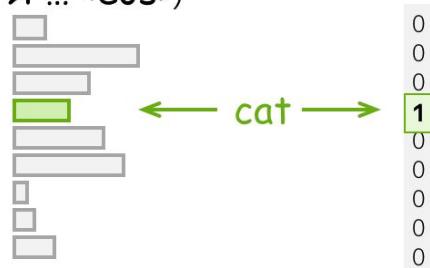
← one training example

← one step for this example

previous tokens

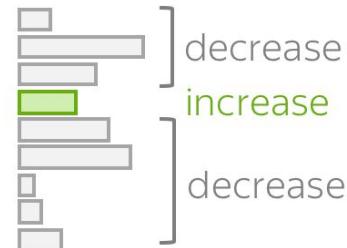
we want the model  
to predict this

Model prediction:  $p(* | I \text{ saw } a,$   
 $\text{я} \dots \text{<eos>})$



Target

$$\text{Loss} = -\log(p(\text{cat})) \rightarrow \min$$



decrease  
increase  
decrease

## Encoder

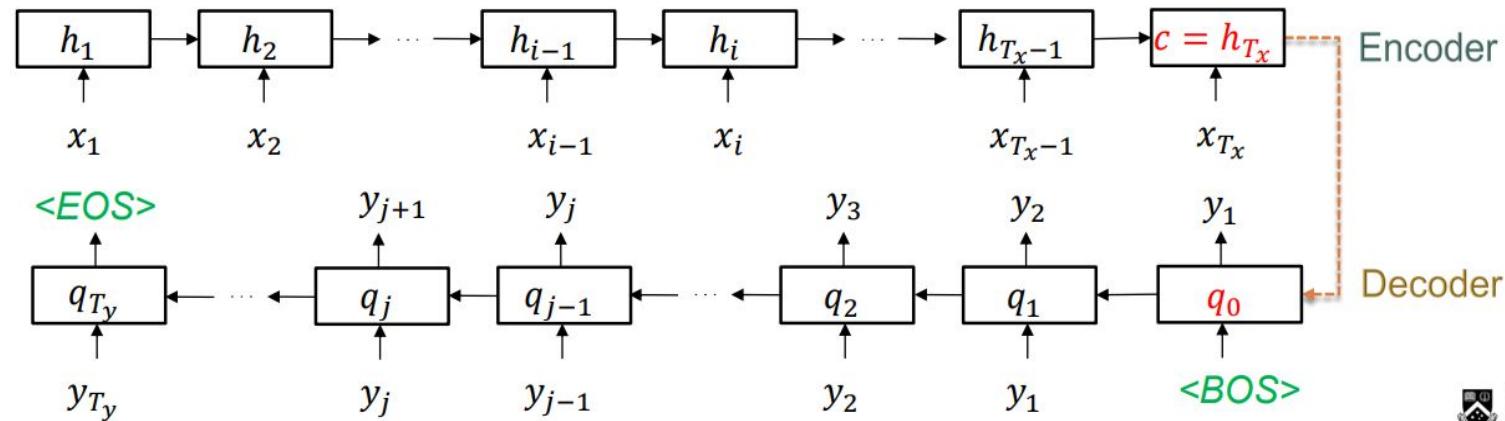
- Produces the context vector  $\mathbf{c} = \mathbf{h}_{T_x}$  of the input sequence
- Context vector  $\mathbf{c}$  summarizes input sequence  $[\mathbf{x}_1, \dots, \mathbf{x}_{T_x}]$ .

## Decoder

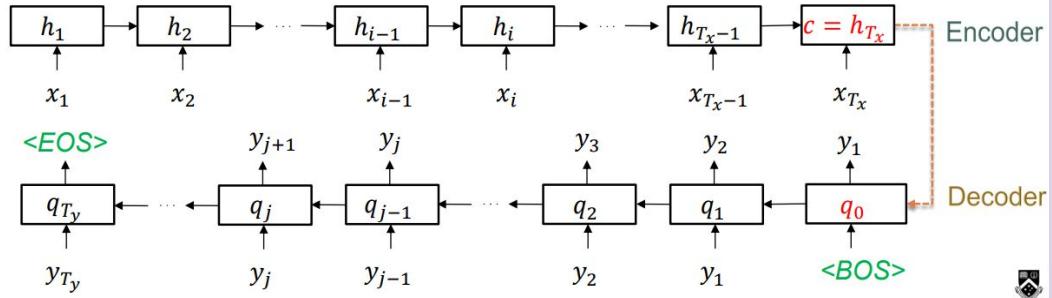
- Decodes the encode  $\mathbf{c}$  to the output sequence

## Special symbols

- **<EOS>** signifies the end of a sequence
- **<BOS>** signifies the beginning of a sequence



# Step 1: Dataset Preparation



```
dataset = datasets.load_dataset ("bentrevett/multi30k")
```

```
{'en': 'Two young, White males are outside near many bushes.'}
```

```
'de': 'Zwei junge weiße Männer sind im Freien in der Nähe vieler Büsche.'}
```

## Statistics

### train

```
(en) 29000 sentences, 377534 words, 13.0 words/sent  
(de) 29000 sentences, 360706 words, 12.4 words/sent  
(fr) 29000 sentences, 409845 words, 14.1 words/sent  
(cs) 29000 sentences, 297212 words, 10.2 words/sent
```

### val

```
(en) 1014 sentences, 13308 words, 13.1 words/sent  
(de) 1014 sentences, 12828 words, 12.7 words/sent  
(fr) 1014 sentences, 14381 words, 14.2 words/sent  
(cs) 1014 sentences, 10342 words, 10.2 words/sent
```

### test\_2016\_flickr

```
(en) 1000 sentences, 12968 words, 13.0 words/sent  
(de) 1000 sentences, 12103 words, 12.1 words/sent  
(fr) 1000 sentences, 13988 words, 14.0 words/sent  
(cs) 1000 sentences, 10497 words, 10.5 words/sent
```

### test\_2017\_flickr

```
(en) 1000 sentences, 11376 words, 11.4 words/sent  
(de) 1000 sentences, 10758 words, 10.8 words/sent  
(fr) 1000 sentences, 12596 words, 12.6 words/sent
```

### test\_2017\_mscoco

```
(en) 461 sentences, 5239 words, 11.4 words/sent  
(de) 461 sentences, 5158 words, 11.2 words/sent  
(fr) 461 sentences, 5710 words, 12.4 words/sent
```

# Step 2: Tokenization

## Text Analytics and Language Models

### Text normalization

#### Expanding contractions

- Contractions are shortened version of words or syllables
  - e.g., isn't → is not, you're → you are
  - Exist extensively and pose a problem to text analytics

#### Lemmatization

- removing word **affixes** to get to a **base form** of the **root word**.
  - e.g. cars → car, running → run, is → be

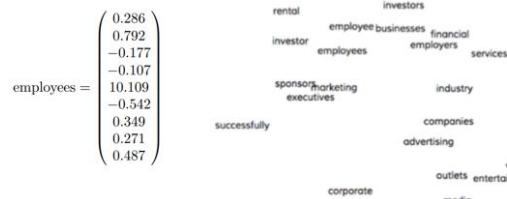
#### Removing special characters and symbols

- e.g. !, .

#### Removing stop words

- e.g., a, and

### Distributed representation of words

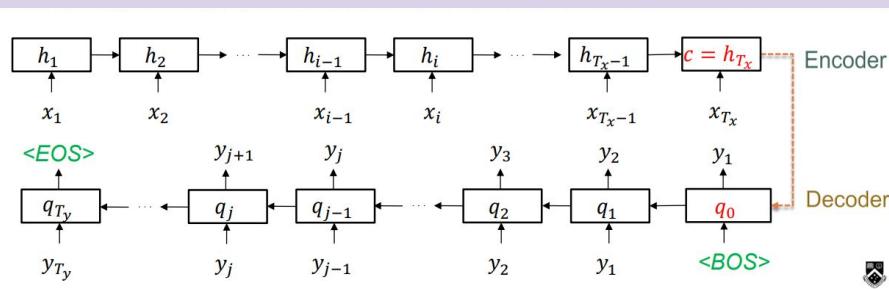


```
import spacy
```

1. A tokenizer is used to turn a string into a list of tokens

2. "good morning!" becomes ["good", "morning", "!"].

# Step 2: Tokenization

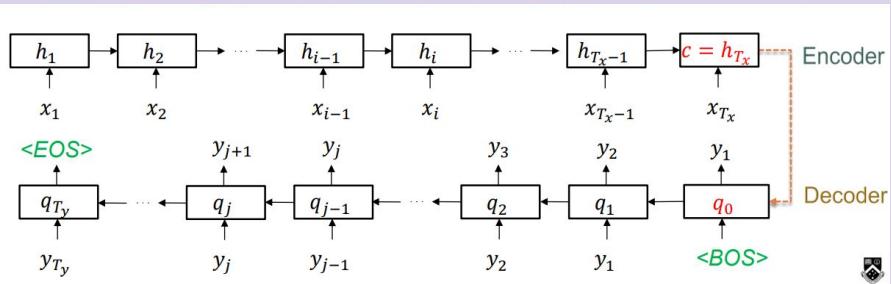


```
import spacy
```

```
1 en_nlp = spacy.load("en_core_web_sm")
2 de_nlp = spacy.load("de_core_news_sm")
```

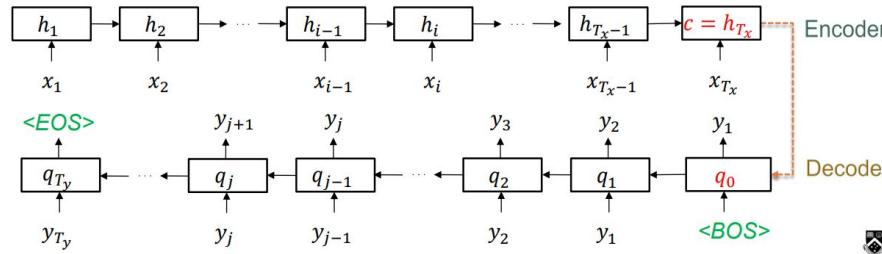
```
1 string = "What a lovely day it is today!"
2
3 [token.text for token in en_nlp.tokenizer(string)]
['What', 'a', 'lovely', 'day', 'it', 'is', 'today', '!']
```

# Step 2: Tokenization



```
1 max_length = 1_000
2 lower = True
3 sos_token = "<sos>"
4 eos_token = "<eos>"
5
6 fn_kwargs = {
7     "en_nlp": en_nlp,
8     "de_nlp": de_nlp,
9     "max_length": max_length,
10    "lower": lower,
11    "sos_token": sos_token,
12    "eos_token": eos_token,
13 }
14
15 train_data = train_data.map(tokenize_example, fn_kwarg=fn_kwarg)
16 valid_data = valid_data.map(tokenize_example, fn_kwarg=fn_kwarg)
17 test_data = test_data.map(tokenize_example, fn_kwarg=fn_kwarg)
```

# Step 2: Tokenization



```
{'en': 'Two young, White males are outside near many bushes.',  
 'de': 'Zwei junge weiße Männer sind im Freien in der Nähe vieler Büsche.',  
 'en_tokens': ['<sos>',  
   'two',  
   'young',  
   ',',  
   'white',  
   'males',  
   'are',  
   'outside',  
   'near',  
   'many',  
   'bushes',  
   '.',  
   '<eos>'],  
 'de_tokens': ['<sos>',  
   'zwei',  
   'junge',  
   'weiße',  
   'männer',  
   'sind',  
   'im',  
   'freien',  
   'in',  
   'der',  
   'nähe',  
   'vieler',  
   'büsche',  
   '.',  
   '<eos>']]}
```

# Step 2.1 : Vocabulary

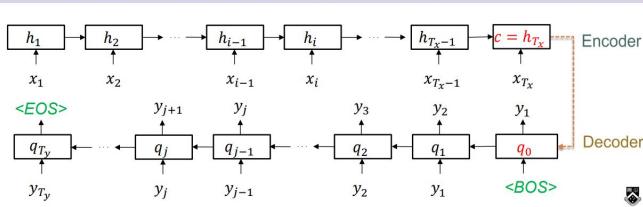
```
{'en': 'Two young, White males are outside near many bushes.',  
 'de': 'Zwei junge weiße Männer sind im Freien in der Nähe vieler Büsche.',  
 'en_tokens': ['<sos>',  
 'two',  
 'young',  
 '.',  
 'white',  
 'males',  
 'are',  
 'outside',  
 'near',  
 'many',  
 'bushes',  
 '.',  
 '<eos>'],  
 'de_tokens': ['<sos>',  
 'zwei',  
 'junge',  
 'weiße',  
 'männer',  
 'sind',  
 'im',  
 'freien',  
 'in',  
 'der',  
 'nähe',  
 'vieler',  
 'büsche',  
 '.',  
 '<eos>']]}
```

## Vocabularies

Next, we'll build the *vocabulary* for the source and target languages. The vocabulary is used to associate each unique token in our dataset with an index (an integer), e.g. "hello" = 1, "world" = 2, "bye" = 3, "hates" = 4, etc

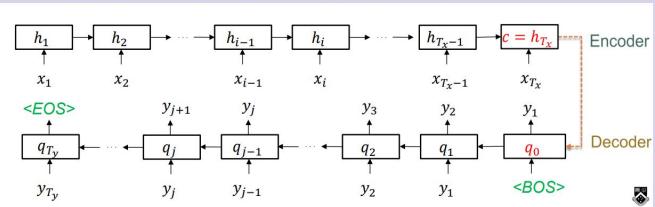
```
1 min_freq = 2  
2 unk_token = "<unk>"  
3 pad_token = "<pad>"  
4  
5 special_tokens = [  
6     unk_token,  
7     pad_token,  
8     sos_token,  
9     eos_token,  
10 ]  
11  
12 en_vocab = torchtext.vocab.build_vocab_from_iterator(  
13     train_data["en_tokens"],  
14     min_freq=min_freq,  
15     specials=special_tokens,  
16 )  
17  
18 de_vocab = torchtext.vocab.build_vocab_from_iterator(  
19     train_data["de_tokens"],  
20     min_freq=min_freq,  
21     specials=special_tokens,  
22 )
```

## Step 2.2: Encoder



```
1 class Encoder(nn.Module):
2     def __init__(self, input_dim, embedding_dim, hidden_dim, n_layers, dropout):
3         super().__init__()
4         self.hidden_dim = hidden_dim
5         self.n_layers = n_layers
6         self.embedding = nn.Embedding(input_dim, embedding_dim)
7         self.rnn = nn.LSTM(embedding_dim, hidden_dim, n_layers, dropout=dropout)
8         self.dropout = nn.Dropout(dropout)
9
10    def forward(self, src):
11        # src = [src length, batch size]
12        embedded = self.dropout(self.embedding(src))
13        # embedded = [src length, batch size, embedding dim]
14        outputs, (hidden, cell) = self.rnn(embedded)
15        # outputs = [src length, batch size, hidden dim * n directions]
16        # hidden = [n layers * n directions, batch size, hidden dim]
17        # cell = [n layers * n directions, batch size, hidden dim]
18        # outputs are always from the top hidden layer
19        return hidden, cell
```

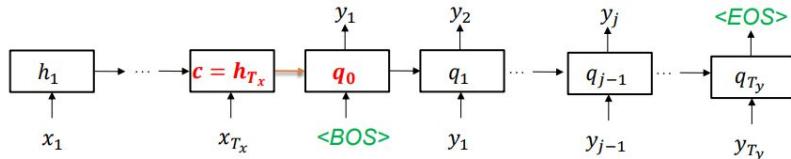
# Step 2.4: Decoder



```
1 class Decoder(nn.Module):
2     def __init__(self, output_dim, embedding_dim, hidden_dim, n_layers, dropout):
3         super().__init__()
4         self.output_dim = output_dim
5         self.hidden_dim = hidden_dim
6         self.n_layers = n_layers
7         self.embedding = nn.Embedding(output_dim, embedding_dim)
8         self.rnn = nn.LSTM(embedding_dim, hidden_dim, n_layers, dropout=dropout)
9         self.fc_out = nn.Linear(hidden_dim, output_dim)
10        self.dropout = nn.Dropout(dropout)
11
12    def forward(self, input, hidden, cell):
13        # input = [batch size]
14        # hidden = [n layers * n directions, batch size, hidden dim]
15        # cell = [n layers * n directions, batch size, hidden dim]
16        # n directions in the decoder will both always be 1, therefore:
17        # hidden = [n layers, batch size, hidden dim]
18        # context = [n layers, batch size, hidden dim]
19        input = input.unsqueeze(0)
20        # input = [1, batch size]
21        embedded = self.dropout(self.embedding(input))
22        # embedded = [1, batch size, embedding dim]
23        output, (hidden, cell) = self.rnn(embedded, (hidden, cell))
24        # output = [seq length, batch size, hidden dim * n directions]
25        # hidden = [n layers * n directions, batch size, hidden dim]
26        # cell = [n layers * n directions, batch size, hidden dim]
27        # seq length and n directions will always be 1 in this decoder, therefore:
28        # output = [1, batch size, hidden dim]
29        # hidden = [n layers, batch size, hidden dim]
30        # cell = [n layers, batch size, hidden dim]
31        prediction = self.fc_out(output.squeeze(0))
32        # prediction = [batch size, output dim]
33        return prediction, hidden, cell
```

# Step 3: Training

## Training of seq2seq



- We need to maximize the log-likelihood:

$$\max_{\theta} J(\theta) = \sum_{(x,y) \in \mathcal{D}} \log P(y|x, \theta)$$

where  $\theta = [\theta_e, \theta_d]$  and  $\theta_e, \theta_d$  are encoding and decoding parameters respectively.

- Product rule:

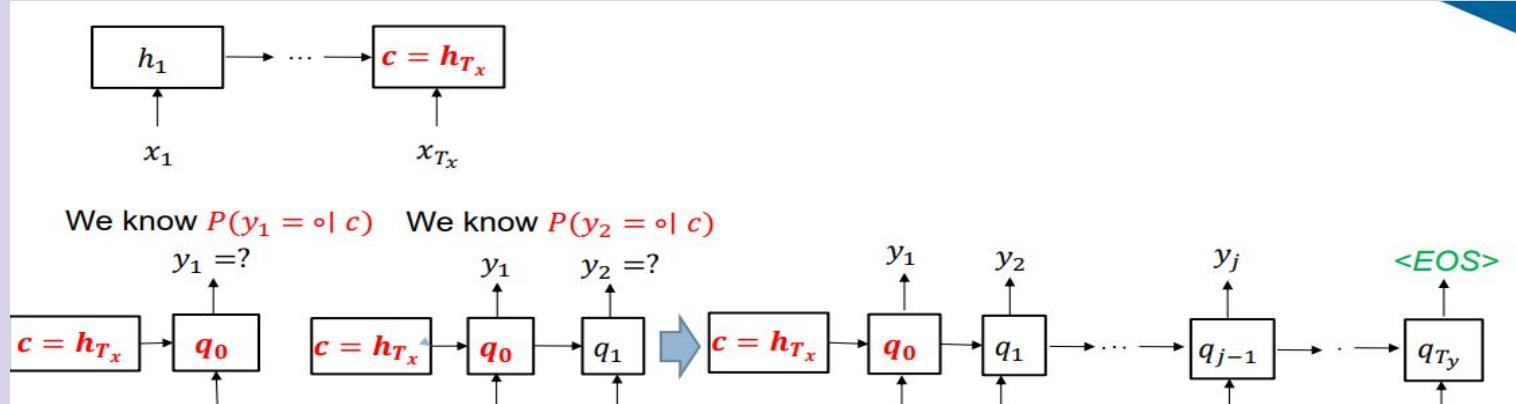
$$\begin{aligned} P(y|\mathbf{x}, \theta) &= P(y_{1:T_y} | \mathbf{x}_{1:T_x}, \theta) = P(y_{1:T_y} | \mathbf{c}, \theta) \\ &= P(y_1 | \mathbf{c}, \theta) P(y_2 | y_1, \mathbf{c}, \theta) \dots P(y_j | y_{1:j-1}, \mathbf{c}, \theta) \dots P(y_{T_y} | y_{1:T_y-1}, \mathbf{c}, \theta) = \prod_{j=1}^{T_y} P(y_j | y_{1:j-1}, \mathbf{c}, \theta) \\ \log P(y|\mathbf{x}, \theta) &= \log P(y|\mathbf{c}, \theta) = \sum_{j=1}^{T_y} \log P(y_j | y_{1:j-1}, \mathbf{c}, \theta) = \sum_{j=1}^{T_y} \log P(y_j | \mathbf{q}_{j-1}, \mathbf{c}, \theta) \end{aligned}$$

- We can compute  $P(y_j | \mathbf{q}_{j-1}, \mathbf{c}) = g(y_j, \mathbf{q}_{j-1}, \mathbf{c})$  where  $g$  is a nonlinear, potentially multi-layered NN that outputs the probability of  $y_j$ .

- Pay attention on how  $\mathbf{c}$  is used in every step during decoding

```
1 input_dim = len(de_vocab)
2 output_dim = len(en_vocab)
3 encoder_embedding_dim = 256
4 decoder_embedding_dim = 256
5 hidden_dim = 512
6 n_layers = 2
7 encoder_dropout = 0.5
8 decoder_dropout = 0.5
9 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
10
11 encoder = Encoder(
12     input_dim,
13     encoder_embedding_dim,
14     hidden_dim,
15     n_layers,
16     encoder_dropout,
17 )
18
19 decoder = Decoder(
20     output_dim,
21     decoder_embedding_dim,
22     hidden_dim,
23     n_layers,
24     decoder_dropout,
25 )
26
27 model = Seq2Seq(encoder, decoder, device).to(device)
```

# Step 4: Inference



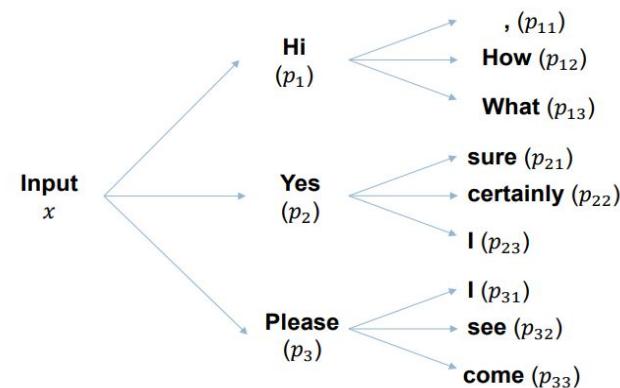
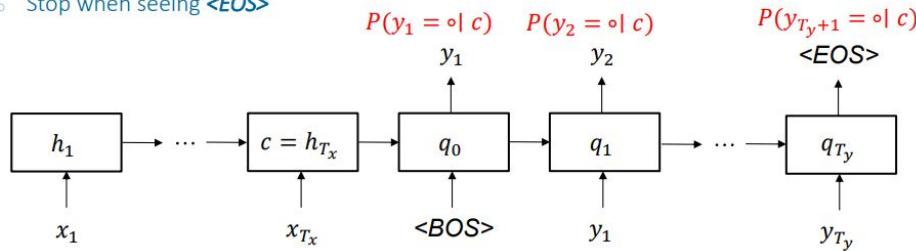
```
1 model.load_state_dict(torch.load("tut1-model.pt"))
2
3 test_loss = evaluate_fn(model, test_data_loader, criterion, device)
4
5 print(f" | Test Loss: {test_loss:.3f} | Test PPL: {np.exp(test_loss):.3f} |")
```

Test Loss: 3.780 | Test PPL: 43.833 |

# Step 4.1: Inference; Greedy Vs Beam

## Greedy Decoding

- Given  $x$ , find word  $y_1$  with highest probability
- Given  $y_1$  and  $x$ , find word  $y_2$  with highest probability
- ...
- Stop when seeing  $\langle EOS \rangle$



**Beam width = 3**  
- We always choose three sentences with highest joint probabilities

# Drawback of fixed context

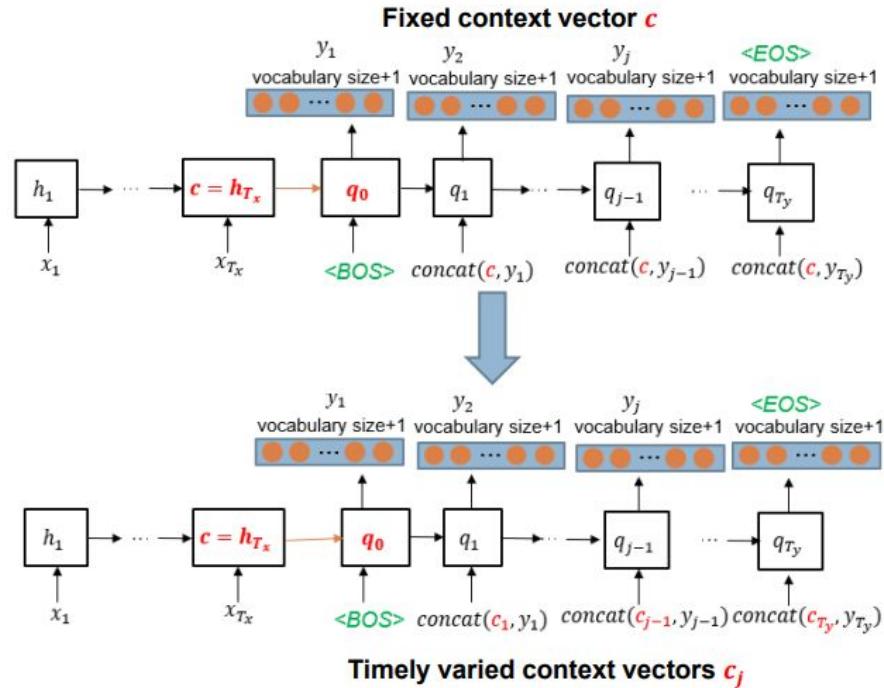
- Fixed context vector  $c$  is easily overwhelmed by long inputs or long outputs.
- At a specific timestep  $j$ , some words or items in the input sequence might possibly contribute more to the generation of next item or word in the output sequence.
  - I want to see you every day → Je veux te voir chaque jour
  - I want to see you every day → Je veux te ? (voir) ....
- How to timely adapt the context vector  $c_j$ ?
  - $c_j = \alpha(h_1, \dots, h_{T_x}, q_{j-1})$
  - Computed using attention mechanism

## NEURAL MACHINE TRANSLATION BY JOINTLY LEARNING TO ALIGN AND TRANSLATE

Dzmitry Bahdanau  
Jacobs University Bremen, Germany

KyungHyun Cho   Yoshua Bengio<sup>\*</sup>  
Université de Montréal

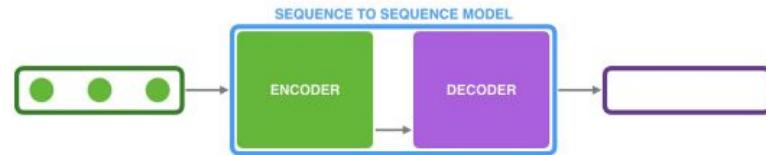
Paper: [paper link](#)



# Attention is all you need

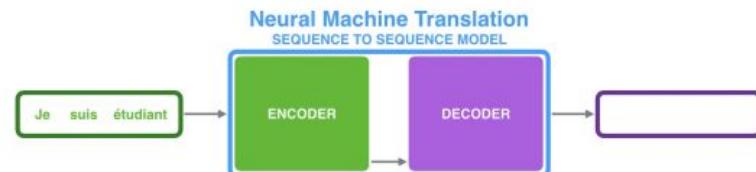
# Attention mechanism

- So far, the input sequence is **summarised** by a **single** context **c** !

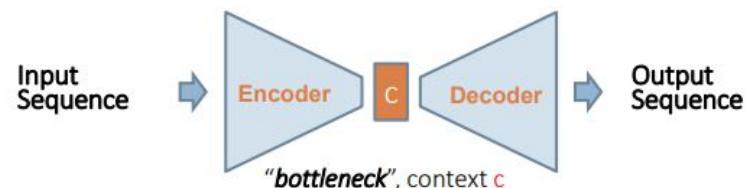


- Fixed-length context** could be problematic as it is **easily overwhelmed** by long inputs or long outputs

- The **fixed-length** context **c** might **not be powerful enough** to capture long input sequences



- Some **specific items** in an input sequence might be **more relevant** and **contributing in generating** a given item in output sequence.



- Gratefully acknowledge the excellent visualizations used from Jay Alammar blog at:

- <https://jalammar.github.io/>
- <https://jalammar.github.io/visualizing-neural-machine-translation-mechanics-of-seq2seq-models-with-attention/>

# Attention mechanism

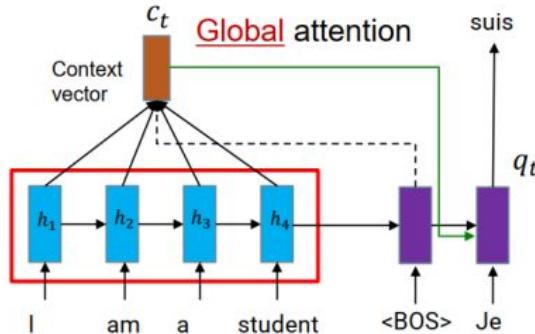
Published as a conference paper at ICLR 2015

## NEURAL MACHINE TRANSLATION BY JOINTLY LEARNING TO ALIGN AND TRANSLATE

Dzmitry Bahdanau  
Jacobs University Bremen, Germany

KyungHyun Cho Yoshua Bengio<sup>\*</sup>  
Université de Montréal

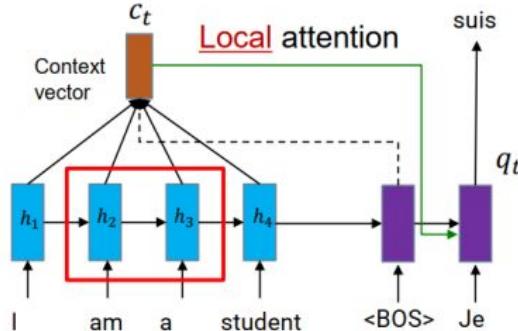
Bahdanau, Cho, Bengio, [Neural Machine Translation by Jointly Learning to Align and Translate](#), ICLR 2015



## Effective Approaches to Attention-based Neural Machine Translation

Minh-Thang Luong Hieu Pham Christopher D. Manning  
Computer Science Department, Stanford University, Stanford, CA 94305  
[{lmthang,hyhieu,manning}@stanford.edu](mailto:{lmthang,hyhieu,manning}@stanford.edu)

Luong, Pham, Manning, [Effective Approach Attention-based Neural Machine Translation](#), EMNLP, 2015



- Attention mechanism allows the decoding network to refer to the input.

- Global attention

- Use all input hidden states of the encoder when deriving the context  $c_t$ .

- Local attention

- Use a selective window of input hidden states of the encoder when deriving the context  $c_t$ .