

FIT3181/5215 Deep Learning

Generative Adversarial Networks and Diffusion Models

Trung Le and Teaching team

Department of Data Science and AI
Faculty of Information Technology, Monash University
Email: trunglm@monash.edu

Outline

❑ Tutorial 11a

- Overview of GANs (*****)
- Implementation of GANs (*****)
- Implementation of WGAN with gradient penalty (***)

❑ Tutorial 11b

- Diffusion models (***)

Overview of GANs

□ Min-max game between two players


- $\min_G \max_D J(G, D) = \mathbb{E}_{x \sim p_d(x)} [\log D(x)] + \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$
- The discriminator D tries to distinguish real and fake samples, while the generator G tries to fool D by generating fake samples indistinguishable from real samples
- $D(x) \in [0,1]$ is the probability that x is a real sample

□ Issues and challenges

- Mode collapsing problems, generating unrealistic images for complex dataset, no single loss to train

□ WS GAN with penalty

- $\min_G \max_D (\mathbb{E}_{x \sim p_d}[D(x)] - \mathbb{E}_{z \sim p_z}[D(G(z))] - \alpha \mathbb{E}_{\hat{x} \sim \hat{p}} [(\|\nabla_{\hat{x}} D(\hat{x})\| - 1)^2])$
- $D(x) \in \mathbb{R}$ where D is called 'critic'



Implementation of GANs

```
class GAN(nn.Module):
    def __init__(self, data_loader=None, batch_size=32, epochs=10, optimizer=None,
                 noise_dim=30, device='cpu'):
        super(GAN, self).__init__()
        self.batch_size = batch_size # Batch size for training
        self.epochs = epochs # Number of training epochs
        self.optimizer = optimizer # Optimizer class for training
        self.noise_dim = noise_dim # Dimension of noise vector
        self.device = device # Device to run the model (CPU/GPU)
        self.generator = None # Placeholder for generator model
        self.discriminator = None # Placeholder for discriminator model
        self.gen_optimizer = None # Placeholder for generator optimizer
        self.dis_optimizer = None # Placeholder for discriminator optimizer
        self.data_loader = data_loader # Data loader for training dataset
```

Abstract GAN class

StandardGAN class

```
class StandardGAN(GAN):
    def __init__(self, data_loader=None, batch_size=32, epochs=10, optimizer=None, noise_dim=30, device='cpu'):
        super(StandardGAN, self).__init__(data_loader, batch_size, epochs, optimizer, noise_dim, device)

    def build(self):
        # Generator network definition
        self.generator = nn.Sequential(
            nn.Linear(self.noise_dim, 256), # Maps noise to 256 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Linear(256, 512), # Upsamples to 512 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Linear(512, 1024), # Upsamples to 1024 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Linear(1024, 784), # Outputs 784 dimensions (28x28 images)
            nn.Tanh(),
            nn.Unflatten(1, torch.Size([1, 28, 28])) # Reshapes to image format
        ).to(self.device)

        # Discriminator network definition
        self.discriminator = nn.Sequential(
            nn.Flatten(), # flattens input images
            nn.Linear(784, 1024), # Processes 784-dimensional input
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.3),
            nn.Linear(1024, 512), # Reduces to 512 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.3),
            nn.Linear(512, 256), # Further reduces to 256 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.3),
            nn.Linear(256, 128), # Reduces to 128 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.3),
            nn.Linear(128, 1), # Outputs a single score (real or fake)
            nn.Sigmoid() # Probability output
        ).to(self.device)
```

```
def train(self):
    self.history = {'disc-loss': [], 'gen-loss': []} # Track losses
    criterion = nn.BCELoss() # Binary Cross Entropy loss

    for epoch in range(self.epochs):
        dis_losses = [] # List to store discriminator losses
        gen_losses = [] # List to store generator losses
        for batch_idx, (X_batch, y_batch) in enumerate(self.data_loader):
            X_batch = X_batch.to(self.device) # Move batch to device

            # PHASE 1 - Train Discriminator
            z = torch.randn(X_batch.size(0), self.noise_dim, device=self.device) # Random noise
            X_gen = self.generator(z) # Generate fake images

            X_fake_and_real = torch.cat([X_gen, X_batch], dim=0) # Combine real and fake images
            y1 = torch.cat([torch.zeros(X_gen.size(0), 1), torch.ones(X_batch.size(0), 1)], dim=0).to(self.device)

            self.discriminator.train() # Set discriminator to training mode
            self.dis_optimizer.zero_grad() # Zero gradients
            dis_loss = criterion(self.discriminator(X_fake_and_real), y1) # Compute discriminator loss
            dis_loss.backward() # Backpropagation
            self.dis_optimizer.step() # Update discriminator weights
            dis_losses.append(dis_loss.item()) # Store loss

            # PHASE 2 - Train Generator
            z = torch.randn(self.batch_size, self.noise_dim, device=self.device) # New random noise
            y2 = torch.ones(self.batch_size, 1, device=self.device) # Labels for generator

            self.discriminator.eval() # Freeze discriminator weights
            self.gen_optimizer.zero_grad() # Zero gradients
            gen_loss = criterion(self.discriminator(self.generator(z)), y2) # Compute generator loss
            gen_loss.backward() # Backpropagation
            self.gen_optimizer.step() # Update generator weights
            gen_losses.append(gen_loss.item()) # Store loss
```

```
class DCGAN(GAN):
    def __init__(self, data_loader=None, batch_size=32, epochs=10, optimizer=None, noise_dim=30, device='cpu'):
        super(DCGAN, self).__init__(data_loader, batch_size, epochs, optimizer, noise_dim, device)

    def build(self):
        # Generator: Upscales noise to images
        self.generator = nn.Sequential(
            nn.Linear(self.noise_dim, 7 * 7 * 128), # Upscale noise
            nn.BatchNorm2d(7 * 7 * 128),
            nn.ReLU(inplace=True),
            nn.Unflatten(1, (128, 7, 7)),
            nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.ConvTranspose2d(64, 1, kernel_size=4, stride=2, padding=1),
            nn.Tanh() # Output between -1 and 1
        ).to(self.device)

        # Discriminator: Classifies images as real or fake
        self.discriminator = nn.Sequential(
            nn.Conv2d(1, 64, kernel_size=4, stride=2, padding=1), # Downsample
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.4),
            nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1), # Further downsample
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.4),
            nn.Flatten(),
            nn.Linear(128 * 7 * 7, 1), # Single output
            nn.Sigmoid() # Probability output
        ).to(self.device)
```

DCGAN class

Abstract GAN class

```
class GAN(nn.Module):
    def __init__(self, data_loader=None, batch_size=32, epochs=10, optimizer=None,
                 noise_dim=30, device='cpu'):
        super(GAN, self).__init__()
        self.batch_size = batch_size # Batch size for training
        self.epochs = epochs # Number of training epochs
        self.optimizer = optimizer # Optimizer class for training
        self.noise_dim = noise_dim # Dimension of noise vector
        self.device = device # Device to run the model (CPU/GPU)
        self.generator = None # Placeholder for generator model
        self.discriminator = None # Placeholder for discriminator model
        self.gen_optimizer = None # Placeholder for generator optimizer
        self.dis_optimizer = None # Placeholder for discriminator optimizer
        self.data_loader = data_loader # Data loader for training dataset

    @staticmethod
    def plot_multiple_images(images, n_cols=None):
        # Plot a grid of images
        n_cols = n_cols or len(images)
        n_rows = (len(images) - 1) // n_cols + 1
        if images.shape[1] == 1:
            images = images.squeeze(1) # Remove channel dimension if grayscale
        plt.figure(figsize=(n_cols, n_rows))
        for index, image in enumerate(images):
            plt.subplot(n_rows, n_cols, index + 1)
            plt.imshow(image.cpu().detach().numpy(), cmap="gray") # Display image
            plt.axis("off") # Hide axes
        plt.show()

    def build(self):
        pass # To be implemented in subclasses

    def generate(self, num_examples=None):
        # Generate fake images from random noise
        num_examples = self.batch_size if num_examples is None else num_examples
        z = torch.randn(num_examples, self.noise_dim, device=self.device) # Random noise
        X_gen = self.generator(z) # Generate images
        return X_gen
```

```
def train(self):
    self.history = {'disc-loss': [], 'gen-loss': []} # Track losses
    criterion = nn.BCELoss() # Binary Cross Entropy loss

    for epoch in range(self.epochs):
        dis_losses = [] # List to store discriminator losses
        gen_losses = [] # List to store generator losses
        for batch_idx, (X_batch, y_batch) in enumerate(self.data_loader):
            X_batch = X_batch.to(self.device) # Move batch to device

            # PHASE 1 - Train Discriminator
            z = torch.randn(X_batch.size(0), self.noise_dim, device=self.device) # Random noise
            X_gen = self.generator(z) # Generate fake images

            X_fake_and_real = torch.cat([X_gen, X_batch], dim=0) # Combine real and fake images
            y1 = torch.cat([torch.zeros(X_gen.size(0), 1), torch.ones(X_batch.size(0), 1)], dim=0).to(self.device) # Labels for discriminator

            self.discriminator.train() # Set discriminator to training mode
            self.dis_optimizer.zero_grad() # Zero gradients
            dis_loss = criterion(self.discriminator(X_fake_and_real), y1) # Compute discriminator loss
            dis_loss.backward() # Backpropagation
            self.dis_optimizer.step() # Update discriminator weights
            dis_losses.append(dis_loss.item()) # Store loss

            # PHASE 2 - Train Generator
            z = torch.randn(self.batch_size, self.noise_dim, device=self.device) # New random noise
            y2 = torch.ones(self.batch_size, 1, device=self.device) # Labels for generator

            self.discriminator.eval() # Freeze discriminator weights
            self.gen_optimizer.zero_grad() # Zero gradients
            gen_loss = criterion(self.discriminator(self.generator(z)), y2) # Compute generator loss
            gen_loss.backward() # Backpropagation
            self.gen_optimizer.step() # Update generator weights
            gen_losses.append(gen_loss.item()) # Store loss

        # Log and plot losses
        print(f"Epoch {epoch + 1}/{self.epochs}: disc-loss = {np.mean(dis_losses)}, gen-loss = {np.mean(gen_losses)}")
        self.history['disc-loss'].append(np.mean(dis_losses)) # Record average loss
        self.history['gen-loss'].append(np.mean(gen_losses))
```

StandardGAN

```
class StandardGAN(GAN):
    def __init__(self, data_loader=None, batch_size=32, epochs=10, optimizer=None, noise_dim=30, device='cpu'):
        super(StandardGAN, self).__init__(data_loader, batch_size, epochs, optimizer, noise_dim, device)

    def build(self):
        # Generator network definition
        self.generator = nn.Sequential(
            nn.Linear(self.noise_dim, 256), # Maps noise to 256 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Linear(256, 512), # Upsamples to 512 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Linear(512, 1024), # Upsamples to 1024 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Linear(1024, 784), # Outputs 784 dimensions (28x28 images)
            nn.Tanh(),
            nn.Unflatten(1, torch.Size([1, 28, 28])) # Reshapes to image format
        ).to(self.device)


        # Discriminator network definition
        self.discriminator = nn.Sequential(
            nn.Flatten(), # Flattens input images
            nn.Linear(784, 1024), # Processes 784-dimensional input
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.3),
            nn.Linear(1024, 512), # Reduces to 512 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.3),
            nn.Linear(512, 256), # Further reduces to 256 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.3),
            nn.Linear(256, 128), # Reduces to 128 dimensions
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.3),
            nn.Linear(128, 1), # Outputs a single score (real or fake)
            nn.Sigmoid() # Probability output
        ).to(self.device)

        # Set optimizers for generator and discriminator
        self.dis_optimizer = self.optimizer(self.discriminator.parameters(), lr=2e-4, betas=(0.5, 0.999))
        self.gen_optimizer = self.optimizer(self.generator.parameters(), lr=2e-4, betas=(0.5, 0.999))
```

Introducing Conv2DTranspose

□ Standard Conv2D

- Conv2D: input = [3,3], kernel = [2,2], stride = [1,1] → output = [2,2]



Introducing Conv2DTranspose

□ Standard Conv2D

- Conv2D: input = [3,3], kernel = [2,2], stride = [1,1] → output = [2,2]

0	0	1
0	1	2
1	2	0

input

0	1
1	1

filter

1	4
4	4

output

0	0
0	0

□ Conv2DTranspose: it is opposite of conv2D

- Input =[2,2], kernel = [2,2], stride = [1,1] → output = [3,3]

0	1
2	3

input


0	1
2	3

filter

Introducing Conv2DTranspose


□ Standard Conv2D

- Conv2D: input = [3,3], kernel = [2,2], stride = [1,1] → output = [2,2]



□ Conv2DTranspose: it is opposite of conv2D

- Input =[2,2], kernel = [2,2], stride = [1,1] → output = [3,3]



Introducing Conv2DTranspose

□ Standard Conv2D

- Conv2D: input = [3,3], kernel = [2,2], stride = [1,1] → output = [2,2]

0	0	1
0	1	2
1	2	0

input

0	1
1	1

filter

1	4
4	4

output

0	0	
0	0	

□ Conv2DTranspose: it is opposite of conv2D

- Input =[2,2], kernel = [2,2], stride = [1,1] → output = [3,3]

0	1
2	3

input

0	1
2	3

filter

0	1
2	3

Introducing Conv2DTranspose

□ Standard Conv2D

- Conv2D: input = [3,3], kernel = [2,2], stride = [1,1] → output = [2,2]

0	0	1
0	1	2
1	2	0

input

0	1
1	1

filter

1	4
4	4

output

0	0	
0	0	

□ Conv2DTranspose: it is opposite of conv2D

- Input =[2,2], kernel = [2,2], stride = [1,1] → output = [3,3]

0	1
2	3

0	1
2	3

input

X

0	1
2	3

filter


0	
2	
4	

0	3
6	9

Introducing Conv2DTranspose


□ Standard Conv2D

- Conv2D: input = [3,3], kernel = [2,2], stride = [1,1] → output = [2,2]



□ Conv2DTranspose: it is opposite of conv2D

- Input =[2,2], kernel = [2,2], stride = [1,1] → output = [3,3]



Conv2DTranspose

□ Basic syntax

python

 Copy code

```
torch.nn.ConvTranspose2d(in_channels, out_channels, kernel_size, stride, padding=0)
```

□ Arguments

- **in_channels**: Number of input channels (depth of the input feature map).
- **out_channels**: Number of output channels (depth of the output feature map).
- **kernel_size**: Size of the convolving kernel (e.g., 3 for a 3x3 filter).
- **stride**: Step size of the transposed convolution (determines the upsampling factor).
- **padding**: Implicit padding added on both sides of the input.

□ Output sizes

- **Output Size=(I-1)×S-2P+K+output padding**
- Input size (I), stride (S), padding (P), kernel size (K), output padding is optional

DCGAN

```

class DCGAN(GAN):
    def __init__(self, data_loader=None, batch_size=32, epochs=10, optimizer=None, noise_dim=30, device='cpu'):
        super(DCGAN, self).__init__(data_loader, batch_size, epochs, optimizer, noise_dim, device)


    def build(self):
        # Generator: Upscales noise to images
        self.generator = nn.Sequential(
            nn.Linear(self.noise_dim, 7 * 7 * 128), # Upscale noise
            nn.BatchNorm1d(7 * 7 * 128),
            nn.ReLU(inplace=True),
            nn.Unflatten(1, (128, 7, 7)),
            nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.ConvTranspose2d(64, 1, kernel_size=4, stride=2, padding=1),
            nn.Tanh() # Output between -1 and 1
        ).to(self.device)

        # Discriminator: Classifies images as real or fake
        self.discriminator = nn.Sequential(
            nn.Conv2d(1, 64, kernel_size=4, stride=2, padding=1), # Downsample
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.4),
            nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1), # Further downsample
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout(0.4),
            nn.Flatten(),
            nn.Linear(128 * 7 * 7, 1), # Single output
            nn.Sigmoid() # Probability output
        ).to(self.device)

        # Optimizers for both networks
        self.dis_optimizer = self.optimizer(self.discriminator.parameters(), lr=2e-4, betas=(0.5, 0.999))
        self.gen_optimizer = self.optimizer(self.generator.parameters(), lr=2e-4, betas=(0.5, 0.999))


```

Diffusion Model: Probabilistic Viewpoint



- **Forward process**
 - Shrinking and adding Gaussian noises
 - $\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}$ where $\beta_t \uparrow 1$
- **Backward process**
 - Learning to generate data by denoising
 - $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$

Overview of Diffusion Models



Reverse diffusion process. Image modified by Ho et al. 2020


$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon} \text{ and } \mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon} \text{ where } \alpha_t = 1 - \beta_t \text{ and } \bar{\alpha}_t = \prod_{s=1}^t \alpha_s$$

$$\diamond \quad \alpha_t \downarrow 0 \text{ and } \bar{\alpha}_t \downarrow 0 \rightarrow \mathbf{x}_T = \sqrt{\bar{\alpha}_T} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_T} \boldsymbol{\epsilon} = \sqrt{\prod_{s=1}^T \alpha_s} \mathbf{x}_0 + \sqrt{1 - \prod_{s=1}^T \alpha_s} \boldsymbol{\epsilon} \xrightarrow{T \rightarrow \infty} \mathcal{N}(\mathbf{0}, \mathbf{I})$$

$$\diamond \quad q(x_{t-1}|x_t, x_0) = N\left(\frac{1}{\sqrt{\alpha_t}}\left(\mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \boldsymbol{\epsilon}\right), \sigma_t^2 \mathbf{I}\right) \text{ with } \sigma_t = \sqrt{\frac{\beta_t(1-\bar{\alpha}_t)}{1-\alpha_t}}$$

$$\diamond \quad p_\theta(x_{t-1} | x_t) = N(\mu_t(x_t, t; \theta), \sigma_t^2 \mathbf{I})$$

$$\text{with } \mu_t(x_t, t; \theta) = \frac{1}{\sqrt{\alpha_t}}\left(\mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)\right) \text{ and } \sigma_t = \sqrt{\frac{\beta_t(1-\bar{\alpha}_t)}{1-\alpha_t}}$$



Algorithm 1 Training

```

1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
       $\nabla_\theta \|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t)\|^2$ 
6: until converged
  
```

Use $\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)$ to predict
the noise $\boldsymbol{\epsilon}$ injected by
minimizing

$$\|\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t) - \boldsymbol{\epsilon}\|^2$$

Algorithm 2 Sampling

```

1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
  
```

Implementation of Diffusion Models

- Step 1: Data Preparation
- Step 2: Forward Diffusion Process
- Step 3: Building a U-net Model for $\epsilon_\theta(\mathbf{x}_t, t)$
- Step 4: Model Training

$$L = \mathbb{E}_{\mathbf{x}_0 \sim q(\mathbf{x}), \epsilon \sim N(0, I), \mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_0 + \sqrt{1 - \alpha_t} \epsilon} [||\epsilon - \epsilon_\theta(\mathbf{x}_t, t)||^2]$$

- Step 5: Generation by Sampling

Thanks for your attention!