

**COPYRIGHT WARNING:** Copyright in these original lectures is owned by Monash University. You may transcribe, take notes, download or stream lectures for the purpose of your research and study only. If used for any other purpose, (excluding exceptions in the Copyright Act 1969 (Cth)) the University may take legal action for infringement of copyright.

Do not share, redistribute, or upload the lecture to a third party without a written permission!

# **FIT3181/5215 Deep Learning**

**Week 04: Stochastic Gradient Descent and Optimization**

**Lecturer: Dr Trung Le**

**Trung Le and Tutor team**

Email: [trunglm@monash.edu](mailto:trunglm@monash.edu)



# Introduction

- The **purpose** of this tutorial is to practise:
  1. Logistic Regression with GD and SGD
  2. AutoGrad in PyTorch
  3. Use AutoGrad to implement Logistic Regression

# Logistic Regression

- A shallow feed-forward neural network
  - Binary classifier with **one layer** and **softmax function**

## Forward propagation

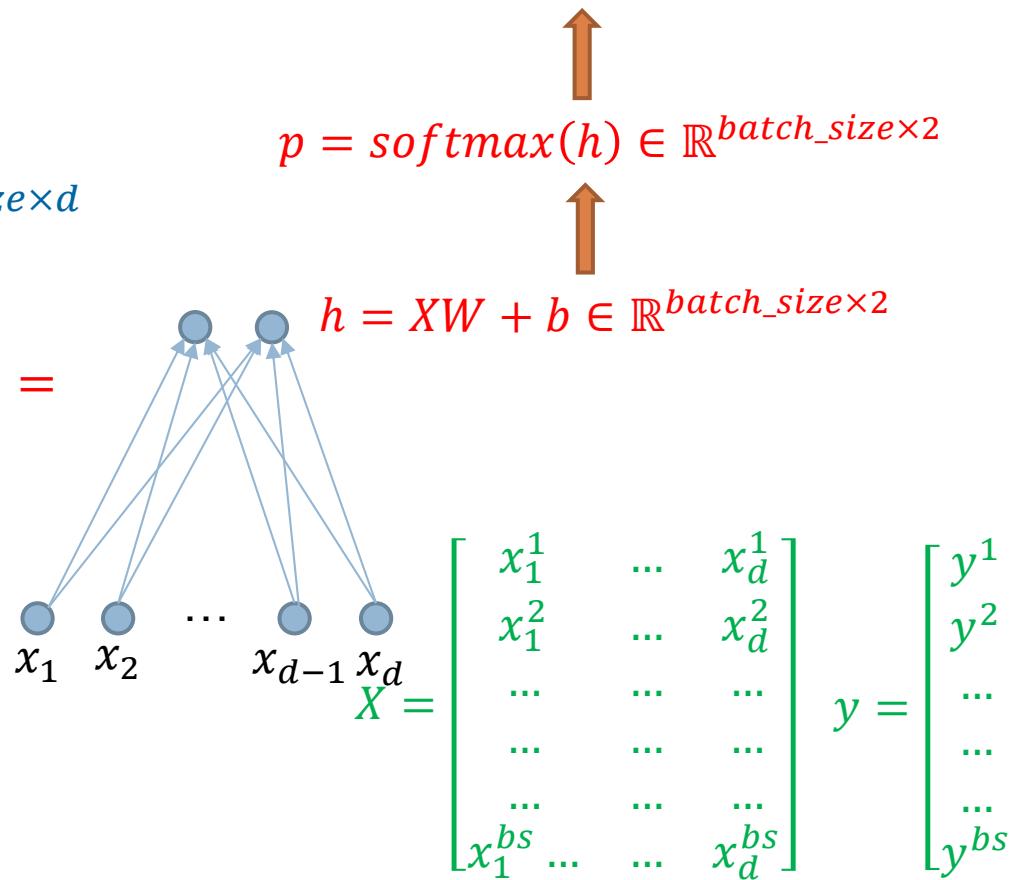
- $h = XW + b \in \mathbb{R}^{batch\_size \times 2}$  with  $X \in \mathbb{R}^{batch\_size \times d}$  and  $W \in \mathbb{R}^{d \times 2}$  and  $b \in \mathbb{R}^{1 \times 2}$

- $p = softmax(h, \text{dim} = 1) = \left[ p_1 = \frac{e^{h_1}}{e^{h_1} + e^{h_2}}, p_2 = \frac{e^{h_2}}{e^{h_1} + e^{h_2}} \right] \in \mathbb{R}^{batch\_size \times 2}$

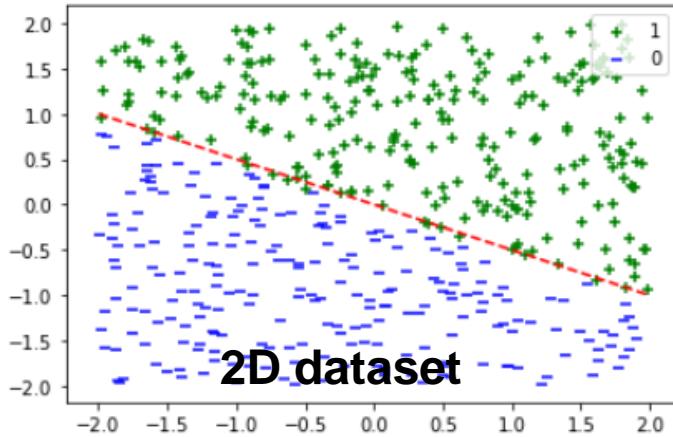
- $batch\_loss = \frac{1}{batch\_size} \sum_{i=1}^{batch\_size} CE(1_{y^i}, p^i) = -\frac{1}{batch\_size} \sum_{i=1}^{batch\_size} \log p_{y^i}^i$

$$batch\_loss = \frac{1}{batch\_size} \sum_{i=1}^{batch\_size} CE(1_{y^i}, p^i)$$

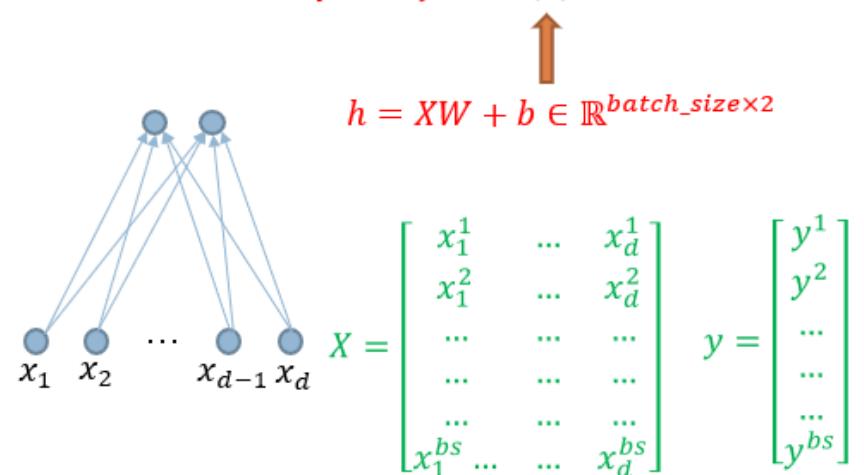
$$= -\frac{1}{batch\_size} \sum_{i=1}^{batch\_size} \log p_{y^i}^i$$



# Logistic Regression



$$p = \text{softmax}(h) \in \mathbb{R}^{\text{batch\_size} \times 2}$$



## Forward propagation

- $h = XW + b \in \mathbb{R}^{\text{batch\_size} \times 2}$  with  $X \in \mathbb{R}^{\text{batch\_size} \times d}$  and  $W \in \mathbb{R}^{d \times 2}$  and  $b \in \mathbb{R}^{1 \times 2}$
- $p = \text{softmax}(h, \text{dim} = 1) = \left[ p_1 = \frac{e^{h_1}}{e^{h_1} + e^{h_2}}, p_2 = \frac{e^{h_2}}{e^{h_1} + e^{h_2}} \right] \in \mathbb{R}^{\text{batch\_size} \times 2}$
- $\text{batch\_loss} = \frac{1}{\text{batch\_size}} \sum_{i=1}^{\text{batch\_size}} CE(1_{y^i}, p^i) = -\frac{1}{\text{batch\_size}} \sum_{i=1}^{\text{batch\_size}} \log p_{y^i}^i$

## Backward propagation

- $\frac{\partial l}{\partial p} = p - \begin{bmatrix} 1_{y^1} \\ \vdots \\ 1_{y^{bs}} \end{bmatrix} \in \mathbb{R}^{\text{batch\_size} \times 2}$
- $\frac{\partial l}{\partial W} = \frac{1}{\text{batch\_size}} X^T \frac{\partial l}{\partial p} \in \mathbb{R}^{d \times 2}$
- $\frac{\partial l}{\partial b} = \frac{1}{\text{batch\_size}} \text{sum}(\frac{\partial l}{\partial p}, \text{dim} = 0)$

## SGD update

- $W = W - \eta \frac{\partial l}{\partial W}$
- $b = b - \eta \frac{\partial l}{\partial b}$

# AutoGrad in PyTorch

## □ Declare a computational graph

```
x = torch.tensor([1.0, 2.0, 3.0], requires_grad=False)
y = torch.tensor([4.0])
W = torch.rand(3,1, requires_grad=True)
b = torch.rand(1, requires_grad=True)
y_hat = torch.matmul(x,W) + b
l = (y_hat - y)**2
```

## □ Do backpropagation and compute the gradients

```
l.backward(retain_graph=True)
#l.backward() raises RuntimeError if we try to through the graph a "second time"
#suggests that you might be attempting to call the .backward() function
#on a tensor that has already had its gradients calculated and the underlying computation graph has been freed.
print(f'W_grad = {W.grad}')
print(f'W_value={W.data}')
print(f'b_grad = {b.grad}')
print(f'b_value={b.data}')

W_grad = tensor([-0.6609],
               [-1.3217],
               [-1.9826]])
W_value=tensor([[0.6718],
               [0.2688],
               [0.6840]])
b_grad = tensor([-0.6609])
b_value=tensor([0.4080])
```

# Implement Logistic Regression using AutoGrad

```
class LogisticRegression(torch.nn.Module):
    def __init__(self):
        super(LogisticRegression, self).__init__()
        self.W = torch.nn.Parameter(torch.randn(2, 2, requires_grad= True)) → Declare W with requires_grad = True
        self.b = torch.nn.Parameter(torch.randn(1,2, requires_grad= True)) → Declare b with requires_grad = True

    def forward(self, x):
        return torch.matmul(x, self.W) + self.b

    def compute_loss(self, x, y):
        eps = 1E-10
        h = self.forward(x)
        #p = torch.nn.functional.softmax(h, dim=1)
        # Convert y to LongTensor for indexing
        #losses = [-torch.log(p[i, y[i].long()]+eps) for i in range(len(y))]
        #loss = torch.mean(torch.tensor(losses, requires_grad=True)) # Enable gradients for the loss tensor
        loss_fn = torch.nn.CrossEntropyLoss()
        loss = loss_fn(h, y.long().squeeze())
        return loss

    def update_SGD(self, x, y, eta):
        loss = self.compute_loss(x, y) → Compute the batch loss
        loss.backward() → Do backpropagation
        self.W.data = self.W.data - eta*self.W.grad → Use SGD to update W and b
        self.b.data = self.b.data - eta*self.b.grad
        if self.W.grad is not None:
            self.W.grad.zero_()
        if self.b.grad is not None:
            self.b.grad.zero_()

    def fit(self, train_loader, epochs, eta):
        plt.ion()
        losses = []
        for epoch in range(epochs):
            for x, y in train_loader:
                self.update_SGD(x, y, eta)
            train_loss = self.compute_loss(train_dataset.data, train_dataset.labels).detach().numpy()
            losses.append(train_loss)
            visualize_hyperplane_loss(X_train, y_train, self.W.data, self.b.data, losses)
            plt.pause(0.5)
        plt.ioff()
```

Thanks for your attention!  
Enjoy the tutorial 4!