


# FIT5230 **MALICIOUS AI**


S2 2025

Week 6:  
Generative Adversarial Networks (GANs)



# Overview

- Recap on Classification
- Generative Models
- Discriminative Models
- Generative Adversarial Networks
  - GAN & Game Theory
  - Game vs Optimisation Problem





# Generative Adversarial Networks

Security Gaming meets AI



# Classification

- Recap: Classification part of Pattern Recognition
  - pre-processing, Feature Extraction, Classification
- Classification: given observed sample  $x$ , predict its class / category label  $y$





# Classification: using statistics

- Classification: given sample  $x$ , predict its class label  $y \in \{1,2\}$
- Training:
  - given  $(x_1, y_1), \dots, (x_n, y_n)$
  - compute features of  $x_i$ 
    - $(f_1^i, f_2^i) = FE(x_i)$
  - plot how frequently the features of a class  $y_i$  have what values  
 $\Rightarrow P(Y | (f_1^i, f_2^i))$   
so given an  $x$ , can know most likely  $y$




# Classification: using statistics

- Classification (check during testing): given unknown sample  $x$ , predict its class label  $y$
- Training:  
 $\Rightarrow P(Y | (f_1, f_2))$
- Testing: given unknown sample  $x$ 
  - compute  $(f_1, f_2) = \text{FE}(x)$
  - output  $y$  s.t.  $\max[P(Y | (f_1, f_2))]$




# Classification Models

- Classification: given  $x$ , predict its class label  $y$ 
  - Discriminative
    - given  $x$ , compute  $P(Y | F=(f_1, f_2))$
    - output  $y$  s.t.  $\max P(Y|F)$
  - Generative
    - given  $x$ , find  $P(F, Y)$  or  $P(F|Y)$  for each class  $y$
    - then compute  $P(Y|F)$  using Bayes' theorem




# Classification Models\*

- Classification: given  $x$ , predict its class label  $y$ 
  - Discriminative
    - given  $x$ , compute  $P(Y | F=(f_1, f_2))$
    - output  $y$  s.t.  $\max P(Y|F)$
  - Generative
    - given  $x$ , find  $P(F, Y)$  or  $P(F|Y)$  for each class  $y$
    - compute  $P(Y|F)$   
$$= \frac{P(F|Y)P(Y)}{P(F)} \quad [\text{Bayes}']$$




# Generative Models




- Case (I):
  - given samples (the points)  $x_i$  from distribution  $p_{\text{real}}$ , obtain an estimation  $p_{\text{model}}$  of that distribution
  - use  $p_{\text{model}}(x)$  to find the probability that  $x$  occurs
- Case (II):
  - given samples (the points)  $x_i$  from distribution  $p_{\text{real}}$ , estimate more examples  $x'$  from the same distribution
  - e.g. GANs

# Generative vs Discriminative Models

- Generative model



- Discriminative model



# Generative Adversarial Networks

## Generative adversarial nets

I Goodfellow, J Pouget-Abadie... - Advances in neural ..., 2014 - proceedings.neurips.cc

... We propose a new framework for estimating **generative** models via **adversarial nets**, in which we simultaneously train two models: a **generative** model  $G$  that captures the data ...

☆ Save ⚡ Cite Cited by 71604 Related articles All 67 versions ☰



## Ian Goodfellow

Unknown affiliation

Verified email at cs.stanford.edu

Deep Learning

@ Aug 2024

# Generative Adversarial Networks

## Generative adversarial nets

I Goodfellow, J Pouget-Abadie ... - Advances in neural ..., 2014 - proceedings.neurips.cc

... We propose a new framework for estimating **generative** models via **adversarial nets**, in which we simultaneously train two models: a **generative** model  $G$  that captures the data ...

☆ Save 99 Cite Cited by 59517 Related articles All 61 versions »



## Ian Goodfellow

Unknown affiliation

Verified email at cs.stanford.edu

Deep Learning

@ Aug 2023

# Generative Adversarial Networks

## Generative adversarial nets

I Goodfellow, J Pouget-Abadie... - Advances in neural ..., 2014 - proceedings.neurips.cc

... We propose a new framework for estimating **generative** models via **adversarial nets**, in which we simultaneously train two models: a **generative** model  $G$  that captures the data ...

☆ Save 99 Cite Cited by 71604 Related articles All 67 versions »



## Ian Goodfellow

Unknown affiliation

Verified email at cs.stanford.edu

## Deep Learning

- + > 12000 citations in a year
- 7100++ citations/year

# GAN & Generative Models

- GAN is a type of generative model
  - obtain more samples from the same distribution



# Generative Adversarial Networks

- a model like security game:
  - 2 opposing players, G vs D

- Generator G



- creates samples  $x'$  from the same distribution as the real samples  $x$   
s.t.  $D(x) = D(x')$ , i.e. D can't tell the difference
- goal: INDistinguishability of generated samples
  - vs IND for crypto

# Generative Adversarial Networks

- a model like security game: 2 opposing players, G vs D

- Generator G

- goal: INDistinguishability of generated samples




- Discriminator D


- checks samples to determine if real or fake
  - goal: break IND
  - outputs probability estimation  $D(\cdot)$  that i/p is real



# Generative Adversarial Networks



# Generative Adversarial Networks



# GAN & Game Theory

- GAN Training: 2-player minimax game
  - each player's cost/loss function  $J_D, J_G$   
 $= f(\text{both players' params } \theta_D, \theta_G)$


G

D

- D  $\xrightarrow{x} \boxed{D} \longrightarrow D(x \text{ or } G(z))$ 
  - $x' = G(z)$
  - maximize the probability  $D(x)$  where  $x \sim p_{\text{data}}$
  - minimize the probability  $D(G(z))$  where  $G(z) \sim p_{\text{model}}$
  - by controlling only  $\theta_D$ , taking turns with G

# GAN & Game Theory

- GAN Training: 2-player minimax game



- maximize the probability  $D(G(z))$  where  $G(z) \sim p_{\text{model}}$ , or
- minimize  $1 - D(G(z))$
- by controlling only  $\theta_G$ , taking turns with D

# GAN & Game Theory

- D:



- max the probability  $D(x)$ ,
- min the probability  $D(G(z))$ , or max  $1-D(G(z))$

- G:



- max the probability  $D(G(z))$ , or min  $1-D(G(z))$

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

real

fake


# Game vs Optimisation Problem

- Each player's cost  $J_D$  ( $J_G$ ) depends on the other player's parameters  $\theta_G$  ( $\theta_D$ ), but cannot control that, so is not optimisation but a game
- Solution to optimisation problem is a minimum point
  - vs
- Solution to a game is a Nash equilibrium


$$\arg \min_{\theta_G} \max_{\theta_D} J(\theta_D, \theta_G)$$

where arg max just means arguments of the maxima i.e. values of  $\theta$   
s.t. the expression has a max value

# GAN Discriminator D



# GAN Generator G



G benefits from D's outputs

# DCGAN @ICLR 2016

## UNSUPERVISED REPRESENTATION LEARNING WITH DEEP CONVOLUTIONAL GENERATIVE ADVERSARIAL NETWORKS

**Alec Radford & Luke Metz**

indico Research

Boston, MA

{alec, luke}@indico.io

**Soumith Chintala**

Facebook AI Research

New York, NY

soumith@fb.com


cited > 18,000 times

### ABSTRACT

In recent years, supervised learning with convolutional networks (CNNs) has seen huge adoption in computer vision applications. Comparatively, unsupervised learning with CNNs has received less attention. In this work we hope to help bridge the gap between the success of CNNs for supervised learning and unsupervised learning. We introduce a class of CNNs called deep convolutional generative adversarial networks (DCGANs), that have certain architectural constraints, and


# DCGAN

- overall network structure based on all convolutional net
  - conventional: conv, pooling, fully connected
  - AllConvNet: replace pooling with conv, no fully connected layer
- batch normalisation layers in most layers of G and D
- ReLU for G except Tanh for last layer, LeakyReLU for D
- does not have pooling nor unpooling layers
  - upsampling done via transposed convolution with stride  $> 1$
- uses Adam optimiser instead of SGD to update the weights




# DCGAN

- Tanh activation function: range –1 to 1



- LeakyReLU:

- non-zero for –ve input




- Adam optimiser:


- alternative optimization to stochastic/random gradient descent, to update the neural network weights

(a)


# DCGAN: G & D



(b)



# DCGAN: G



# CycleGAN@ ICCV 2017

## Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks

Jun-Yan Zhu\*      Taesung Park\*      Phillip Isola      Alexei A. Efros  
Berkeley AI Research (BAIR) laboratory, UC Berkeley

### Abstract


*Image-to-image translation is a class of vision and graphics problems where the goal is to learn the mapping between an input image and an output image using a training set of aligned image pairs. However, for many tasks, paired training data will not be available. We present an approach for learning to translate an image from a source domain  $X$  to a target domain  $Y$  in the absence of paired examples. Our goal is to learn a mapping  $G : X \rightarrow Y$  such that the distribution of images from  $G(X)$  is indistinguishable from the distribution  $Y$  using an adversarial loss. Because this mapping is highly under-constrained, we couple it with an inverse mapping  $F : Y \rightarrow X$  and introduce a cycle consistency loss to push  $F(G(X)) \approx X$  (and vice versa). Qualitative results are presented on several tasks where paired training data does not exist, including collection style transfer, object transfiguration, season transfer, photo enhancement, etc. Quantitative comparisons against several prior methods demonstrate the superiority of our approach.*

cited > 28,000 times

# CycleGAN

- Learns to translate between domains using **unpaired datasets**.
- CycleGAN consists of:
  - Two Generators:
    - $G: X \rightarrow Y$  (e.g., horse  $\rightarrow$  zebra)
    - $F: Y \rightarrow X$  (e.g., zebra  $\rightarrow$  horse)
  - Two Discriminators:
    - $D_Y$ : distinguishes real vs fake images in domain Y
    - $D_X$ : distinguishes real vs fake images in domain X


# CycleGAN



**Figure 1:** Given any two unordered image collections  $X$  and  $Y$ , our algorithm learns to automatically “translate” an image from one into the other and vice versa. Example application (*bottom*): using a collection of paintings of a famous artist, learn to render a user’s photograph into their style.

# CycleGAN

- Cycle consistency loss:
  - If you translate an image from domain  $X$  to  $Y$  using  $G$ , and then back to  $X$  using  $F$ , you should get the original image:  $F(G(x)) \approx x$
  - Similarly:  $G(F(y)) \approx y$
- Prevents the generators from producing arbitrary outputs and enforces semantic preservation.



# CycleGAN

- Adversarial Loss encourages generators to produce realistic images that fool the discriminators:

$$\begin{aligned}\mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) = & \mathbb{E}_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] \\ & + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log(1 - D_Y(G(x)))]\end{aligned}$$


- Cycle Consistency Loss ensures that translating back and forth yields the original image:

$$\begin{aligned}\mathcal{L}_{\text{cyc}}(G, F) = & \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] \\ & + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1].\end{aligned}$$

# Evaluation Metrics for GANs

- **Inception Score (IS)**

- Measures image quality and diversity.
  - Image Quality: A good image should be confidently classified into a specific class.
  - Image Diversity: The generated images should span across many different classes.
- Uses a pretrained classifier (e.g., Inception v3).
- High score = sharp, diverse images.



Ideal marginal distribution

[Image source](#)

# Inception Score (IS)\*

- The Inception Score is defined as:

$$\exp(\mathbb{E}_{\mathbf{x}} \text{KL}(p(y|\mathbf{x}) || p(y)))$$

where:

- $\mathbf{x}$  be a generated image.
- $p(y|\mathbf{x})$  be the conditional label distribution predicted by the Inception model.
- $p(y)$  be the marginal distribution over all generated images.
- $D_{\text{KL}}$  is the Kullback-Leibler divergence.
- $\mathbb{E}_{\mathbf{x}}$  is the expectation over all generated images.

[Image source](#)

# Evaluation Metrics for GANs

- **Fréchet Inception Distance (FID)**
  - Compares statistics (e.g., mean and covariance) of real vs generated images.
  - Lower FID = closer to real distribution.
  - More robust than IS for detecting mode collapse.

# Fréchet Inception Distance (FID)\*

- FID is computed as:

$$\|\mathbf{m} - \mathbf{m}_w\|_2^2 + \text{Tr}(\mathbf{C} + \mathbf{C}_w - 2(\mathbf{C}\mathbf{C}_w)^{1/2})$$

Where:


- $\|\mathbf{m} - \mathbf{m}_w\|_2^2$  measures the difference in means
- $\text{Tr}$  is the trace of a matrix (sum of diagonal elements)
- $(\mathbf{C}\mathbf{C}_w)^{1/2}$  is the matrix square root of the product of covariances.

# Evaluation Metrics for GANs

- **Perceptual Path Length (PPL)**

- Measures smoothness of latent space interpolation
- Reflects how well the generator maps small changes in latent vectors to perceptually smooth changes in the generated images.
- Lower PPL = more consistent transitions

Image source



# Perceptual Path Length (PPL)\*

- PPL is computed as:
  - Sample two latent vectors  $z_1$  and  $z_2$
  - Interpolate between them using a small step  $\epsilon$ :  $z(t) = (1 - t)z_1 + tz_2$
  - Generate images from interpolated latent vectors:  $I(t) = G(z(t))$
  - Compute PPL:

$$\mathbb{E}_{z_1, z_2, t} \left[ \frac{d(I(t), I(t+\epsilon))}{\epsilon^2} \right]$$

Where:

- $d(x)$  is a perceptual distance (not pixel-wise)
- The expectation is taken over many random pairs and interpolation points

# Evaluation Metrics for GANs

- **Precision & Recall**
  - Precision: How realistic are the generated samples?
  - Recall: How diverse are the generated samples?
- Embed images (real and generated) into a feature space using a pretrained network.
- Model the distributions of real and generated images in this space.
- For each generated image:
  - Check if it lies within the manifold of real images → contributes to precision.
- For each real image:
  - Check if it lies within the manifold of generated images → contributes to recall.

# Precision & Recall

- Let:
  - $G$ : set of generated samples
  - $R$ : set of real samples
  - $\phi(x)$ : feature embedding of image  $x$
- Precision: Fraction of  $\phi(g) \in G$  that are close to some  $\phi(r) \in R$ .
- Recall: Fraction of  $\phi(r) \in R$  that are close to some  $\phi(g) \in G$ .
- “Closeness” is typically defined using a k-nearest neighbor or radius-based threshold in feature space.



# Appendix

# Cross Entropy Loss\*

- Loss? Common in machine learning
- Cross-entropy loss / log loss used in GAN
  - for classification where output = probability
- Information vs Entropy

# Cross Entropy Loss\*

- Cross-entropy loss / log loss
- Information  $h(x) = -\log( p(x) )$  , –ve is to ensure the result is +ve
- Measure of Information:
  - to quantify how much info is available
  - how much surprise when something happens?
    - low probability event: high info / surprising / uncertain
      - if  $p(x) \approx 0$ , info  $h(x) = -\log( p(x) ) \rightarrow \infty$  , very surprising
    - high probability event: low info / ho-hum

# Cross Entropy Loss\*


- Information  $h(x) = -\log(p(x))$ , -ve is to ensure the result is +ve
- Entropy  $H(x)$ :
  - average number of bits/info/surprise/uncertainty to transmit a random event from a probability distribution

$$H(X) := - \sum_{x \in \mathcal{X}} p(x) \log p(x) = \mathbb{E}[-\log p(X)],$$

- expected value of self-information of a random variable

# Cross Entropy Loss\*

- Loss? Common in machine learning
- GAN loss? cross-entropy loss



$$H(p, q) = \mathbb{E}_{x \sim p(x)}[-\log q(x)]$$

- where label is predicted by D