

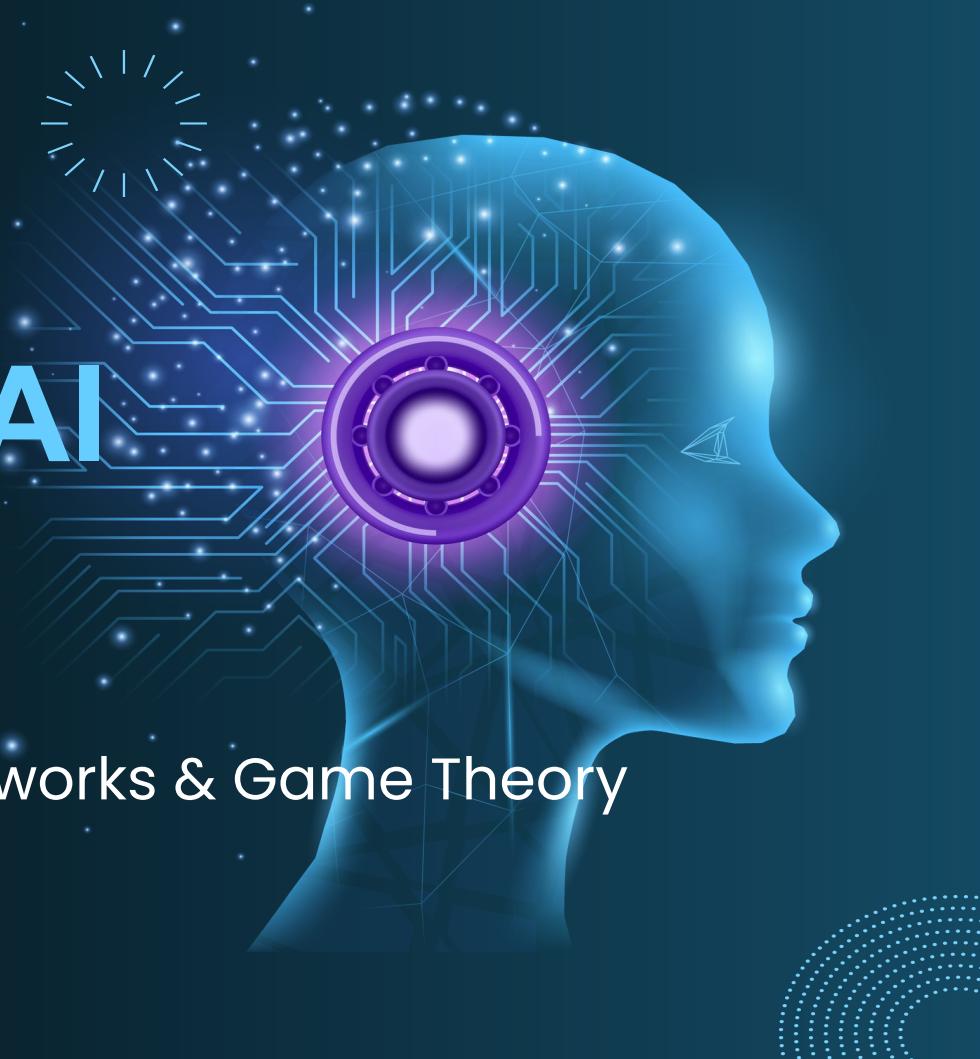


FIT5230 **MALICIOUS AI**

S2 2025

Week 7:
Generative Adversarial Networks & Game Theory

•
•
•



Overview

- GAN Quick Recap
- G/D Min-Max Game
- Game Theory
- Prisoner's Dilemma
- Nash Equilibrium
- Matching Coins Game
- GAN Motivation
- GAN for AI Security



Generative Adversarial Network **Quick Recap**

Generative Adversarial Networks

Generative Adversarial Nets

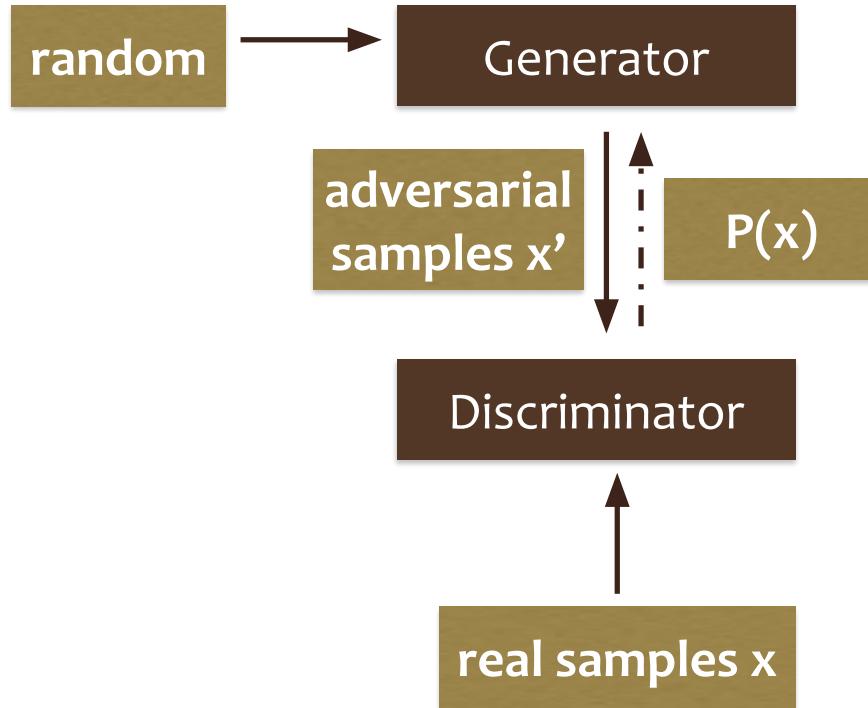
**Ian J. Goodfellow, Jean Pouget-Abadie*, Mehdi Mirza, Bing Xu, David Warde-Farley,
Sherjil Ozair†, Aaron Courville, Yoshua Bengio‡**

Département d'informatique et de recherche opérationnelle
Université de Montréal
Montréal, QC H3C 3J7

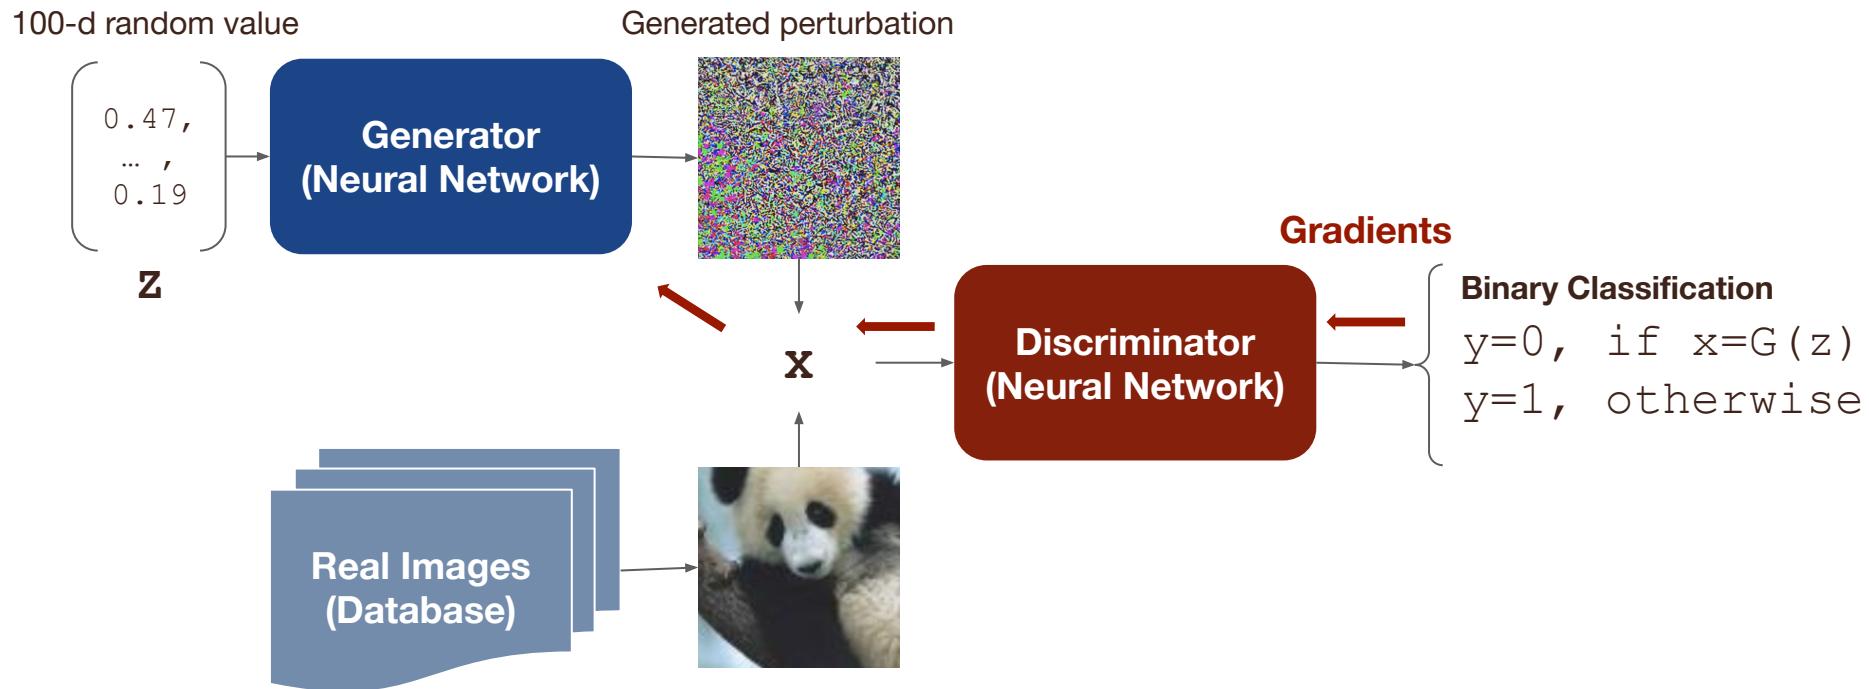
Abstract

We propose a new framework for estimating generative models via an adversarial process, in which we simultaneously train two models: a generative model G that captures the data distribution, and a discriminative model D that estimates the probability that a sample came from the training data rather than G . The training procedure for G is to maximize the probability of D making a mistake. This framework corresponds to a minimax two-player game. In the space of arbitrary functions G and D , a unique solution exists, with G recovering the training data

Generative Adversarial Networks



G/D Game



G/D Min-Max Game

The adversarial nature of GANs creates a Min-Max Game

- leading to the final loss function:

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

- The **discriminator** tries to maximize this function by correctly classifying real and fake data.
- The **generator** tries to minimize it by generating realistic samples that fool the discriminator.

GANs Summary

- First presented by Ian Goodfellow and associates in 2014
- Composed of two neural networks that cooperate through a competitive process
 - **Generator:** To generate human imperceptible adversarial examples
 - **Discriminator:** To determine real vs fake data

The outcome of this adversarial process is extremely advanced models that can recognize and produce synthetic media that is remarkably realistic.



Game Theory **Prisoner's Dilemma**

GAN & Game Theory

- Two players G and D in a game
- Shown that Nash equilibrium exists when function space corresponds to $p_{\text{real}} = p_{\text{model}}$
- i.e. distribution of real = distribution of G's outputs

Theory of Games

Game:

- players indexed by $i \in \{1, \dots, \ell\}$: P_i

Strategies S_i = set of available actions for player i

- $s_i \in S_i$ = action by player i

Payoffs $u_i(S)$ = payoff/utility function of player i
where S is list of actions of player i

Q: payoff depends on whose strategies?

Theory of Games: the Matrix

Matrix game / strategic form game

- 2-player game with identical rounds
- in each round: each player simultaneously makes a move
- outcome: tie or win/lose (loser pays winner)

	Rock	Paper	Scissors
Rock	0, 0	-1, 1	1, -1
Paper	1, -1	0, 0	-1, 1
Scissors	-1, 1	1, -1	0, 0

Theory of Games: the Matrix

Payoff matrix

- a_{ij} is payoff in a round
- P1 makes move i, P2 makes move j
- +ve if P1 wins, 0 if tie, -ve if P1 loses

Zero-sum game

- sum of winnings of all players = 0

	Rock	Paper	Scissors
Rock	0, 0	-1, 1	1, -1
Paper	1, -1	0, 0	-1, 1
Scissors	-1, 1	1, -1	0, 0

Game Theory: Prisoner's Dilemma

Two suspects 呂 and 吳 arrested, questioned by police

1. If no one confesses: 1 year jail each
2. If one confesses, the other keeps quiet:
 - leniency for the rat: 0 years
 - severe punishment for silence: 20 years
3. If both confess:
 - both punished equally: 5 years

Best strategy? To confess or not to confess?

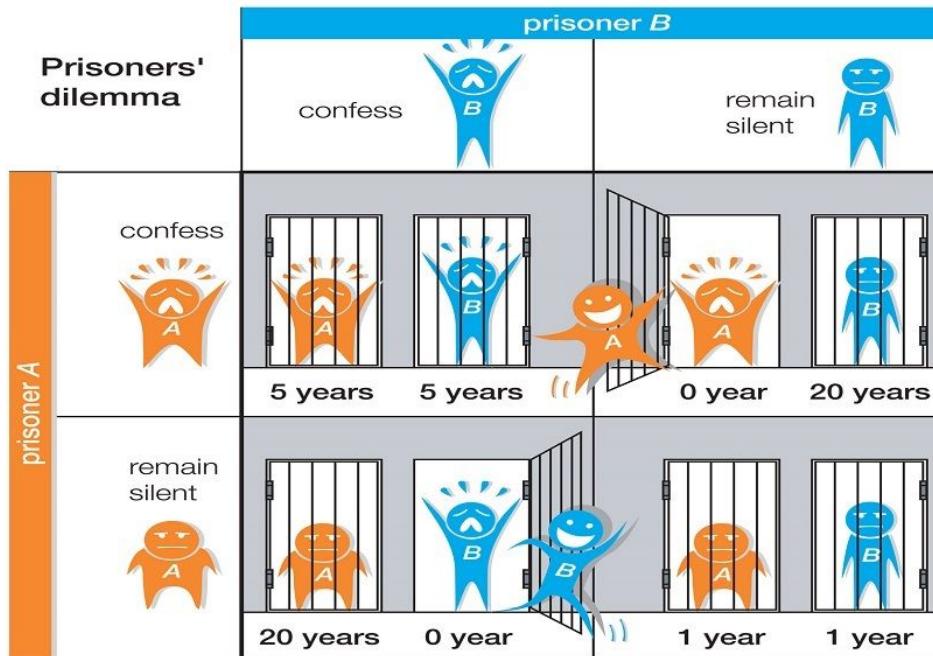
Game Theory: Prisoner's Dilemma



Market Realist^Q

Source: Encyclopedia Britannica

Game Theory: Prisoner's Dilemma



Market Realist^Q

Source: Encyclopedia Britannica

웃\웃	Confess	Defect
Confess	(-5,-5)	(0,-20)
Defect	(-20,0)	(-1,-1)

Game Theory: Prisoner's Dilemma

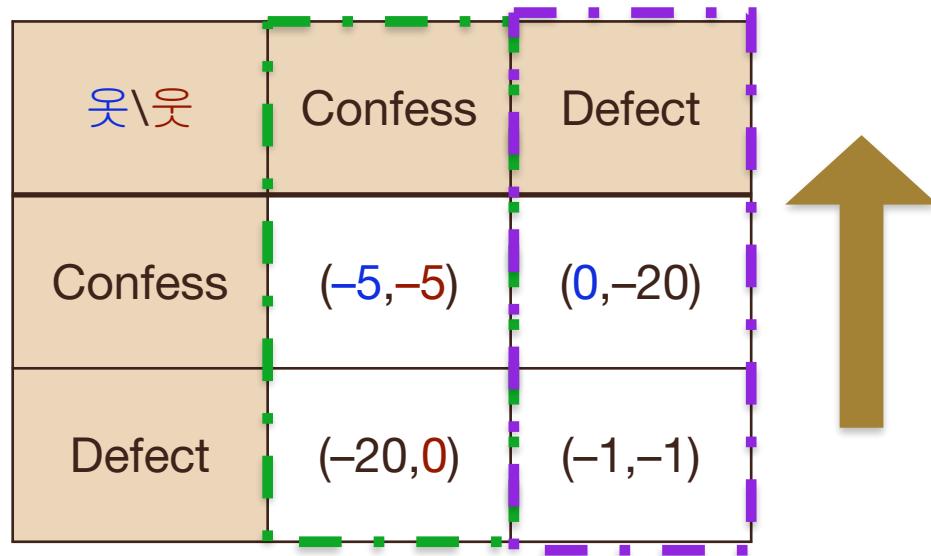
Payoff matrix:

1. if 㠭 confesses:
confess gets 㠭 less years

2. if 㠭 defects
confess gets 㠭 less years

'confess' vs 'defect'
the only possible outcome is: ???

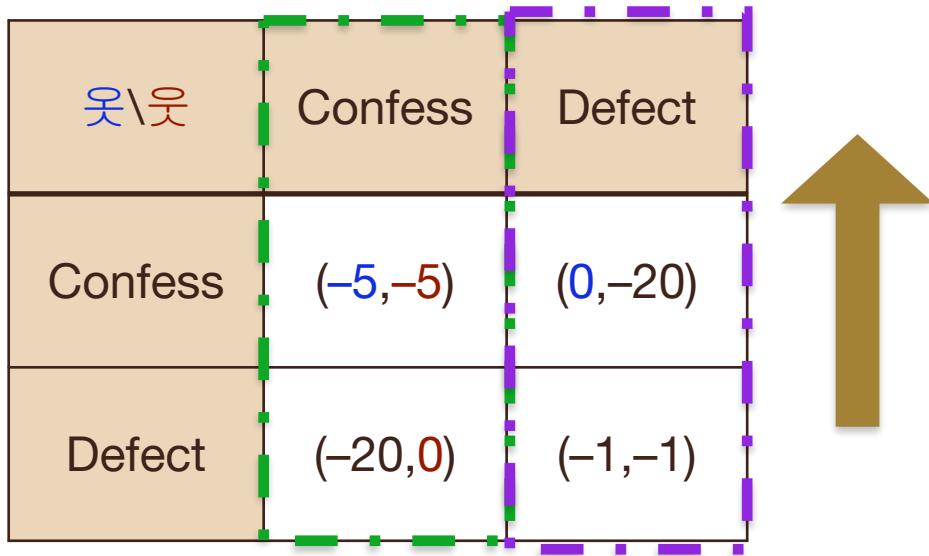
Q: Why didn't they go for ???



Game Theory: Prisoner's Dilemma

- 'confess' strictly dominates 'defect'
- **rational** players never play strictly dominated strategies
- the only possible outcome is
 - $\langle \text{Confess}, \text{Confess} \rangle \rightarrow (-5, -5)$

Why didnt they go for
 $\langle \text{Defect}, \text{Defect} \rangle \rightarrow (-1, -1)$?



Game Theory: Prisoner's Dilemma

Why not <Defect,Defect>?

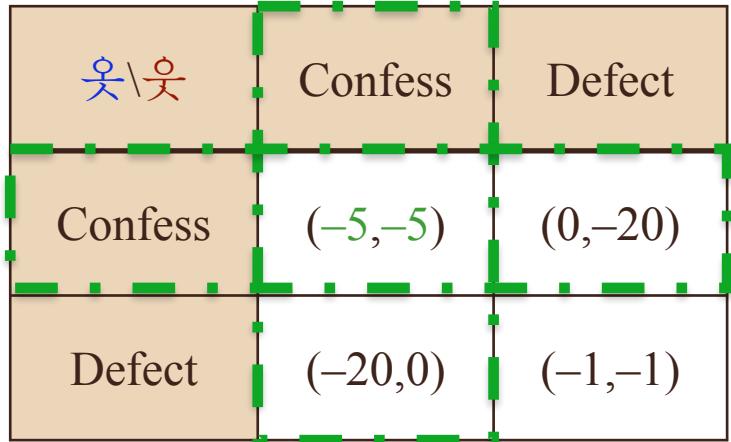
- even if they had agreed to both not confess (i.e. defect) once in interrogation room, its in their best interest to confess irrespective of what the other party does
- so, its **not** a **stable** outcome, each of them can get better payoff by flipping their strategies, so **not equilibrium**

웃\웃	Confess	Defect
Confess	(-5,-5)	(0,-20)
Defect	(-20,0)	(-1,-1)

Nash Equilibrium

Stable outcome:

- the only possible outcome: <Confess, Confess> → (-5,-5)
- For this outcome,
- each player can't do any better by changing his/her strategy/action, so no incentive to change \Rightarrow Nash equilibrium
 - since 웃 confesses:
changing from confess to defect ...
 - since 웃 confesses:
changing from confess to defect ...



Nash Equilibrium

Stable outcome/state,

- where no player can gain any better payoff by changing his/her strategy
- given that other players have chosen their action

Pure strategy Nash Equilibrium

- each player only one (pure) strategy:
{Confess,Defect}

웃\웃	Confess	Defect
Confess	(-5,-5)	(0,-20)
Defect	(-20,0)	(-1,-1)



Non-cooperative games

[J Nash - Annals of mathematics, 1951 - JSTOR](#)

... prove that a finite **non-cooperative game** always has at least ... **non-cooperative game** and prove a theorem on the geometrical structure of the set of equilibrium points of a solvable **game**. ...

[☆ Save](#) [⤒ Cite](#) [Cited by 14413](#) [Related articles](#) [All 35 versions](#) [⤓](#)

Matching Coins game

Matrix game: the Matching Coins game

- 2 players 옷 and 웃
- actions: {Heads, Tails}
- Payoff matrix:

웃\웃	Heads	Tails
Heads	(1,-1)	(-1,1)
Tails	(-1,1)	(1,-1)

Best strategy: Heads or Tails?

Penalty Kick game

Matrix game: the Penalty Kick game

- 2 players Goalie 웃 and Striker 웃
- actions: {Left, Right}
- Payoff matrix:

웃\웃	Left	Right
Left	(1,-1)	(-1,1)
Right	(-1,1)	(1,-1)

Best strategy: Left or Right?

Matching Coins game

1. <Heads,Heads>: 웃 chose Heads
웃 prefers to have chosen Tails
so <Heads,Heads> not equilibrium

2. <Heads,Tails>: 웃 chose Tails
웃 prefers to have chosen Tails
so <Heads,Tails> not equilibrium

Recall... Payoff matrix:

웃\웃	Heads	Tails
Heads	1.(1,-1)	2.(-1,1)

웃\웃	Tails
Heads	2.(-1,1)
Tails	3.(1,-1)

웃\웃	Heads	Tails
Heads	1.(1,-1)	2.(-1,1)
Tails	4.(-1,1)	3.(1,-1)

Matching Coins game



3. < Tails,Tails >: 웃 chose Tails
웃 prefers to have chosen Heads
so <Heads,Tails> not equilibrium

4. < Tails,Heads >: 웃 chose Heads
웃 prefers to have chosen Heads
so <Tails,Heads> not equilibrium

No pure strategy equilibrium
Q: can equilibrium be achieved?

웃\웃	Heads	Tails
Tails	4.(-1,1)	3.(1,-1)

웃\웃	Heads	
Heads	1.(1,-1)	
Tails	4.(-1,1)	



웃\웃	Heads	Tails
Heads	1.(1,-1)	2.(-1,1)
Tails	4.(-1,1)	3.(1,-1)

Nash Equilibrium

Stable outcome/state,

- where no player can gain any better payoff by changing his/her strategy
- given that other players have chosen their action

Pure strategy Nash Equilibrium

- each player only one (pure) strategy:
{Confess,Defect}

웃\웃	Confess	Defect
Confess	(-5,-5)	(0,-20)
Defect	(-20,0)	(-1,-1)



Non-cooperative games

[J Nash - Annals of mathematics, 1951 - JSTOR](#)

... prove that a finite **non-cooperative game** always has at least ... **non-cooperative game** and prove a theorem on the geometrical structure of the set of equilibrium points of a solvable **game**. ...

[☆ Save](#) [⤒ Cite](#) [Cited by 14413](#) [Related articles](#) [All 35 versions](#) [⤓](#)

Nash Equilibrium

Pure strategy Nash Equilibrium

- each player only one (pure) strategy
- e.g. H or T, L or R, Go or Stop, Confess or Defect

Q: What if can't achieve pure strategy Nash Equilibrium? but one must exist ...

- **Nash's Existence Theorem**
 - every game with a finite number of players who can choose from finitely many pure strategies, has at least one Nash equilibrium

Mixed Strategy Nash Equilibrium

Based on Nash's Existence Theorem, if no equilibrium for pure strategies, must have for mixed

Mixed strategy

- probability distribution over multiple pure strategies: each one may be chosen based on some probability $P()$

Matching Coins game

Q: what if playing against a mind reader? how to not always lose?

A: ?

웃\웃 0.5	Heads 0.5	Tails 0.5
Heads 0.5	(1,-1)	(-1,1)
Tails 0.5	(-1,1)	(1,-1)

Matching Coins game

Q: what if playing against a mind reader? how to not always lose?

A: just **flip the coin** (i.e. random)

- at best, opponent wins only half the time
- he can't do anything to change outcome
- neither player can change strategy & expect to do better
- mixed strategy Nash equilibrium

웃\웃	Heads 0.5	Tails 0.5
Heads 0.5	(1,-1)	(-1,1)
Tails 0.5	(-1,1)	(1,-1)

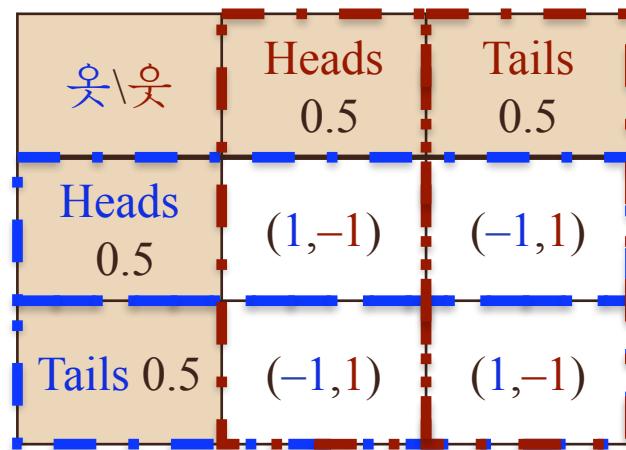
Mixed Strategy Nash Equilibrium

Mixed strategy

- each player chooses an action with probability 0.5
- Payoff for each player for choosing an action:
 - payoff $\pi(H) = 0.5(1) + 0.5(-1) = 0$
 - payoff $\pi(T) = 0.5(-1) + 0.5(1) = 0$
 - payoff $\pi(H) = 0.5(-1) + 0.5(1) = 0$
 - payoff $\pi(T) = 0.5(1) + 0.5(-1) = 0$

vs pure strategy

- payoff either 1 or -1



Battle of the Sexes game

- 2 players: man 옷 and woman 옷
 - get together for night for entertainment, but no communication
 - actions: choose ballet or watch fight
 - man 옷 prefers fight, woman 옷 prefers ballet
 - both prefer together vs alone

Payoff matrix

옷\웃	Ballet	Fight
Ballet	(1,2)	(0,0)
Fight	(0,0)	(2,1)

Q: any Nash equilibrium?

Battle of the Sexes game

Pure strategy Nash equilibrium **exists**

- since 웃 chose Ballet
웃: no point to change from Ballet to Fight
 - since 웃 chose Ballet
웃: no point to change from Ballet to Fight
- ...

웃\웃	Ballet	Fight
Ballet	(1,2)	(0,0)
Fight	(0,0)	(2,1)

Pure strategy Nash equilibrium:

- <Ballet, Ballet> or
- <Fight, Fight>

Battle of the Sexes game

Pure strategy Nash equilibrium exists

But how to decide? Each has his/her own preference, how to coordinate even though it can be done? Not sure which one?

Check if **mixed strategy Nash equilibrium** exists using the **Mixed Strategy Algorithm**

웃\웃	Ballet	Fight
Ballet	(1,2)	(0,0)
Fight	(0,0)	(2,1)

Mixed Strategy Algorithm

Solve for 옷's mixed strategy:

① Target: payoff 옷(Ballet) = payoff 옷(Fight)

- so it won't matter to 옷 which one she chooses

② payoff 옷(L)= $P_{\text{웃}}(U)(2) + (1-P_{\text{웃}}(U))(0)$

③ payoff 옷(R)= $P_{\text{웃}}(U)(0) + (1-P_{\text{웃}}(U))(1)$

- where $P_{\text{웃}}(U)$ is probably of 옷 choosing U

웃 \ 웃	L	R
U	(1,2)	(0,0)
D	(0,0)	(2,1)

Mixed Strategy Algorithm

Solve for 웃's mixed strategy:

① Target: payoff 웃(Балет) = payoff 웃(Бой)

- so it won't matter to 웃 which one she chooses

② payoff 웃(U)= $P_{웃}(L)(1) + (1-P_{웃}(L))(0)$

③ payoff 웃(D)= $P_{웃}(L)(0) + (1-P_{웃}(L))(2)$

- where $P_{웃}(L)$ is probably of 웃 choosing L

웃 \ 웃	L	R
U	(1,2)	(0,0)
D	(0,0)	(2,1)

Mixed Strategy Algorithm

Payoff for each player for choosing an action:

- ① compute the P of each outcome, using $P_{\text{웃}}$ and $P_{\text{웃}}$
- ② for each outcome, multiply P by player's payoff, & sum these up for all outcomes

$$\text{payoff } \text{웃} = 2/9(1) + 1/9(0) + 4/9(0) + 2/9(2) = 2/3$$

$$\text{payoff } \text{웃} = 2/9(2) + 1/9(0) + 4/9(0) + 2/9(1) = 2/3$$

Pure strategy payoffs are better; just give in & follow

웃\웃	Ballet ($\frac{2}{3}$)	Fight ($\frac{1}{3}$)
Ballet ($\frac{1}{3}$)	(1, 2)	(0, 0)
Fight ($\frac{2}{3}$)	(0, 0)	(2, 1)



Generative Adversarial Network **AI Security**

GANs for Security

The Threat Landscape

- Deepfakes → AI-generated media (video, audio, text, images) that depict fabricated events or speech
- Forged Content → Broader manipulation of digital media beyond deepfakes

Key Risks:

- Misinformation: Undermines trust in news and media.
- Identity Theft & Fraud: Impersonation for financial or personal gain.
- Political Manipulation: Influences elections and public opinion.
- Defamation: Damages reputations through falsified portrayals.

GANs for Security: Threat Detection and Prevention

GANs as a Defense Mechanism

- GANs can be trained on real vs. fake datasets to identify manipulation artifacts.
- Detects subtle inconsistencies in lighting, audio, facial expressions, motion patterns.

Application in Content Moderation

- Automatic Flagging: AI detection for suspicious media.
- Real-Time Analysis: Prevents rapid spread of harmful content.
- User Reporting Integration: Enhances system accuracy with human input.

The Dark Side: GANs as a Security Threat

- Realistic Malware Creation:
 - GANs generate polymorphic malware that evades signature-based antivirus.
 - Subtle perturbations → difficult to classify as malicious.
- Adversarial Evasion Attacks:
 - GANs create adversarial inputs that mislead ML-based security systems.
 - Exploits classifier vulnerabilities.
- Data Poisoning:
 - GAN-synthetic data injected into training pipelines.
 - Corrupts learning process, weakening long-term detection accuracy.

Reference: [Bringing a GAN to a Knife-Fight: Adapting Malware Communication to Avoid Detection](#), [Generating Adversarial Malware Examples for Black-Box Attacks Based on GAN](#), [Mal-D2GAN: Double-Detector based GAN for Malware Generation](#), [VagueGAN: A GAN-Based Data Poisoning Attack Against Federated Learning Systems](#)